Polynomial Upper Bounds on Characteristic-Data Length for Multiple Context-Free Languages under Fixed Finite-Semigroup Typing

Takayuki Kuriyama
Independent Researcher, Tokyo, Japan
growup.kuriyama@gmail.com

Abstract

We study positive-data learning of multiple context-free languages that are tuple-substitutable under a fixed finite-semigroup typing and bound the total encoded length of characteristic samples relative to their MCFG presentations. Qualitative identification follows Yoshinaka's distributional framework; our contribution is quantitative. With the alphabet, fan-out, rule rank, and typing fixed, a reduced good grammar G of rule thickness $\bar{\theta}_G$ admits characteristic data of length $O(\|G\|^2\bar{\theta}_G)$ for prefix–suffix-window typings combined with any finite, possibly noncommutative, group, and of length $O(\|G\|^3\bar{\theta}_G)$ when the positive typing image is a local group. The proofs shorten interleaved tuple components simultaneously while preserving ordered group products. A further refinement-thickness dichotomy characterizes local groups as precisely the finite-semigroup regime admitting polynomially bounded full yield-type refinement thickness at every fixed fan-out; this is distinct from a lower bound on characteristic-data length. General input MCFGs are covered by parameterized bounds using component-pattern thickness, which may be exponential in grammar size. A separate conservative learner gives Gold identification. We exhibit a class-wide finite-feature bias for cross-serial languages and explicit characteristic samples for that family. Whether the local-group condition is necessary for polynomial characteristic-data length at fan-out two or more remains open.

Keywords: grammatical inference; multiple context-free grammars; positive data; finite-semigroup typing; multidimensional substitutability; Gold identification; rule rank.

1 Introduction

Positive-data grammatical inference asks when a language can be identified in the limit from positive examples alone [1, 2]. Distributional learning answers this question by treating words, or tuples of words, that occur in common contexts as candidate constituents. The approach began with substitutable CFGs [3, 12] and was extended to multiple context-free grammars (MCFGs), whose tuple-valued nonterminals express discontinuous and cross-serial dependencies [13, 14, 29, 30]. Yoshinaka's finite-signature generalization [15] provides positive-data identification and, in the corresponding bounded-rank specialization, characteristic samples consisting of polynomially many words.

The question: how long must characteristic data be? A sample of polynomial *cardinality* may still contain exponentially long words. We therefore measure a finite sample K by its total encoded length $\|K\|_+ = \sum_{w\in K}(|w| + 1)$, not merely by $|K|$. This follows the polynomial *thick-data*

perspective reviewed by Eyraud, Heinz and Yoshinaka [10, Section 2.5.2]: the size of a grammar and the length needed to exercise its productive rules are distinct resources. Our aim is to bound the total encoded length of characteristic data relative to an MCFG presentation of the target. Qualitative identification is inherited from Yoshinaka's framework; the contribution of this paper is quantitative.

The setting: a fixed finite-semigroup typing. We study this question for learners equipped with a finite bias. Before any data are seen, we fix a fan-out bound f and a surjective finite-semigroup homomorphism $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$, called a *typing*. The learner compares observed *nonempty* tuple components only when their s-values agree; the target class $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ consists of the f-MCFLs that are tuple-substitutable under this restriction (Definition 2.8). The canonical fresh-identity completion $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$ is used to type grammar yields and evaluate rule templates, including empty terminal factors. The substitutability condition itself compares nonempty tuple components using s. Fixing the typing in advance separates the choice of a finite bias from learning under that bias [5]. Prefix–suffix windows, values in a finite group, their products, and DFA transition morphisms are typical instances. Other biases restrict derivations by regular control sets [6, 7] or restrict the evidence used to compare factors [9, 11].

The typing is a *class-level* bias: it is supplied once, before the target is known, and is not chosen to recognize each target separately. Finite features, such as morphological compatibility or case licensing, can constrain substitution between discontinuous constituents while the MCFG accounts for unbounded dependencies. For example, fix one window–group typing $s = (s_{1,1}, g)|_{\Sigma^+_{\mathrm{feat}}}$ for *every* $R \subseteq \{0,1\}^2$, and consider

$$L_R = \{\mathtt{a}^m \mathtt{p}_i \mathtt{b}^n \mathtt{c}^m \mathtt{q}_j \mathtt{d}^n : m, n \geq 1,\ (i,j) \in R\}.$$

Every L_R belongs to $\mathcal{C}^{\mathrm{mcf}}_{2,s}$, and for nonempty R these are non-context-free cross-serial languages (Corollary 6.4). The group component g records the two feature *values*, not the compatibility relation R; neither R nor a target grammar is an input to the learner, which must recover both the licensing information and the unbounded crossing dependencies from positive examples. Explicit characteristic samples of encoded length at most $36|R| + 22$ exist for $R \neq \emptyset$, while every sample characteristic for a single set-driven operator on the whole family must exhibit every licensed pair (Proposition 6.5). The nonrectangular instance $R = \{(0,0), (0,1), (1,0)\}$ shows that untyped and boundary-only substitutability both fail here (Proposition 6.3).

Why bound the original presentation? A natural route is to refine the grammar by types and apply untyped estimates to the refinement. This route can be expensive. The full typed construction may contain as many as $|V||\mathsf{T}^\sharp|^f$ nonterminals and $|P||\mathsf{T}^\sharp|^{fr}$ rules, and its thickness can explode. This is a *potential* cost, not a monotonicity theorem for minimal characteristic-data length, but it can be exponential: for a fixed two-element typing $s_{\mathtt{c}}$ there are rank-two CFG presentations G_n of $\{\mathtt{a}, \mathtt{c}\}^+$ with $\|G_n\| = O(n)$ and rule thickness $\theta_{G_n} \leq 2$, whose trimmed full yield-type refinement has thickness $\theta_{\widetilde{G}_{n,0}} \geq 2^n$ (Proposition 4.5). We therefore bound characteristic data directly in terms of the original presentation, and ask when the algebra of the fixed typing permits short witnesses *before* this cost is paid. For the same reason we treat the typing as given rather than something to infer or optimize: a finer typing $s \preceq t$ (that is, $\ker t \subseteq \ker s$ on Σ^+) enlarges the class, $\mathcal{C}^{\mathrm{mcf}}_{f,s} \subseteq \mathcal{C}^{\mathrm{mcf}}_{f,t}$, but it also changes the reconstruction operator and may leave the algebraic regimes covered by our bounds. We ask when a fixed, externally given typing has class-uniform short witnesses.

The technical obstacle. To obtain short representatives of prescribed root types, we replace selected long subderivations by shorter derivations that need not preserve the componentwise group values of the replaced subtrees. In a multi-component MCFG derivation, one subtree may contribute several blocks to different root components, interleaved with blocks from other subtrees. For noncommutative groups, the effect of simultaneous replacements depends on the order of these blocks. The key quantitative step is an *owner-wise simultaneous selection* of replacements that preserves the ordered group product of every root component, even though individual replacement blocks may change their group values (Proposition 4.9).

Main results. Two input-sensitive thickness scales appear. For a reduced good MCFG G, the primary scale is its *rule thickness* $\bar{\theta}_G = \max\{1, \theta_G\}$: the maximum, over rules, of the shortest accepted sentence using that rule (Section 4.1). For an arbitrary standard MCFG H, the normalization bounds use instead the *component-pattern thickness* $\bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})$, which measures the shortest way to realize each feasible pattern of nonempty tuple components (Definition 4.23). The two scales are not interchangeable for inputs with empty components; Table 1 compares them with the auxiliary quantities used in the proofs. With the alphabet, fan-out f, rule-rank bound $r_{\max}$ and typing fixed, write $\bar{\theta} = \bar{\theta}_G$ for a reduced good grammar G, and $n = \max\{1, \|H\|\}$, $t = \bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})$ for a standard input grammar H. The principal bounds on total encoded characteristic-data length for the set-driven Const/Link/Comp reconstruction operator are:

Fixed typing / rank	Reduced good $(\|G\|, \bar{\theta})$	Any standard input (n, t)	Source
Rule rank ≤ 1; any typing	—	$O(n^3)$	Thm. 5.1
Prefix–suffix window	$O(\|G\|^2\bar{\theta})$	$O(n^3 t)$	Thm. 4.7
Window with finite group	$O(\|G\|^2\bar{\theta})$	$O(n^3 t)$	Thm. 4.15
Positive image in **LG**	$O(\|G\|^3\bar{\theta})$	$O(n^4 t)$	Thm. 4.20

The finite group may be noncommutative. The arbitrary-input column follows from Corollary 4.25 via normalization controlled by component-pattern thickness (Proposition 4.24), and the rank-one row is a separate normalization argument. The bounds are *class-uniform*: they apply to every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ admitting a presentation of the stated form, with constants independent of the particular target. The hidden constants depend on f, $r_{\max}$, $|\Sigma|$ and the typing, and may grow rapidly with them; for noncommutative-group shortening they contain $M_{A,f} = U_f \max\{3, |A|^f\}^{2^{f^2}}$, which may grow doubly exponentially with f. The abelian case has sharper constants (Theorem 4.16). A separate conservative learner gives Gold identification.

The obstruction of Proposition 4.5 has an exact algebraic boundary. Corollary 4.21 shows that the typings whose positive image is a local group, and only those, give uniformly polynomial full-refinement thickness at every fan-out; unlike the characteristic-data theorems, this needs no target-specific substitutability hypothesis. It concerns refinement thickness, and is *not* a lower bound on characteristic-data length.

Further results include a local-group sufficient condition strictly beyond $\mathbf{LI} \vee \mathbf{Grp}$ (Proposition B.5), a rule-only computable witness budget (Proposition 5.3), and an intrinsically rank-two example (Theorem 6.13). Theorem 5.5 complements the class-uniform bounds: it gives a polynomial bound for targets that *have* a representation $L = \widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathsf{H})$ for an abelian-group morphism φ and a subgroup H, even when the auxiliary abelian group is infinite. This is not a bound for every member of $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, and neither the window–group nor the local-group bound assumes such a representation. A refinement s' of s may put an individual example into this form, but it changes the supplied typing and hence the reconstruction operator. The typing-class separations in Section 6 do not assert that either shortening technique is necessary for any single example.

Relation to prior work. The following comparison concerns the guarantees used in this paper; it does not attempt an exhaustive classification of prior learning models.

Source	Established scope	Present additional question
Yoshinaka (2008) [12]	(k,ℓ)-substitutable CFGs; characteristic-data size with a grammar-thickness parameter	Higher-fan-out, fixed-group, and LG-typed shortening
Yoshinaka (2009, 2011) [13, 14]	Multidimensionally substitutable MCFGs and positive-data learning; polynomial identification in the corresponding grammar regime	Original-thickness bounds under supplied finite-semigroup typing
Yoshinaka (2015) [15]	Finite-signature and disjoint-sort learning principles, including finite witness cardinality under fixed bounded-rank signatures	Encoded-length upper bounds that survive the typing refinement

The Const/Link/Comp construction specializes Yoshinaka's qualitative learning principle rather than claiming it as new. The finite-semigroup structure theory and local-group/delay representations used here are classical [36–38]. The qualitative learning principles and finite witness-cardinality results are due to Yoshinaka [14, 15]. Classical sliding-factor coordinates serve as input to our new shortening step rather than a separate novelty claim; the rank-one obstruction uses [21, 22]. Our independent contribution is the noncommutative owner-wise shortening, the resulting window–group and local-group bounds (Theorems 4.15 and 4.20), and the normalization controlled by component-pattern thickness. A related, separately submitted manuscript addresses finite-semigroup-typed CFG learning [25]; no lower-bound construction or sharpness theorem from that manuscript is reproduced or used as a premise here. Positive-data learning of other mildly context-sensitive formalisms is studied in [19], and MCFG learning with membership queries or dual context sets in [16, 18, 20]. A recent substitution lemma for MCFLs [23] is a necessary condition for membership, not a learning criterion.

Scope and organization. The characteristic-data guarantees apply to the specified reconstruction operator and the stated grammar presentations; no presentation-independent stage of Gold convergence is claimed. The bounds measure total encoded length, not an intrinsic minimum grammar size; for general standard inputs, component-pattern thickness may be exponential in grammar size. The polynomial exponents are not established as optimal. The local-group condition is sufficient at every fixed fan-out; whether it is necessary for polynomial characteristic-data length at fan-out two or more is open (Open Problem 7.1). Section 2 fixes the setting, Section 3 defines the reconstruction operator and its qualitative guarantees, Sections 4 and 5 prove the quantitative bounds, and Section 6 compares the typed classes with multidimensional substitutability and develops the examples. The appendices prove the technical shortening, normalization and reconstruction statements.

2 Preliminaries

Throughout, Σ is a finite alphabet, ε the empty word and $\Sigma^+ := \Sigma^* \setminus \{\varepsilon\}$. If $\Sigma = \emptyset$, the only languages are $\emptyset$ and $\{\varepsilon\}$ and are handled directly. For a finite $K \subseteq \Sigma^*$ put $\|K\|_+ := \sum_{w\in K}(|w|+1)$, and for a tuple $\vec{u} = (u_1, \ldots, u_d)$ put $\|\vec{u}\|_1 := \sum_i |u_i|$. Every polynomial bound below states whether it concerns cardinality or $\|\cdot\|_+$. The table fixes symbols used across sections; all others (sliding-factor labellings, auxiliary morphisms, padding lengths) are defined where they occur.

Symbol	Meaning
$G,\ H$	a (reduced good) MCFG; an arbitrary standard input MCFG
B, C, X, Y	nonterminals; X, Y also denote typed nonterminals where specified
$\|G\|$, θ_G, t_G	encoding size (Definition 2.5), rule thickness, nonterminal thickness
s, $\widehat{s}$, S	positive typing, its fresh-identity completion, its positive image
$\preceq$	kernel refinement (Definition 2.3), also between monoid morphisms on Σ^*
A	a finite group in group-theoretic lemmas (otherwise locally defined); A an abelian group in Section 5.3
g	a monoid morphism $\Sigma^* \to A$ into a finite group
T	a derivation tree

2.1 Positive-Data Learning

Definition 2.1 (Text, TxtEx identification, conservativeness)**.** Fix a pause symbol $\perp \notin \Sigma$. A *text* for $L \subseteq \Sigma^*$ is an infinite sequence over $L \cup \{\perp\}$ in which every element of L occurs; the *content* of a finite prefix is the set of elements of Σ^* occurring in it. A *learner* is a computable map from finite prefixes to grammars. It *TxtEx-identifies* L [1] if on every text for L its outputs are, from some stage on, one fixed grammar H with $L(H) = L$; it TxtEx-identifies a class if it does so for each member. It is *conservative* if it changes its hypothesis only on a datum $w \in \Sigma^*$ not generated by the current hypothesis. We fix a hypothesis H_0 with $L(H_0) = \emptyset$.

Definition 2.2 (Set-driven operator and characteristic sample)**.** A *set-driven reconstruction operator* $\mathcal{B}$ maps each finite $K \subseteq \Sigma^*$ to a grammar $\mathcal{B}(K)$ depending only on the set K; it is *sample-consistent* if $K \subseteq L(\mathcal{B}(K))$ for all K. A finite $\mathcal{F} \subseteq L$ is a *characteristic sample for $\mathcal{B}$ and L* if $L(\mathcal{B}(K)) = L$ for every finite K with $\mathcal{F} \subseteq K \subseteq L$ (cf. [4]). This constrains generated languages only: $\mathcal{B}(K)$ need not stabilize syntactically as K grows.

2.2 Fixed Finite-Semigroup Typing

Definition 2.3 (Typing, fresh completion, refinement)**.** A *typing* is a surjective semigroup morphism $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ onto a finite semigroup, given by the multiplication table of S and the images of the letters, so that $s(w)$ is computable in time $O(|w|)$; S is its *positive image*. (A morphism into a larger finite semigroup is replaced by its corestriction to the image.) Let $\mathsf{S}^\sharp := \mathsf{S} \uplus \{\mathbf{1}_\sharp\}$ be obtained by adjoining a *fresh* identity, even if S already has one, and let $\widehat{s} : \Sigma^* \twoheadrightarrow \mathsf{S}^\sharp$ be the monoid morphism with $\widehat{s}(\varepsilon) = \mathbf{1}_\sharp$ and $\widehat{s}|_{\Sigma^+} = s$. For $\vec{w} \in (\Sigma^*)^d$ put $\widehat{s}^{(d)}(\vec{w}) := (\widehat{s}(w_1), \dots, \widehat{s}(w_d))$; on $(\Sigma^+)^d$ this is $s^{(d)}(\vec{w})$. For typings s, t write $s \preceq t$ if $\ker t \subseteq \ker s$ on Σ^+, and let (s, t) be the typing $w \mapsto (s(w), t(w))$ onto its image in $\mathsf{S} \times \mathsf{T}$; thus $s, t \preceq (s, t)$. Write s_{triv} for the one-element typing.

The fresh completion evaluates grammar templates (including empty terminal gaps) and types tuple yields in the refinement of Section 3.2; its new identity is not a positive-word type. Algebraic conditions such as **LG** concern S, not $\mathsf{S}^\sharp$. A finite-monoid morphism $q : \Sigma^* \to N$ determines the positive typing $q|_{\Sigma^+}$: by Lemma E.1(i), it gives the same target class and set-driven reconstruction rules as the fresh completion of that typing. Lemma E.1(ii) additionally identifies their trimmed yield-type refinements for *good* grammars. This last identification is not asserted for arbitrary standard grammars with empty-component yields.

Definition 2.4 (Prefix–suffix typing)**.** For $k, \ell \geq 0$ put $m_0 := \max\{1, k + \ell\}$, and for $|w| \geq m_0$ let $\mathrm{pref}_k(w)$, $\mathrm{suf}_\ell(w)$ be the prefix and suffix of the indicated lengths (length-zero windows are ε). Define

$$\widehat{s}_{k,\ell}(w) = \begin{cases} w, & |w| < m_0, \\ \langle \mathrm{pref}_k(w), \mathrm{suf}_\ell(w) \rangle, & |w| \geq m_0, \end{cases}$$

with disjoint tags for short words and pairs, and put $M_{k,\ell} := \widehat{s}_{k,\ell}(\Sigma^*)$. For short x, y and pairs $\langle u, v\rangle, \langle u', v'\rangle$ define

$$x \cdot y := \widehat{s}_{k,\ell}(xy), \qquad x \cdot \langle u, v\rangle := \langle \mathrm{pref}_k(xu), v\rangle,$$
$$\langle u, v\rangle \cdot y := \langle u, \mathrm{suf}_\ell(vy)\rangle, \quad \langle u, v\rangle \cdot \langle u', v'\rangle := \langle u, v'\rangle.$$

The *boundary typing* is $s_{k,\ell} := \widehat{s}_{k,\ell}|_{\Sigma^+}$; a typing s with $s \preceq s_{k,\ell}$ is (k,ℓ)-*boundary-determined*, or briefly a *boundary typing*.

A case check on the lengths of $x, y \in \Sigma^*$ gives $\widehat{s}_{k,\ell}(xy) = \widehat{s}_{k,\ell}(x) \cdot \widehat{s}_{k,\ell}(y)$. Since $\widehat{s}_{k,\ell}$ is onto $M_{k,\ell}$, the product is associative and $\widehat{s}_{k,\ell}(\varepsilon)$ is an identity, so $\widehat{s}_{k,\ell}$ is a finite-monoid morphism. The short-word tag separates ε from $\widehat{s}_{k,\ell}(\Sigma^+)$, so $\widehat{s}_{k,\ell}$ is isomorphic to the completion $\widehat{s_{k,\ell}}$.

2.3 Multiple Context-Free Grammars

Definition 2.5 (Standard MCFG presentation and size)**.** An MCFG [29, 30] is $G = (V, \Sigma, P, S, \mathrm{fan})$ with finite V, $S \in V$, $\mathrm{fan}(S) = 1$ and $\mathrm{fan}(A) \geq 1$. A rule of rank $r \geq 0$ is $A \to \rho(B_1, \ldots, B_r)$ with template $\rho = (\alpha_1, \ldots, \alpha_{d_0})$, $d_0 = \mathrm{fan}(A)$, $d_j = \mathrm{fan}(B_j)$, where each α_i is a word over $\Gamma_\rho := \Sigma \uplus \Xi_1 \uplus \cdots \uplus \Xi_r$, $\Xi_j := \{\xi_j^1, \ldots, \xi_j^{d_j}\}$, and every variable occurs at most once in $\alpha_1 \cdots \alpha_{d_0}$ (*linearity*). For $\vec{u}_j \in (\Sigma^*)^{d_j}$, $\rho(\vec{u}_1, \ldots, \vec{u}_r)$ substitutes $u_{j,\kappa}$ for ξ_j^κ; for $r = 0$ the rule is a *constant rule*. The tuple languages $L_A(G) \subseteq (\Sigma^*)^{\mathrm{fan}(A)}$ form the least family closed under the rules, and $L(G) := \{w : (w) \in L_S(G)\}$. We write $A \Rightarrow_G^* \vec{u}$ for a *terminal derivation* (a finite derivation tree rooted at A, evaluated bottom-up); one from S is *successful*. G is an f-MCFG if $\mathrm{fan}(A) \leq f$ for all A, and has rule-rank cap $r_{\max}$ if all ranks are at most $r_{\max}$. Such a presentation, with no further restriction, is called *standard*. Its *size* is

$$\|G\| := |V| + \sum_{(A \to \rho(B_1, \ldots, B_r)) \in P} \Big(1 + r + \textstyle\sum_i |\alpha_i|\Big),$$

where $|\alpha_i|$ counts terminal and variable occurrences. For fixed fan-out and rule rank, this is the symbol-level grammar-size measure used in the quantitative bounds; it does not count binary encodings of nonterminal names.

$\mathsf{MCFL}(f, r)$ is the class of languages of MCFGs of fan-out $\leq f$ and rule rank $\leq r$; f-MCFL $= \bigcup_{r<\infty} \mathsf{MCFL}(f, r)$, and the minimum rule rank of a language is taken over all its fan-out-$\leq f$ presentations. For nonterminals of equal fan-out d, $A \to B$ abbreviates $A \to (\xi_1^1, \ldots, \xi_1^d)(B)$, and $A \to \vec{c}$ a constant rule. Rule variables are always ξ_j^κ; terminals in examples are in `typewriter` font. Fan-out and rank are distinct resources: no rank reduction preserves an arbitrary fixed fan-out [24, 32], so the reconstruction below assumes neither a binary normal form nor a known rank cap.

Definition 2.6 (Grammar properties)**.** A rule is *nondeleting* if every variable occurs exactly once in its template; a nondeleting rule is *nonpermuting* if for each j the variables $\xi_j^1, \ldots, \xi_j^{d_j}$ occur in this order in $\alpha_1 \cdots \alpha_{d_0}$, and *nonmerging* if moreover no α_i contains a factor $\xi_j^\kappa \xi_j^{\kappa+1}$. Following [14, Section 2.3], a rule is λ-*free* if no α_i is empty. A rule, template or grammar is *good* if it is λ-free, nondeleting, nonpermuting and nonmerging. A nonterminal A is *reachable* if there is a chain $S = A_0, \ldots, A_m = A$ in which each A_{i+1} is a child in a rule for A_i, and *productive* if $L_A(G) \neq \emptyset$; G is *reduced* if every nonterminal is both.

In a λ-free grammar every component of every derived tuple is nonempty (induction on derivations). λ-freeness concerns template components and is distinct from $\varepsilon \notin L(G)$.

2.4 Tuple Contexts and Substitutability

A *sentence context of arity d* is a word $E = v_0\Box_1 v_1 \cdots v_{d-1}\Box_d v_d$ with each hole occurring once, $v_0, v_d \in \Sigma^*$ and $v_1, \ldots, v_{d-1} \in \Sigma^+$; $\mathcal{E}_d$ is the set of these *admissible* contexts [14, Section 3.1]. $E[\vec{w}]$ fills $\Box_i$ by w_i, $|E|_{\text{fix}} := \sum_{i=0}^{d} |v_i|$, and E *accepts* $\vec{w}$ in L if $E[\vec{w}] \in L$. The *d-tuple distribution* of $\vec{x} \in (\Sigma^*)^d$ is $\mathcal{D}_L^{(d)}(\vec{x}) := \{E \in \mathcal{E}_d : E[\vec{x}] \in L\}$; at fan-out one it is the two-sided distribution $\{(u, v) : uxv \in L\}$.

Definition 2.7 (Tuple occurrence)**.** For $w = a_1 \cdots a_n$ number the cut positions $⓪a_1① \cdots a_n⓷$ and write $w[i : i'] := a_{i+1} \cdots a_{i'}$. An *arity-d tuple occurrence* in w is a sequence $0 \le i_1 < i'_1 < i_2 < i'_2 < \cdots < i_d < i'_d \le n$; it determines $\vec{x} = (w[i_1 : i'_1], \ldots, w[i_d : i'_d]) \in (\Sigma^+)^d$ and the context $E \in \mathcal{E}_d$ obtained by replacing these intervals by $\Box_1, \ldots, \Box_d$, so that $E[\vec{x}] = w$. For finite K, $\text{Obs}_d(K) := \{\vec{x} \in (\Sigma^+)^d : E[\vec{x}] \in K \text{ for some } E \in \mathcal{E}_d\}$ is the set of tuples arising from such occurrences, and $\text{Obs}(K) := \bigcup_{1 \le d \le f} \text{Obs}_d(K)$.

Definition 2.8 ((f, s)-tuple substitutability and the target class)**.** Let $f \ge 1$ and let s be a typing. A language $L \subseteq \Sigma^*$ is (f, s)*-tuple-substitutable* if for every $1 \le d \le f$ and all $\vec{x}, \vec{y} \in (\Sigma^+)^d$,

$$s^{(d)}(\vec{x}) = s^{(d)}(\vec{y}) \;\text{ and }\; \mathcal{D}_L^{(d)}(\vec{x}) \cap \mathcal{D}_L^{(d)}(\vec{y}) \neq \emptyset \implies \mathcal{D}_L^{(d)}(\vec{x}) = \mathcal{D}_L^{(d)}(\vec{y}).$$

$\mathcal{C}_{f,s}^{\text{mcf}}$ is the class of (f, s)-tuple-substitutable f-MCFLs; a member of $\mathcal{C}_{f,s}^{\text{mcf}}$ for a boundary typing s is *boundary-typed*. For $\vec{x}, \vec{y} \in (\Sigma^+)^d$ write $\vec{x} \equiv_{L,s}^{d} \vec{y}$ (or $\vec{x} \equiv_{L,s} \vec{y}$) if $s^{(d)}(\vec{x}) = s^{(d)}(\vec{y})$ and $\mathcal{D}_L^{(d)}(\vec{x}) = \mathcal{D}_L^{(d)}(\vec{y})$; this is an equivalence relation.

Empty components are never compared, following [14, Section 3.1 and footnote 3] (see [8] for CFGs), and $\varepsilon \in L$ is handled at the start symbol; otherwise already $\mathtt{a}^+$ would fail, since $\mathtt{a}$ and ε share a context but not their distributions. The bias consists of both s and f:

(a) If $s \preceq t$, equal t-types are equal s-types, so $\mathcal{C}_{f,s}^{\text{mcf}} \subseteq \mathcal{C}_{f,t}^{\text{mcf}}$; in particular $\mathcal{C}_{f,s_{\text{triv}}}^{\text{mcf}} \subseteq \mathcal{C}_{f,s}^{\text{mcf}}$. This compares classes only: the operators $\mathcal{B}_{f,s}$ and $\mathcal{B}_{f,t}$ of Section 3 differ, and no monotonicity of characteristic-data length in the typing is claimed.

(b) $\mathcal{C}_{f+1,s}^{\text{mcf}} \cap f\text{-MCFL} \subseteq \mathcal{C}_{f,s}^{\text{mcf}}$, but the class is not monotone in f: for trivial typing $\{w\#w^R\}$ is 1- but not 2-dimensionally substitutable [14, Section 3.1, Example 5].

(c) For every $\mathsf{Acc} \subseteq \mathsf{S}^\sharp$, $\widehat{s}^{-1}(\mathsf{Acc}) \in \mathcal{C}_{f,s}^{\text{mcf}}$: it is regular, and componentwise equal types give equal total types in every context.

Proposition 2.9 (Intersection and abelian filtering)**.** *Fix f and s.*

(i) If L_1, L_2 are (f, s)-tuple-substitutable, so is $L_1 \cap L_2$.

(ii) If L is (f, s)-tuple-substitutable, A is an abelian group, $\varphi : \Sigma^* \to \mathsf{A}$ a monoid morphism, $\mathsf{H} \le \mathsf{A}$ and $\gamma \in \mathsf{A}$, then $L \cap \varphi^{-1}(\gamma + \mathsf{H})$ is (f, s)-tuple-substitutable.

Whenever such an intersection is an f-MCFL, it belongs to $\mathcal{C}_{f,s}^{\text{mcf}}$; in particular so does every f-MCFL $\widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathsf{H})$. Such factorizations are not supplied to the learner.

Proof. (i) A context accepting equal-type $\vec{x}, \vec{y}$ in $L_1 \cap L_2$ accepts both in each L_i, so their distributions agree in each L_i, hence in the intersection. (ii) If some context E accepts $\vec{x}, \vec{y}$ in $\varphi^{-1}(\gamma + \mathsf{H})$, then $\delta := \sum_i (\varphi(y_i) - \varphi(x_i)) \in \mathsf{H}$, and by commutativity $\varphi(E'[\vec{y}]) - \varphi(E'[\vec{x}]) = \delta$ for every $E' \in \mathcal{E}_d$, so coset membership transfers in both directions. Thus $\varphi^{-1}(\gamma + \mathsf{H})$ is (f, s_{triv})-, hence (f, s)-tuple-substitutable; apply (i). The final membership claim also uses the type-recognizable-language observation in item (c) immediately before the proposition, and the stated f-MCFL assumption. □

3 Reconstruction and Qualitative Learning

This section fixes the reconstruction operator analysed later and records its qualitative guarantees, which are not claimed as independent novelty. For fixed f, rank cap and typing, the operator is a concrete form of Yoshinaka's learner for substitutable languages over a finite signature [15, Theorem 1], for the MCFG signature of [15, Example 3] with tuple sorts refined by s-types (cf. [15, Section 8]). The correspondence is

Here	Antecedent	Addition needed later
Prop. 3.10 (soundness)	[15, Lemma 2]	uncapped rank
Prop. 3.11 (completeness)	[15, Lemma 3]	arbitrary same-type representatives
Thm. 3.13, Cor. 3.16	[15, Theorem 1]	uncapped rank

The additions are why the statements are given for the concrete grammar, with routine proofs in Appendices D and E; no embedding into the signature framework is used as a premise. The uncapped operator has sample-dependent rank, outside the bounded-arity assumption of [15]. Representatives that are not derivable from their nonterminal are used in Propositions 5.4 and 5.3. Finally, [15] bounds the *number* of characteristic words by the number of target rules and nonterminals; their *length*, which this paper bounds, is not addressed there.

3.1 Set-Driven Reconstruction

All interaction between templates and contexts rests on one observation.

Lemma 3.1 (Induced child contexts)**.** *Let $\rho = (\alpha_1, \dots, \alpha_e)$ be a good template of rank r with child arities $d_1, \dots, d_r$, let $E \in \mathcal{E}_e$, let $1 \le j \le r$, and fix $\vec{u}_i \in (\Sigma^+)^{d_i}$ for $i \neq j$. Let E_j be the word obtained from $E[(\alpha_1, \dots, \alpha_e)]$ by replacing ξ_j^κ by $\square_\kappa$ and every other variable ξ_i^ι by $u_{i,\iota}$. Then $E_j \in \mathcal{E}_{d_j}$ and $E_j[\vec{x}] = E[\rho(\vec{u}_1, \dots, \vec{u}_{j-1}, \vec{x}, \vec{u}_{j+1}, \dots, \vec{u}_r)]$ for every $\vec{x} \in (\Sigma^*)^{d_j}$.*

Proof. The identity holds because substitutions commute. By linearity and nondeletion each $\square_\kappa$ occurs exactly once in E_j, and by nonpermutation in the order $\kappa = 1, \dots, d_j$, since E places $\alpha_1, \dots, \alpha_e$ from left to right. Consider ξ_j^κ and $\xi_j^{\kappa+1}$. If both lie in one α_i, nonmerging puts a terminal or a variable ξ_i^ι with $i \neq j$ between them, which becomes a nonempty word. Otherwise they lie in $\alpha_i, \alpha_{i'}$ with $i < i'$, and the material between them contains the separator $v_i \in \Sigma^+$ of E. So every internal separator of E_j is nonempty. □

Lemma 3.2 (Composition preserves equivalence)**.** *Let $L \subseteq \Sigma^*$, let ρ be a good template of rank r, and let $\vec{x}_j, \vec{u}_j \in (\Sigma^+)^{d_j}$ satisfy $\vec{u}_j \equiv_{L,s} \vec{x}_j$ for every j. Then $\rho(\vec{u}_1, \dots, \vec{u}_r) \equiv_{L,s} \rho(\vec{x}_1, \dots, \vec{x}_r)$.*

Definition 3.3 (Composition witness)**.** Let K be finite and $r \ge 1$. A *composition witness of rank r* in K is a triple $(\vec{z}, \rho, (\vec{u}_1, \dots, \vec{u}_r))$ in which ρ is a good template of rank r with output arity $e \le f$ and child arities $d_1, \dots, d_r \le f$, $\vec{z} \in \mathrm{Obs}_e(K)$, $\vec{u}_j \in (\Sigma^+)^{d_j}$, and $\vec{z} = \rho(\vec{u}_1, \dots, \vec{u}_r)$.

Equivalently, a composition witness is a segmentation of the components of an observed $\vec{z}$ into terminal gaps and nonempty intervals labelled by variables, subject to the good-template conditions. If $E[\vec{z}] \in K$, Lemma 3.1 gives $E_j[\vec{u}_j] = E[\vec{z}] \in K$ with E_j admissible, so every $\vec{u}_j \in \mathrm{Obs}(K)$. The children occupy disjoint, nonempty parts of $\vec{z}$, so

$$r \le \textstyle\sum_j \|\vec{u}_j\|_1 \le \|\vec{z}\|_1 \le \max_{w \in K} |w| \le \|K\|_+. \tag{3.1}$$

Definition 3.4 (Capped and uncapped reconstruction operators)**.** Fix f and a typing s. For $r_{\max} \ge 1$ let $\mathcal{B}_{f,s}^{\le r_{\max}}(K) = \widehat{G}_{\le r_{\max}}(K)$ be the grammar with start symbol $\widehat{S}$, a nonterminal $[\vec{x}]$ of arity d for each $\vec{x} \in \mathrm{Obs}_d(K)$, $1 \le d \le f$, and the rules

(Start) $\widehat{S} \to [(w)]$ for $w \in K \cap \Sigma^+$, and $\widehat{S} \to \varepsilon$ if $\varepsilon \in K$;

(Const) $[\vec{x}] \to \vec{x}$ for every $\vec{x} \in \mathrm{Obs}(K)$;

(Comp) $[\vec{z}] \to \rho([\vec{u}_1], \ldots, [\vec{u}_r])$ for every composition witness of rank $1 \le r \le r_{\max}$;

(Link) $[\vec{x}] \to [\vec{y}]$ for $\vec{x}, \vec{y} \in \mathrm{Obs}_d(K)$ with $s^{(d)}(\vec{x}) = s^{(d)}(\vec{y})$ and $E[\vec{x}], E[\vec{y}] \in K$ for some $E \in \mathcal{E}_d$.

For $K \neq \emptyset$ put $\mathcal{B}_{f,s}(K) := \widehat{G}(K) := \widehat{G}_{\le \|K\|_+}(K)$; $\mathcal{B}_{f,s}(\emptyset)$ has no rules. Fixed encoding orders make both operators deterministic and set-driven.

By (3.1), $\mathcal{B}_{f,s}(K)$ contains a Comp rule for every composition witness in K. Both operators are sample-consistent, via $\widehat{S} \to [(w)] \to (w)$ and $\widehat{S} \to \varepsilon$. All non-start templates are λ-free, so $\widehat{S} \Rightarrow^* \varepsilon$ only by the rule $\widehat{S} \to \varepsilon$. For trivial typing and $\varepsilon \notin K$, Comp/Const, Link and Start are the rule Types I–III of [14, Section 4.1], and $\mathcal{B}_{f,s}$ is the operator of the unbounded-rank learner $\mathcal{A}(p, *)$ of [14, Corollary 2].

3.2 Yield-Typed Refinement and Witnesses

Definition 3.5 (Yield-type refinement)**.** Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a finite nondeleting MCFG and s a typing. For a template ρ and $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{d_j}$, let $\mathrm{out}_\rho(\mathbf{t}_1, \ldots, \mathbf{t}_r) \in (\mathsf{S}^\sharp)^{d_0}$ have ith component the image of α_i under the monoid morphism $\Gamma_\rho^* \to \mathsf{S}^\sharp$ extending $a \mapsto \widehat{s}(a)$ and $\xi_j^\kappa \mapsto (\mathbf{t}_j)_\kappa$. The *full yield-type refinement* $G^{\widehat{s}}$ has a fresh start symbol $\widetilde{S}$, nonterminals $A_{\mathbf{t}}$ for $A \in V$ and $\mathbf{t} \in (\mathsf{S}^\sharp)^{\mathrm{fan}(A)}$, start rules $\widetilde{S} \to S_{(\eta)}$ for $\eta \in \mathsf{S}^\sharp$, and for every rule $A \to \rho(B_1, \ldots, B_r)$ (including $r = 0$) and all $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{\mathrm{fan}(B_j)}$ the rule $A_{\mathrm{out}_\rho(\mathbf{t}_1,\ldots,\mathbf{t}_r)} \to \rho(B_{1,\mathbf{t}_1}, \ldots, B_{r,\mathbf{t}_r})$. The *trimmed refinement* $\widetilde{G}_0$ keeps exactly the nonterminals and rules occurring in some successful derivation from $\widetilde{S}$. Unless $G^{\widehat{s}}$ is named, *typed nonterminals* and *typed rules* are the non-start nonterminals and rules of $\widetilde{G}_0$, and $L_X := \{\vec{u} : X \Rightarrow^*_{\widetilde{G}_0} \vec{u}\}$.

Typed rules have the templates of their sources, so $\widetilde{G}_0$ is good whenever G is; each L_X is nonempty, and its tuples have nonempty components if G is λ-free.

Proposition 3.6 (Yield-type invariants)**.** *(i) If $A_{\mathbf{t}} \Rightarrow^* \vec{u}$ in $G^{\widehat{s}}$, then $\widehat{s}^{(\mathrm{fan}(A))}(\vec{u}) = \mathbf{t}$, and erasing indices gives a G-derivation of $\vec{u}$.*

(ii) Every terminal derivation $A \Rightarrow^*_G \vec{u}$ lifts uniquely to $A_{\widehat{s}^{(\mathrm{fan}(A))}(\vec{u})} \Rightarrow^*_{G^{\widehat{s}}} \vec{u}$; a successful derivation of w lifts uniquely through $\widetilde{S} \to S_{(\widehat{s}(w))}$.

(iii) $L(\widetilde{G}_0) = L(G^{\widehat{s}}) = L(G)$.

Lemma 3.7 (Canonical contexts for typed nonterminals)**.** *Let G be good. Every typed nonterminal X of arity d has a context $\chi(X) \in \mathcal{E}_d$ with $\chi(X)[\vec{u}] \in L(G)$ for all $\vec{u} \in L_X$; if $\widetilde{S} \to X$ is in $\widetilde{G}_0$, one may take $\chi(X) = \Box_1$.*

Lemma 3.8 (Typed nonterminals are distributionally homogeneous)**.** *Let G be good with $L := L(G) \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. For every typed nonterminal X of arity d, L_X lies in a single $\equiv^d_{L,s}$-class, denoted $[X]$.*

The proofs of these three statements are in Appendix E. The homogeneity of Lemma 3.8 is weaker than that of congruential MCFGs [17], whose contexts may have empty separators.

Definition 3.9 (Canonical tuples, contexts and the witness set)**.** Let G be good. For every typed nonterminal X fix a *canonical tuple* $\omega(X) \in L_X$ and a *canonical context* $\chi(X)$ as in Lemma 3.7, with $\chi(X) = \square_1$ whenever $\widetilde{S} \to X$ is in $\widetilde{G}_0$. Let $\mathsf{W}(\widetilde{G}_0)$ consist of the words

$$\chi(X)[\omega(X)] \ \ (X \text{ typed}), \qquad \chi(X)[\rho(\omega(Y_1), \ldots, \omega(Y_q))] \ \ (X \to \rho(Y_1, \ldots, Y_q) \text{ typed}, \ q \geq 0).$$

Since $\omega(X)$ and $\rho(\omega(Y_1), \ldots, \omega(Y_q))$ belong to L_X, $\mathsf{W}(\widetilde{G}_0)$ is a finite subset of $L(G)$, nonempty if $L(G) \neq \emptyset$, and $|\mathsf{W}(\widetilde{G}_0)|$ is at most the number of typed nonterminals and rules. Canonical objects are properties of the target, not inputs of the operator: qualitatively any choice works, while Section 4 chooses short ones.

3.3 Qualitative Guarantees

Proposition 3.10 (Soundness of every hypothesis)**.** *Let L be (f,s)-tuple-substitutable and $K \subseteq L$ finite. For every $r_{\max} \geq 1$, if $[\vec{x}] \Rightarrow^* \vec{u}$ in $\widehat{G}_{\leq r_{\max}}(K)$, then $\vec{u} \equiv_{L,s} \vec{x}$. Consequently $L(\mathcal{B}_{f,s}^{\leq r_{\max}}(K)) \subseteq L$ and $L(\mathcal{B}_{f,s}(K)) \subseteq L$.*

Proposition 3.11 (Completeness from a witness assignment)**.** *Let G be a finite good MCFG of fan-out $\leq f$, let r_G be the maximum rank of a typed rule of $\widetilde{G}_0$ ($r_G := 0$ if none), and assign to each typed nonterminal $X = A_{\mathbf{t}}$ a tuple $\iota(X) \in (\Sigma^+)^{\mathrm{fan}(A)}$ with $\widehat{s}^{(\mathrm{fan}(A))}(\iota(X)) = \mathbf{t}$ and a context $\chi(X) \in \mathcal{E}_{\mathrm{fan}(A)}$, with $\chi(X) = \square_1$ if $\widetilde{S} \to X$ is in $\widetilde{G}_0$. Suppose K contains $\chi(X)[\iota(X)]$ for each X and $\chi(X)[\rho(\iota(Y_1), \ldots, \iota(Y_q))]$ for each typed rule $X \to \rho(Y_1, \ldots, Y_q)$, $q \geq 0$. Then*

$$L(G) \subseteq L(\mathcal{B}_{f,s}^{\leq r_{\max}}(K)) \quad (r_{\max} \geq \max\{1, r_G\}), \qquad L(G) \subseteq L(\mathcal{B}_{f,s}(K)).$$

Neither substitutability nor $\iota(X) \in L_X$ is assumed; $\iota = \omega$ is the case $\mathsf{W}(\widetilde{G}_0) \subseteq K$.

Exact reconstruction needs a good presentation of the target; the following normalization controls size, fan-out and rank, but not thickness (for reduced λ-free inputs, see Proposition F.2).

Proposition 3.12 (Good normalization with rank preservation)**.** *Let G_{in} be a finite standard f-MCFG of rule rank at most $r_{\max}$. There is a finite good f-MCFG G_{good} of rule rank at most $r_{\max}$ with $L(G_{\mathrm{good}}) = L(G_{\mathrm{in}}) \setminus \{\varepsilon\}$ and, for fixed $f, r_{\max}$, $\|G_{\mathrm{good}}\| = O_{f,r_{\max}}(\|G_{\mathrm{in}}\|)$. If $L(G_{\mathrm{in}}) \cap \Sigma^+ \neq \emptyset$, G_{good} is reduced.*

Theorem 3.13 (Exact reconstruction for a fixed typing)**.** *Fix $f \geq 1$ and a typing s, and let $L \in \mathcal{C}_{f,s}^{\mathrm{mcf}}$. If $L \cap \Sigma^+ = \emptyset$ put $\mathcal{F}_L := L$ and $r_G := 0$. Otherwise let G be a reduced good fan-out-$\leq f$ presentation of $L \cap \Sigma^+$ (Proposition 3.12), let r_G be as in Proposition 3.11, and put $\mathcal{F}_L := \mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$. Then $\mathcal{F}_L \subseteq L$ is finite, and every finite K with $\mathcal{F}_L \subseteq K \subseteq L$ satisfies*

$$L(\mathcal{B}_{f,s}(K)) = L, \qquad L(\mathcal{B}_{f,s}^{\leq r_{\max}}(K)) = L \quad \text{for every } r_{\max} \geq \max\{1, r_G\}.$$

If $L \in \mathsf{MCFL}(f, r)$, G can be chosen of rank $\leq r$, so $r_{\max} = \max\{1, r\}$ suffices (also when $r = 0$).

Theorem 3.14 (Polynomial construction under a rule-rank cap)**.** *Fix $f, r_{\max} \geq 1$ and s. The encoding of $\widehat{G}_{\leq r_{\max}}(K)$ is computable in time $(1 + \|K\|_+)^{O(f r_{\max})}$; its fan-out is at most f and its rank at most $r_{\max}$.*

Definition 3.15 (Conservative sequential learner)**.** The learner $\mathcal{A}_{f,s}$ (respectively $\mathcal{A}_{f,s}^{\leq r_{\max}}$) starts with H_0, keeps its hypothesis on $\perp$ and on every datum the hypothesis generates, and otherwise replaces it by $\mathcal{B}_{f,s}(K)$ (respectively $\mathcal{B}_{f,s}^{\leq r_{\max}}(K)$), where K is the content read so far. This is the policy of [14, Algorithm 1].

Corollary 3.16 (Gold identification)**.** *Fix $f \geq 1$ and s.*

(i) *$\mathcal{A}_{f,s}$ TxtEx-identifies $\mathcal{C}^{\mathrm{mcf}}_{f,s}$.*

(ii) *For fixed $r_{\max}$, $\mathcal{A}^{\leq r_{\max}}_{f,s}$ TxtEx-identifies $\mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathsf{MCFL}(f, r_{\max})$, and processes each datum in time polynomial in $\|K\|_+$ for the content K read so far.*

Proofs are in Appendix D (Proposition 3.12: Appendix F). None of these statements bounds the length of characteristic data or the stage of convergence.

4 Characteristic-Data Bounds

Under a fixed rule-rank cap the typed witness construction already gives polynomially many characteristic words; the problem is to bound their total encoded length. Full yield-type refinement can have exponentially greater rule thickness than the original presentation (Proposition 4.5). The proofs instead work on the original grammar. Window types are preserved using a bounded derivation skeleton (Lemma 4.6); possibly noncommutative group products are preserved by simultaneous heavy-subtree replacement (Proposition 4.9 and Lemmas 4.10–4.11). Sliding-factor coordinates extend these arguments to local-group typings (Subsection 4.5), while component-pattern thickness controls normalizations of general input grammars (Subsection 4.6).

4.1 Primary thickness scales and the witness-length template

For the quantitative results, presentations are reduced and *good*, as in Yoshinaka's fixed-dimension, fixed-rank good form [14, Section 2.3, after Example 2]. For a rule R of G, define its *rule thickness* by

$$\theta_G(R) := \min\{\, |w| : w \in L(G) \text{ has a successful derivation using } R \,\}.$$

Thus length means the number of terminal letters in the *derived sentence*, not the number of steps of its derivation. Put $\theta_G := \max_{R \in P} \theta_G(R)$, $\bar{\theta}_G := \max\{1, \theta_G\}$, and $t_G := \max_A \min\{\|\vec{w}\|_1 : \vec{w} \in L_A(G)\}$. Thus θ_G is rule thickness and t_G nonterminal thickness; Yoshinaka uses the notation τ_G for our θ_G [14, Section 4.3]. We keep θ_G throughout this paper. The two primary parameters and three auxiliary measures are organized below; all comparisons involving $\widetilde{G}_0$ refer to its trimmed, nonempty yield-type refinement of a reduced good grammar.

Table 1: Thickness scales: two theorem parameters and three auxiliary quantities.

Scale	What it measures / role	Comparison or obstruction
θ_G	Shortest accepted sentence using each rule, maximized over rules; *primary for reduced good G.*	$t_G \leq \theta_G \leq 2\|G\| \max\{1, t_G\}$ (Lemma 4.1).
t_G	Largest minimum total yield length of a nonterminal; proof-scale for good forms.	Polynomially equivalent to θ_G on reduced good presentations.
$t_{\mathrm{pat}}(H)$	Largest minimum yield length over *realized nonempty-component support patterns*; *primary for arbitrary H.*	$t_H \leq t_{\mathrm{pat}}(H)$; the family E_n has $\theta_{E_n} = O(n)$ but $t_{\mathrm{pat}}(E_n) \geq 2^n$ (Prop. F.3).
θ_{typ}	Rule thickness of the trimmed full yield-type refinement $\widetilde{G}_0$; baseline only.	$\theta_{\mathrm{typ}} \geq \theta_G$; it can be $\geq 2^n$ with $\theta_G \leq 2$ (Prop. 4.5).
κ_G	Shortest representative of each typed distribution class, maximized over classes; auxiliary for abelian-filtered targets.	$\kappa_G \leq \theta_{\mathrm{typ}}$; representatives need not be derivable from the typed nonterminal (Prop. 5.4).

Here t_H uses the same nonterminal-minimum formula as t_G, without a goodness assumption; the formal definition of t_{pat} appears in Section 4.6. In a reduced good (nonempty-component) grammar there is only the full support pattern for each productive nonterminal, so $t_{\mathrm{pat}}(G) = t_G$. The refined-thickness and distributional inequalities compare the same reduced good grammar with its own trimmed refinement, not two unrelated presentations. In particular, neither θ_{typ} nor κ_G is an alternative input parameter in the main theorems. If a target contains ε, every statement below is applied to its λ-free core and the single datum ε is added, exactly as in Theorem 3.13; this adds one to the encoded length.

Derivational sentence contexts. A sentence context for a typed nonterminal X is *derivational* if it results from a successful typed derivation by deleting one occurrence of an X-subtree. Lemma 3.1 ensures that these contexts are admissible.

Lemma 4.1 (Thickness lemma)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG.*

(i) *Every rule writes at most θ_G terminal symbols explicitly.*

(ii) *$t_G \le \theta_G$; that is, every $A \in V$ has a tuple $\vec{u} \in L_A(G)$ with $\|\vec{u}\|_1 \le \theta_G$.*

(iii) *$\theta_G \le 2\|G\| \max\{1, t_G\}$.*

(iv) *Every $A \in V$ of fan-out d has a context $E_A \in \mathcal{E}_d$ with $|E_A|_{\mathrm{fix}} \le \theta_G$ and $E_A[\vec{u}] \in L(G)$ for every $\vec{u} \in L_A(G)$, and $E_S = \square_1$. For every finite typing $\widehat{s}$ and every non-start typed nonterminal $X = A_{\mathbf{t}}$ of $\widetilde{G}_0$, the choice $\chi(X) := E_A$ meets Definition 3.9, and E_A arises from a successful typed derivation by deleting an occurrence of an X-subtree.*

Proof. Fix a rule R with left-hand side A and a successful derivation D_R using R of length $\theta_G(R)$; it exists because G is reduced. By nondeletion, every terminal written by R and every component of the tuple produced at that occurrence of A survives exactly once in the final sentence. This gives (i), and also (ii) because every productive A has a rule. For (iv) with $A \neq S$, delete the subtree rooted at that occurrence of A and put ordered holes in place of its components; iterating Lemma 3.1 along the path to the root shows that the result E_A lies in $\mathcal{E}_d$, and its fixed material is part of the selected sentence. Substituting any $\vec{u} \in L_A(G)$ yields another successful derivation. Since $L_X \subseteq L_A(G)$ by Proposition 3.6(i), E_A accepts L_X; a typed nonterminal with a start link has the form $S_{\mathbf{t}}$, for which $E_S = \square_1$. Filling E_A's derivation with an X-derivation and lifting by Proposition 3.6(ii) places X at the hole, so E_A is derivational for X.

For (iii), follow a shortest dependency path from S to A; it repeats no nonterminal, so its rules are distinct and differ from R. Keep the path child open, complete every sibling and every child of R by a shortest terminal tuple (of length at most t_G), and apply R. The distinct rules used write at most $\|G\|$ terminals and have at most $\|G\|$ child occurrences, so the resulting sentence has length at most $\|G\| + \|G\| t_G \le 2\|G\| \max\{1, t_G\}$. □

Hence the two thickness parameters are polynomially equivalent on reduced good presentations. Yoshinaka's efficiency criterion asks for polynomial hypothesis construction, characteristic-set cardinality, and total positive data in target size and rule thickness [14, Definition 1]; see also [10]. Proposition F.2 transfers the size–thickness scale to reduced λ-free inputs; genuine empty components may instead cause the exponential growth of Proposition F.3.

All higher-rank bounds below are assembled by the following two lemmas.

Lemma 4.2 (Witness-length template)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \geq 1$, let $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$ be a finite typing, and let $\widetilde{G}_0$ be the trimmed $\widehat{s}$-refinement. If every non-start typed nonterminal has a canonical tuple with $\|\omega(X)\|_1 \leq b_{\mathrm{can}}$ and $\chi(X) := E_A$ is taken from Lemma 4.1(iv), then every word of $\mathsf{W}(\widetilde{G}_0)$ has length at most $2\bar{\theta}_G + r_{\max} b_{\mathrm{can}}$, and*

$$|\mathsf{W}(\widetilde{G}_0)| \leq |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}}, \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ \leq \big(|V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}}\big)\big(1 + 2\bar{\theta}_G + r_{\max} b_{\mathrm{can}}\big).$$

If moreover $L(G) = L \cap \Sigma^+$ for some $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, then $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}^{\leq r_{\max}}_{f,s}$.

Proof. A witness is $\chi(X)[\omega(X)]$, $\chi(X)[\vec{c}]$ or $\chi(X)[\rho(\omega(Y_1), \ldots, \omega(Y_q))]$. The context contributes at most $\bar{\theta}_G$ terminals, the source template at most $\bar{\theta}_G$ by Lemma 4.1(i), and there are at most $r_{\max}$ child tuples. The count follows because there is one witness per typed nonterminal and per typed rule, and each rule is determined by its source and the type vectors of its at most $r_{\max}$ children; multiply. The characteristic property is Proposition 3.11 (the cap $r_{\max}$ bounds every typed rank) together with Proposition 3.10 and the Start rule for ε. □

Lemma 4.3 (Type-preserving replacement)**.** *Let G be reduced and good, let $s : \Sigma^+ \to \mathsf{S}$ be a finite-semigroup typing and $\lambda : \Sigma^* \to N$ an auxiliary morphism with $s \preceq \lambda|_{\Sigma^+}$, and let $X = B_{\mathbf{t}}$ be a non-start typed nonterminal of the trimmed $\widehat{s}$-refinement $\widetilde{G}_0$. If $\vec{u} \in L_X$, $\vec{u}' \in L_B(G)$ and $\lambda^{(d)}(\vec{u}') = \lambda^{(d)}(\vec{u})$, then $\vec{u}' \in L_X$; in particular $\vec{u}'$ may be chosen as $\omega(X)$.*

Proof. Goodness makes all components nonempty, so $\widehat{s} \preceq \lambda$ gives $\widehat{s}^{(d)}(\vec{u}') = \mathbf{t}$. Proposition 3.6(ii) lifts a G-derivation of $\vec{u}'$ to a $G^{\widehat{s}}$-derivation rooted at X. Inserting it at an occurrence of X in a successful derivation of $\widetilde{G}_0$ gives a successful derivation, because typed rules depend only on typed symbols; hence all its rules survive trimming. □

4.2 Refined-Thickness Baseline

For a reduced good presentation G of a λ-free target, let $\widetilde{G}_0$ be its trimmed full yield-type refinement and put $\bar{\theta}_{\mathrm{typ}} := \max\{1, \theta_{\widetilde{G}_0}\}$, where $\theta_{\widetilde{G}_0}$ is rule thickness computed in the refined grammar. No comparison with the original untyped thickness is assumed.

Lemma 4.4 (Refined-thickness baseline)**.** *Fix $f \geq 1$, $r_{\max} \geq 1$, a finite alphabet Σ, and a finite typing $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$. Let G be a reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ presentation of a λ-free language in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, and let $\widetilde{G}_0$ be its trimmed yield-type refinement. Then the target witness set has cardinality polynomial in $\|G\|$ and total positive length polynomial in $\|G\|$ and $\bar{\theta}_{\mathrm{typ}} := \max\{1, \theta_{\widetilde{G}_0}\}$.*

Proof. $\widetilde{G}_0$ is reduced and good, so Lemma 4.1(i),(ii),(iv) applies to it with $\bar{\theta}_{\mathrm{typ}}$ in place of $\bar{\theta}_G$: canonical tuples and (derivational) canonical contexts of length at most $\bar{\theta}_{\mathrm{typ}}$ exist, with $\chi(X) = \Box_1$ for start links. Each typed rule writes at most $\bar{\theta}_{\mathrm{typ}}$ terminals and has at most $r_{\max}$ children, so every witness has length at most $(r_{\max} + 2)\bar{\theta}_{\mathrm{typ}}$, and

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \leq (|V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}})(1 + (r_{\max} + 2)\bar{\theta}_{\mathrm{typ}})$$

by counting the typed nonterminals and rules. □

For trivial typing no productive nonterminal has more than one yield type, so the original thickness scale is retained; in our notation this recovers Yoshinaka's untyped estimates $|K_G| \leq |P|$ and $\|K_G\| \leq |P|\theta_G$ [14, Lemma 7 and Theorem 1]. The quantitative claims below therefore concern typed, multi-component settings. Without uniqueness the refinement may be much thicker:

Proposition 4.5 (Exponential blow-up of full yield-type refinement)**.** *There are a fixed two-element typing $s_{\mathtt{c}}$, reduced good fan-out-one rank-two grammars G_n of size $O(n)$, and a constant bound $\theta_{G_n} \le 2$, such that $L(G_n) = \{\mathtt{a}, \mathtt{c}\}^+$ and the trimmed full yield-type refinement $\widetilde{G}_{n,0}$ satisfies $\theta_{\widetilde{G}_{n,0}} \ge 2^n$.*

Proof. Let $M_{\mathtt{c}} = \{1, \zeta\}$ with $\zeta^2 = \zeta$, $s_{\mathtt{c}}(\mathtt{a}) = 1$, and $s_{\mathtt{c}}(\mathtt{c}) = \zeta$. For $n \ge 1$, view

$$S \to SS \mid \mathtt{a} \mid \mathtt{c} \mid A_0 \mid \cdots \mid A_n, \qquad A_0 \to \mathtt{a}, \qquad A_i \to A_{i-1}A_{i-1} \mid \mathtt{c} \quad (1 \le i \le n)$$

as a fan-out-one MCFG. It is reduced and good, has size $O(n)$ and rule rank two, generates $\{\mathtt{a}, \mathtt{c}\}^+$, and every rule has a successful use of length at most two, so $\theta_{G_n} \le 2$. In the full refinement, the retained copy $(A_n)_1$ can derive only words without $\mathtt{c}$; recursively the unique such yield is $\mathtt{a}^{2^n}$. Hence the typed rule $S_{(1)} \to (A_n)_1$ has successful thickness at least 2^n. □

This is the refined-presentation obstruction summarized in Table 1, not a lower bound on characteristic-data length for $\{\mathtt{a}, \mathtt{c}\}^+$ or for arbitrary reconstruction operators. The later proofs therefore work directly with the original presentation.

4.3 Boundary Windows and Auxiliary Invariants

The refined thickness need not be polynomially controlled by the original good presentation (Proposition 4.5). A transfer is nevertheless possible whenever the typing is determined by a fixed prefix–suffix window together with an auxiliary finite invariant that can itself be preserved by short replacements. We use the thickness-sensitive representation model of Yoshinaka [14, Section 4.3].

Lemma 4.6 (Window skeletons)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \ge 1$. Fix $k, \ell \ge 0$, a morphism $\kappa : \Sigma^* \to N$ into a finite monoid, and a finite typing $\widehat{s} \preceq (\widehat{s}_{k,\ell}, \kappa)$. Let $\widetilde{G}_0$ be the trimmed $\widehat{s}$-refinement and n_{typ} its number of non-start typed nonterminals. Suppose that $\beta \ge 0$ is such that for every $C \in V$ and $\vec{x} \in L_C(G)$ there is $\vec{x}' \in L_C(G)$ with $\kappa^{(\mathrm{fan}(C))}(\vec{x}') = \kappa^{(\mathrm{fan}(C))}(\vec{x})$ and $\|\vec{x}'\|_1 \le \beta$. Then every non-start typed nonterminal X has a canonical tuple with*

$$\|\omega(X)\|_1 \le f(k+\ell+1)\, n_{\mathrm{typ}}\, (\bar{\theta}_G + r_{\max}\beta).$$

For trivial κ, $\beta = \bar{\theta}_G$ is admissible by Lemma 4.1(ii), so $\|\omega(X)\|_1 \le f(k+\ell+1)n_{\mathrm{typ}}(r_{\max}+1)\bar{\theta}_G$ whenever $\widehat{s} \preceq \widehat{s}_{k,\ell}$.

Proof. Choose a terminal derivation D from X in $\widetilde{G}_0$ with the minimum number of grammar-rule nodes, and let $\vec{u}$ be its root tuple. No typed nonterminal occurs twice on a root-to-leaf path: replacing the upper subtree by the lower one would give a derivation from X with fewer nodes. Hence every path has at most n_{typ} grammar nodes.

Mark in each component of $\vec{u}$ the terminals that determine its $\widehat{s}_{k,\ell}$-type (Definition 2.4): all terminals of a component shorter than $m_0 = \max\{1, k+\ell\}$, and otherwise its first k and last ℓ terminals (one arbitrary terminal when $k = \ell = 0$). There are at most $f(k+\ell+1)$ marks. Each terminal is written explicitly by exactly one node, so the union $\mathcal{K}$ of the root paths to the marked terminals has at most $f(k+\ell+1)n_{\mathrm{typ}}$ nodes. Call a child of a node of $\mathcal{K}$ that is not itself in $\mathcal{K}$ *omitted*. By linearity and nondeletion, each component of an omitted child is one contiguous block of exactly one root component, and this block contains no mark.

Replace each omitted subtree, with underlying nonterminal C and yield $\vec{x}$, by a G-derivation of a tuple $\vec{x}'$ as in the hypothesis, and forget all typed labels; this gives a G-derivation from the

underlying nonterminal B of X with root tuple $\vec{u}'$. A short root component consists only of marked terminals, so it receives no replaced block and is unchanged. In a long component every replaced block lies strictly between the marked first k and last ℓ terminals, and goodness keeps it nonempty; the marked prefix and suffix remain literal, and the component stays on the long side of the threshold. Hence $\widehat{s}_{k,\ell}$-vectors agree. Each root component is the ordered product of explicit terminals and child components, and every replaced block keeps its κ-value; so κ-vectors agree, with no commutativity used. Thus $(\widehat{s}_{k,\ell},\kappa)^{(d)}(\vec{u}') = (\widehat{s}_{k,\ell},\kappa)^{(d)}(\vec{u})$, and Lemma 4.3 gives $\vec{u}' \in L_X$. Each node of $\mathcal{K}$ writes at most $\bar{\theta}_G$ terminals (Lemma 4.1(i)) and has at most $r_{\max}$ omitted children, each of length at most β; the bound follows. □

Theorem 4.7 (Polynomial data for prefix–suffix-determined typings)**.** *Fix $f \geq 1$, $r_{\max} \geq 1$, k, ℓ, a finite alphabet Σ, and a finite-semigroup typing $s : \Sigma^+ \to \mathsf{S}$ satisfying $s \preceq s_{k,\ell}$; write $\widehat{s}$ for its canonical completion. Let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and let G be a reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar with $L(G) = L \cap \Sigma^+$. Then $\mathcal{B}^{\leq r_{\max}}_{f,s}$ is polynomial-time in the positive sample. A characteristic sample for L is $K = \mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$, with*

$$|\mathsf{W}(\widetilde{G}_0)| = O_{f,r_{\max},s}(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},k,\ell,s,\Sigma}(\|G\|^2 \bar{\theta}_G),$$

where $\bar{\theta}_G = \max\{1, \theta_G\}$. For a reduced λ-free input presentation G_{in} (deleting rules allowed), the structural good-form conversion of Proposition F.2 yields the explicit presentation-relative bound

$$O_{f,r_{\max},k,\ell,s,\Sigma}\Big(\|G_{\mathrm{in}}\|^3 \max\{1, t_{G_{\mathrm{in}}}\}\Big).$$

Proof. By Lemma E.1 we may use $\widehat{s}$. Lemma 4.6 with trivial κ gives $b_{\mathrm{can}} = f(k{+}\ell{+}1)n_{\mathrm{typ}}(r_{\max}{+}1)\bar{\theta}_G$, and $n_{\mathrm{typ}} \leq |V||\mathsf{S}^\sharp|^f = O(\|G\|)$. Lemma 4.2 bounds every witness by $O(\|G\|\bar{\theta}_G)$ and their number by $O(\|G\|)$, and gives the characteristic property; the datum ε adds at most one. Theorem 3.14 gives the construction time. For the input clause, Proposition F.2 gives $\|G\| = O(\|G_{\mathrm{in}}\|)$ and $\bar{\theta}_G = O(\|G_{\mathrm{in}}\| \max\{1, t_{G_{\mathrm{in}}}\})$; nonterminal thickness is used because deleting source rules may erase material before good-form conversion. □

Corollary 3.16 gives polynomial time per datum for the capped conservative learner. At fan-out one, polynomial time-and-thick-data learning for fixed windows is already due to Yoshinaka [12, Theorem 1]. As explained in the Introduction, our quadratic total-encoded-length bound is measured relative to a reduced good presentation; the degree in that presentation size is uniform for fixed window parameters, without claiming an improvement over Yoshinaka's bounds after normalization. The contribution of Lemma 4.6 is to preserve boundary information simultaneously across *multiple* tuple components at every fixed $f \geq 2$. Theorem 6.13 shows that this higher-rank regime is needed at the language level: it gives a non-context-free $(1,1)$-typed target of minimum fan-out two and minimum rule rank two over all finite fan-outs, whereas the rank-one examples are covered by Theorem 5.1.

4.4 Finite-Group Invariants

Window–group typings require short tuples with *prescribed* componentwise group values. The length bound is obtained from a single count of heavy nodes, which we state for an arbitrary finite-monoid invariant because Appendix B reuses it.

Realizable yield-type vectors. For a nonterminal B of arity d and a monoid morphism $\mu : \Sigma^* \to M$, a vector $\mathbf{t} \in M^d$ is *realizable at* B if $\mu^{(d)}(\vec{u}) = \mathbf{t}$ for some $\vec{u} \in L_B(G)$. For a positive-word typing we use its canonical fresh-identity completion; realizability does not mean that a representative has already been computed.

Lemma 4.8 (Heavy-node count)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \geq 1$, let $\lambda : \Sigma^* \to N$ be a morphism into a finite monoid, and let $b \geq \bar{\theta}_G$ and $c \geq 0$. Let T be a terminal derivation from $B \in V$ that minimizes first the total yield length and then the number of nodes among derivations from B with the same root λ-vector. Call a node* heavy *if its subtree yield has total length greater than b. If every set of pairwise incomparable heavy nodes of T has at most c elements, then*

$$\|\textit{root tuple of } T\|_1 \leq (r_{\max} + 1)\, b\, \max\{1, c\,|V|\,|N|^f\}.$$

Proof. Label each node by its nonterminal and the λ-vector of its yield; there are at most $|V|\,|N|^f$ labels. If an ancestor and a proper descendant had the same label, replacing the former subtree by the latter would keep the root λ-vector (templates evaluate homomorphically, as in Proposition 3.6), would not increase the yield (nondeletion) and would remove nodes. So every root-to-leaf path has at most $|V|\,|N|^f$ nodes. If the root is light, its yield has length at most b. Otherwise the heavy nodes form an ancestor-closed subtree, whose leaves are pairwise incomparable; hence there are at most $c|V|\,|N|^f$ heavy nodes. Every terminal of the root yield is written either by a heavy node (at most $\bar{\theta}_G \leq b$ per node, Lemma 4.1(i)) or inside a light child of a heavy node (at most $r_{\max}$ children of length at most b each). □

For a finite abelian group H write $\mathsf{d}(H)$ for the maximal length of a zero-sum-free sequence over H, so that $\mathsf{D}(H) = \mathsf{d}(H) + 1$ is its Davenport constant. For a finite group A and $f \geq 1$ put

$$N_{A,f} = \max\{3, |A|^f\}, \quad U_f = f^{f+1} f!, \quad M_{A,f} = U_f\, N_{A,f}^{2^{f^2}}, \qquad c_{A,f} := \begin{cases} \mathsf{d}(A^f), & A \text{ abelian}, \\ M_{A,f} - 1, & \text{otherwise}. \end{cases}$$

The lower floor $N_{A,f} \geq 3$ is required for the mixed-direction cases in the block-order decomposition; the $|A|^f$ term supplies repeated group boundary states.

Proposition 4.9 (Simultaneous block replacement: the exact interface)**.** *Let A be a finite group and $f \geq 1$. There are d ordered output rows ($1 \leq d \leq f$). Row c is a fixed-order product of q_c blocks, interspersed with unaltered constants $a_{c,0}, \ldots, a_{c,q_c} \in A$. A block $\beta_{c,j}$ has an owner $o(c,j) \in [m]$ and an original and a replacement value $x_{c,j}, x'_{c,j} \in A$. Every owner has at most f blocks over all rows. Different owners' blocks may interleave arbitrarily; all material not replaced is absorbed into the intervening constants. For $J \subseteq [m]$ put*

$$\mathrm{val}_c(J) = a_{c,0} y_{c,1} a_{c,1} \cdots y_{c,q_c} a_{c,q_c}, \qquad y_{c,j} = \begin{cases} x'_{c,j}, & o(c,j) \in J, \\ x_{c,j}, & o(c,j) \notin J. \end{cases}$$

With $c_{A,f}$ as defined above, if $m > c_{A,f}$ there exists a nonempty $J \subseteq [m]$ such that $\mathrm{val}_c(J) = \mathrm{val}_c(\varnothing)$ $(c = 1, \ldots, d)$. For abelian A, one may take $c_{A,f} = \mathsf{d}(A^f)$; for an arbitrary finite group, $c_{A,f} = M_{A,f} - 1$ suffices. Individual block or owner values need not be preserved.

Proof. This is Lemma A.2 of Appendix A, written with root-component notation. For abelian A, associate with each owner its difference vector in A^d and select a nonempty zero-sum subsequence. In a nonabelian group one must not commute factors. Instead fix a common owner pattern and

homogenize pairwise order comparisons by Lemma A.5; Lemma 4.10 then decomposes the output rows into consecutive rank chunks, and Lemma 4.11 selects a nonempty interval of owners preserving all ordered row products. The threshold is at most $M_{A,f}$. The detailed extraction and constants appear in Appendix A. □

Two noncommutative ingredients. The following combinatorial statements are used in the proof of Proposition 4.9. Here a block-replacement family is as formally specified in Definition A.1; the proofs, including opposite owner orders on different output components, are in Appendix A.1.

Lemma 4.10 (A homogeneous block-order lemma)**.** *Fix positive integers e, n with $n \geq 3$. For each $t \in [e]$ and $i \in [n]$, let (t, i) be a distinct point on a line, with a total left-to-right order. Suppose that (i) the order of (t, i) and (u, i) is independent of i; and (ii) for each ordered pair (t, u), including $t = u$, the order of (t, i) and (u, j) for $i < j$ is independent of the choice of $i < j$. Then the types at the points have a partition into classes, each occupying a consecutive region of the line. Every class has one direction (increasing or decreasing in i); inside its region the points occur in n successive* chunks*, one per rank in that direction. Each chunk consists of exactly one point of each type in the class (in a fixed within-chunk order).*

Lemma 4.11 (Simultaneous finite-group interval shortening)**.** *Let A be any finite group, and consider a block-replacement family of width f over A with $n \geq \max\{3, |A|^f\}$ indices in which all indices have the same number of blocks, the same lines for them and the same within-index order, and in which every cross-index comparison between the t-th block of index i and the u-th block of index j on a common line has the same truth value for all $i < j$. Then a nonempty interval J of indices satisfies $\mathrm{val}_c(J) = \mathrm{val}_c(\emptyset)$ for every line c.*

Theorem 4.12 (Short tuples with prescribed finite-group values)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and maximum rule rank $r_{\max} \geq 1$. Let $g : \Sigma^* \to A$ be a homomorphism into* any *finite group A. For every $B \in V$ and every realizable $\mathbf{s} \in A^{\mathrm{fan}(B)}$ there is $\vec{u} \in L_B(G)$ with $g^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{s}$ and*

$$\|\vec{u}\|_1 \leq \beta_g := (r_{\max} + 1)\bar{\theta}_G \max\{1, c_{A,f}|V||A|^f\}.$$

In particular $\beta_g = O_{A,f,r_{\max}}(\|G\|\bar{\theta}_G)$; when A is abelian the antichain constant is the zero-sum constant $\mathsf{d}(A^f)$.

Proof. Let T be minimal as in Lemma 4.8 with $\lambda = g$ and $b = \bar{\theta}_G$, and let $v_1, \ldots, v_m$ be pairwise incomparable heavy nodes. For each v_i, labelled C_i, fix a C_i-derivation of length at most $\bar{\theta}_G$ (Lemma 4.1(ii)); for $J \subseteq [m]$ let T_J replace the subtrees at v_i, $i \in J$. By linearity and nondeletion each component of v_i is a contiguous block of one root component, and blocks of different v_i are disjoint. The g-value of a root component of T_J is the ordered product of the g-values of its fixed segments and of its blocks, each block contributing its original or its replacement value according to whether its owner lies outside or inside J. This is a block-replacement family over A in the sense of Definition A.1. If $m > c_{A,f}$, Proposition 4.9 yields a nonempty J for which T_J has the root g-vector of T; but T_J is strictly shorter, since every replaced subtree has length greater than $\bar{\theta}_G$. This contradicts minimality, so Lemma 4.8 applies with $c = c_{A,f}$ and $N = A$. □

Lemma 4.13 (Marked heavy nodes for simultaneous window–group shortening)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \geq 1$. Fix $k, \ell \geq 0$, a finite group A, and a homomorphism $g : \Sigma^* \to A$. Put*

$\lambda = (\widehat{s}_{k,\ell}, g) : \Sigma^* \to N := M_{k,\ell} \times A$, $h := f(k+\ell+1)$, *and* $n_\lambda := |V||N|^f$. *For every* $B \in V$ *and realizable vector* $\mathbf{t} \in N^{\mathrm{fan}(B)}$, *there exists* $\vec{u} \in L_B(G)$ *with* $\lambda^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{t}$ *and*

$$\|\vec{u}\|_1 \le b_{\mathrm{mark}} := (r_{\max}+1)\bar{\theta}_G \max\{1, (c_{A,f}+h)n_\lambda\}.$$

Hence $b_{\mathrm{mark}} = O(\|G\|\bar{\theta}_G)$ *with* $f, r_{\max}, k, \ell, A, g, \Sigma$ *fixed. If* A *is abelian,* $c_{A,f}$ *can be taken as* $\mathsf{d}(A^f)$.

Proof. Among B-derivations with root componentwise λ-value $\mathbf{t}$, choose T minimizing total yield length, then node count. As in Lemma 4.8, a root–leaf path repeats no pair of a nonterminal and its componentwise λ-value; its number of nodes is at most n_λ.

Mark the first k and last ℓ terminal positions in each root component of length at least $m_0 := \max\{1, k+\ell\}$, and all positions in a shorter component. If $k = \ell = 0$, mark one position per component to ensure that nonemptiness is explicitly preserved. There are at most $h = f(k+\ell+1)$ marks in total.

Call a node heavy if its subtree yields more than $\bar{\theta}_G$ terminals. Fix an arbitrary antichain Q of heavy nodes. Since incomparable subtrees occupy disjoint sets of terminal positions, at most h nodes of Q contain marked terminals. Let Q_0 consist of all other nodes. For each $v \in Q_0$ with nonterminal C_v, fix a derivation producing a tuple of total length at most $\bar{\theta}_G$, which exists by Lemma 4.1(ii). Each replacement strictly shortens the original heavy subtree.

Suppose $|Q_0| > c_{A,f}$. By linearity and nondeletion, every node in Q_0 contributes at most f nonoverlapping blocks distributed among at most f ordered root components. Everything outside these subtrees remains fixed. Since g is a homomorphism, each root group value is the ordered product of the unchanged pieces and the original or replacement block values. Proposition 4.9 thus selects a nonempty set of owners whose simultaneous replacement preserves all root group values, without assuming commutativity.

No selected subtree contains a mark. Short root components are untouched, while the marked prefixes and suffixes of long components retain their letters and order. Every replacement component is nonempty by goodness. Consequently all root $\widehat{s}_{k,\ell}$-values are preserved. The root λ-vector is unchanged, while the total yield length strictly decreases, a contradiction to the choice of T. Every heavy antichain therefore has at most $c_{A,f}+h$ members. Lemma 4.8, with $b = \bar{\theta}_G$, $c = c_{A,f}+h$ and $N = M_{k,\ell} \times A$, gives the displayed bound. For abelian A use Lemma A.2(a). □

Example 4.14 (Two components with crossed owner order). Two owners with two blocks each may appear in the root tuple in the order $(\beta_{1,1}\beta_{2,1},\ \beta_{2,2}\beta_{1,2})$. Take $A = S_3$, put $\sigma = (12)$ and $\tau = (23)$, and assign value σ to the blocks of owner 1 and τ to those of owner 2. The root group vector is $(\sigma\tau, \tau\sigma)$, with $\sigma\tau \neq \tau\sigma$. Thus commuting owner contributions would be invalid. This illustrates a block-replacement arrangement, not a proof that any particular target language requires a nonabelian typing. Proposition 4.9 preserves the two ordered root products simultaneously.

Theorem 4.15 (Polynomial characteristic data for window–finite-group typing)**.** *Fix* $f, r_{\max}, k, \ell, \Sigma$, *a finite group* A, *a homomorphism* $g : \Sigma^* \to A$, *and a fixed positive-word typing* $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ *whose fresh-identity completion satisfies* $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$. *For every target* $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ *and reduced good fan-out-*$\le f$, *rank-*$\le r_{\max}$ *grammar* G *with* $L(G) = L \cap \Sigma^+$, *the capped reconstruction operator* $\mathcal{B}^{\le r_{\max}}_{f,s}$ *has a characteristic sample with*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},k,\ell,\widehat{s},g,\Sigma}(\|G\|^2\bar{\theta}_G).$$

Proof. Apply Lemma 4.13 to $\lambda = (\widehat{s}_{k,\ell}, g)$. Since $\widehat{s} \preceq \lambda$, every realizable $\widehat{s}$-typed non-start nonterminal $X = B_{\mathbf{s}}$ has an untyped B-tuple of the same root λ-value and total length at most

$b_{\mathrm{mark}} = O(\|G\|\bar{\theta}_G)$. Lemma 4.3 lifts its derivation to X, so $b_{\mathrm{can}} \le b_{\mathrm{mark}}$. The witness template (Lemma 4.2) gives $O(\|G\|)$ witnesses, each of length $O(\|G\|\bar{\theta}_G)$; the total encoded length is $O(\|G\|^2\bar{\theta}_G)$. Handle ε by its direct start rule. □

Theorem 4.16 (Polynomial data for window–abelian typing)**.** *Fix $f, r_{\max}, k, \ell, \Sigma$, a finite abelian group A, a homomorphism $g : \Sigma^* \to A$, and a fixed positive-word typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ whose fresh-identity completion satisfies $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$. For every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and reduced good fan-out-$\le f$, rank-$\le r_{\max}$ grammar G with $L(G) = L \cap \Sigma^+$, the capped operator $\mathcal{B}^{\le r_{\max}}_{f,s}$ has a characteristic sample satisfying*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},k,\ell,\widehat{s},g,\Sigma}(\|G\|^2\bar{\theta}_G).$$

In particular the bound uses the marked-node constant

$$b_{\mathrm{mark}} = (r_{\max}+1)\bar{\theta}_G \max\{1, (\mathsf{d}(A^f) + f(k+\ell+1))|V|(|M_{k,\ell}||A|)^f\}.$$

Proof. Apply the proof of Theorem 4.15 using the abelian case of Lemma 4.13. Lemma A.2(a) permits $c_{A,f} = \mathsf{d}(A^f)$ rather than $M_{A,f} - 1$, with $\mathsf{d}(A^f) = O(\exp(A) f \log|A|)$ (Remark A.3). □

Commutativity is used only in the zero-sum case of Lemma A.2. For nonabelian A, the order-homogenization proof of Appendix A.1 controls blocks interleaving across multiple root components. The finite-group block-replacement lemma and its full proof, including the nonabelian order-homogenization and interval-state arguments, are in Appendix A.

4.5 An Algebraic Criterion for Polynomial Characteristic Data

Write **LI** for the pseudovariety of finite locally trivial semigroups and **Grp** for the pseudovariety of finite groups. Their join $\mathbf{LI} \vee \mathbf{Grp}$ consists of divisors of finite products $S \times A$, with $S \in \mathbf{LI}$ and A a finite group [27, 28]. The following criterion makes the window–group hypothesis intrinsic to the positive image of the fixed typing. The semigroup identities used in its proof are classical [26].

Lemma 4.17 (Locally trivial–group typings admit window–group refinements)**.** *Let Σ be a nonempty finite alphabet and let $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ be a finite-semigroup typing; write $\widehat{s}$ for its canonical completion. If $\mathsf{S} \in \mathbf{LI} \vee \mathbf{Grp}$, then there exist $k, \ell \ge 0$, a finite group A, and a morphism $g : \Sigma^* \to A$ such that $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$ on all of Σ^*. If $\mathsf{S} \in \mathbf{LI}$, the group factor can be taken trivial, so $\widehat{s} \preceq \widehat{s}_{k,\ell}$.*

Proof. Choose a finite locally trivial semigroup S, a finite group A, a subsemigroup $T \subseteq S \times A$, and an onto semigroup morphism $\pi : T \to \mathsf{S}$. Lifting each $s(a)$ to some $(s_a, g_a) \in T$ defines a morphism $\rho : \Sigma^+ \to T$ with $\pi\rho = s$; let $\sigma : \Sigma^+ \to S$ and $g : \Sigma^* \to A$ (with $g(\varepsilon) = 1_A$) be its coordinates. The positive image $S_0 = \sigma(\Sigma^+)$ is locally trivial. Put $q = |S_0| + 1$. The finite locally trivial identities $S_0^q = S_0 E(S_0) S_0$ and $euf = ef$ for idempotents e, f and $u \in S_0$ [26] show that $\sigma(w)$ is determined by the length-q prefix and suffix of w when $|w| \ge 2q$: writing $w = pxz$ with $|p| = |z| = q$, $\sigma(p) = aeb$, $\sigma(z) = cfd$ and $\sigma^1(\varepsilon) = 1$, $\sigma(pxz) = ae\,(b\,\sigma^1(x)\,c)\,fd = aefd = \sigma(pz)$, where the parenthesized factor lies in S_0 because $b, c \in S_0$. Words shorter than $2q$ are recorded exactly by $\widehat{s}_{q,q}$. Hence equal $(\widehat{s}_{q,q}, g)$-values on Σ^+ give equal ρ- and $\widehat{s}$-values; the window typing separates ε from Σ^+. Take $k = \ell = q$. If $\mathsf{S} \in \mathbf{LI}$, use $S = \mathsf{S}$, $\pi = \mathrm{id}$ and the trivial group. □

Corollary 4.18 (Algebraic sufficient conditions for polynomial data)**.** *Fix a nonempty finite Σ, a fan-out cap $f \ge 1$, a rule-rank cap $r_{\max} \ge 1$, and a finite-semigroup typing $s : \Sigma^+ \to \mathsf{S}$; write $\widehat{s}$ for its canonical completion. Let G be any reduced good fan-out-$\le f$, rank-$\le r_{\max}$ presentation of a*

target in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, and put $\bar\theta_G = \max\{1, \theta_G\}$. The capped reconstruction operator has a characteristic sample of cardinality $O_{f,r_{\max},s}(\|G\|)$ with:

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},s,\Sigma}(\|G\|^2\bar\theta_G) \qquad \text{if } \mathsf{S} \in \mathbf{LI} \vee \mathbf{Grp}.$$

The locally trivial case **LI** *is included, with a trivial group factor.*

Proof. Combine Lemma 4.17 with Theorem 4.7 or Theorem 4.15; the refining parameters depend only on s and Σ, not on G. □

The local-group characterization and the detailed shortening argument appear in Appendix B. A finite semigroup belongs to **LG** (the *local groups*) if eSe is a group for every idempotent $e \in S$. This condition concerns the *positive image*, never the fresh-identity completion.

Definition 4.19 (Local-group typing)**.** Fix $k \geq 1$ and a finite group A. Let $\nu : \Sigma^{k+1} \to A$ be any map, the *sliding-factor labelling*. For $|w| \geq 2k$ put

$$\gamma_{k,\nu}(w) := \nu(F_1)\nu(F_2)\cdots\nu(F_{|w|-k}), \qquad \sigma_k(w) = (F_1, \ldots, F_{|w|-k}),$$

the ordered product in A, and put $\gamma_{k,\nu}(w) := 1_A$ for $|w| < 2k$. The *local-group typing* is $\lambda_{k,\nu}(w) := \big(\widehat{s}_{k,k}(w), \gamma_{k,\nu}(w)\big)$. A finite typing $\widehat{s}$ is *local-group-determined* if $\widehat{s} \preceq \lambda_{k,\nu}$ for some k, A and ν.

Theorem 4.20 (Polynomial characteristic data for local-group typing)**.** *Fix $f, r_{\max}, \Sigma$ and a finite-semigroup typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ with $\mathsf{S} \in \mathbf{LG}$; write $\widehat{s}$ for its canonical fresh-identity completion. For every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar G with $L(G) = L \cap \Sigma^+$, the capped operator $\mathcal{B}^{\leq r_{\max}}_{f,s}$ has a characteristic sample with*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},s,\Sigma}(\|G\|^3\bar\theta_G).$$

Proof. By the classical finite local-group representation, implemented in Theorem B.4 of Appendix B, there are fixed k, a finite group A and a sliding-factor labelling ν such that $\widehat{s} \preceq \lambda_{k,\nu}$. Lemma B.6 constructs, for each realizable $\lambda_{k,\nu}$-tuple type, a representative of length $O(\|G\|^2\bar\theta_G)$: first replace heavy subtrees by short window-compatible tuples, then select a nonempty simultaneous family whose ordered group products remain unchanged (Lemmas B.2 and A.2). Type-preserving replacement (Lemma 4.3) transfers these representatives to the typed refinement, which has $O(\|G\|)$ relevant rules. The witness-length template (Lemma 4.2) gives the cubic total-length bound. The coordinate construction and all intermediate constants are proved in Appendix B. □

The strict inclusion $\mathbf{LI} \vee \mathbf{Grp} \subsetneq \mathbf{LG}$ is proved in Proposition B.5. The Rees/sliding-factor coordinates are classical; their combination with owner-wise MCFG block shortening is the quantitative step here. Window–group typings retain the sharper quadratic bound of Theorem 4.15.

The algebraic boundary for full refinement thickness. For a surjective typing s and reduced good G, set $\theta_{\mathrm{typ}}(G, s) := \max\{1, \theta_{\widetilde{G}_0}\}$, where $\widetilde{G}_0$ is the trimmed full yield-type refinement. For a fixed $f \geq 1$, say a pseudovariety $\mathbf{V}$ has *polynomial refinement thickness at fan-out f* if, for every nonempty finite alphabet Σ, every rank cap $r_{\max} \geq 1$, and every surjective $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ with $\mathsf{S} \in \mathbf{V}$, a polynomial $q = q_{f,r_{\max},s,\Sigma}$ exists such that every reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar G over Σ satisfies

$$\theta_{\mathrm{typ}}(G, s) \leq q(\|G\|, \bar\theta_G).$$

The polynomial may depend on the fixed typing, but not on G or its language. Its degree in the following result is uniform for fixed $f, r_{\max}$.

Corollary 4.21 (Refinement-thickness dichotomy)**.** *For every $f \geq 1$ and every pseudovariety $\mathbf{V}$ of finite semigroups, polynomial refinement thickness at fan-out f holds if and only if $\mathbf{V} \subseteq \mathbf{LG}$. For $\mathbf{V} \subseteq \mathbf{LG}$, one may take*

$$\theta_{\mathrm{typ}}(G, s) = O_{f,r_{\max},s,\Sigma}(\|G\|^2 \bar{\theta}_G).$$

The equivalence also holds when the defining property is restricted to grammars G presenting $L \cap \Sigma^+$ for some $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. In fact, for each individual surjective typing whose positive image is not a local group, the polynomial bound already fails at fan-out one and rule rank two.

Proof. Suppose first that $\mathsf{S} \in \mathbf{LG}$. Theorem B.4 gives a fixed sliding-factor morphism $\lambda = \lambda_{k,\nu}$ with $\widehat{s} \preceq \lambda$. Let $X = B_{\mathbf{t}}$ be any retained non-start typed nonterminal. By productivity, some $\vec{u} \in L_X$ realizes $\mathbf{t}$. Its λ-vector is realizable at B, so Lemma B.6 supplies $\vec{u}' \in L_B(G)$ with the same λ-vector and $\|\vec{u}'\|_1 \leq \beta_\lambda = O(\|G\|^2 \bar{\theta}_G)$. Since $\widehat{s} \preceq \lambda$, its $\widehat{s}$-vector is also $\mathbf{t}$. The unique-lifting property (Proposition 3.6(ii)) therefore gives $\vec{u}' \in L_X$. This applies to every retained non-start typed nonterminal, including those with the original start symbol on the left-hand side.

Now take any retained non-start typed rule $\widetilde{R} : X \to \rho(Y_1, \ldots, Y_q)$ with source rule $R : B \to \rho(C_1, \ldots, C_q)$. Choose the preceding short terminal derivations for its children Y_i and apply $\widetilde{R}$ to obtain a derivation D from X with root yield $\vec{z}$ satisfying

$$\|\vec{z}\|_1 \leq \bar{\theta}_G + r_{\max}\beta_\lambda$$

by Lemma 4.1(i). If $\widetilde{S} \to X$ is a retained start link, this immediately gives a successful derivation using $\widetilde{R}$. Otherwise choose the context E_B of Lemma 4.1(iv), obtained by deleting a B-subtree from a successful original derivation, with $|E_B|_{\mathrm{fix}} \leq \bar{\theta}_G$. Filling this hole with the erasure of D gives a successful derivation in G. Its unique yield-typed lift restricts to precisely D at the inserted occurrence, because the typed labels and rules are determined by the terminal yields of their subtrees (Proposition 3.6(ii)). Thus the successful typed derivation uses $\widetilde{R}$ and has length at most $2\bar{\theta}_G + r_{\max}\beta_\lambda$. Every retained start-link rule similarly has a successful derivation of length at most β_λ. Therefore every rule of $\widetilde{G}_0$ has rule thickness $O_{f,r_{\max},s,\Sigma}(\|G\|^2 \bar{\theta}_G)$. No tuple-substitutability hypothesis entered this argument.

Conversely, the obstruction is stronger than required for the pseudovariety assertion: it occurs for *each fixed surjective* typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ with $\mathsf{S} \notin \mathbf{LG}$. By the classical U_1-exclusion criterion (Theorem B.4(i)–(ii)), there are distinct idempotents $e, z \in \mathsf{S}$ with $e^2 = e$, $ez = ze = z$, and $z^2 = z$. Choose fixed nonempty words $u, v \in \Sigma^+$ with $s(u) = e$ and $s(v) = z$; these exist by surjectivity. For each $n \geq 1$ consider the fan-out-one, rank-two presentation

$$S \to SS \mid a\ (a \in \Sigma) \mid A_0 \mid \cdots \mid A_n, \qquad A_0 \to u, \qquad A_i \to A_{i-1}A_{i-1} \mid v \quad (1 \leq i \leq n).$$

Here u and v are fixed terminal strings in the respective constant templates, not extra alphabet symbols. The grammar is reduced and good, has $\|G_n\| = O_s(n)$, generates Σ^+, and has $\bar{\theta}_{G_n} = O_s(1)$: every rule has a successful use in which each A_i is completed by v (or u for A_0), and a binary rule needs only two such bounded completions. On the other hand, $s(w) = e$ for a yield $w \in L_{A_n}(G_n)$ requires every leaf at which a constant rule is applied to use u, because any occurrence of v changes the value of the total product to the absorbing element z in $\{e, z\}$. Consequently the sole yield of $(A_n)_{(e)}$ is u^{2^n}, and the retained typed rule $S_{(e)} \to (A_n)_{(e)}$ can occur in a successful typed derivation only at sentence length at least $|u|2^n$. Hence $\theta_{\mathrm{typ}}(G_n, s) \geq |u|2^n$, contradicting every polynomial bound already at rank cap two. Since Σ^+ belongs to $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ for every $f \geq 1$, the same construction proves the version restricted to target presentations. Finally, if $\mathbf{V} \not\subseteq \mathbf{LG}$, choose any $\mathsf{S} \in \mathbf{V} \setminus \mathbf{LG}$ and apply this pointwise counterexample. The fixed- U_1 instance of Proposition 4.5 is the case of one-letter choices u, v. □

Remark 4.22 (Refinement thickness is not a data lower bound). The corollary concerns the thickness of a particular full yield-type refinement, not the minimal encoded length of characteristic positive data and not a lower bound for arbitrary Gold learners. The negative direction uses rank two; it does not assert failure in the rank-one fragment, where typed derivations have bounded-length acyclic chains. Polynomial refinement thickness gives a polynomial upper bound on data through Lemma 4.4, but the converse does not follow: short distributional representatives may exist even when full refinement is exponentially thick (Proposition 5.4). At fan-out two and above, **LG** remains a sufficient regime for short data, while its necessity for characteristic-data length is open (Open Problem 7.1).

4.6 Unrestricted Standard Input MCFGs

The preceding bounds use a reduced good presentation to make the quantitative witness argument transparent. Goodness, nondeletion, nonpermutation, nonmerging, and nonemptiness of tuple components need *not* be imposed on the grammar supplied for the target language. What matters when empty components are permitted is the cost of realizing each nonempty-component pattern; this is the second primary scale in Table 1, not a renaming of the source's rule thickness.

Definition 4.23 (Component-pattern thickness)**.** For a productive nonterminal A of a standard MCFG H, let $\mathrm{supp}(\vec{u}) := \{i : u_i \neq \varepsilon\}$ and $\mathsf{Pat}_H(A) := \{\mathrm{supp}(\vec{u}) : \vec{u} \in L_A(H)\}$. Put

$$t_{\mathrm{pat}}(H) := \max_{\substack{A \text{ productive} \\ I \in \mathsf{Pat}_H(A)}} \min_{\substack{\vec{u} \in L_A(H) \\ \mathrm{supp}(\vec{u}) = I}} \|\vec{u}\|_1, \qquad \bar{t}_{\mathrm{pat}}(H) := \max\{1, t_{\mathrm{pat}}(H)\},$$

with the maximum over an empty set taken as zero. In particular, the all-empty pattern costs zero.

Proposition 4.24 (Good normalization controlled by pattern thickness)**.** *Fix $f, r \geq 1$. Let H be any standard fan-out-$\leq f$, rank-$\leq r$ MCFG with $L(H) \cap \Sigma^+ \neq \emptyset$, and let H_{red} be its trimmed presentation. There exists a reduced good MCFG G for $L(H) \cap \Sigma^+$ with the same fan-out and rank bounds and*

$$\|G\| = O_{f,r}(\|H\|), \qquad t_G \leq t_{\mathrm{pat}}(H_{\mathrm{red}}), \qquad \bar{\theta}_G = O_{f,r}(\|H\| \bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})).$$

Proof. Use Proposition 3.12 and the stagewise tuple-language identities of Lemma F.1. Every retained nonterminal language is an image of a union of original support-pattern fibres under component deletion, permutation and concatenation. These maps do not increase total yield length, so a shortest representative of one contributing original pattern bounds each retained nonterminal's shortest yield by $t_{\mathrm{pat}}(H_{\mathrm{red}})$. Lemma 4.1(iii) then gives the thickness estimate. The precise fibre-transfer argument is in Appendix C. □

Corollary 4.25 (Total characteristic-data length from any standard MCFG)**.** *Fix $f, r \geq 1$, a finite alphabet Σ, and a fixed finite typing s. Let H be* any *finite standard fan-out-$\leq f$, rank-$\leq r$ MCFG (deleting rules, permutations, merging and empty components permitted) for $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. Put $n := \max\{1, \|H\|\}$ and $t := \bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})$. The capped reconstruction operator admits a characteristic sample $K \subseteq L$ of cardinality $O_{f,r,s}(n)$ with*

$$\|K\|_+ = \begin{cases} O(n^3 t), & \text{if } s \preceq s_{k,\ell} \text{ or } \widehat{s} \preceq (\widehat{s}_{k,\ell}, g) \text{ for a fixed finite-group morphism } g, \\ O(n^4 t), & \text{if } \mathsf{S} \in \mathbf{LG}. \end{cases}$$

The hidden constants depend on the fixed alphabet, typing, rank, fan-out, and (where applicable) window and group, not on H; an optional datum ε costs one.

Proof. For a nonempty positive core apply Proposition 4.24 and then Theorem 4.7, 4.15, or 4.20, respectively. The size is $O(n)$ and the new thickness is $O(nt)$. If $L \subseteq \{\varepsilon\}$, use $K = L$. □

5 Further Quantitative Results

This section collects quantitative results that complement the window, finite-group, and local-group regimes but are not used in their proofs: a typing-independent rank-one upper bound, effective construction of the short witnessed samples, and abelian-filtered distributional representatives.

5.1 Rule Rank One

Theorem 5.1 (Rank-one characteristic-data bound)**.** *Fix f, Σ, and a finite-semigroup typing s. For any standard rank-≤ 1, fan-out-$\leq f$ grammar G_{in} presenting a target in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, the operator $\mathcal{B}^{\leq 1}_{f,s}$ admits a characteristic sample K with*

$$\|K\|_+ = O_{f,s,\Sigma}(\|G_{\mathrm{in}}\|^3).$$

Proof. Normalize without increasing rank (Proposition 3.12) and trim, obtaining a reduced good rank-$\leq$ 1 grammar G with $\|G\| = O_f(\|G_{\mathrm{in}}\|)$, and trim its typed refinement. Let $n_{\mathrm{typ}} \leq |V||\mathsf{S}^\sharp|^f$ be the number of non-start typed nonterminals and $b_G := \max\{1, \text{largest explicit terminal count of a rule}\} \leq \max\{1, \|G\|\}$. A terminal derivation from X is a chain of unary rules ending in a constant rule. Choose one minimizing first the yield length and then the number of rules. No typed nonterminal repeats: splicing in the later suffix would not increase the yield (nondeletion) and would shorten the chain. Hence $\|\omega(X)\|_1 \leq n_{\mathrm{typ}} b_G$. For the context, take $\Box_1$ if $\widetilde{S} \to X$ is present; otherwise follow a shortest path $\widetilde{S} \to X_1 \to \cdots \to X_p = X$ in the typed dependency graph, so $p \leq n_{\mathrm{typ}}$. Composing its unary templates and putting ordered holes for the variables of X gives, by Lemma 3.1 iterated with no siblings, a derivational context $\chi(X) \in \mathcal{E}_{\mathrm{fan}(X)}$ with $|\chi(X)|_{\mathrm{fix}} \leq n_{\mathrm{typ}} b_G$ accepting all of L_X. Every witness word therefore has length at most $n_{\mathrm{typ}} b_G + b_G + n_{\mathrm{typ}} b_G$, and counting the typed rules and nonterminals gives

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \leq (n_{\mathrm{typ}} + |P||\mathsf{S}^\sharp|^f)(1 + (2n_{\mathrm{typ}} + 1)b_G) = O_{f,s}(\|G_{\mathrm{in}}\|^3).$$

This set, with ε when $\varepsilon \in L$, is characteristic for $\mathcal{B}^{\leq 1}_{f,s}$ by Propositions 3.11 and 3.10. □

Polynomial-time construction and Gold convergence follow separately from Theorem 3.14 and Corollary 3.16.

This extends the rank-one untyped scale [14, Lemma 7] to arbitrary fixed finite typing at every fixed fan-out. Rank two cannot generally yield a grammar-size-only bound: the size-$O(n)$ grammar $A_0 \to \mathsf{c}$, $A_i \to A_{i-1}A_{i-1}$, $S \to A_n$ generates $\{\mathsf{c}^{2^n}\}$, which is (f,s)-tuple-substitutable for every fixed f, s (over a one-letter alphabet, sharing a sentence context forces equal total component lengths), and every characteristic sample for it contains a word of length at least 2^n.

5.2 Computing Short Characteristic Data

The bounds above exhibit short canonical tuples and contexts. Here *derivational* has the meaning fixed in Subsection 4.1; all canonical contexts used in the main shortening results have this property (Lemma 4.1(iv)).

Proposition 5.2 (Effective short characteristic data)**.** *Fix $f, r_{\max}$ and a finite typing $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$. Given a reduced good MCFG G of fan-out at most f and rule rank at most $r_{\max}$, the following can be computed in time polynomial in $\|G\|$, with lengths written in binary:*

(i) $\theta_G(R)$ for every rule R, and hence θ_G and t_G;

(ii) the trimmed refinement $\widetilde{G}_0$ and, for each non-start typed nonterminal X, a canonical tuple $\omega(X) \in L_X$ of minimum total length and a derivational context $\chi(X)$ of minimum fixed length, with $\chi(X) = \square_1$ whenever $\widetilde{S} \to X$ belongs to $\widetilde{G}_0$;

(iii) derivation DAGs of size polynomial in $\|G\|$ representing these tuples and contexts, and hence every word of the witness set $\mathsf{W}(\widetilde{G}_0)$ of Definition 3.9 built from them.

Every word-length and total-length bound proved for $\mathsf{W}(\widetilde{G}_0)$ in Sections 4.1–4.5 and in Appendix B holds for this computed witness set. This does not apply to the relaxed samples $K^\dagger$ of Section 5.3, whose representatives need not be generated by their nonterminals. The words can be written out explicitly in time polynomial in $\|G\| + \|\mathsf{W}(\widetilde{G}_0)\|_+$.

Proof. For fixed $f, r_{\max}, s$ the full refinement $G^{\widehat{s}}$ has at most $|V||\mathsf{S}^\sharp|^f$ typed nonterminals and $|P||\mathsf{S}^\sharp|^{fr_{\max}}$ typed rules, and trimming is a linear-time productivity and reachability computation. View the typed rules as weighted deduction steps. A rule $X \to \rho(Y_1, \dots, Y_r)$ maps child weights to $(c_\rho, 1) + \sum_j (x_j, n_j)$, where c_ρ is the number of terminals written by ρ and pairs are compared lexicographically (terminal length, derivation-node count). This function is strictly superior to each argument even when no terminal is written, so the shortest-derivation algorithm in the weighted deductive form of Nederhof [35], building on Knuth's generalization of Dijkstra's method (see also Ramalingam and Reps [34]), computes $\ell_{\mathrm{in}}(X) = \min\{\|\vec{u}\|_1 : \vec{u} \in L_X\}$ together with an acyclic choice of minimizing rules, that is, a derivation DAG for $\omega(X)$. A second Dijkstra pass computes outside lengths, breaking equal-length ties by the number of edges: it starts with $\ell_{\mathrm{out}}(X) = 0$ when $\widetilde{S} \to X$ is present, and for every typed rule $X \to \rho(Y_1, \dots, Y_r)$ and every j adds an edge from X to Y_j of weight $c_\rho + \sum_{i \neq j} \ell_{\mathrm{in}}(Y_i)$. Then $\ell_{\mathrm{out}}(X)$ is the minimum fixed length of a derivational context of X, and the shortest-path tree supplies its DAG. The same two passes on G give $t_G = \max_A \ell_{\mathrm{in}}(A)$ and $\theta_G(R) = \ell_{\mathrm{out}}(A) + c_\rho + \sum_j \ell_{\mathrm{in}}(B_j)$ for $R : A \to \rho(B_1, \dots, B_r)$.

Each of the cited proofs bounds a witness word through the length of its context, of its rule template and of its canonical child tuples, and does so by exhibiting *some* tuple of L_X and *some* derivational context. The computed minima are no longer, and the number of witnesses is unchanged, so every stated bound persists. Expanding a DAG takes time linear in the length of the word it represents. □

Proposition 5.3 (Rule-only witnesses and a computable data budget)**.** *In the setting of Proposition 5.2, choose all canonical tuples by the consistent acyclic minimum-weight derivations described there, and let $\mathcal{R}$ be the non-start typed rules of $\widetilde{G}_0$ (rank-zero rules included).*

(i) For each non-start typed nonterminal X, let $R_X : X \to \rho_X(Y_1, \dots, Y_q)$ be the first rule of its chosen minimizing derivation (including $q = 0$). Then $\omega(X) = \rho_X(\omega(Y_1), \dots, \omega(Y_q))$. Consequently

$$\mathsf{W}(\widetilde{G}_0) = \{\chi(\mathrm{lhs}\, R)[\rho_R(\omega(Y_1), \dots, \omega(Y_q))] : R \in \mathcal{R}\}, \qquad |\mathsf{W}(\widetilde{G}_0)| \le |\mathcal{R}| \le |P||\mathsf{S}^\sharp|^{fr_{\max}},$$

and this set remains characteristic for $\mathcal{B}_{f,s}^{\le r_{\max}}$ whenever the target is in $\mathcal{C}_{f,s}^{\mathrm{mcf}}$.

(ii) Let $\ell_{\mathrm{in}}(X)$ be the minimum total yield length from X and $\ell_{\mathrm{out}}(X)$ the minimum fixed length of a derivational context of X (with $\ell_{\mathrm{out}}(X) = 0$ for a direct start link). For $R : X \to \rho(Y_1, \dots, Y_q)$ put

$$b(R) := c_\rho + \sum_{j=1}^{q} \ell_{\mathrm{in}}(Y_j), \qquad \mathsf{B}_{\mathrm{rule}}(G, \widehat{s}) := \sum_{R \in \mathcal{R}} (1 + \ell_{\mathrm{out}}(\mathrm{lhs}\, R) + b(R)).$$

Then $\|\mathsf{W}(\widetilde{G}_0)\|_+ \le \mathsf{B}_{\mathrm{rule}}(G,\widehat{s})$*, and the integer* $\mathsf{B}_{\mathrm{rule}}(G,\widehat{s})$ *is computable in time polynomial in* $\|G\|$ *for fixed* $f, r_{\max}, s$*, its value being represented in binary. If* $L(G)$ *also contains* ε *and its nonempty core is treated separately, the* ε *datum adds at most one to the budget.*

Proof. (i) An optimal derivation of X starts with R_X, and each of its child derivations is optimal, since replacing a nonoptimal child by the chosen minimum would decrease the lexicographic weight. The node-count tie-breaker makes the simultaneous choice acyclic, which gives the identity, also for rank zero. Hence $\chi(X)[\omega(X)]$ is *identical* to the rule word of R_X, and every other term of Definition 3.9 is a rule word; the two sets coincide as sets of terminal words, so completeness and the characteristic-sample property are unchanged. There is at most one word per typed rule. (ii) The rule word of $R : X \to \rho(Y_1, \dots, Y_q)$ has length exactly $\ell_{\mathrm{out}}(X) + b(R)$. Summing $|w| + 1$ over the indexed family gives $\mathsf{B}_{\mathrm{rule}}(G,\widehat{s})$; passing to distinct words can only lower it. The values $\ell_{\mathrm{in}}, \ell_{\mathrm{out}}$ and the number of typed rules are computed, with polynomial binary length, by Proposition 5.2. □

This rule-only budget is presentation-sensitive and can be much smaller than the worst-case polynomial bounds; it is an efficiently computable *upper certificate*, not a claim of a minimum-size characteristic sample or an improved worst-case exponent. For the singleton $\{\mathsf{c}^{2^n}\}$ after Theorem 5.1, explicit words are exponentially longer than $\|G\|$, but the derivation DAGs remain of polynomial size.

5.3 Distributional Representatives and Abelian-Filtered Targets

For an existential characteristic-sample bound one may instead choose short representatives of the same *typed distribution classes*, even when they have no derivation from the corresponding nonterminal.

Proposition 5.4 (Relaxed distributional witnesses)**.** *Let* G *be a reduced good fan-out-*$\le f$*, rank-*$\le r$ *presentation of* $L = L(G) \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$*, and let* $\widetilde{G}_0$ *be its trimmed yield-type refinement. For every non-start typed nonterminal* X*, let* $[X]$ *be the* $\equiv_{L,s}$*-class containing* L_X *(Lemma 3.8), and put* $\kappa_G(X) = \min_{\vec{v}\in[X]} \|\vec{v}\|_1$*,* $\kappa_G = \max_X \kappa_G(X)$*. There exists a characteristic positive sample* $K^\dagger$ *for both the capped operator* $\mathcal{B}^{\le r}_{f,s}$ *(with cap at least the rank of* G*) and the uncapped* $\mathcal{B}_{f,s}$ *satisfying*

$$|K^\dagger| \le |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{fr},$$
$$\|K^\dagger\|_+ \le (|V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{fr})(1 + \theta_G + \|G\| + (r+1)\kappa_G).$$

The bound concerns the nonempty core; if $\varepsilon \in L$*, add the sample word* ε *at encoded cost one.*

Proof. For each $X = A_{\mathbf{t}}$ choose $\omega^\dagger(X) \in [X]$ of total length $\kappa_G(X)$, and let $\chi^\dagger(X) := E_A$ from Lemma 4.1(iv), of fixed length at most θ_G ($\square_1$ for a start link). It accepts every tuple of L_X, hence also $\omega^\dagger(X) \equiv_{L,s} \vec{u} \in L_X$. Let $K^\dagger$ consist of the words $\chi^\dagger(X)[\omega^\dagger(X)]$, the rule words $\chi^\dagger(X)[\rho(\omega^\dagger(Y_1), \dots, \omega^\dagger(Y_q))]$ and the rank-zero rule words. By Lemma 3.2 the rule tuple is equivalent to a tuple of L_X, so $K^\dagger \subseteq L$; it also has the $\widehat{s}$-type of X. Apply the generalized Proposition 3.11 with $\iota(X) = \omega^\dagger(X)$ and $\chi(X) = \chi^\dagger(X)$: the first family of sample words is its representative evidence and the rule-word family is precisely its rule evidence. Componentwise types agree by the typed rules and Lemma 3.2; all contexts are admissible and start-linked contexts are identities. The proposition supplies the capped and uncapped derivation simulations. Proposition 3.10 gives the reverse inclusion for all larger positive samples. A nonterminal word has length at most $\theta_G + \kappa_G$ and a rule word at most $\theta_G + \|G\| + r\kappa_G$; sum the encoded costs. Unlike Proposition 5.3, the nonterminal words cannot in general be omitted. □

For fixed $f, r, \widehat{s}$, a family G_n with $\|G_n\| \max\{1, \theta_{G_n}\} = n^{O(1)}$ and superpolynomial characteristic-data requirements must therefore have $\kappa_{G_n} = n^{\omega(1)}$; this is a necessary condition, not a proof that such a family exists.

Theorem 5.5 (Polynomial data for abelian-filtered targets)**.** *Fix a finite nonempty alphabet Σ, a finite semigroup typing $s : \Sigma^+ \to \mathsf{S}$ and bounds $f, r \geq 1$. Write $\widehat{s}$ for the canonical completion and $\mathsf{S}^\sharp$ for its codomain. There are constants C, e depending only on Σ, s, f, r with the following property. Let A be any abelian group, $\varphi : \Sigma^* \to \mathsf{A}$ a morphism, $\mathsf{H} \leq \mathsf{A}$, $\gamma \in \mathsf{A}$, and $\mathsf{Acc} \subseteq \mathsf{S}^\sharp$. If $L = \widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathsf{H})$ has a reduced good fan-out-$\leq f$, rank-$\leq r$ MCFG presentation G of $L \cap \Sigma^+$, then $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and admits a characteristic sample for $\mathcal{B}^{\leq r}_{f,s}$ with*

$$\|K^\dagger\|_+ \leq C(\|G\| \max\{1, \theta_G\})^e.$$

The constants do not depend on $\mathsf{A}, \varphi, \mathsf{H}, \gamma$. The datum ε is included when $\varepsilon \in L$; the degenerate cases $L \subseteq \{\varepsilon\}$ are handled directly.

Proof. Membership in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ follows from Proposition 2.9 and the remark preceding it; the positive part has the assumed presentation, and ε is admitted by Start. By Proposition 5.4 it suffices to bound κ_G by a polynomial in $\|G\|\bar{\theta}_G$ with exponent depending only on $\Sigma, \widehat{s}, f, r$.

Let $\psi : \Sigma^* \to \mathbb{N}^\Sigma$ be the Parikh map, $\Psi(\vec{u}) = \sum_i \psi(u_i)$, and let $\overline{\varphi} : \mathbb{Z}^\Sigma \to \mathsf{A}$ be the group homomorphism through which φ factors; put $\Lambda = \{z \in \mathbb{Z}^\Sigma : \overline{\varphi}(z) \in \mathsf{H}\}$. Fix for each $C \in V$ a tuple $\omega_C \in L_C(G)$ with $\|\omega_C\|_1 \leq t_G \leq \theta_G$. The context E_C accepts every C-tuple, and all accepted words lie in one coset of H, so

$$\Psi(\vec{u}) - \Psi(\omega_C) \in \Lambda \qquad (\vec{u} \in L_C(G)). \tag{$*$}$$

For $R : C \to \rho(B_1, \ldots, B_q)$ let $c_R \in \mathbb{N}^\Sigma$ count its explicit letters and $\delta_R = \Psi(\omega_C) - c_R - \sum_j \Psi(\omega_{B_j})$. By (*) applied to $\rho(\omega_{B_1}, \ldots, \omega_{B_q}) \in L_C(G)$, $\delta_R \in \Lambda$, and $\|\delta_R\|_1 \leq T := (r+1)t_G + \|G\|$. Let $\Gamma_T = \{z \in \Lambda : \|z\|_1 \leq T\}$ and $\Lambda_T = \langle \Gamma_T \rangle$. Induction on derivations, using $\Psi(\vec{u}) - \Psi(\omega_C) = \sum_j (\Psi(\vec{u}_j) - \Psi(\omega_{B_j})) - \delta_R$ for $\vec{u} = \rho(\vec{u}_1, \ldots, \vec{u}_q)$, shows $\Psi(\vec{u}) - \Psi(\omega_C) \in \Lambda_T$ for every $\vec{u} \in L_C(G)$.

Now let $X = C_{\mathbf{t}}$ have arity $d \leq f$ and fix $\vec{u} \in L_X$. For each i choose, from a fixed semilinear representation of the Parikh image of the regular language $\widehat{s}^{-1}(t_i) \cap \Sigma^+$, a linear set $c_i + \sum_{p \in P_i} \mathbb{N}p$ containing $\psi(u_i)$; the number and size of all such generators depend only on Σ and $\widehat{s}$. The nonnegative integer system

$$\sum_{i=1}^{d} \Big(c_i + \sum_{p \in P_i} \eta_{i,p} p \Big) - \sum_{z \in \Gamma_T} (\mu^+_z - \mu^-_z) z = \Psi(\omega_C)$$

is feasible (use the Parikh vectors of the u_i and the Λ_T-membership of $\Psi(\vec{u}) - \Psi(\omega_C)$). It has $|\Sigma|$ equations, $O_{f,s,\Sigma}((T+1)^{|\Sigma|})$ variables, and coefficients and right-hand side of absolute value $O_{f,s,\Sigma}(T+1)$; in particular, each coordinate of $\Psi(\omega_C)$ is at most $t_G \leq \theta_G \leq T$ for the nonempty core. Papadimitriou's bound for small nonnegative integral solutions of equality systems [39] (a solution with all entries at most $n(ma)^{2m+1}$ for m equations, n variables and entries bounded by a) gives a solution whose entries are bounded by a polynomial in $T+1$ with exponent depending only on $|\Sigma|$. Realize each resulting Parikh vector by a nonempty word $v_i \in \widehat{s}^{-1}(t_i)$ and put $\vec{v} = (v_1, \ldots, v_d)$. Then $\|\vec{v}\|_1$ is polynomial in $T+1$, $\widehat{s}^{(d)}(\vec{v}) = \mathbf{t}$, and $\Psi(\vec{v}) - \Psi(\vec{u}) \in \Lambda_T \subseteq \Lambda$. For every sentence context E, the words $E[\vec{v}]$ and $E[\vec{u}]$ have the same $\widehat{s}$-value and $\varphi(E[\vec{v}]) - \varphi(E[\vec{u}]) = \overline{\varphi}(\Psi(\vec{v}) - \Psi(\vec{u})) \in \mathsf{H}$; hence $\vec{v} \equiv_{L,s} \vec{u}$ and $\kappa_G(X) \leq \|\vec{v}\|_1$. Since $T \leq (r+1)\theta_G + \|G\|$, the claim follows, with parameters independent of the auxiliary abelian group. □

Theorem 5.5 applies even when the positive image of $\widehat{s}$ lies *outside* **LG**: it is a sufficient condition on the target language, not on the finite control, and it claims no polynomial-time construction for

ineffectively presented auxiliary groups. The displayed abelian-coset representation is an additional target-side hypothesis. Unlike Theorems 4.15 and 4.20, this theorem does not give a bound for all members of $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ having the prescribed MCFG presentations.

6 Comparison with Multidimensional Substitutability

Yoshinaka's multidimensional substitutability is the direct positive-data baseline [13, 14]. Under trivial typing the present framework reproduces that semantic class; finite typing refines which observed tuples may be identified while preserving compositional rule witnesses, and $\widehat{s}_{k,\ell}$ lifts the boundary information of Yoshinaka's CFG (k,ℓ) hierarchy [12] componentwise. Since every (f,s) condition includes arity one, $\mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathrm{CFL} \subseteq \mathcal{C}^{\mathrm{mcf}}_{1,s}$.

6.1 Yoshinaka's Baseline and Prefix–Suffix Specialization

Definition 6.1 (Yoshinaka multidimensional substitutability)**.** Using the sentence-context sets $\mathcal{E}_d$ defined in Section 2, a language L is *fD-substitutable* if, for every $1 \le d \le f$, all $E, F \in \mathcal{E}_d$, and all $\vec{x}, \vec{y} \in (\Sigma^+)^d$, $E[\vec{x}], \quad E[\vec{y}], \quad F[\vec{x}] \in L \Longrightarrow F[\vec{y}] \in L$. Write $\mathsf{YSub}(f)$ for this semantic class and $\mathsf{YSL}(f,r) := \mathsf{YSub}(f) \cap \mathsf{MCFL}(f,r)$, where r bounds rule rank. We refer to the displayed four-instance implication as the *substitution implication*. Yoshinaka's original definition quantifies over every arity $d \le f$, uses nonempty tuple components, and requires nonempty internal separators in d-contexts; footnote 3 separately discusses the stronger alternative that permits empty components and separators [14, Section 3.1 and footnote 3].

The trivial typing removes the componentwise type premise, so $\mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{triv}}} = \mathsf{YSub}(f) \cap f\text{-MCFL} \subseteq \mathcal{C}^{\mathrm{mcf}}_{f,s}$. In Yoshinaka's word-level (k,ℓ) condition [12], with $u \in \Sigma^k$, $v \in \Sigma^\ell$ and $uy_1v, uy_2v \neq \varepsilon$,

$$w_1uy_1vw_1',\ w_1uy_2vw_1' \in L \quad \Longrightarrow \quad (w_2uy_1vw_2' \in L \iff w_2uy_2vw_2' \in L).$$

The compared fragments uy_iv are nonempty although the y_i may be empty, and equality of their $\widehat{s}_{k,\ell}$-values is exactly the boundary information used here, including exact identity of words shorter than $m_0 = \max\{1, k+\ell\}$. Hence $(1, \widehat{s}_{k,\ell})$-substitutability is precisely Yoshinaka's (k,ℓ) condition, and the analogous componentwise boundary requirement at higher fan-out is a definition of this paper, not a claim that Yoshinaka (2008) formulated an MCFG (k,ℓ) notion.

A language is *strictly 2-local* if membership is determined by permitted adjacent pairs in $\vdash w \dashv$, for fixed endmarkers.

Lemma 6.2 (Strictly 2-local constraints are boundary typed)**.** *If $L_{\mathrm{sl}} \subseteq \Sigma^*$ is strictly 2-local, then for every $f \ge 1$, $L_{\mathrm{sl}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}}$.*

Proof. Let equal-$\widehat{s}_{1,1}$-type nonempty tuples $\vec{x}, \vec{y}$ share an accepting context E_0. Internal adjacent pairs of each y_i are permitted because they occur in $E_0[\vec{y}]$. If $E[\vec{x}] \in L_{\mathrm{sl}}$, every adjacent pair of $E[\vec{y}]$ is either in fixed material, internal to a y_i, or crosses a component boundary. These are permitted respectively by $E[\vec{x}]$, by $E_0[\vec{y}]$, or by equality of the first/last letters encoded by the boundary type; end-marker pairs are unchanged. Thus $E[\vec{y}] \in L_{\mathrm{sl}}$, and symmetry finishes the proof. □

Proposition 6.3 (Cross-serial dependencies with finite feature licensing)**.** *Let* $\Sigma_{\mathrm{feat}} = \{\mathtt{p}_0, \mathtt{p}_1, \mathtt{q}_0, \mathtt{q}_1, \mathtt{a}, \mathtt{b}, \mathtt{c}, \mathtt{d}\}$, $\mathcal{R}_{\mathrm{feat}} = \{(0,0),(0,1),(1,0)\}$, *and*

$$L_{\mathrm{feat}} = \{\mathtt{a}^m\mathtt{p}_i\mathtt{b}^n\mathtt{c}^m\mathtt{q}_j\mathtt{d}^n : m,n \ge 1,\ (i,j) \in \mathcal{R}_{\mathrm{feat}}\}.$$

Let $\mathsf{F} = \mathbb{Z}_2 \times \mathbb{Z}_2$*, and define a morphism* $g : \Sigma^*_{\mathrm{feat}} \to \mathsf{F}$ *by* $g(\mathtt{p}_1) = (1,0)$*,* $g(\mathtt{q}_1) = (0,1)$*, and* $g(z) = (0,0)$ *for all other letters* z*. Fix the positive-image semigroup typing*

$$s = (s_{1,1}, g)|_{\Sigma^+_{\mathrm{feat}}},$$

where $s_{1,1}$ *is the first/last-letter morphism and the codomain is the finite positive image of this product. Then, for every* $f \geq 2$*,*

$$L_{\mathrm{feat}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s} \setminus \mathsf{YSub}(f), \qquad L_{\mathrm{feat}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}.$$

Moreover L_{feat} *is not even* $(1, s_{1,1})$*-substitutable: the finite group component is needed by this example.*

Proof. Two crossed dependencies. Use fan-out-two nonterminals A, B with

$$A \to (\mathtt{a}, \mathtt{c}) \mid (\mathtt{a}\xi^1_1, \mathtt{c}\xi^2_1)(A),$$
$$B \to (\mathtt{b}, \mathtt{d}) \mid (\mathtt{b}\xi^1_1, \mathtt{d}\xi^2_1)(B).$$

For each $(i,j) \in \mathcal{R}_{\mathrm{feat}}$ add the rank-two rule

$$S \to \xi^1_1 \mathtt{p}_i \xi^1_2 \xi^2_1 \mathtt{q}_j \xi^2_2 (A, B).$$

These rules generate exactly L_{feat}, using both components of each child once and ordering them as A_1, B_1, A_2, B_2. The grammar is nondeleting, nonmerging, and has fan-out and rank two. The central crossing-count language $\{\mathtt{a}^m\mathtt{b}^n\mathtt{c}^m\mathtt{d}^n : m, n \geq 1\}$ is not context-free: for a pumping length p, a window of length p in $a^p b^p c^p d^p$ affects at most two adjacent letter blocks, and pumping cannot preserve both a/c and b/d equalities. Intersecting L_{feat} with the regular language $\mathtt{a}^+\mathtt{p}_0\mathtt{b}^+\mathtt{c}^+\mathtt{q}_0\mathtt{d}^+$ and erasing $\mathtt{p}_0, \mathtt{q}_0$ would yield this non-context-free language, so L_{feat} is not context-free.

Finite group-valued feature licensing. Put $Q = \mathtt{a}^+\{\mathtt{p}_0, \mathtt{p}_1\}\mathtt{b}^+\mathtt{c}^+\{\mathtt{q}_0, \mathtt{q}_1\}\mathtt{d}^+$. The language Q is strictly 2-local: its allowed adjacent pairs and endmarkers determine precisely these blocks and their order. Define the regular feature-compatibility filter $K_{\mathrm{feat}} = g^{-1}(\mathcal{R}_{\mathrm{feat}})$, viewing each (i,j) as the corresponding element of F. Because s includes g, this filter is the inverse image of an accepting subset of the finite s-type space. Let $\varphi : \Sigma^*_{\mathrm{feat}} \to \mathbb{Z}^2$ map $\mathtt{a}, \mathtt{c}, \mathtt{b}, \mathtt{d}$ to $(1,0), (-1,0), (0,1), (0,-1)$, respectively, and map the feature markers to $(0,0)$. Then

$$L_{\mathrm{feat}} = Q \cap K_{\mathrm{feat}} \cap \varphi^{-1}(0).$$

Lemma 6.2 applies to Q because s refines $s_{1,1}$. The inverse-image observation after Definition 2.8 applies to K_{feat}, and Proposition 2.9 to the count filter. Their intersection is (f,s)-tuple-substitutable for every f.

Why the untyped and boundary-only implications fail. Take $\vec{x} = \mathtt{a}\mathtt{p}_0\mathtt{b}$, $\vec{y} = \mathtt{a}\mathtt{p}_1\mathtt{b}$, $E = \square_1\mathtt{c}\mathtt{q}_0\mathtt{d}$, $F = \square_1\mathtt{c}\mathtt{q}_1\mathtt{d}$. Then $E[x], E[y], F[x] \in L_{\mathrm{feat}}$, but $F[y] \notin L_{\mathrm{feat}}$, contradicting the arity-one substitution implication. These fragments have the same first/last-letter type $(\mathtt{a}, \mathtt{b})$, but $g(x) = (0,0) \neq (1,0) = g(y)$. Thus neither trivial typing nor the boundary window alone blocks the unsound substitution, whereas the fixed product typing does. □

Corollary 6.4 (Uniform finite feature licensing)**.** *The* same *fixed typing* $s = (s_{1,1}, g)|_{\Sigma^+_{\mathrm{feat}}}$ *from Proposition 6.3 works for every relation* $R \subseteq \{0,1\}^2$*: putting*

$$L_R = \{\mathtt{a}^m\mathtt{p}_i\mathtt{b}^n\mathtt{c}^m\mathtt{q}_j\mathtt{d}^n : m, n \geq 1,\ (i,j) \in R\},$$

we have $L_R \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ *for all* $f \geq 2$*. The typing is chosen once for the entire family, not tailored to* R*.*

Proof. The grammar in Proposition 6.3 retains exactly the start rules indexed by R (or has no productive start rule if $R = \varnothing$). Its proof of typed substitutability applies with $g^{-1}(R)$ in place of $g^{-1}(\mathcal{R}_{\text{feat}})$: each is an accepting union of fibres of the same fixed g. The other two filters do not depend on R. □

Proposition 6.5 (Explicit data budget for a common feature typing)**.** *Fix the typing s from Corollary 6.4. For every nonempty $R \subseteq \{0,1\}^2$ and $f \geq 2$, the capped operator $\mathcal{B}^{\leq 2}_{f,s}$ admits a characteristic sample $K_R \subseteq L_R$ with*

$$\|K_R\|_+ = 36|R| + 22.$$

For $R = \emptyset$, the empty sample is characteristic for the empty language. Conversely, let $\mathcal{D}$ be any set-driven operator admitting characteristic samples for all *L_R in this family. Every characteristic sample K for L_R with respect to $\mathcal{D}$ satisfies*

$$\|K\|_+ \geq 7|R|.$$

The lower bound applies to the common-family learning requirement; it does not assert a lower bound for a learner designed for one known relation R.

Proof. Suppose $R \neq \emptyset$ and select one $(i_0, j_0) \in R$. For $(i,j) \in R$ and $m, n \geq 1$ write

$$w_{ij}^{m,n} = \mathtt{a}^m \mathtt{p}_i \mathtt{b}^n \mathtt{c}^m \mathtt{q}_j \mathtt{d}^n.$$

Put

$$K_R = \{w_{ij}^{m,n} : (i,j) \in R,\ m,n \in \{1,2\}\} \cup \{w_{i_0 j_0}^{3,1}, w_{i_0 j_0}^{1,3}\}.$$

All these words lie in L_R and are distinct. Since $|w_{ij}^{m,n}| = 2m + 2n + 2$, their encoded costs for $(m,n) = (1,1), (2,1), (1,2), (2,2)$ are $7, 9, 9, 11$, respectively; the two extra words have encoded cost 11 each. Thus $\|K_R\|_+ = 36|R| + 22$.

We check the witnesses of Proposition 3.11 for the good fan-out-two, rank-two grammar in Proposition 6.3, retaining precisely the start rules indexed by R. In its trimmed $\widehat{s}$-refinement, each of A, B has *two* yield types, since the first/last-letter typing $s_{1,1}$ distinguishes one-letter components from components of length at least two. Write A_1, A_+ for the types realized by $(\mathtt{a}, \mathtt{c})$ and $(\mathtt{a}^2, \mathtt{c}^2)$, and B_1, B_+ for those realized by $(\mathtt{b}, \mathtt{d})$ and $(\mathtt{b}^2, \mathtt{d}^2)$. These are all the productive child types. There is one typed start nonterminal S_{ij} for each $(i,j) \in R$: its output type distinguishes the feature pair, but is independent of m, n.

Choose representatives

$$\omega(A_1) = (\mathtt{a}, \mathtt{c}), \quad \omega(A_+) = (\mathtt{a}^2, \mathtt{c}^2),$$
$$\omega(B_1) = (\mathtt{b}, \mathtt{d}), \quad \omega(B_+) = (\mathtt{b}^2, \mathtt{d}^2), \quad \omega(S_{ij}) = (w_{ij}^{1,1}).$$

Give S_{ij} the identity context. For both types of A and of B, use the derivational contexts with fixed feature pair (i_0, j_0) and the other child at its one-letter base:

$$\chi(A_\bullet)[(x,y)] = x \mathtt{p}_{i_0} \mathtt{b}\, y \mathtt{q}_{j_0} \mathtt{d},$$
$$\chi(B_\bullet)[(x,y)] = \mathtt{a} \mathtt{p}_{i_0} x \mathtt{c} \mathtt{q}_{j_0} y \qquad (\bullet \in \{1, +\}).$$

The representatives of A_1, B_1 and their rank-zero rules are witnessed by $w_{i_0 j_0}^{1,1}$; the representatives of A_+, B_+ are witnessed by $w_{i_0 j_0}^{2,1}$ and $w_{i_0 j_0}^{1,2}$. Each recursion into the longer type from the one-letter type has one of these latter two witnesses; recursion from a longer type into itself is witnessed by

$w^{3,1}_{i_0 j_0}$ or $w^{1,3}_{i_0 j_0}$. For each $(i,j) \in R$, the four typed start rules with children in $\{A_1, A_+\} \times \{B_1, B_+\}$ are witnessed by $w^{1,1}_{ij}, w^{2,1}_{ij}, w^{1,2}_{ij}, w^{2,2}_{ij}$. These account for every typed nonterminal and typed rule; hence the witness set of Definition 3.9 is contained in K_R. By Theorem 3.13, K_R is characteristic for $\mathcal{B}^{\leq 2}_{f,s}$. For $R = \emptyset$, the only possible positive sample is empty, and the reconstruction grammar has no start rule.

For the lower bound, fix $(i,j) \in R$. If a characteristic sample K for L_R contained no word with markers $\mathtt{p}_i, \mathtt{q}_j$, then $K \subseteq L_{R\setminus\{(i,j)\}}$. Take a characteristic sample F for the latter language under $\mathcal{D}$. The finite set $K \cup F$ is simultaneously a superset of K contained in L_R and a superset of F contained in $L_{R\setminus\{(i,j)\}}$, forcing $L(\mathcal{D}(K \cup F))$ to equal both distinct target languages. This is impossible. Hence K contains at least one word for each $(i,j) \in R$. Distinct pairs require distinct words, each of length at least 6 and encoded cost at least 7. □

Remark 6.6 (More feature values, with one typing per inventory). The construction also works for any *fixed* finite feature inventories $I = \{0, \ldots, p-1\}$ and $J = \{0, \ldots, q-1\}$. Use distinct one-letter markers $\mathtt{p}_i, \mathtt{q}_j$ and let $g : \Sigma^* \to \mathbb{Z}_p \times \mathbb{Z}_q$ send these markers to $(i,0)$ and $(0,j)$, respectively, and all four counting letters to zero. With the same first/last-letter component, the proof of Corollary 6.4 applies to every $R \subseteq I \times J$. Thus one typing for the *chosen* inventories accommodates 2^{pq} distinct target languages, and the preceding explicit witness construction still has cost $36|R| + 22$ for nonempty R. This does not give a uniform bound as p, q vary: changing the inventories changes the alphabet, group and fixed typing.

Remark 6.7 (What the feature example establishes). The markers represent a finite case-licensing or feature-compatibility constraint; no claim about the grammar of a specific natural language is made. Its nonrectangular three-pair relation makes untyped substitution unsound even between equal-boundary fragments. A finite group signature excludes that substitution, while two independent unbounded crossing-count dependencies remain to be learned. Corollary 6.4 and Proposition 6.5 make the class-level interpretation precise: the relation varies with the target, but the morphism and the reconstruction operator do not. This is an example of a meaningful *known bias*, not an algorithm for learning or optimizing the typing itself. It does not establish the necessity of noncommutative or local-group shortening.

6.2 Targets outside window–group-determined typings

The feature-licensing example above is covered by a window and an abelian group. This subsection separates the typing classes behind Section 4 at the level of targets: one fan-out-two target is admitted by a local-group typing but by no window–group-determined typing, in particular by none in $\mathbf{LI} \vee \mathbf{Grp}$; the other is admitted by a window–S_3 typing but by no typing in $\mathbf{LI} \vee \mathbf{Ab}$, where $\mathbf{Ab}$ is the pseudovariety of finite abelian groups. Both use the same template: a flag letter $\mathtt{r}$ leaves the following word unconstrained, a flag letter $\mathtt{l}$ imposes a typing-dependent condition on it, and a cross-serial counting block makes the target non-context-free.

Throughout, $\mathtt{1}, \mathtt{2}, \mathtt{3}, \mathtt{4}, \$, \mathtt{l}, \mathtt{r}$ are letters, and for a word w we write $\#_{\mathtt{ab}}(w)$ for the number of occurrences of the factor $\mathtt{ab}$ in w. Put $C_\times := \{\mathtt{1}^m\mathtt{2}^n\mathtt{3}^m\mathtt{4}^n : m, n \geq 1\}$ and let $\varphi_\times$ map $\mathtt{1}, \mathtt{3}, \mathtt{2}, \mathtt{4}$ to $(1,0), (-1,0), (0,1), (0,-1) \in \mathbb{Z}^2$ and every other letter to $(0,0)$.

Lemma 6.8 (Flag separation). *Let t be a finite typing of Σ^+, let $u, v \in \Sigma^+$ with $t(u) = t(v)$, and let $z \in \Sigma^*$. If L contains $\mathtt{r}uz$, $\mathtt{r}vz$ and $\mathtt{l}uz$ but not $\mathtt{l}vz$, then $L \notin \mathcal{C}^{\mathrm{mcf}}_{f,t}$ for every $f \geq 1$.*

Proof. The arity-one contexts $E = \mathtt{r}\Box_1 z$ and $F = \mathtt{l}\Box_1 z$ belong to $\mathcal{E}_1$. E is a common accepting context of u and v, while F accepts u but not v, so $\mathcal{D}^{(1)}_L(u) \neq \mathcal{D}^{(1)}_L(v)$. Every (f,t) condition includes arity one. □

A separation from window–group typings. Let $\Sigma_{\mathrm{A}} = \{\mathtt{l}, \mathtt{r}, \mathtt{a}, \mathtt{b}, \mathtt{c}, \$, 1, 2, 3, 4\}$ and

$$L_{\mathrm{A}} := \{\sigma w \,\$\, c : \sigma \in \{\mathtt{l}, \mathtt{r}\},\ w \in \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}^+,\ c \in C_\times,\ (\sigma = \mathtt{r} \text{ or } \#_{\mathtt{ab}}(w) \text{ is even})\}.$$

Let $s_{\mathrm{A}} := \lambda_{1,\nu}|_{\Sigma_{\mathrm{A}}^+}$, where $\nu : \Sigma_{\mathrm{A}}^2 \to \mathbb{Z}_2$ sends $\mathtt{ab}$ to 1 and every other factor to 0 (Definition 4.19; this is $h_{\mathtt{ab}}$ of Proposition B.5 on the larger alphabet). Thus s_{A} records the first and last letters and, for words of length at least two, the parity of $\#_{\mathtt{ab}}$.

Proposition 6.9 (A target separating local-group and window–group typings)**.** (i) $L_{\mathrm{A}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}$.

(ii) $L_{\mathrm{A}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{A}}}$ *for every* $f \geq 2$, *and* $s_{\mathrm{A}}(\Sigma_{\mathrm{A}}^+) \in \mathbf{LG}$.

(iii) *For every finite typing* t *of* Σ_{A}^+ *that is window–group-determined, that is,* $t \preceq (s_{k,\ell}, g'|_{\Sigma_{\mathrm{A}}^+})$ *for some* k, ℓ *and some morphism* g' *into a finite group, and in particular for every* t *whose positive image lies in* $\mathbf{LI} \vee \mathbf{Grp}$, *we have* $L_{\mathrm{A}} \notin \mathcal{C}^{\mathrm{mcf}}_{f,t}$ *for all* $f \geq 1$.

Consequently, among the window-type results, L_{A} *is covered by Theorem 4.20 but by none of Theorems 4.7, 4.15 and 4.16 or Corollary 4.18, for any choice of typing (but see Remark 6.11); and* $s_{\mathrm{A}}(\Sigma_{\mathrm{A}}^+) \notin \mathbf{LI} \vee \mathbf{Grp}$.

Proof. (i) Let W be a rank-one, fan-out-one good grammar for the regular language $\{\sigma w : \sigma \in \{\mathtt{l}, \mathtt{r}\},\ w \in \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}^+,\ \sigma = \mathtt{r}$ or $\#_{\mathtt{ab}}(w)$ even$\}$ (a right-linear grammar tracking the last letter and the parity). Add fan-out-two nonterminals C, D and a start symbol S with

$$\begin{aligned} C &\to (\xi_1^1 \$ \mathtt{1},\ \mathtt{3})(W) \mid (\xi_1^1 \mathtt{1},\ \xi_1^2 \mathtt{3})(C), \\ D &\to (\mathtt{2},\ \mathtt{4}) \mid (\xi_1^1 \mathtt{2},\ \xi_1^2 \mathtt{4})(D), \\ S &\to (\xi_1^1 \xi_2^1 \xi_1^2 \xi_2^2)(C, D). \end{aligned}$$

Then $L_C = \{(\sigma w \$ \mathtt{1}^m, \mathtt{3}^m)\}$ with $\sigma w \in L(W)$, $L_D = \{(\mathtt{2}^n, \mathtt{4}^n)\}$, and S generates L_{A}. The start template is nonmerging because adjacent variables belong to different children. Intersecting with the regular set $\mathtt{ra}\$\mathtt{1}^+\mathtt{2}^+\mathtt{3}^+\mathtt{4}^+$ and erasing $\mathtt{ra}\$$ yields $\{\mathtt{1}^m\mathtt{2}^n\mathtt{3}^m\mathtt{4}^n\}$, which is not context-free.

(ii) Write $L_{\mathrm{A}} = Q_{\mathrm{A}} \cap \Pi_{\mathrm{A}} \cap \varphi_\times^{-1}(0)$ with $Q_{\mathrm{A}} = \{\mathtt{l}, \mathtt{r}\}\{\mathtt{a}, \mathtt{b}, \mathtt{c}\}^+ \$\, \mathtt{1}^+\mathtt{2}^+\mathtt{3}^+\mathtt{4}^+$ and Π_{A} the set of nonempty words whose first letter is $\mathtt{r}$, or whose first letter is $\mathtt{l}$ and whose number of $\mathtt{ab}$ factors is even. On Q_{A} no factor $\mathtt{ab}$ meets σ, \$ or the counting block, so the parity is that of w.

- Q_{A} is strictly 2-local, with permitted pairs $\vdash\mathtt{l}$, $\vdash\mathtt{r}$, σx and xy for $\sigma \in \{\mathtt{l}, \mathtt{r}\}$ and $x, y \in \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}$, $x\$$, $\$\mathtt{1}$, $\mathtt{11}$, $\mathtt{12}$, $\mathtt{22}$, $\mathtt{23}$, $\mathtt{33}$, $\mathtt{34}$, $\mathtt{44}$ and $\mathtt{4}\dashv$. By Lemma 6.2 it is $(f, s_{1,1})$-tuple-substitutable, hence (f, s_{A})-tuple-substitutable since $s_{1,1} \preceq s_{\mathrm{A}}$.
- $\Pi_{\mathrm{A}} = \widehat{s}_{\mathrm{A}}^{-1}(\mathsf{Acc})$ for the set Acc of types of long words whose first letter is $\mathtt{r}$, or whose first letter is $\mathtt{l}$ and whose group coordinate is 0, together with the two short types $\mathtt{r}$ and $\mathtt{l}$ (the latter has zero $\mathtt{ab}$-count). By the observation following Definition 2.8 it is (f, s_{A})-tuple-substitutable.
- $\varphi_\times^{-1}(0)$ is handled by Proposition 2.9(ii) with $\mathsf{H} = 0$, $\gamma = 0$.

Proposition 2.9(i) gives the intersection; with (i), $L_{\mathrm{A}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{A}}}$ for $f \geq 2$. The positive image lies in $\mathbf{LG}$ by Theorem B.4, (iv)⇒(i).

(iii) Fix k, ℓ and $g' : \Sigma_{\mathrm{A}}^* \to A'$. Let $\hat{k} = \max\{k, \ell, 1\}$, let $q \geq 2$ be an even multiple of the exponent of A', and put

$$u = \mathtt{a}^{\hat{k}}(\mathtt{abc})^q\mathtt{a}^{\hat{k}}, \qquad v = \mathtt{a}^{\hat{k}}(\mathtt{bac})^q\mathtt{a}^{\hat{k}}.$$

As in Proposition B.5(ii), both are long with prefix $\mathtt{a}^k$ and suffix $\mathtt{a}^\ell$, and $g'(u) = g'(\mathtt{a})^{2\hat{k}} = g'(v)$; hence $(\widehat{s}_{k,\ell}, g')(u) = (\widehat{s}_{k,\ell}, g')(v)$, and also $t(u) = t(v)$ for every coarser t. Now $\#_{\mathtt{ab}}(u) = q$ is even and $\#_{\mathtt{ab}}(v) = 1$. With $z = \$\mathtt{1234}$, the words $\mathtt{r}uz$ and $\mathtt{r}vz$ and $\mathtt{l}uz$ lie in L_{A} and $\mathtt{l}vz$ does not. Lemma 6.8 applies. A typing with positive image in $\mathbf{LI} \vee \mathbf{Grp}$ is window–group-determined by Lemma 4.17.

For the consequences: if $s_{\mathrm{A}}(\Sigma_{\mathrm{A}}^+)$ were in $\mathbf{LI} \vee \mathbf{Grp}$, then by (iii) applied to $t = s_{\mathrm{A}}$ we would contradict (ii). The window and window–group theorems require a window–group-determined typing, and by (iii) no such typing admits L_{A}. □

A typing-class separation from abelian groups. Let $\Sigma_{\mathrm{B}} = \{\mathtt{l}, \mathtt{r}, \mathtt{a}, \mathtt{b}, \$, \mathtt{1}, \mathtt{2}, \mathtt{3}, \mathtt{4}\}$, let $\alpha = (12)$ and $\beta = (23)$ in S_3, and let $g_{\mathrm{B}} : \Sigma_{\mathrm{B}}^* \to S_3$ send $\mathtt{a} \mapsto \alpha$, $\mathtt{b} \mapsto \beta$ and every other letter to 1. Put

$$L_{\mathrm{B}} := \{\sigma w \,\$\, c : \sigma \in \{\mathtt{l}, \mathtt{r}\},\ w \in \{\mathtt{a}, \mathtt{b}\}^+,\ c \in C_\times,\ (\sigma = \mathtt{r} \text{ or } g_{\mathrm{B}}(w) = 1)\},$$

and let $s_{\mathrm{B}} := (s_{1,1}, g_{\mathrm{B}}|_{\Sigma_{\mathrm{B}}^+})$ restricted to its positive image. We write $\mathbf{Ab}$ for the pseudovariety of finite abelian groups; $\mathbf{LI} \vee \mathbf{Ab}$ consists of the divisors of products $T \times A$ with $T \in \mathbf{LI}$ and A finite abelian, since both pseudovarieties are closed under finite products.

Proposition 6.10 (A target separating nonabelian and window–abelian typings)**.** (i) $L_{\mathrm{B}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}$.

(ii) $L_{\mathrm{B}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{B}}}$ *for every* $f \geq 2$*; the typing* s_{B} *is window–group-determined with group* S_3*.*

(iii) *For every finite typing* t *of* Σ_{B}^+ *with* $t \preceq (s_{k,\ell}, g'|_{\Sigma_{\mathrm{B}}^+})$ *for some* k, ℓ *and some morphism* g' *into a finite* abelian *group, and in particular for every* t *whose positive image lies in* $\mathbf{LI} \vee \mathbf{Ab}$*,* $L_{\mathrm{B}} \notin \mathcal{C}^{\mathrm{mcf}}_{f,t}$ *for all* $f \geq 1$*.*

Consequently L_{B} *is covered by Theorem 4.15 with the nonabelian constant* $c_{S_3,f} = M_{S_3,f} - 1$*, but by no instance of Theorem 4.7 or 4.16 (but see Remark 6.11).*

Proof. (i) As for L_{A}, with W a right-linear grammar for $\{\sigma w : \sigma \in \{\mathtt{l}, \mathtt{r}\}, w \in \{\mathtt{a}, \mathtt{b}\}^+, \sigma = \mathtt{r} \text{ or } g_{\mathrm{B}}(w) = 1\}$, which tracks the last letter and the current value in S_3.

(ii) As for L_{A}: $L_{\mathrm{B}} = Q_{\mathrm{B}} \cap \Pi_{\mathrm{B}} \cap \varphi_\times^{-1}(0)$, where $Q_{\mathrm{B}} = \{\mathtt{l}, \mathtt{r}\}\{\mathtt{a}, \mathtt{b}\}^+ \$\, \mathtt{1}^+\mathtt{2}^+\mathtt{3}^+\mathtt{4}^+$ is strictly 2-local and Π_{B} is the set of nonempty words beginning with $\mathtt{r}$, or beginning with $\mathtt{l}$ and having g_{B}-value 1. Since g_{B} is trivial on all letters other than $\mathtt{a}, \mathtt{b}$, on Q_{B} the g_{B}-value of the whole word equals $g_{\mathrm{B}}(w)$, and $\Pi_{\mathrm{B}} = \widehat{s_{\mathrm{B}}}^{-1}(\mathsf{Acc})$ for a suitable Acc.

(iii) Let $g' : \Sigma_{\mathrm{B}}^* \to A'$ with A' abelian, let $\hat{k}$ be even with $\hat{k} \geq \max\{k, \ell, 1\}$, and put

$$u = \mathtt{a}^{\hat{k}}(\mathtt{ab})^3\mathtt{a}^{\hat{k}}, \qquad v = \mathtt{a}^{\hat{k}}(\mathtt{ab})^2\mathtt{ba}\,\mathtt{a}^{\hat{k}}.$$

Both are long with prefix $\mathtt{a}^k$ and suffix $\mathtt{a}^\ell$, and they have the same Parikh vector, so $g'(u) = g'(v)$ by commutativity. In S_3, $\alpha^2 = \beta^2 = 1$, $\alpha\beta$ has order 3 and $\beta\alpha = (\alpha\beta)^{-1}$. Since $\hat{k}$ is even,

$$g_{\mathrm{B}}(u) = (\alpha\beta)^3 = 1, \qquad g_{\mathrm{B}}(v) = (\alpha\beta)^2(\beta\alpha) = \alpha\beta \neq 1.$$

With $z = \$\mathtt{1234}$, Lemma 6.8 applies. For the clause on $\mathbf{LI} \vee \mathbf{Ab}$, the proof of Lemma 4.17 goes through verbatim with an abelian group factor and produces an abelian g'. □

Remark 6.11 (What the two examples do and do not show). Propositions 6.9 and 6.10 separate *typing classes*: no window–group-determined (respectively, window–abelian-determined) typing admits L_{A} (respectively, L_{B}). They do *not* show that the local-group or nonabelian shortening is

needed to obtain polynomial characteristic data for these targets. Indeed, let σ_{Q} be the positive restriction of the syntactic morphism of the regular language Q_{A} and put $s' = (s_{\mathrm{A}}, \sigma_{\mathrm{Q}})$, restricted to its positive image. Then $Q_{\mathrm{A}} \cap \Pi_{\mathrm{A}}$ is recognized by $\widehat{s'}$, so $L_{\mathrm{A}} = \widehat{s'}^{-1}(\mathsf{Acc}') \cap \varphi_{\times}^{-1}(0)$ for some Acc', and $L_{\mathrm{A}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s'}$ because $s_{\mathrm{A}} \preceq s'$. With the grammar of Proposition 6.9(i), Theorem 5.5 gives polynomial characteristic data for $\mathcal{B}^{\leq 2}_{\widetilde{f},s'}$, with no algebraic hypothesis on s'; the same applies to L_{B} after the analogous refinement. These are bounds for the *refined* typing's reconstruction operator, not bounds transferred to the original operators $\mathcal{B}^{\leq 2}_{f,s_{\mathrm{A}}}$ or $\mathcal{B}^{\leq 2}_{f,s_{\mathrm{B}}}$. Moreover the separation happens at arity one (Lemma 6.8): the typing-sensitive material lies in a single fan-out-one constituent, and the crossing block $C_{\times}$ only makes the targets non-context-free. The examples therefore isolate the algebraic expressiveness of the supplied typing, not its interaction with interleaved components, which is what Example A.4 illustrates inside the proof. Neither example asserts a lower bound on characteristic data or necessity of a shortening method for an individual language.

6.3 An Intrinsically Rank-Two Boundary-Typed Target

In this subsection, $s_{1,1}$ denotes the positive-word restriction $\widehat{s}_{1,1}|_{\Sigma^+_{\mathrm{intr}}}$ of the prefix–suffix typing from Definition 2.4, with alphabet Σ_{intr} defined below.

Let

$$\Sigma_{\mathrm{intr}} := \{\mathtt{a}, \mathtt{b}, \mathtt{l}, \mathtt{r}, \mathtt{p}, \mathtt{q}, \mathtt{c}, \mathtt{d}, \#, \$, \mathtt{u}, \mathtt{v}, \mathtt{w}\}, \qquad P_{\mathrm{intr}} := \{\mathtt{lpc}, \mathtt{lpd}, \mathtt{lqc}, \mathtt{rpc}, \mathtt{rpd}\},$$

$$\mathrm{MIX}_2^+ := \{x \in \{\mathtt{a}, \mathtt{b}\}^+ : |x|_{\mathtt{a}} = |x|_{\mathtt{b}}\}, \qquad T_3^+ := \{\mathtt{u}^n \mathtt{v}^n \mathtt{w}^n : n \geq 1\}, \qquad L_{\mathrm{intr}} := P_{\mathrm{intr}} \,\#\, \mathrm{MIX}_2^+ \,\$\, T_3^+.$$

The MIX_2^+ factor carries the rule-rank obstruction, which is not created by finite typing: $\mathrm{MIX}_2^+ = \{\mathtt{a}, \mathtt{b}\}^+ \cap \varphi_{\mathrm{mix}}^{-1}(0)$ for $\varphi_{\mathrm{mix}}(\mathtt{a}) = 1$, $\varphi_{\mathrm{mix}}(\mathtt{b}) = -1$ in $\mathbb{Z}$, so Proposition 2.9(ii) already places it in $\mathsf{YSub}(f)$ for every f. The T_3^+ factor forces minimum fan-out two, and the wrapper P_{intr} separates the target from $\mathsf{YSub}(f)$.

Lemma 6.12 (Rank-one MCFGs have finite-index EDT0L representations)**.** *For every fixed finite fan-out bound f', each language in $\mathsf{MCFL}(f', 1)$ has a finite-index EDT0L representation whose index is at most f'.*

Proof. Reduce the rank-one grammar H to productive reachable symbols. For each nonterminal B and component j introduce one EDT0L symbol $[B, j]$. Start with $[S, 1]$. For each rank-one rule $A \to \rho(B)$ take a deterministic table that replaces every $[A, j]$ simultaneously by the jth component of ρ, with ξ_1^k represented by $[B, k]$. Every other symbol is fixed. Rank-zero rules have the corresponding terminal-only tables. Because the rule has at most one child and is linear, at every reachable nonterminal-bearing sentential form the active markers belong to a single nonterminal and appear at most once per component. Thus no more than f' such markers occur at a time. Every MCFG derivation chooses a corresponding sequence of tables. Conversely, a successful table sequence acts on the unique active nonterminal; unobserved child components can be supplied with productive derivations in H. It therefore corresponds to a valid MCFG derivation with exactly the same terminal word. Hence the obtained EDT0L system generates $L(H)$ and has finite index at most f'. □

Theorem 6.13 (A non-context-free intrinsic-rank-two typed witness)**.** *The language L_{intr} satisfies $L_{\mathrm{intr}} \in \mathsf{MCFL}(2, 2) \setminus \mathrm{CFL}$, and, for every $f \geq 2$, $L_{\mathrm{intr}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}} \setminus \mathsf{YSub}(f)$. Moreover, for every finite $f' \geq 1$, $L_{\mathrm{intr}} \notin \mathsf{MCFL}(f', 1)$. Consequently the minimum fan-out of L_{intr} is exactly two, and its minimum rule rank among all finite-fan-out MCFG presentations is exactly two. In particular, Theorem 4.7 applies to a non-context-free language that cannot be covered by Theorem 5.1 at any finite fan-out.*

Proof. Grammar. Let B have the rules

$$B \to \mathtt{a}B\mathtt{b}B \mid \mathtt{b}B\mathtt{a}B \mid \mathtt{a}B\mathtt{b} \mid \mathtt{ab}B \mid \mathtt{ab} \mid \mathtt{b}B\mathtt{a} \mid \mathtt{ba}B \mid \mathtt{ba};$$

the first-return induction on $|x|_{\mathtt{a}} - |x|_{\mathtt{b}}$ gives $L_B = \mathrm{MIX}_2^+$. Let the fan-out-two nonterminal A have rules $A \to (\mathtt{u}, \mathtt{w}) \mid (\mathtt{u}\xi_1^1\mathtt{v}, \mathtt{w}\xi_1^2)(A)$, so $L_A = \{(\mathtt{u}^n\mathtt{v}^{n-1}, \mathtt{w}^n) : n \geq 1\}$, and for every $\alpha \in P_{\mathrm{intr}}$ add $S \to \alpha\#\xi_1^1\$\xi_2^1\mathtt{v}\xi_2^2(B, A)$. The grammar generates L_{intr}, is reduced, λ-free, linear, nondeleting and nonpermuting, and is nonmerging because the two components of A are separated by the literal $\mathtt{v}$. It is good, of fan-out two and rule rank two.

Non-context-freeness. With $R := \mathtt{lpc}\#\mathtt{ab}\$\,\mathtt{u}^+\mathtt{v}^+\mathtt{w}^+$, $L_{\mathrm{intr}} \cap R = \{\mathtt{lpc}\#\mathtt{ab}\$\mathtt{u}^n\mathtt{v}^n\mathtt{w}^n : n \geq 1\}$, which a homomorphism erasing the fixed prefix maps onto the non-context-free $\{\mathtt{u}^n\mathtt{v}^n\mathtt{w}^n : n \geq 1\}$. By closure of CFL under intersection with regular languages and homomorphism, L_{intr} is not context-free; fan-out-one MCFGs are exactly CFGs here, so the minimum fan-out is two.

Typed membership. Let $Q_{\mathrm{intr}} := P_{\mathrm{intr}} \# \{\mathtt{a}, \mathtt{b}\}^+ \$\, \mathtt{u}^+\mathtt{v}^+\mathtt{w}^+$. It is strictly 2-local, with the allowed adjacent pairs

$$\vdash \mathtt{l},\ \vdash \mathtt{r},\ \mathtt{lp},\ \mathtt{lq},\ \mathtt{rp},\ \mathtt{pc},\ \mathtt{pd},\ \mathtt{qc},\ \mathtt{c}\#,\ \mathtt{d}\#,\ \#\mathtt{a},\ \#\mathtt{b},\ \mathtt{aa},\ \mathtt{ab},\ \mathtt{ba},\ \mathtt{bb},$$
$$\mathtt{a}\$,\ \mathtt{b}\$,\ \$\mathtt{u},\ \mathtt{uu},\ \mathtt{uv},\ \mathtt{vv},\ \mathtt{vw},\ \mathtt{ww},\ \mathtt{w} \dashv .$$

Let $\Phi_{\mathrm{intr}} : \Sigma_{\mathrm{intr}}^* \to \mathbb{Z}^3$ send $\mathtt{a}, \mathtt{b}, \mathtt{u}, \mathtt{v}, \mathtt{w}$ to $(1,0,0)$, $(-1,0,0)$, $(0,1,0)$, $(0,-1,1)$, $(0,0,-1)$ and all other letters to 0. Its three coordinates impose $|x|_{\mathtt{a}} = |x|_{\mathtt{b}}$, $|x|_{\mathtt{u}} = |x|_{\mathtt{v}}$ and $|x|_{\mathtt{v}} = |x|_{\mathtt{w}}$, so $L_{\mathrm{intr}} = Q_{\mathrm{intr}} \cap \Phi_{\mathrm{intr}}^{-1}(0)$. Lemma 6.2 and Proposition 2.9(ii) make this intersection $(f, s_{1,1})$-tuple-substitutable for every f; with the grammar above, $L_{\mathrm{intr}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}}$ for every $f \geq 2$.

Separation. For $x = \mathtt{p}$, $y = \mathtt{q}$, $E = \mathtt{l}\Box_1\mathtt{c}\#\mathtt{ab}\$\mathtt{uvw}$ and $F = \mathtt{r}\Box_1\mathtt{d}\#\mathtt{ab}\$\mathtt{uvw}$, we have $E[x], E[y], F[x] \in L_{\mathrm{intr}}$ but $F[y] \notin L_{\mathrm{intr}}$, so the arity-one substitution implication fails and $L_{\mathrm{intr}} \notin \mathsf{YSub}(f)$ for every $f \geq 1$; the boundary typing separates the offending factors, $\widehat{s}_{1,1}(\mathtt{p}) \neq \widehat{s}_{1,1}(\mathtt{q})$.

Intrinsic rule rank. Lemma 6.12 shows that every finite-fan-out rank-one MCFG language is EDT0L. This agrees with Gallot's non-branching characterization [21, Definition 38 and Theorems 28–29]. Independently, Bishop et al. show that the word problem of the infinite cyclic group is not EDT0L [22]. Over $\{\mathtt{a}, \mathtt{b}\}$ that word problem is precisely $\mathrm{MIX}_2 = \{x : |x|_{\mathtt{a}} = |x|_{\mathtt{b}}\}$. Let π preserve $\mathtt{a}, \mathtt{b}$ and erase the other letters. If L_{intr} had a finite-fan-out rank-one presentation, its homomorphic image $\pi(L_{\mathrm{intr}}) = \mathrm{MIX}_2^+$ would be EDT0L, since the table construction commutes with terminal homomorphisms. Adding the singleton $\{\varepsilon\}$ preserves EDT0L, giving MIX_2, a contradiction. Hence $L_{\mathrm{intr}} \notin \mathsf{MCFL}(f', 1)$ for every finite f', and the displayed grammar proves the claimed intrinsic rule rank. □

The two lower bounds arise from separate factors rather than from an irreducible interaction between the resources, and the non-branching lower bound itself is prior work. What is specific here is that a strictly local wrapper and the $(1,1)$ boundary typing place the target inside the typed class while the substitution implication keeps it outside Yoshinaka's trivial-typing baseline.

7 Conclusion

This paper converts the qualitative positive-data learning principle for typed MCFGs into presentation-sensitive upper bounds on the *total encoded length* of characteristic samples. For reduced good grammars, window and window–finite-group typings yield $O(\|G\|^2\bar{\theta}_G)$, including noncommutative groups and interleaved tuple components. Positive-image local-group typings yield $O(\|G\|^3\bar{\theta}_G)$ by coupling classical sliding-factor coordinates with the MCFG ordered-block

decomposition and shortening. For arbitrary standard input grammars, component-pattern thickness controls the transfer of these bounds through normalization. Rank-one presentations admit a polynomial bound for every fixed typing. The refinement-thickness dichotomy (Corollary 4.21) places the exact boundary for polynomial full yield-type refinement thickness at **LG** at every fan-out; it does not establish necessity for characteristic-data length. The internal feature-licensing example in Proposition 6.3 is a non-context-free window–group typed target for which the untyped distributional substitution rule is unsound even with the boundary-window condition alone. No lower-bound theorem or separating family from the related CFG manuscript is required for these results.

Open Problem 7.1 (Characteristic-data length at higher fan-out). Fix $f \geq 2$. Is there a fixed finite-semigroup typing whose positive image lies outside **LG**, for which every target in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ with a reduced good rank-bounded presentation still admits characteristic data polynomial in the original grammar size and rule thickness? Conversely, can the two-element semilattice U_1 already force a superpolynomial characteristic-data lower bound at higher fan-out? Corollary 4.21 answers the analogous *refinement-thickness* question, but it does not answer either question about characteristic data.

The key unresolved issue is whether higher-arity distributional constraints can be turned into short typed representatives for every fixed finite semigroup, or whether a fixed-typing counterexample exists.

Acknowledgments

The author thanks Makoto Kanazawa and Ryo Yoshinaka for valuable discussions and encouragement.

Declaration of competing interest

The author declares that he has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data availability

No empirical dataset is associated with this theoretical paper.

A Simultaneous Block Replacement in Finite Groups

Semigroup convention. The finite group A and the morphisms into it are auxiliary to the shortening proofs. The semantic typing is always the positive-image restriction of the window–group (or local-group) product, with the fresh identity used only to evaluate empty template factors; by Lemma E.1 this changes neither nonempty tuple comparisons nor the selected witnesses.

We now prove the purely group-theoretic statement used in Theorem 4.12 and Lemma B.6. The commutative case is a zero-sum argument; the general case needs an order-homogenization argument

because, for $f \geq 2$, blocks of different indices can interleave differently on different lines. The proof has three logically separate parts: the classification of homogeneous point orders (Lemma 4.10), the group-valued interval invariant (Lemma 4.11), and the polynomial-size extraction of homogeneous indices (Lemma A.5). Neither the order classification nor the selection of a replacement interval assumes that the group is abelian.

Definition A.1 (Block-replacement family)**.** Let A be a finite group and $f \geq 1$. A *block-replacement family* of width f over A with index set $[m]$ consists of $d \leq f$ *lines*; on line c, a left-to-right sequence of blocks $\beta_{c,1}, \ldots, \beta_{c,q_c}$ and constants $a_{c,0}, \ldots, a_{c,q_c} \in A$; for every block β an *owner* $o(\beta) \in [m]$ and two values $x_\beta, x'_\beta \in A$; every index owns at most f blocks. For $J \subseteq [m]$ put

$$\mathrm{val}_c(J) := a_{c,0}\, y_{c,1}\, a_{c,1} \cdots y_{c,q_c}\, a_{c,q_c}, \qquad y_{c,j} := \begin{cases} x'_{\beta_{c,j}}, & o(\beta_{c,j}) \in J, \\ x_{\beta_{c,j}}, & o(\beta_{c,j}) \notin J. \end{cases}$$

Lemma A.2 (Block replacement)**.** *In a block-replacement family of width f over A there is a nonempty $J \subseteq [m]$ with $\mathrm{val}_c(J) = \mathrm{val}_c(\emptyset)$ for every line c, provided that either* (a) *A is abelian and $m > \mathsf{d}(A^f)$, or* (b) *A is arbitrary and $m \geq M_{A,f}$, where d and $M_{A,f}$ are as in Section 4.4.*

Part (a) is proved below under commutativity, and part (b) under the arbitrary-group order homogenization argument.

Commutative groups: zero sums

Proof of Lemma A.2(a). Write A additively. By commutativity $\mathrm{val}_c(J) - \mathrm{val}_c(\emptyset) = \sum_{i \in J} \delta_{i,c}$, where $\delta_{i,c}$ is the sum of $x'_\beta - x_\beta$ over the blocks β of line c owned by i. Put $\delta_i = (\delta_{i,c})_c \in A^d$. A sequence over A^d of length $m > \mathsf{d}(A^f) \geq \mathsf{d}(A^d)$ has a nonempty zero-sum subsequence $(\delta_i)_{i \in J}$; the inequality holds because A^d embeds in A^f and a zero-sum-free sequence in a subgroup is zero-sum free in the group. □

Remark A.3 (Size of the zero-sum constant). The prefix-sum pigeonhole argument gives $\mathsf{d}(A^f) \leq |A|^f - 1$; the zero-sum bound is usually much stronger. Write $A \cong C_{n_1} \oplus \cdots \oplus C_{n_s}$ and $n = \exp(A) \geq 2$. Then

$$f \sum_{i=1}^{s} (n_i - 1) \;\leq\; \mathsf{d}(A^f) \;\leq\; (n-1) + n \log \frac{|A|^f}{n},$$

where the lower bound is attained by a basis sequence and the upper bound is due to van Emde Boas and Kruyswijk [33, Proposition 5.1.8 and Theorem 5.5.5]; for a p-group the lower bound is an equality [33, Theorem 5.5.9] (Olson). Hence $\mathsf{d}(A^f) = O(\exp(A)\, f \log |A|)$ is linear in f, whereas $|A|^f - 1$ is exponential; for $A = C_2$ the antichain bound drops from $2^f - 1$ to f. The value $\mathsf{d}(A^d)$ is also the limit of the argument as stated, since a zero-sum-free sequence of that length exists; we do not claim that a grammar realizes such a configuration.

A.1 Arbitrary Finite Groups: Order Homogenization

Proof of Lemma 4.10.

Proof. By (ii) with $t = u$, each type's points are monotone. For distinct types t, u write X for $(t,i) < (u,j)$ and Y for $(u,i) < (t,j)$, $i < j$; both are uniform by (ii). Suppose first that both types increase. $X \wedge Y$ puts all points of rank i before all points of rank $j > i$. Under $X \wedge \neg Y$ all off-diagonal comparisons put t left of u, and so do equal ranks: $(u,i) < (t,i)$ would give $(t,1) < (t,2) < (u,1) < (t,1)$; thus the types are separated. The case $\neg X \wedge Y$ is symmetric, and

$\neg X \wedge \neg Y$ gives $(t,1) < (t,2) < (u,1) < (u,2) < (t,1)$, a contradiction. If t increases and u decreases, $X \wedge Y$ gives $(t,2) < (u,3) < (u,1) < (t,2)$ and $\neg X \wedge \neg Y$ gives $(t,3) < (u,2) < (t,1) < (t,3)$; under $X \wedge \neg Y$ all off-diagonal comparisons put t left of u, and a contrary equal-rank comparison would give $(t,1) < (u,2) < (u,1) < (t,1)$, contrary to (i), so the types are separated; $\neg X \wedge Y$ is symmetric. Reversing indices covers two decreasing types. Thus any two types are either *separated*, or have the same direction and are *interleaved by rank*.

Interleaving is transitive: if t interleaves with u and u with v, all three have one direction, say increasing; if t, v were separated with t to the left we would have $(v,1) < (u,2) < (t,3) < (v,1)$, and with v to the left $(t,1) < (u,2) < (v,3) < (t,1)$. Hence interleaving, with equality, partitions the types into classes. A type outside a class is separated from each type inside it and lies on the same side of all of them, since otherwise two interleaved types would be separated. So class regions are consecutive, the rank comparisons give successive chunks, and the within-chunk order is independent of rank by (i). □

Proof of Lemma 4.11.

Proof. Apply Lemma 4.10 on each line to the block types going to that line. Across all lines there are $q \le f$ class regions. For such a region $\mathcal{Q}$, enumerate its chunks $1, \dots, n$ from left to right. Assign any constant between two successive chunks to the chunk on its left; keep the initial and terminal constants outside the chunks. Since those two outer constants never change, equality of the ordered product between them is sufficient. In the following, P_p and Q_p refer to this interior product (so both begin at 1_A). Let P_p be the original value of its prefix ending after chunk p and Q_p the value of that prefix when chunks $1, \dots, p$ are replaced; set $P_0 = Q_0 = 1_A$ and $\kappa_p^{\mathcal{Q}} = P_p Q_p^{-1}$. For $0 \le p < p' \le n$, the *ordered* value after replacing exactly chunks $p+1, \dots, p'$ is

$$\underbrace{P_p}_{\text{old prefix}} \underbrace{(Q_p^{-1} Q_{p'})}_{\text{new middle}} \underbrace{(P_{p'}^{-1} P_n)}_{\text{old suffix}}.$$

The middle expression is correct even with constants between chunks: $Q_{p'} = Q_p$ (new middle) by the definition of prefix products. Therefore the displayed product equals P_n exactly when

$$P_p Q_p^{-1} Q_{p'} P_{p'}^{-1} = 1_A, \quad \text{equivalently} \quad \kappa_p^{\mathcal{Q}} = P_p Q_p^{-1} = P_{p'} Q_{p'}^{-1} = \kappa_{p'}^{\mathcal{Q}}.$$

No factors have been commuted. Note that this requires equality of *two boundary states*; it does not require an individual block's original and replacement values to agree.

If $\mathcal{Q}$ increases with the index order, replacing owners $(a,b]$ replaces chunks $(a,b]$ in the region. If the order decreases, the same owners occupy chunks $(n-b, n-a]$ (counted from the left). This reversal is why the two orientations must use distinct state indices. For $j = 0, \dots, n$ define a state in A^q by taking $\kappa_j^{\mathcal{Q}}$ from each increasing and $\kappa_{n-j}^{\mathcal{Q}}$ from each decreasing class. There are $n+1 > |A|^q$ states, so two coincide, at $a < b$. Replacing the indices $a+1, \dots, b$ then preserves every region value; everything else on every line is unchanged. □

Example A.4 (Two rows, three owners, reversed order, and noncommuting values). This finite example illustrates the *boundary-state identity*, not an invocation of the numerical pigeonhole threshold (which can be much larger than three). Work in the nonabelian group $A = S_3$ and put $\alpha = (12)$, $\beta = (23)$, so that $\alpha^2 = \beta^2 = 1$ but $\alpha\beta \ne \beta\alpha$. There are three owners $i = 1, 2, 3$ with three blocks each: two blocks U_i, V_i on row 1, and one block W_i on row 2. Reading left to right, with the fixed group-valued material *shown explicitly*, the rows are

$$\begin{aligned} &\text{row 1:} \quad (U_1 \,\beta\, V_1)\, \alpha\, (U_2 \,\beta\, V_2)\, \alpha\, (U_3 \,\beta\, V_3), \\ &\text{row 2:} \quad W_3 \,\beta\, W_2 \,\beta\, W_1. \end{aligned}$$

Thus row 1 has two owner types occurring in increasing rank chunks, while row 2 has one type in *decreasing* rank chunks. Originally take $U_i = \beta$, $V_i = 1$, $W_i = 1$ for all i; both complete rows have value 1. Now replace *exactly* owners $J = \{2, 3\}$ using the table

owner	U_i'	V_i'	W_i'	replaced?
1	β	1	1	no
2	1	1	$\beta\alpha\beta$	yes
3	$\alpha\beta\alpha\beta$	1	α	yes

(the unchanged entries in row 1 for owner 1 are listed for comparison). The row 1 chunk values were $(1, 1, 1)$ and become $(1, \beta, \alpha\beta\alpha)$; its complete new value is

$$1 \cdot \alpha \cdot \beta \cdot \alpha \cdot (\alpha\beta\alpha) = 1.$$

In the reverse-order row 2 the new value is

$$\alpha \cdot \beta \cdot (\beta\alpha\beta) \cdot \beta \cdot 1 = 1.$$

These are *ordered* calculations in S_3; neither replaces $\alpha\beta$ by $\beta\alpha$. More precisely, for the interval of owner indices $(a, b] = (1, 3]$, the increasing region of row 1 has $\kappa_1 = \kappa_3 = 1$, whereas the decreasing region of row 2 has $\kappa_{3-b} = \kappa_0 = \kappa_2 = \kappa_{3-a} = 1$. The same owner interval therefore matches the two *different* boundary-index orientations. Note that no individual replacement owner need preserve its original row contribution; only both *full row products* are fixed. In the general proof there can be up to f such regions on multiple rows, and their states are compared simultaneously.

Lemma A.5 (Polynomial order homogenization)**.** *Let $N \geq 3$, and define $m_0 = N$ and $m_{s+1} = (m_s - 1)^2 + 1$, so that $m_k \leq N^{2^k}$. Let T be a finite set of point types, each assigned to one of several lines. For every index $i \in [n]$ and $t \in T$ let $p_t(i)$ be a point on the line of t, the points on each line being distinct. Let Π be the set of ordered pairs (t, u), including $t = u$, of types on a common line, and put $k = |\Pi|$. If $n \geq m_k$, there is a set $I \subseteq [n]$ of N indices such that, for every $(t, u) \in \Pi$, the truth value of $p_t(i) < p_u(j)$ is the same for all $i < j$ in I.*

Proof. Treat the pairs one at a time, the pairs (t, t) first. Each step passes from a current index set of size at least $(m-1)^2 + 1$ to a subset of size at least m; homogeneity already obtained, and monotonicity of each type, persist in subsets. For (t, t), homogeneity means that p_t is monotone, and the Erdős–Szekeres theorem gives a monotone subsequence of length m.

For $t \neq u$, both p_t and p_u are already monotone. Relabel the current indices $1, \ldots, m'$ in order, with $m' \geq (m-1)^2 + 1$, and put $X(i, j) = [p_t(i) < p_u(j)]$ for $i < j$. Because $p_u(j)$ is monotone in j, every row $X(i, \cdot)$ is monotone; because $p_t(i)$ is monotone in i, every column $X(\cdot, j)$ is monotone. The respective directions are fixed over all rows and columns by the directions of the two point types. Replacing X by $1 - X$ if necessary, assume each row goes from 0 to 1.

If the columns go from 1 to 0, put $J(i) = \min\{j > i : X(i, j) = 1\}$ (or ∞); J is nondecreasing. Set $i_1 = 1$ and $i_{s+1} = J(i_s)$ while this is at most m'. For $a < b$, $i_b \geq J(i_a)$ and so $X(i_a, i_b) = 1$; if the chain has m terms we are done. Otherwise its at most $m - 1$ terms cut $[m']$ into at most $m - 1$ intervals $[i_s, i_{s+1})$ and $[i_L, m']$, one of which has more than $m'/(m-1) > m - 1$ indices; inside it every $i < j$ satisfies $j < J(i_s) \leq J(i)$, so $X(i, j) = 0$ (and $J \equiv \infty$ in the last interval).

If the columns go from 0 to 1, the set $\{X = 1\}$ is closed upward in both coordinates. Let κ be the least index with $X(\kappa, \kappa + 1) = 1$, or $\kappa = m'$. For $\kappa \leq i < j$, $X(i, j) \geq X(\kappa, \kappa + 1) = 1$; for $i < j \leq \kappa$, $X(i, j) \leq X(j - 1, j) = 0$. One of $[1, \kappa]$ and $[\kappa, m']$ has at least $(m' + 1)/2 \geq m$ indices, because $(m-1)^2 + 1 \geq 2m - 1$.

After all k pairs at least $m_0 = N$ indices remain, and $m_{s+1} \leq m_s^2$ gives $m_k \leq N^{2^k}$. □

Proof of Lemma A.2(b). If an index owns no blocks, choosing that index alone preserves every line, so assume each index owns at least one block. The *unary pattern* of an index records the number of its blocks, the line of each, and their order within any shared line; there are at most $f \cdot f^f \cdot f! = U_f$ patterns. Hence at least $N_{A,f}^{2^{f^2}}$ indices share one pattern; the blocks of all other indices are never replaced and are absorbed into the constants. Number the blocks of the common pattern as types $1, \dots, f'$ ($f' \le f$) and use their positions as points on their lines. There are at most f^2 ordered pairs of types on a common line, including $t = u$, and $m_k \le N^{2^k} \le N^{2^{f^2}}$ for $k \le f^2$, so Lemma A.5 with $N = N_{A,f}$ yields $N_{A,f}$ indices with constant cross-index comparisons for all $i < j$. These satisfy the hypotheses of Lemma 4.11, which supplies the required nonempty interval. □

For fixed f, the constant $M_{A,f} = U_f \max\{3, |A|^f\}^{2^{f^2}}$ is polynomial in $|A|$; its exponent depends doubly exponentially on f.

B Local-Group Coordinates and MCFG Block Shortening

Semigroup convention. Throughout this subsection **LG** is tested on the *positive image* $S = \widehat{s}(\Sigma^+)$ of the typing, not on the canonical completion $S^\sharp$. Auxiliary group-valued maps and the displayed monoid morphisms are implementations whose restrictions to nonempty words define the semantic bias (Lemma E.1).

LG is the pseudovariety of finite semigroups S in which every local monoid eSe, with e idempotent, is a group; the classical relation $\mathbf{LG} = \mathbf{Grp} * \mathbf{D}$ places it within the sliding-window and delay-theorem theory [37, 38]. The equivalent exclusion of the two-element semilattice $U_1 = \{1, \zeta\}$ ($\zeta^2 = \zeta$) is given by [36, Proposition 2.3]. The sliding-factor coordinates below are a self-contained implementation of this established algebra, tailored to the precise boundary and group values used here; they are not claimed as a new algebraic characterization. The contribution to MCFG learning is the factorization of group values into fixed separators and variable blocks in Lemma B.2, followed by simultaneous replacement of those blocks.

Throughout, Σ is nonempty, $k \ge 1$, and A is a finite group. For a word w with $|w| \ge k+1$ let $\sigma_k(w)$ be the sequence of its factors of length $k+1$, ordered by their starting positions.

Recall the local-group typing of Definition 4.19.

For $\ell = k$, words of length $< 2k$ are *short* and recorded exactly, while a *long* w is recorded by $\langle \mathrm{pref}_k(w), \mathrm{suf}_k(w) \rangle$ (Definition 2.4). For $s, p \in \Sigma^k$ write $J_\nu(s,p) := \prod_{i=1}^{k} \nu((sp)[i..i+k])$, the ordered product over the k factors of length $k+1$ of sp that start in s.

Lemma B.1 (Local-group typings are morphisms)**.** *The kernel of $\lambda_{k,\nu}$ is a congruence of finite index on Σ^*. Hence $\lambda_{k,\nu}$ is a homomorphism onto a finite monoid $M_{k,\nu}$. For long u, v,*

$$\lambda_{k,\nu}(uv) = \big(\langle \mathrm{pref}_k u, \mathrm{suf}_k v \rangle,\ \gamma_{k,\nu}(u)\, J_\nu(\mathrm{suf}_k u, \mathrm{pref}_k v)\, \gamma_{k,\nu}(v)\big).$$

Proof. The window component is a morphism by Definition 2.4. For long u, v, the factors of length $k+1$ of uv are, in order, those of u, the k factors starting in the last k positions of u (inside $\mathrm{suf}_k(u)\,\mathrm{pref}_k(v)$), and those of v. If u is short and v long, the factors starting in u lie inside $u\,\mathrm{pref}_k(v)$ and the rest are those of v, so $\gamma_{k,\nu}(uv)$ is determined by u, $\mathrm{pref}_k(v)$ and $\gamma_{k,\nu}(v)$; the case of long u and short v is symmetric, and if both are short uv is determined exactly. Thus $\lambda_{k,\nu}(uv)$ is a function of $\lambda_{k,\nu}(u)$ and $\lambda_{k,\nu}(v)$, and the image has at most $|\Sigma|^{<2k} + |\Sigma|^{2k}|A|$ elements. □

Lemma B.2 (Block decomposition)**.** *Let $w = x_0\beta_1 x_1 \cdots \beta_q x_q$ be a factorization of a word into designated blocks β_j and fixed words x_j. Let $w' = x_0\beta'_1 x_1 \cdots \beta'_q x_q$, where $\widehat{s}_{k,k}(\beta'_j) = \widehat{s}_{k,k}(\beta_j)$ for*

every j. Then $\widehat{s}_{k,k}(w') = \widehat{s}_{k,k}(w)$. If w is long, there are $c_0, \dots, c_q \in A$, depending only on the x_j and the window types of the β_j, such that

$$\gamma_{k,\nu}(w) = c_0\, \gamma_{k,\nu}(\beta_1)\, c_1 \cdots \gamma_{k,\nu}(\beta_q)\, c_q, \qquad \gamma_{k,\nu}(w') = c_0\, \gamma_{k,\nu}(\beta'_1)\, c_1 \cdots \gamma_{k,\nu}(\beta'_q)\, c_q.$$

Proof. A short block is recorded exactly, so $\beta'_j = \beta_j$, and we absorb it into the fixed words. Every remaining block is long and its replacement is long with the same first and last k letters; hence w' is long exactly when w is, with the same first and last k letters, which gives the window claim. A factor F of length $k+1$ of w either lies inside a long block β_j or meets each long block it touches only in a prefix or a suffix of length at most k (meeting both ends would force $|\beta_j| \le k$). Factors of the second kind are spelled by the fixed words and the windows of the blocks, and occur in the same order in w'. Ordering all factors by starting position, those inside β_j form one consecutive run contributing $\gamma_{k,\nu}(\beta_j)$; let c_j be the ordered product of the remaining factors between consecutive runs. □

Lemma B.3 (Explicit sliding-factor coordinates for block shortening)**.** *Let S be a finite semigroup with minimal ideal K and suppose every idempotent of S lies in K. Put $n = |S|$ and fix any morphism $h : \Sigma^+ \to S$. In a Rees representation $\mathsf{K} = \mathcal{M}(H; I, \Lambda; P)$ there is a map $\nu : \Sigma^{n+1} \to H$ such that, for every word w with $|w| \ge 2n$,*

$$h(w) = \big(\mathrm{row}(h(\mathrm{pref}_n\, w)),\ \mathrm{grp}(h(\mathrm{pref}_n\, w))\, \gamma_{n,\nu}(w),\ \mathrm{col}(h(\mathrm{suf}_n\, w))\big). \tag{$\star$}$$

Here row*,* grp *and* col *send $(i, a, \lambda) \in \mathsf{K}$ to $i \in I$, $a \in H$ and $\lambda \in \Lambda$, and $\gamma_{n,\nu}(w)$ is the* ordered *product of the $|w| - n$ sliding-factor labels from Definition 4.19. Consequently $h(w)$ is determined by the first and last n letters and this ordered product, with no commutativity hypothesis on H.*

Proof. (1) Uniform entry into K. The Rees quotient S/K has no idempotent except zero. Every element therefore has a power equal to zero. If a product of $m = |S/\mathsf{K}|$ factors were nonzero, its m nonempty prefix products would all be nonzero, so two, say $p_i = p_j$ with $i < j$, would coincide. With $x = s_{i+1} \cdots s_j$ we would have $p_i = p_i x^q$ for every q; taking q with $x^q = 0$ is a contradiction. Thus every product of $n \ge |S/\mathsf{K}|$ letters belongs to K.

(2) Outer coordinates and the right action. Write $(i, a, \lambda)(j, b, \mu) = (i, a p_{\lambda j} b, \mu)$. Fix $i_0 \in I$ and set $e_\lambda = (i_0, p^{-1}_{\lambda i_0}, \lambda) \in \mathsf{K}$. For all $(i, a, \lambda) \in \mathsf{K}$, multiplication on the right by e_λ is the identity. For $s \in S$ write $e_\lambda s = (j, c, \mu) \in \mathsf{K}$ and put $\rho(\lambda, s) = p_{\lambda j} c$. Associativity then gives the exact identity

$$(i, a, \lambda)s = (i, a\rho(\lambda, s), \mu(\lambda, s)).$$

In particular the I-coordinate is unchanged by right multiplication, even if $s \notin \mathsf{K}$. Dually, left multiplication preserves the Λ-coordinate of a product already in K: choose an idempotent $f_i \in \mathsf{K}$ with $f_i(i, a, \lambda) = (i, a, \lambda)$; then $s(i, a, \lambda) = (sf_i)(i, a, \lambda)$, and the Rees rule forces the same Λ-coordinate. The I-coordinate of a long word is therefore fixed by its first n letters, and its Λ-coordinate by its last n letters.

(3) The sliding-factor group coordinate. Write $w = a_1 \cdots a_m$, $m \ge 2n$, and set $h(a_1 \cdots a_n) = (i, a, \lambda_n)$. For $j \ge n$ let λ_j denote the Λ-coordinate of $h(a_1 \cdots a_j)$. Since left multiplication preserves the Λ-coordinate on K, this equals the Λ-coordinate of $h(a_{j-n+1} \cdots a_j)$; both products have length at least n. Define for each $F \in \Sigma^{n+1}$

$$\nu(F) = \rho\big(\mathrm{col}(h(F[1..n])),\ h(F[n+1])\big).$$

Successive applications of the right-action identity, in their actual left-to-right order, yield

$$h(w) = \Big(i,\ a \prod_{j=n}^{m-1} \nu(a_{j-n+1} \cdots a_{j+1}),\ \lambda_m\Big).$$

The displayed product is precisely $\gamma_{n,\nu}(w)$, and $\lambda_m = \mathrm{col}(h(\mathrm{suf}_n\, w))$. This proves $(\star)$. The case $m = 2n$ is included: it has exactly n sliding factors and the same boundary-coordinate formula. □

Theorem B.4 (Classical local-group criteria in sliding-factor coordinates)**.** *For a finite typing $s : \Sigma^+ \to \mathsf{S}$ put $S := \widehat{s}(\Sigma^+)$. The following are equivalent.*

(i) *$S \in \mathbf{LG}$; that is, eSe is a group for every idempotent e of S.*

(ii) *S has no subsemigroup isomorphic to U_1.*

(iii) *Every idempotent of S lies in the minimal ideal of S.*

(iv) *$\widehat{s}$ is local-group-determined. One may take $k = |S|$.*

Proof. The equivalence of (i)–(iii) is classical. More precisely, [36, Proposition 2.3] proves that (i) and (ii) are equivalent to $S^m = \mathsf{K}$ for some m, where K is the minimal ideal. By [36, Lemma 2.2], for $n = |S|$ one has $S^n = SE(S)S$. Thus $S^m = \mathsf{K}$ for some m if and only if $E(S) \subseteq \mathsf{K}$: the forward implication uses $e = e^m$ for every idempotent e, while the converse follows from $S^n = SE(S)S \subseteq \mathsf{K}$ and the fact that S^n is a nonempty ideal containing K.

(iii)⇒(iv). Apply Lemma B.3 to $h = \widehat{s}|_{\Sigma^+}$ with $n = |S|$. Equation $(\star)$ gives the value of $\widehat{s}(w)$ for every $|w| \geq 2n$ from its n-letter prefix, n-letter suffix, and the ordered sliding-factor product $\gamma_{n,\nu}(w)$. The window congruence $\widehat{s}_{n,n}$ records all shorter words exactly and separates ε by the fresh identity. Hence $\widehat{s} \preceq \lambda_{n,\nu}$.

(iv)⇒(i). Since **LG** is closed under homomorphic images, it suffices to show that the positive image $T := \lambda_{k,\nu}(\Sigma^+)$ belongs to **LG**. No short element is idempotent: squaring a short word increases its length, which the window component records exactly up to the long threshold. Let $e = (\langle p, s\rangle, \gamma_e)$ be idempotent. By Lemma B.1, $\gamma_e J_\nu(s,p)\gamma_e = \gamma_e$; moreover, every exe $(x \in T)$ has window $\langle p, s\rangle$. Thus eTe lies in the group $H_e := \{(\langle p, s\rangle, \alpha) : \alpha \in A\}$, whose multiplication is $\alpha * \beta = \alpha J_\nu(s,p)\beta$ and which is isomorphic to A via $\alpha \mapsto \alpha J_\nu(s,p)$. The finite subsemigroup eTe is therefore a group. □

Proposition B.5 (Window–group typings are local-group typings, strictly)**.** (i) *If $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$ for a homomorphism g into a finite group, then $\widehat{s} \preceq \lambda_{\hat{k},\nu}$ with $\hat{k} = \max\{k, \ell, 1\}$ and $\nu(F) := g(\text{last letter of } F)$. Hence $\mathbf{LI} \vee \mathbf{Grp} \subseteq \mathbf{LG}$.*

(ii) *Let $\Sigma = \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}$. Let $h_{\mathtt{ab}} := \lambda_{1,\nu}$ with $\nu(\mathtt{ab}) = 1 \in \mathbb{Z}_2$ and $\nu(F) = 0$ otherwise; it records the parity of the number of occurrences of $\mathtt{ab}$, together with the window. Then $h_{\mathtt{ab}}$ is not window–group-determined. Hence $h_{\mathtt{ab}}(\Sigma^+) \in \mathbf{LG} \setminus (\mathbf{LI} \vee \mathbf{Grp})$.*

Proof. (i) For $|w| \geq 2\hat{k}$, $g(w) = g(\mathrm{pref}_{\hat{k}}\, w)\, \gamma_{\hat{k},\nu}(w)$, and shorter words are recorded exactly. For the inclusion, let $S \in \mathbf{LI} \vee \mathbf{Grp}$ and take $\Sigma = S$ with the evaluation typing $s : \Sigma^+ \twoheadrightarrow S$. Lemma 4.17 makes $\widehat{s}$ window–group determined, hence local-group-determined by the first part, so $S \in \mathbf{LG}$ by Theorem B.4.

(ii) Suppose $h_{\mathtt{ab}} \preceq (\widehat{s}_{k,\ell}, g')$ with $g' : \Sigma^* \to A'$. Let $\hat{k} = \max\{k, \ell, 1\}$ and let q be an even multiple of the exponent of A'. Put $u = \mathtt{a}^{\hat{k}}(\mathtt{abc})^q\mathtt{a}^{\hat{k}}$ and $v = \mathtt{a}^{\hat{k}}(\mathtt{bac})^q\mathtt{a}^{\hat{k}}$. Both are long, with prefix $\mathtt{a}^k$ and suffix $\mathtt{a}^\ell$, and $g'(u) = g'(\mathtt{a})^{2\hat{k}} = g'(v)$. But u contains q occurrences of $\mathtt{ab}$, while v contains exactly one, at $\mathtt{a}^{\hat{k}}\mathtt{b}$, so $h_{\mathtt{ab}}(u) \neq h_{\mathtt{ab}}(v)$. Membership in **LG** follows from Theorem B.4; if the positive image were in $\mathbf{LI} \vee \mathbf{Grp}$, Lemma 4.17 would make $h_{\mathtt{ab}}$ window–group-determined. □

Lemma B.6 (Short tuples with prescribed local-group type)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \geq 1$, and let $\lambda = \lambda_{k,\nu}$. Let $b_{\mathrm{w}} := f(2k+1)n_{\mathrm{w}}(r_{\max}+1)\bar{\theta}_G$, where $n_{\mathrm{w}} \leq |V||M_{k,k}|^f$ is the number of non-start typed nonterminals of the trimmed $\widehat{s}_{k,k}$-refinement. Put $n_\lambda := |V||M_{k,\nu}|^f$, and let $M_{A,f}$ be the constant of Section 4.4. For every $B \in V$ and every realizable $\mathbf{t} \in M_{k,\nu}^{\mathrm{fan}(B)}$ there is $\vec{u} \in L_B(G)$ with $\lambda^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{t}$ and*

$$\|\vec{u}\|_1 \leq \beta_\lambda := (r_{\max}+1)\, b_{\mathrm{w}} \max\{1, (M_{A,f}-1)\, n_\lambda\}.$$

If A is abelian, $M_{A,f} - 1$ may be replaced by $\mathsf{d}(A^f)$.

Proof. Let T be minimal as in Lemma 4.8 with root λ-vector $\mathbf{t}$ and threshold $b = b_{\mathrm{w}} \geq \bar{\theta}_G$. For a heavy node v labelled C, with yield $\vec{x}$ and window vector $\boldsymbol{\omega} = \widehat{s}_{k,k}(\vec{x})$, the typed nonterminal $C_{\boldsymbol{\omega}}$ is retained in the trimmed $\widehat{s}_{k,k}$-refinement, because T can be inserted into a successful derivation (G is reduced). Lemma 4.6 (trivial κ) gives a replacement tuple of C with the same window vector and length at most b_{w}. For pairwise incomparable heavy nodes $v_1, \ldots, v_m$ and $J \subseteq [m]$, let T_J replace the subtrees at v_i, $i \in J$. The components of the v_i are disjoint contiguous blocks of the root components. By Lemma B.2, every T_J has the root window vector of T; a short root component is unchanged, and on each long root component $\gamma_{k,\nu}$ is an ordered product $c_0 y_1 c_1 \cdots y_q c_q$ of block values between constants that do not depend on J. The long root components thus form a block-replacement family over A (Definition A.1). If $m \geq M_{A,f}$, or $m > \mathsf{d}(A^f)$ for abelian A, Lemma A.2 gives a nonempty J with the root λ-vector of T, while T_J is strictly shorter. This contradicts minimality, and Lemma 4.8 with $N = M_{k,\nu}$ gives the bound. □

By Proposition B.5(i), this theorem also applies to window and window–group typings, but its general cubic bound is weaker than the quadratic bounds proved for those subclasses in Theorems 4.7 and 4.15.

C Pattern-Thickness Transfer for General MCFG Presentations

Here are the support-fibre details behind Proposition 4.24. Write $F_{A,I} := \{\vec{u} \in L_A(H_{\mathrm{red}}) : \mathrm{supp}(\vec{u}) = I\}$. Every realized fibre has a shortest tuple, of length at most $t_{\mathrm{pat}}(H_{\mathrm{red}})$.

For clarity, the source-language invariant is preserved *through all four stages*: each intermediate decorated nonterminal is the image of a union of realizable original fibres $F_{A,I}$ under one composition of projection, permutation and consecutive-component merging; where an intermediate support is prescribed, retain only the original fibres whose image has that support. The map is length-nonincreasing, and every such fibre has an actual derivation in the trimmed input. The stagewise bookkeeping is:

Stage	Tuple-language operation	Support/length information
(a) Nondeletion	$\pi_J(L_A)$, union over source fibres	Omitted children have derivations; projection does not increase length
(b) Nonpermutation	$\sigma(\pi_J(L_A))$	Support is permuted; length is unchanged
(c) Empty-pattern removal	Restrict to a realized support I, then π_I	Select preimage source fibres with the required support; length does not increase
(d) Nonmerging	$\mu(\sigma\pi_J(L_A))$, with the support restriction inherited from (c)	Merged support is a function of previous support; length is unchanged

In this table, L_A refers to the original input nonterminal and the operations are understood as applying to the appropriate union of $F_{A,I}$, not to an unrestricted L_A after stage (c). The original rank

and fan-out bounds survive because the four constructions introduce only finitely many component decorations per rule for fixed f, r.

At stage (a) of Lemma F.1, a new nonterminal has language $\pi_J(L_A(H_{\mathrm{red}}))$, a union of the images $\pi_J(F_{A,I})$. Stage (b) permutes components and thus permutes support patterns without changing total length. At stage (c), restricting to a realized output pattern selects the images of those original support fibres whose projected/permuted support is exactly that pattern. Because each such original fibre is nonempty, selecting its shortest tuple yields an output tuple of total length at most $t_{\mathrm{pat}}(H_{\mathrm{red}})$. Stage (d) merges consecutive components; concatenation cannot increase total length, and its output support is determined by the input support, so the same argument applies after merging. The stagewise equalities are equalities of tuple languages, not merely inclusions; hence each retained nonterminal has an actual bounded-length terminal derivation. Fresh nonterminal names are only the finite decorators of these source nonterminals. Final trimming removes symbols without changing the tuple languages of the surviving ones. This proves $t_G \le t_{\mathrm{pat}}(H_{\mathrm{red}})$.

The normal form has $\|G\| = O_{f,r}(\|H\|)$ by Proposition 3.12; consequently $\bar{\theta}_G \le 2\max\{1, \|G\|\}\max\{1, t_G\} = O_{f,r}(\|H\|\bar{t}_{\mathrm{pat}}(H_{\mathrm{red}}))$ by Lemma 4.1(iii). This argument does not assert polynomial control by the *ordinary* untyped rule thickness of H: Proposition F.3 shows why elimination of empty components can be expensive.

Example C.1 (An explicit four-stage fan-out-two transfer). Here is a finite example in which *deletion, permutation, an empty component, and a merge* all occur. It illustrates the exact source-fibre invariant used above, rather than replacing the proof of Lemma F.1. Let the terminal alphabet contain $a, b, c, d, z, \#$ and take fan-out-one start symbol S, fan-out-two symbols A, B, C, and the rank-≤ 2 grammar

$$\begin{aligned} B &\to (a,\varepsilon) \mid (a,b),\\ C &\to (c,d),\\ A &\to (\xi_1^2\xi_1^1, \xi_2^1)(B,C) \mid (\varepsilon, z),\\ S &\to (\xi_1^1 \# \xi_1^2)(A) \mid (\varepsilon). \end{aligned}$$

The first A-rule *deletes* C's second component d, uses B's components in *reverse* order, and puts the adjacent pieces B^2B^1 together. The terminal tuple (ε, z) realizes the support pattern $\{2\}$, whereas the other A-derivations realize $\{1,2\}$. Direct evaluation gives

X	realized tuple fibres (by support)	minimum lengths per fibre
B	$F_{B,\{1\}} = \{(a,\varepsilon)\}, \quad F_{B,\{1,2\}} = \{(a,b)\}$	$1, 2$
C	$F_{C,\{1,2\}} = \{(c,d)\}$	2
A	$F_{A,\{2\}} = \{(\varepsilon,z)\}, \quad F_{A,\{1,2\}} = \{(a,c),(ba,c)\}$	$1, 2$
S	$F_{S,\emptyset} = \{(\varepsilon)\}, \quad F_{S,\{1\}} = \{(\#z),(a\#c),(ba\#c)\}$	$0, 2$

The last column records the minimum length *within each realized support fibre*, rather than the lengths of all its tuples. For example, $\|(a,c)\|_1 = 2$ and $\|(ba,c)\|_1 = 3$, so the minimum for $F_{A,\{1,2\}}$ is 2. Thus $t_{\mathrm{pat}}(H) = 2$ for this H and $L(H) = \{\varepsilon, \#z, a\#c, ba\#c\}$.

Stage (a): projecting C to its used component gives $C^{\{1\}} \to (c)$ with language $\pi_{\{1\}}(L_C) = \{(c)\}$; the deleted d is restored by the omitted component in the original C-derivation. All derived A-tuples are unchanged.

Stage (b): the decorated reverse-order copy of B has language $\{(\varepsilon,a),(b,a)\}$. Reading this copy in its own component order makes the first A-rule nonpermuting. This is precisely a permutation image of L_B, and neither support nor length is silently changed (the support is permuted).

Stage (c): the two nonempty support fibres of this reordered B become $\{(a)\}$ from $\{2\}$ and $\{(b,a)\}$ from $\{1,2\}$. The A-copy of support $\{2\}$ generates (z), whereas the $\{1,2\}$-copy generates

(a, c) and (ba, c). The positive start language is $\{\#z, a\#c, ba\#c\}$; its separate ε-rule is removed. In particular, the all-empty output is not accidentally reconstructed by a good-rule template.

Stage (d): the contiguous pair (b, a) in the reordered B-copy is merged to (ba); the singleton fibre remains (a). The corresponding A-copy still generates exactly (a, c) or (ba, c) and S retains exactly the three positive words. After collapsing duplicate names, one possible *final good grammar* for the nonempty core is the following fully explicit one:

$$\begin{array}{lcllcl} D_0 & \to & (a), & & D_1 & \to & (ba), \\ C_0 & \to & (c), & & A_2 & \to & (z), \\ A_{12} & \to & (\xi_1^1, \xi_2^1)(D_0, C_0) \mid (\xi_1^1, \xi_2^1)(D_1, C_0), & & & & \\ S_+ & \to & (\#\,\xi_1^1)(A_2) \mid (\xi_1^1\,\#\,\xi_1^2)(A_{12}). & & & & \end{array}$$

Here D_0 and D_1 are exactly the two transformed B fibres: $L_{D_0} = \{(a)\}$ and $L_{D_1} = \{(ba)\}$; the omitted d from C contributes no root output. Each displayed component is nonempty, no child component is deleted, permuted, or separately merged at a rule, and the fan-out and rank remain 2. For each output copy, a representative is the image of a tuple in one of the displayed *original* fibres under a composition $\mu\sigma\pi_J$, possibly after conditioning on a realized support pattern. Each map is length-nonincreasing, and the shortest chosen original-fibre representative has length at most two. Thus every surviving decorated nonterminal has a terminal yield of length at most $2 = t_{\mathrm{pat}}(H)$, as asserted by the general transfer. In this example the difference between preserving the *whole source language* and preserving a *support fibre* is visible at stage (c).

D Auxiliary Verification of Qualitative Reconstruction

D.1 Soundness

Proof of Lemma 3.2. Replace the children one at a time, keeping all other children fixed and nonempty. Given an accepting outer context E for the result of any intermediate replacement, Lemma 3.1 induces an admissible context E_j of the child being replaced. Since $\vec{u}_j \equiv_{L,s} \vec{x}_j$, membership in L is unchanged under that replacement, in either direction. The same argument applies at each step even if earlier children have already been replaced. Finally, the componentwise s-types of the output tuples agree because the template is fixed and $\widehat{s}$ is multiplicative on concatenation, including the fresh identity for empty terminal gaps. Therefore the output distributions coincide. □

Proof of Proposition 3.10. Proceed by induction on the height of a derivation rooted at $[\vec{x}]$. A Const rule derives exactly $\vec{x}$. For Link, $[\vec{x}] \to [\vec{y}]$ is witnessed by a common admissible context of same-type tuples in $K \subseteq L$, hence tuple substitutability gives $\vec{x} \equiv_{L,s} \vec{y}$; compose this equivalence with the inductive hypothesis for $[\vec{y}]$. For Comp, each child derivation produces a tuple equivalent to its witnessed child, and Lemma 3.2 transfers those equivalences through the witnessed template. Its original output is exactly the observed parent $\vec{x}$, establishing the invariant. Every positive Start rule points to an observed sentence (w), whose derivational context is the identity; the invariant implies $w \in L$. The empty Start rule is included only for $\varepsilon \in K \subseteq L$. This argument holds at every finite rank cap; for nonempty K the uncapped construction has cap $\|K\|_+$, and the empty case is immediate. □

D.2 Effective Reconstruction and Conservative Learning

Proof of Theorem 3.14. Put $n_K = \max\{1, \|K\|_+\}$. An occurrence is a sample word with at most $2f$ cut positions, so there are $n_K^{O(f)}$ occurrences, typed by fixed-monoid products; grouping occurrences

by context and type lists the Link rules. A composition witness of rank $r \le r_{\max}$ is determined by a parent occurrence and at most $2fr_{\max}$ further cut positions inside it, with a labelling whose number depends only on $f, r_{\max}$; goodness is checked directly. This gives $n_K^{O(fr_{\max})}$ candidate Comp rules; the remaining rules have rank at most one, and deduplication leaves an encoding of the same size. □

Proof of Corollary 3.16. For fixed rule rank, qualitative identification follows from the direct soundness and exact-reconstruction results (Proposition 3.10 and Theorem 3.13) with the conservative policy of [14, Algorithm 1]; the finite-signature antecedent is Yoshinaka [15, Theorem 1]. For unbounded rank, Theorem 3.13 provides a finite characteristic sample, while Proposition 3.10 gives soundness. Membership is decidable [29]. After a text has presented the characteristic sample, the next rebuild (if any) is correct and permanent; if no rebuild follows, sample consistency and conservativeness imply that the current hypothesis already generates every text word and hence L. For fixed rank, the per-datum bound follows from Theorem 3.14 and [14, Proposition 4]. □

Update complexity versus convergence stage. The polynomial-time guarantee of Theorem 3.14 and Corollary 3.16 concerns hypothesis reconstruction and membership testing at each finite stage, measured against the data seen so far. TxtEx convergence in Corollary 3.16 does not bound the first stage of stabilization or the number of mind changes polynomially: a text may postpone the first presentation of a particular positive example arbitrarily long. Short characteristic data by itself gives no presentation-independent convergence-stage bound.

Every CFG has a rank-two presentation, so $\mathcal{A}_{1,s}^{\le 2}$ covers $\mathcal{C}_{1,s}^{\mathrm{mcf}}$; this does not imply polynomial characteristic-data length at rank two (Section 4).

Proof of Proposition 3.11. This is the rule-witness argument of [15, Lemma 3], applied to the concrete Link/Const/Comp presentation. For each typed rule put $\vec{z} = \rho(\iota(Y_1), \dots, \iota(Y_q))$. The two witnesses supply Link $[\iota(X)] \to [\vec{z}]$, followed by Const if $q = 0$ or Comp of rank $q \le r_G$ otherwise. The latter also belongs to the uncapped grammar, since $q \le \|\vec{z}\|_1 \le \|K\|_+$. Induction on a source derivation thus simulates $X \Rightarrow^*_{\widetilde{G}_0} \vec{u}$ from $[\iota(X)]$; $\chi(X) = \square_1$ at a successful start link supplies Start. Apply Proposition 3.6(iii). □

Proof of Theorem 3.13. If $L \cap \Sigma^+ = \emptyset$ then $K = L \in \{\emptyset, \{\varepsilon\}\}$ and the claim is immediate from Definition 3.4. Otherwise $\mathcal{F}_L \subseteq L$ is finite by the remark after Definition 3.9. Proposition 3.11 with $\iota = \omega$ (admissible by Proposition 3.6(i)) gives $L \cap \Sigma^+ \subseteq L(\mathcal{B}(K))$, and Proposition 3.10 gives $L(\mathcal{B}(K)) \subseteq L$, for both operators. The hypothesis contains ε iff $\varepsilon \in K$, iff $\varepsilon \in L$, because $\mathcal{F}_L \subseteq K \subseteq L$. Rank-preserving normalization gives the last sentence. □

E Auxiliary Proofs for Typed Witness Construction

The statements used for quantitative witness bounds remain in Section 3.2; the routine verification details are recorded here.

Proof of Proposition 3.6. At every rule application, the monoid morphism property gives

$$\widehat{s}^{(\mathrm{fan}(A))}(\rho(\vec{u}_1, \dots, \vec{u}_r)) = \mathrm{out}_\rho(\widehat{s}^{(d_1)}(\vec{u}_1), \dots, \widehat{s}^{(d_r)}(\vec{u}_r)).$$

Induction on derivation height gives (i) and the unique lifting (ii). Erasing indices gives $L(\widehat{G^s}) \subseteq L(G)$, (ii) the converse, and trimming keeps every rule of every successful derivation. □

Lemma E.1 (Monoid implementations)**.** *Let $h_0 : \Sigma^* \to M_0$ be a finite-monoid morphism, $S_0 := h_0(\Sigma^+)$ and $s_0 := h_0|_{\Sigma^+}$.*

(i) *For $1 \le d \le f$ and $\vec{x}, \vec{y} \in (\Sigma^+)^d$, the equalities $h_0^{(d)}(\vec{x}) = h_0^{(d)}(\vec{y})$, $s_0^{(d)}(\vec{x}) = s_0^{(d)}(\vec{y})$ and $\widehat{s_0}^{(d)}(\vec{x}) = \widehat{s_0}^{(d)}(\vec{y})$ are equivalent. Hence $\mathcal{C}^{\mathrm{mcf}}_{f,h_0} = \mathcal{C}^{\mathrm{mcf}}_{f,s_0}$, and the reconstruction operators built with h_0 and with $\widehat{s_0}$ return the same grammar on every sample.*

(ii) *For every good MCFG G, the trimmed refinements of G under h_0 and under $\widehat{s_0}$ coincide (the untrimmed ones need not); hence the same canonical tuples, contexts and witness words may be chosen for both.*

Proof of Lemma E.1. (i) The three maps agree on Σ^+ as maps into S_0, and the class and the Link guard compare types of nonempty components only. (ii) S_0 is a subsemigroup of M_0 and of $S_0^\sharp$ with the same multiplication. By λ-freeness and Proposition 3.6(i), every typed nonterminal occurring in a successful derivation under either morphism is some $A_{\mathbf{t}}$ with $\mathbf{t} \in S_0^{\mathrm{fan}(A)}$, and the start links used are $\widetilde{S} \to S_{(\eta)}$ with $\eta \in S_0$. For child indices in S_0, every component α_i of a good template is a nonempty word whose letters and variables are evaluated in S_0 identically under both morphisms, so both give the same out_ρ. Thus the two full refinements contain the same rules among S_0-indexed symbols, which carry all successful derivations. □

Proof of Lemma 3.7. Take a successful typed derivation containing X, and follow the path from the root to an occurrence of X, starting with $\square_1$. At each rule on the path apply Lemma 3.1 with the (nonempty) yields of the sibling subderivations fixed; this gives $\chi(X) \in \mathcal{E}_d$. The rules on the path depend only on the label X, so replacing the subtree of X by any X-derivation yields a successful derivation of $\chi(X)[\vec{u}]$. If $\widetilde{S} \to X$ is a rule, use the path of length zero. □

Proof of Lemma 3.8. For $\vec{u}, \vec{v} \in L_X$, Proposition 3.6(i) gives equal types and $\chi(X)$ accepts both; substitutability gives $\mathcal{D}_L(\vec{u}) = \mathcal{D}_L(\vec{v})$. □

F Details of the Good Normalization

The normalization uses the four stagewise constructions in Lemma F.1; their semantic and size invariants are verified below.

We record the four conversions used in Proposition 3.12 once, together with the tuple-language identity that each satisfies. Both the fixed-parameter size bound of that proposition and the thickness transfer of Proposition F.2 are read off from these identities. Semantic correctness of the nondeleting and nonmerging conversions is standard [14, 29, 31]; Kanazawa's existence statement [31, Section 2.1] does not by itself give the size or thickness bounds, which are therefore derived here. For $\vec{w} = (w_1, \dots, w_d)$ and nonempty $I \subseteq [d]$, $\pi_I(\vec{w})$ retains the components indexed by I in increasing order, $\sigma\vec{w}$ permutes components, and $\mu(\vec{w})$ concatenates consecutive blocks of components; thus $\|\pi_I(\vec{w})\|_1 \le \|\vec{w}\|_1$ and $\|\sigma\vec{w}\|_1 = \|\mu(\vec{w})\|_1 = \|\vec{w}\|_1$.

Lemma F.1 (Stagewise tuple languages)**.** *Fix f and $r_{\max}$, and let G_0 be a reduced f-MCFG of rule rank at most $r_{\max}$. Each of the following conversions preserves linearity, does not increase fan-out or rule rank, and multiplies the number of nonterminal and rule copies by a factor depending only on f and $r_{\max}$, without enlarging any template.*

(a) Nondeletion *(the construction (f3) in the proof of [29, Lemma 2.2]). Introduce A^I for nonempty $I \subseteq [\mathrm{fan}(A)]$. For $R : A \to \rho(B_1, \dots, B_q)$ let J_j be the set of components of B_j used by the*

I-indexed outputs of ρ; keep these outputs, rename their variables, use $B_j^{J_j}$ when $J_j \neq \emptyset$ and omit B_j otherwise. The resulting rules are nondeleting, and $L_{A^I} = \pi_I(L_A(G_0))$.

(b) Nonpermutation. *Decorate a fan-out-d nonterminal by $\sigma \in \mathfrak{S}_d$; the order in which a permuted parent template uses a child's variables determines the child's permutation. The decorated nonterminal generates σL_A.*

(c) Empty components. *Let $\mathsf{Pat}(A)$ be the set of nonempty-component patterns of tuples in L_A, computed by finite saturation. For realizable child patterns, replace variables of absent components by ε, delete empty parent components and omit all-empty children; introduce A^I for each nonempty $I \in \mathsf{Pat}(A)$. Then L_{A^I} is the set of π_I-projections of the tuples of L_A with pattern I, the result is λ-free, and $S^{\{1\}}$ generates the nonempty part of the language. If every derived component is already nonempty, this stage is the identity.*

(d) Nonmerging *(the construction in the proof of [14, Section 2.3, Lemma 1]). Decorate a nonterminal by a composition μ of its fan-out; in a merged parent template, maximal adjacent runs of consecutive variables of a child determine that child's merge. The merge-decorated nonterminal generates the* image $L_{A^\mu} = \mu(L_A)$.

Stages (b)–(d) preserve the conditions established by earlier stages.

Proof. (a) For arbitrary child tuples, $\pi_I(\rho(\vec{v}_1, \ldots, \vec{v}_q)) = \rho^I\big((\pi_{J_j}\vec{v}_j)_{J_j \neq \emptyset}\big)$, so induction on derivations gives $\subseteq$. For $\supseteq$, lift every retained child tuple by the induction hypothesis and supply a terminal derivation of each omitted child; this uses that G_0 is reduced, hence every omitted child is productive. Without prior trimming only $\pi_I(L_A) \subseteq L_{A^I}$ would be automatic. There are at most 2^f copies of each nonterminal, and I determines the J_j, so at most 2^f copies of each rule; omitting children cannot increase rank.

(b) Linearity and nondeletion make the induced child permutations well defined; induction in both directions identifies the language with σL_A. There are at most $f!$ copies, no child is duplicated, and total length, λ-freeness, fan-out and rank are preserved.

(c) Retained variables still occur once and in order, so nondeletion, nonpermutation, fan-out and rank are preserved. Induction on derivations gives the projection identity; omitted all-empty children are restored by terminal derivations realizing their all-empty pattern. There are at most 2^f patterns per nonterminal and $2^{f r_{\max}}$ choices per rule, and projection only shortens templates.

(d) Nonpermutation makes the maximal runs compatible with each child's component order, and nondeletion makes them cover all its components. The merged template satisfies $\rho^\mu(\mu_1(\vec{v}_1), \ldots, \mu_q(\vec{v}_q)) = \mu(\rho(\vec{v}_1, \ldots, \vec{v}_q))$, and induction gives the image equality. No inverse-image claim is made, because the merge map need not be injective. Concatenation preserves total length, λ-freeness, nonpermutation, fan-out and rank; for fixed dimension and rank the conversion has linear size [14, Section 2.3, after Example 2]. □

Proof of Proposition 3.12. Apply (a)–(d) to the trimmed input and trim again. Each stage multiplies the encoding by a factor depending only on $f, r_{\max}$ and trimming can only decrease it, so $\|G_{\text{good}}\| = O_{f,r_{\max}}(\|G_{\text{in}}\|)$, and $L(G_{\text{good}}) = L(G_{\text{in}}) \setminus \{\varepsilon\}$ by (c). If this language is nonempty, the start symbol is productive and the final trimming leaves a reduced grammar; trimming preserves goodness. □

Thickness transfer for λ-free presentations

Proposition F.2 (Thickness under good-form conversion)**.** *Fix f and $r_{\max}$. Let G_{in} be a reduced λ-free f-MCFG of maximum rule rank at most $r_{\max}$, possibly with deleting rules, and suppose*

$L(G_{\mathrm{in}}) \cap \Sigma^+ \neq \emptyset$. *The good normalization of Proposition 3.12 yields an equivalent reduced good grammar G with the same fan-out and rule-rank bounds such that*

$$\|G\| = O_{f,r_{\max}}(\|G_{\mathrm{in}}\|), \qquad t_G \leq t_{G_{\mathrm{in}}}, \qquad \theta_G = O_{f,r_{\max}}(\|G_{\mathrm{in}}\| \max\{1, t_{G_{\mathrm{in}}}\}).$$

Proof. Because G_{in} is λ-free, every component of every derived tuple is nonempty, by induction on derivations; this persists through (a) and (b), so stage (c) is the identity. By Lemma F.1, every nonterminal X of the final grammar is associated with an original nonterminal A such that

$$L_X(G) = \mu\big(\sigma\pi_I(L_A(G_{\mathrm{in}}))\big), \qquad \min_{\vec{u} \in L_X(G)} \|\vec{u}\|_1 \leq \min_{\vec{w} \in L_A(G_{\mathrm{in}})} \|\vec{w}\|_1.$$

Trimming leaves the tuple languages of surviving nonterminals unchanged, so $t_G \leq t_{G_{\mathrm{in}}}$, and the size bound was shown above. Finally Lemma 4.1(iii) gives $\theta_G \leq 2\|G\| \max\{1, t_G\} = O_{f,r_{\max}}(\|G_{\mathrm{in}}\| \max\{1, t_{G_{\mathrm{in}}}\})$. □

Thus deleting rules do not by themselves obstruct polynomial thickness transfer from a reduced λ-free presentation; genuinely empty components are the problematic step in the scale comparison of Table 1.

Proposition F.3 (Empty-component elimination can blow up thickness)**.** *For every $n \geq 1$ there is a reduced fan-out-two, rule-rank-two MCFG E_n of size $O(n)$ with $t_{E_n} \leq 2$ and $\theta_{E_n} = O(n)$ such that the empty-component elimination stage of Proposition 3.12 contains a retained nonterminal and a retained rule of thickness at least 2^n.*

Proof. Let E_n have fan-out-one S, fan-out-two $A_0, \ldots, A_n$, and

$$A_0 \to (\mathtt{a}, \varepsilon), \qquad A_i \to (\varepsilon, \mathtt{c}) \mid (\xi_1^1\xi_2^1, \xi_1^2\xi_2^2)(A_{i-1}, A_{i-1}) \quad (1 \leq i \leq n), \qquad S \to (\xi_1^1 \mathtt{d}\, \xi_1^2)(A_n).$$

It is reduced, nondeleting, nonpermuting, and nonmerging. Short side derivations give $t_{E_n} \leq 2$, and placing any chosen rule on one branch of the binary spine while completing siblings minimally gives $\theta_{E_n} = O(n)$. If an A_i-tuple has nonempty first and empty second component, the rule $(\varepsilon, \mathtt{c})$ cannot occur in it; inductively the shortest such tuple is $(\mathtt{a}^{2^i}, \varepsilon)$. Thus pattern $\{1\}$ for A_n survives, and the projected start rule $S \to \xi_1^1 \mathtt{d}\, (A_n^{\{1\}})$ succeeds only with a child of length at least 2^n. □

This is a distinct obstruction from Proposition 4.5: here a *support pattern* forces an exponentially long tuple already in the source presentation, while the earlier example concerns the thickness of a *typed refinement*. Neither assertion alone is an operator-independent characteristic-data lower bound.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During preparation of this work, the author used ChatGPT (OpenAI) and Claude (Anthropic) to assist with organization of the manuscript, expression and translation of text, cross-checks of references and notation, and editing of LaTeX source. After using these tools, the author reviewed and edited the material as needed and takes full responsibility for the content of the publication.

References

[1] E. M. Gold. Language identification in the limit. *Information and Control*, 10(5):447–474, 1967.

[2] D. Angluin. Inductive inference of formal languages from positive data. *Information and Control*, 45(2):117–135, 1980.

[3] A. Clark and R. Eyraud. Polynomial identification in the limit of substitutable context-free languages. *Journal of Machine Learning Research*, 8:1725–1745, 2007.

[4] C. de la Higuera. Characteristic sets for polynomial grammatical inference. *Machine Learning*, 27:125–138, 1997.

[5] F. Coste, D. Fredouille, C. Kermorvant, and C. de la Higuera. Introducing domain and typing bias in automata inference. In *Grammatical Inference: Algorithms and Applications – ICGI 2004*, LNCS 3264, pp. 115–126, Springer, 2004.

[6] Y. Takada. Grammatical inference for even linear languages based on control sets. *Information Processing Letters*, 28(4):193–199, 1988. doi:10.1016/0020-0190(88)90208-6.

[7] Y. Takada. Learning formal languages based on control sets. In K. P. Jantke and S. Lange (eds.), *Algorithmic Learning for Knowledge-Based Systems*, LNCS 961, pp. 316–339, Springer, 1995.

[8] A. Clark. Learning trees from strings: A strong learning algorithm for some context-free grammars. *Journal of Machine Learning Research*, 14:3537–3559, 2013.

[9] F. Coste, G. Garet, and J. Nicolas. Local substitutability for sequence generalization. In *Proceedings of the 11th International Conference on Grammatical Inference*, JMLR Workshop and Conference Proceedings 21, pp. 97–111, 2012.

[10] R. Eyraud, J. Heinz, and R. Yoshinaka. Efficiency in the identification in the limit learning paradigm. In J. Heinz and J. M. Sempere (eds.), *Topics in Grammatical Inference*, pp. 25–46. Springer, 2016. doi:10.1007/978-3-662-48395-4_2.

[11] F. Coste and J. Nicolas. Learning local substitutable context-free languages from positive examples in polynomial time and data by reduction. In *Proceedings of the 14th International Conference on Grammatical Inference*, Proceedings of Machine Learning Research 93, pp. 155–168, 2019.

[12] R. Yoshinaka. Identification in the limit of k, l-substitutable context-free languages. In *Grammatical Inference: Algorithms and Applications*, LNCS 5278, pp. 266–279. Springer, 2008.

[13] R. Yoshinaka. Learning mildly context-sensitive languages with multidimensional substitutability from positive data. In *Algorithmic Learning Theory*, LNCS 5809, pp. 278–292. Springer, 2009. doi:10.1007/978-3-642-04414-4_24.

[14] R. Yoshinaka. Efficient learning of multiple context-free languages with multidimensional substitutability from positive data. *Theoretical Computer Science*, 412(19):1821–1831, 2011. doi:10.1016/j.tcs.2010.12.058.

[15] R. Yoshinaka. General perspective on distributionally learnable classes. In *Proceedings of the 14th Meeting on the Mathematics of Language (MoL 2015)*, pp. 87–98. Association for Computational Linguistics, 2015. doi:10.3115/v1/W15-2308.

[16] R. Yoshinaka. Polynomial-time identification of multiple context-free languages from positive data and membership queries. In *Grammatical Inference: Theoretical Results and Applications*, LNCS 6339, pp. 230–244. Springer, 2010. doi:10.1007/978-3-642-15488-1_19.

[17] A. Clark and R. Yoshinaka. Distributional learning of context-free and multiple context-free grammars. In J. Heinz and J. M. Sempere (eds.), *Topics in Grammatical Inference*, pp. 143–172. Springer, 2016. doi:10.1007/978-3-662-48395-4_6.

[18] A. Clark and R. Yoshinaka. Distributional learning of parallel multiple context-free grammars. *Machine Learning*, 96(1–2):5–31, 2014. doi:10.1007/s10994-013-5403-2.

[19] T. Oates, T. Armstrong, L. Becerra-Bonache, and M. Atamas. Inferring grammars for mildly context-sensitive languages in polynomial-time. In *Grammatical Inference: Algorithms and Applications – ICGI 2006*, LNCS 4201, pp. 137–147. Springer, 2006.

[20] M. Kanazawa and R. Yoshinaka. Extending distributional learning from positive data and membership queries. In *Proceedings of the 16th International Conference on Grammatical Inference*, Proceedings of Machine Learning Research 217, pp. 8–22, 2023.

[21] P. D. Gallot. *Safety of transformations of data trees: Tree transducer theory applied to a verification problem on shell scripts*. PhD thesis, Université de Lille, 2021. `https://pepite-depot.univ-lille.fr/LIBRE/EDMADIS/2021/2021LILUB030.pdf`.

[22] A. Bishop, M. Elder, A. Evetts, P. Gallot, and A. Levine. On groups with EDT0L word problem. *International Journal of Algebra and Computation*, 36(4):425–486, 2026. doi:10.1142/S0218196726500165.

[23] A. Duncan, M. Elder, L. Frenkel, and M. Lyu. A substitution lemma for multiple context-free languages. *International Journal of Algebra and Computation*, 2026 (online publication). doi:10.1142/S0218196726500463.

[24] O. Rambow and G. Satta. Independent parallelism in finite copying parallel rewriting systems. *Theoretical Computer Science*, 223(1–2):87–120, 1999. doi:10.1016/S0304-3975(97)00190-4.

[25] T. Kuriyama. Distributional learning of context-free languages under fixed finite-semigroup typing. Revised manuscript submitted to *Theoretical Computer Science*, October 2026 (TCS-D-26-00494R1). The cited Proposition 6.4 is in this revision; the earlier, substantively different public preprint is arXiv:1409.6247v4. `https://arxiv.org/abs/1409.6247`.

[26] J.-E. Pin. *Varieties of Formal Languages*. North Oxford Academic, London, and Plenum, New York, 1986.

[27] J. C. Costa. Some pseudovariety joins involving locally trivial semigroups. *Semigroup Forum*, 64(1):12–28, 2002. doi:10.1007/s002330010060.

[28] J. C. Costa. Some pseudovariety joins involving groups and locally trivial semigroups. In *Semigroups, Algorithms, Automata and Languages*, pp. 341–348. World Scientific, 2002.

[29] H. Seki, T. Matsumura, M. Fujii, and T. Kasami. On multiple context-free grammars. *Theoretical Computer Science*, 88(2):191–229, 1991. doi:10.1016/0304-3975(91)90374-B.

[30] L. Kallmeyer. *Parsing Beyond Context-Free Grammars*. Springer, Berlin, Heidelberg, 2010. doi:10.1007/978-3-642-14846-0.

[31] M. Kanazawa. Ogden's lemma, multiple context-free grammars, and the control language hierarchy. *Information and Computation*, 269:104449, 2019. doi:10.1016/j.ic.2019.104449.

[32] C. Gómez-Rodríguez, M. Kuhlmann, G. Satta, and D. Weir. Optimal reduction of rule length in linear context-free rewriting systems. In *Proceedings of Human Language Technologies: The 2009 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 539–547. Association for Computational Linguistics, 2009.

[33] A. Geroldinger and F. Halter-Koch. *Non-Unique Factorizations: Algebraic, Combinatorial and Analytic Theory*, volume 278 of *Pure and Applied Mathematics*. Chapman & Hall/CRC, Boca Raton, 2006.

[34] G. Ramalingam and T. Reps. An incremental algorithm for a generalization of the shortest-path problem. *Journal of Algorithms*, 21(2):267–305, 1996. doi:10.1006/jagm.1996.0046.

[35] M.-J. Nederhof. Weighted deductive parsing and Knuth's algorithm. *Computational Linguistics*, 29(1):135–143, 2003.

[36] J. Almeida, S. Margolis, B. Steinberg, and M. Volkov. Representation theory of finite semigroups, semigroup radicals and formal language theory. *Transactions of the American Mathematical Society*, 361(3):1429–1461, 2009. doi:10.1090/S0002-9947-08-04712-0.

[37] H. Straubing. Finite semigroup varieties of the form $\mathbf{V} * \mathbf{D}$. *Journal of Pure and Applied Algebra*, 36:53–94, 1985. doi:10.1016/0022-4049(85)90062-3.

[38] B. Tilson. Categories as algebra: an essential ingredient in the theory of monoids. *Journal of Pure and Applied Algebra*, 48:83–198, 1987. doi:10.1016/0022-4049(87)90108-3.

[39] C. H. Papadimitriou. On the complexity of integer programming. *Journal of the ACM*, 28(4):765–768, 1981. doi:10.1145/322276.322287.