Characteristic-Data Bounds for Multiple Context-Free Languages under Fixed Finite-Semigroup Typing

Takayuki Kuriyama
Independent Researcher, Tokyo, Japan
growup.kuriyama@gmail.com

Abstract

We bound the *total encoded length* of characteristic positive samples, $\|K\|_+ = \sum_{w\in K}(|w|+1)$, for multiple context-free languages that are tuple-substitutable relative to a fixed finite-semigroup typing. The qualitative reconstruction belongs to Yoshinaka's distributional learning framework. Under fixed fan-out, rule rank and typing, hypothesis construction is polynomial-time in the encoded length of the observed finite sample, and a separate conservative sequential learner provides Gold identification.

For a reduced good presentation G, write $\|G\|$ for its size and $\bar{\theta}_G = \max\{1, \theta_G\}$ for its nonzero rule-thickness parameter. Both prefix–suffix window typings and windows combined with *any* finite group, including noncommutative groups, admit characteristic samples of total encoded length $O(\|G\|^2\bar{\theta}_G)$. The group result uses marked heavy nodes and owner-wise simultaneous replacement while preserving ordered products across output components. A typing whose positive-image semigroup is a local group admits the bound $O(\|G\|^3\bar{\theta}_G)$. We also establish further polynomial bounds for rank-one grammars and abelian-filtered targets.

A boundary-typed non-context-free witness has intrinsic fan-out and rule rank two. The companion CFG exponential lower-bound family fails 2-dimensional substitutability under every fixed finite typing. These are sufficient upper-bound conditions; their optimality is not claimed. Superpolynomial characteristic-data lower bounds at fan-out two or higher remain open.

Keywords: grammatical inference; multiple context-free grammars; positive data; finite-semigroup typing; multidimensional substitutability; Gold identification; rule rank.

1 Introduction

Positive-data grammatical inference asks when a language can be recovered from an enumeration of its positive examples alone [1, 2]. Distributional approaches make this possible for restricted grammar classes by turning observed interchangeability of substrings or tuples into grammar nonterminals. For context-free grammars (CFGs) this includes substitutable languages [3] and Yoshinaka's (k, ℓ) hierarchy [12]. Yoshinaka extended the idea to multiple context-free grammars (MCFGs), whose nonterminals generate tuples and whose substitutions are tested by multi-hole sentence contexts [13,14]. MCFGs are a standard mildly context-sensitive formalism for discontinuous and cross-serial dependencies [30, 31].

This paper fixes, as part of the learning bias, a surjective finite-semigroup homomorphism $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ and compares observed *nonempty* tuple components only when their s-values agree. The canonical fresh-identity completion $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$ is used only to evaluate grammar templates, including empty terminal gaps. The typing and the fan-out bound are fixed before any data are

seen, separating the choice of a finite bias from learning under that bias [5]; prefix–suffix windows and DFA transition morphisms are typical instances. Other biases restrict derivations by regular control sets [6, 7] or restrict the evidence used to compare factors [9, 11]; positive-data learning of other mildly context-sensitive formalisms is studied in [19], and MCFG learning with membership queries or dual context sets in [16, 18, 20]. A recent substitution lemma for MCFLs [23] is a necessary condition for membership, not a learning criterion.

Prior work and scope. The qualitative observed-tuple grammar and identification mechanism derive from Yoshinaka's MCFG work [13,14]; his finite-signature framework also covers bounded-rank learning [15]. Its main proof concerns the context-indexed $\mathrm{SUBST}_{\mathrm{D}}$ constructor; our soundness and completeness arguments establish the corresponding substructure-indexed construction (Lemma 3.18). Polynomial characteristic data for prefix–suffix typed CFGs are due to Yoshinaka [12], while the companion CFG paper supplies the fan-out-one $\mathbf{LI} \vee \mathbf{Grp}$ bound and the semilattice lower bound [25]. There is also a quantitative distinction in the *degree* of the stated bounds: Yoshinaka's (k, ℓ)-Greibach-normal-form construction, for $k + \ell \geq 1$, bounds the number of witnesses by $|P'|^{7(k+\ell)+1}$ for the converted grammar G' with rules P' [12, Lemmas 2, 6 and 7]. Our $O(\|G\|^2 \bar{\theta}_G)$ bound is relative to a reduced good presentation and has grammar-size degree two independent of fixed k, ℓ; the window parameters affect the hidden constant. Since the presentations and thickness measures differ, this compares the *displayed upper bounds*, not their optimality or a lower bound for other learning algorithms. The rank-one separation uses Gallot's result [21]. The present contribution is quantitative control of *total* witness length for typed MCFG tuples, particularly simultaneous shortening that preserves noncommutative ordered products across output components and the local-group extension.

Contributions. Our object is the *total encoded length* of characteristic positive data relative to an MCFG presentation, for typed *tuples* at higher rule rank, where several component boundaries must be preserved simultaneously. An untyped pruning argument cannot simply shorten child derivations while retaining their types: in noncommutative finite groups, blocks from one subtree may appear in different orders across different output components. The new quantitative step is an *owner-wise simultaneous* selection of replacements preserving all ordered root products, rather than preserving each omitted subtree's group value separately. The bounds below hold with $f, r_{\max}, |\Sigma|$ and the typing fixed; their hidden constants can grow rapidly with these fixed parameters. For a reduced good MCFG G put $\bar{\theta}_G = \max\{1, \theta_G\}$, where θ_G is rule thickness (Section 4). With alphabet, fan-out, rule-rank bound and typing fixed:

Fixed typing regime	Positive-data length	Source
Rule rank one; arbitrary typing	$\mathrm{poly}(\|G_{\mathrm{in}}\|)$	Thm. 4.6
Prefix–suffix window	$O(\|G\|^2 \bar{\theta}_G)$	Thm. 4.8
Window and any finite group	$O(\|G\|^2 \bar{\theta}_G)$	Thm. 4.14
Positive image in **LG**	$O(\|G\|^3 \bar{\theta}_G)$	Thm. 4.18

The group result includes noncommutative groups; the abelian case has sharper constants (Theorem 4.15). The **LG** hypothesis concerns the positive-image semigroup; its known algebraic characterization [37] is distinct from the learning bound. Even at fan-out one, our local-group bound strictly extends the companion CFG paper's sufficient condition $\mathbf{LI} \vee \mathbf{Grp}$ (Proposition C.5). This is a consequence of the new MCFG shortening argument, not a claim that the companion already established the local-group theorem. The matching pseudovariety-level lower-bound direction is currently established only at fan-out one (Corollary C.7). We also give a computable rule-only data

budget (Proposition 4.20), a language-based bound for abelian-filtered targets (Theorem 4.22), and a boundary-typed non-context-free target of intrinsic fan-out and rule rank two (Theorem 5.3). Characteristic data concern language-level exactness of a set-driven operator; syntactic stabilization is supplied by a separate conservative learner.

Section 2 fixes the model, Section 3 the reconstruction and its exactness, Section 4 the quantitative bounds, Section 5 the separations, and Section 6 the higher-arity obstruction.

2 Preliminaries

Throughout, Σ is a finite alphabet, ε the empty word, and $\Sigma^+ := \Sigma^* \setminus \{\varepsilon\}$. If $\Sigma = \emptyset$, the only languages are $\emptyset$ and $\{\varepsilon\}$ and are handled directly.

2.1 Positive-Data Learning

Definition 2.1 (Text, TxtEx identification, conservativeness)**.** Fix a pause symbol $\perp \notin \Sigma$. A *text* for $L \subseteq \Sigma^*$ is an infinite sequence over $L \cup \{\perp\}$ in which every element of L occurs; thus $\perp, \perp, \ldots$ is the text for $\emptyset$. The *content* of a finite prefix is the set of elements of Σ^* occurring in it. A *learner* is a computable map from finite prefixes to grammars. It *TxtEx-identifies* L [1] if on every text for L its outputs are, from some stage on, one fixed grammar H with $L(H) = L$, and it TxtEx-identifies a class if it does so for each member. A learner is *conservative* if it changes its hypothesis only on a datum $w \in \Sigma^*$ not generated by the current hypothesis. We fix a hypothesis H_0 with $L(H_0) = \emptyset$.

Definition 2.2 (Set-driven operator and characteristic sample)**.** A *set-driven reconstruction operator* $\mathcal{B}$ maps each finite $K \subseteq \Sigma^*$ to a grammar $\mathcal{B}(K)$ depending only on the set K; it is *sample-consistent* if $K \subseteq L(\mathcal{B}(K))$. For a language L, a finite $\mathcal{F} \subseteq L$ is a *characteristic sample for* $\mathcal{B}$ if $\mathcal{F} \subseteq K \subseteq L$ implies $L(\mathcal{B}(K)) = L$ for every finite K (cf. [4]). Unlike TxtEx identification this is a statement about generated languages: $\mathcal{B}(K)$ need not stabilize syntactically as K grows.

For finite $K \subseteq \Sigma^*$ put $\|K\|_+ := \sum_{w \in K}(|w| + 1)$, and for a tuple $\vec{u} = (u_1, \ldots, u_d)$ put $\|\vec{u}\|_1 := \sum_i |u_i|$. Every polynomial bound below states whether it concerns cardinality or $\|\cdot\|_+$.

2.2 Fixed Finite-Semigroup Typing

Following the companion CFG paper [25], the learning bias is a surjective semigroup morphism $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ onto a finite semigroup, given by its multiplication table and the letter images (so $s(w)$ is computable in linear time). A finite codomain can always be replaced by the positive image, so surjectivity loses no generality.

Definition 2.3 (Fresh completion, tuple types, refinement)**.** Let $\mathsf{S}^\sharp = \mathsf{S} \uplus \{\mathbf{1}_\sharp\}$ be obtained by adjoining a *fresh* identity (even if S already has one). The canonical monoid extension $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$ satisfies $\widehat{s}(\varepsilon) = \mathbf{1}_\sharp$ and $\widehat{s}(w) = s(w)$ for $w \in \Sigma^+$. For $\vec{w} = (w_1, \ldots, w_d)$ write $\widehat{s}^{(d)}(\vec{w}) := (\widehat{s}(w_1), \ldots, \widehat{s}(w_d))$; on $(\Sigma^+)^d$ it agrees with $s^{(d)}$. The fresh identity evaluates empty template factors but does not create an additional positive type. For typings $s, t : \Sigma^+ \to$ finite semigroups write $s \preceq t$ if $\ker t \subseteq \ker s$ on Σ^+, and (s, t) for their product restricted to its positive image.

Positive typing and auxiliary morphisms. Empty terminal gaps in rule templates are evaluated by $\widehat{s}$ with values in $\mathsf{S}^\sharp$. An auxiliary finite-monoid morphism $q : \Sigma^* \to N$ represents the positive typing $q|_{\Sigma^+}$, and Lemma 3.7 proves equivalence with its fresh-identity completion on all nonempty components and all successful trimmed refinements. In particular, algebraic conditions such as **LG** concern the positive image S, not $\mathsf{S}^\sharp$. Write s_{triv} for the one-element positive typing.

Definition 2.4 (Prefix–suffix typing)**.** For $k, \ell \geq 0$ put $m_0 := \max\{1, k+\ell\}$, and for $|w| \geq m_0$ let $\mathrm{pref}_k(w)$, $\mathrm{suf}_\ell(w)$ be the prefix and suffix of the indicated lengths (length-zero windows are ε). Define

$$\widehat{s}_{k,\ell}(w) = \begin{cases} w, & |w| < m_0, \\ \langle \mathrm{pref}_k(w), \mathrm{suf}_\ell(w) \rangle, & |w| \geq m_0, \end{cases}$$

with disjoint tags for short words and pairs, and put $M_{k,\ell} := \widehat{s}_{k,\ell}(\Sigma^*)$. For short x, y and pairs $\langle u, v\rangle, \langle u', v'\rangle$ define

$$x \cdot y := \widehat{s}_{k,\ell}(xy), \qquad x \cdot \langle u, v\rangle := \langle \mathrm{pref}_k(xu), v\rangle,$$
$$\langle u, v\rangle \cdot y := \langle u, \mathrm{suf}_\ell(vy)\rangle, \quad \langle u, v\rangle \cdot \langle u', v'\rangle := \langle u, v'\rangle.$$

A case check gives $\widehat{s}_{k,\ell}(xy) = \widehat{s}_{k,\ell}(x) \cdot \widehat{s}_{k,\ell}(y)$; since $\widehat{s}_{k,\ell}$ is onto $M_{k,\ell}$, associativity is inherited from concatenation and $\widehat{s}_{k,\ell}(\varepsilon)$ is the identity. Thus $\widehat{s}_{k,\ell}$ is a finite-monoid morphism; its positive restriction $s_{k,\ell}$ is the boundary typing, and since the short-word tag separates ε from Σ^+, $\widehat{s}_{k,\ell}$ is isomorphic to $\widehat{s_{k,\ell}}$. A positive typing s with $s \preceq s_{k,\ell}$ is called (k,ℓ)-*boundary-determined*, or briefly a *boundary typing*.

2.3 Multiple Context-Free Grammars

Definition 2.5 (Standard MCFG presentation)**.** An MCFG [30, 31] is $G = (V, \Sigma, P, S, \mathrm{fan})$ with finite V, $S \in V$, $\mathrm{fan}(S) = 1$ and $\mathrm{fan}(A) \geq 1$. A rule of rank $r \geq 0$ is $A \to \rho(B_1, \ldots, B_r)$ with template $\rho = (\alpha_1, \ldots, \alpha_{d_0})$, $d_0 = \mathrm{fan}(A)$, $d_j = \mathrm{fan}(B_j)$, where each α_i is a word over $\Gamma_\rho := \Sigma \uplus \Xi_1 \uplus \cdots \uplus \Xi_r$, $\Xi_j := \{\xi_j^1, \ldots, \xi_j^{d_j}\}$, and every variable occurs at most once in $\alpha_1 \cdots \alpha_{d_0}$ (*linearity*); the rule is *nondeleting* if every variable occurs exactly once (derived components may still be empty). For $\vec{u}_j \in (\Sigma^*)^{d_j}$, $\rho(\vec{u}_1, \ldots, \vec{u}_r)$ substitutes $u_{j,\kappa}$ for ξ_j^κ; for $r = 0$ the rule is a constant rule. The tuple languages $L_A(G)$ form the least family closed under the rules, and $L(G) := \{w : (w) \in L_S(G)\}$. We write $A \Rightarrow_G^* \vec{u}$ for a *terminal derivation* (a finite derivation tree rooted at A evaluated bottom-up); one from S is *successful*. G is an f-MCFG if $\mathrm{fan}(A) \leq f$ for all A, and has rule-rank cap $r_{\max}$ if all ranks are at most $r_{\max}$.

We write $\mathsf{MCFL}(f, r)$ for the languages of MCFGs of fan-out $\leq f$ and rule rank $\leq r$, so that f-MCFL $= \bigcup_{r<\infty} \mathsf{MCFL}(f, r)$. Minimum rule rank of a language refers to all presentations, not to a given grammar. $\|G\|$ is the length of a fixed symbol-level encoding counting nonterminal and rule records and every terminal or variable occurrence in templates. For nonterminals of equal fan-out d, $A \to B$ abbreviates the identity-template rule $A \to (\xi_1^1, \ldots, \xi_1^d)(B)$, and $A \to \vec{c}$ a constant rule. Rule variables are always ξ_j^κ; terminals in examples are in `typewriter` font.

A nondeleting rule is *nonpermuting* if, for each j, $\xi_j^1, \ldots, \xi_j^{d_j}$ occur in this order in $\alpha_1 \cdots \alpha_{d_0}$; following [14, Section 2.3] it is λ-*free* if no α_i is empty, and a nonpermuting rule is *nonmerging* if no α_i contains a factor $\xi_j^\kappa \xi_j^{\kappa+1}$. A rule, template or presentation is *good* if it is λ-free, nondeleting, nonpermuting and nonmerging. By induction on derivations, every component of every tuple derived from a nonterminal of a λ-free grammar is nonempty. A nonterminal is *reachable* if it is reachable from S in the graph with an edge $A \to B$ whenever B is a child in a rule for A, and *productive* if $L_A(G) \neq \emptyset$; G is *reduced* if all nonterminals are both.

Proposition 2.6 (Good normalization with rank preservation)**.** *Let G_{in} be a finite f-MCFG of rule rank at most $r_{\max}$. There is a finite good f-MCFG G_{good} of rule rank at most $r_{\max}$ with $L(G_{\mathrm{good}}) = L(G_{\mathrm{in}}) \setminus \{\varepsilon\}$ and, for fixed $f, r_{\max}$, $\|G_{\mathrm{good}}\| = O_{f,r_{\max}}(\|G_{\mathrm{in}}\|)$.*

Proof. Trim; apply the nondeleting construction (f3) of [30, proof of Lemma 2.2] by retaining subsets of output components; decorate nonterminals by component permutations; decorate them by realizable nonempty-component patterns and project away empty components (this removes ε and makes all templates λ-free); finally apply Yoshinaka's nonmerging conversion [14, Section 2.3, Lemma 1]. No step increases fan-out or rule rank, and with $f, r_{\max}$ fixed each step creates $O_{f,r_{\max}}(1)$ copies of each nonterminal and rule (at most 2^f, $f!$ and $2^{fr_{\max}}$ decorations). The derivation correspondences and the size bookkeeping are given in Appendix A; cf. also [32, Section 2.1] and [14, Section 2.3, after Example 2]. □

The proposition controls size, fan-out and rank, not thickness; thickness transfer for reduced λ-free inputs is Proposition A.2. Fan-out and rule rank are distinct resources: there is no rank reduction preserving an arbitrary fixed fan-out [24, 33]. The reconstruction below therefore imposes neither a binary normal form nor a known rank cap.

2.4 Tuple Contexts and Substitutability

A *sentence context of arity d* is a word $E = v_0\Box_1 v_1 \cdots v_{d-1}\Box_d v_d$ with each hole occurring once, $v_0, v_d \in \Sigma^*$ and $v_1, \ldots, v_{d-1} \in \Sigma^+$; such contexts are *admissible*, and $\mathcal{E}_d$ denotes their set ([14, Section 3.1]). $E[\vec{w}]$ fills $\Box_i$ by w_i, $|E|_{\text{fix}} := \sum_{i=0}^{d} |v_i|$, and E *accepts* $\vec{w}$ in L if $E[\vec{w}] \in L$.

Definition 2.7 (Tuple occurrence)**.** For $w = a_1 \cdots a_n$ number the cut positions $\textcircled{0}a_1\textcircled{1}\cdots a_n\textcircled{n}$ and write $w[i : i'] := a_{i+1} \cdots a_{i'}$. An *arity-d tuple occurrence* in w is a sequence $0 \le i_1 < i'_1 < i_2 < i'_2 < \cdots < i_d < i'_d \le n$; it determines $\vec{x} = (w[i_1 : i'_1], \ldots, w[i_d : i'_d]) \in (\Sigma^+)^d$ and the context $E \in \mathcal{E}_d$ obtained by replacing these intervals by $\Box_1, \ldots, \Box_d$, so that $E[\vec{x}] = w$. For a finite K, $\mathrm{Obs}_d(K)$ is the set of tuples $\vec{x}$ arising from occurrences in words of K, i.e. $\vec{x} \in (\Sigma^+)^d$ with $E[\vec{x}] \in K$ for some $E \in \mathcal{E}_d$, and $\mathrm{Obs}(K) := \bigcup_{d \le f} \mathrm{Obs}_d(K)$.

For $L \subseteq \Sigma^*$ and $\vec{x} \in (\Sigma^*)^d$ the *d-tuple distribution* is $\mathcal{D}_L^{(d)}(\vec{x}) := \{E \in \mathcal{E}_d : E[\vec{x}] \in L\}$.

Definition 2.8 ((f,s)-tuple substitutability)**.** Let $f \ge 1$ and $s : \Sigma^+ \to \mathsf{S}$ a finite-semigroup morphism. A language $L \subseteq \Sigma^*$ is (f,s)*-tuple-substitutable* if for every $1 \le d \le f$ and all $\vec{x}, \vec{y} \in (\Sigma^+)^d$,

$$s^{(d)}(\vec{x}) = s^{(d)}(\vec{y}) \ \text{ and } \ \mathcal{D}_L^{(d)}(\vec{x}) \cap \mathcal{D}_L^{(d)}(\vec{y}) \neq \emptyset \implies \mathcal{D}_L^{(d)}(\vec{x}) = \mathcal{D}_L^{(d)}(\vec{y}).$$

Empty components are never compared (Yoshinaka's convention [14, Section 3.1 and footnote 3]; see [8] for CFGs); $\varepsilon \in L$ is handled at the start symbol. The restriction matters already at fan-out one: $\mathtt{a}$ and ε share a context in $\mathtt{a}^+$ but have different distributions. At fan-out one, $\mathcal{D}_L^{(1)}(x)$ is the two-sided distribution $\{(u,v) : uxv \in L\}$. For $\vec{x}, \vec{y} \in (\Sigma^+)^d$ write $\vec{x} \equiv_{L,s}^{d} \vec{y}$ (or $\vec{x} \equiv_{L,s} \vec{y}$) if $\widehat{s}^{(d)}(\vec{x}) = \widehat{s}^{(d)}(\vec{y})$ and $\mathcal{D}_L^{(d)}(\vec{x}) = \mathcal{D}_L^{(d)}(\vec{y})$; this is an equivalence relation.

Definition 2.9 (The target class)**.** $\mathcal{C}_{f,s}^{\mathrm{mcf}}$ is the class of (f,s)-tuple-substitutable f-MCFLs. A target in $\mathcal{C}_{f,s}^{\mathrm{mcf}}$ for a boundary typing $\widehat{s}$ is called *boundary-typed*.

If $s \preceq t$, then $\mathcal{C}_{f,s}^{\mathrm{mcf}} \subseteq \mathcal{C}_{f,t}^{\mathrm{mcf}}$, since equal finer types are equal coarser types; in particular $\mathcal{C}_{f,s_{\mathrm{triv}}}^{\mathrm{mcf}} \subseteq \mathcal{C}_{f,s}^{\mathrm{mcf}}$. The class is not monotone in f: only $\mathcal{C}_{f+1,s}^{\mathrm{mcf}} \cap f\text{-MCFL} \subseteq \mathcal{C}_{f,s}^{\mathrm{mcf}}$ holds, and already for trivial typing $\{w\#w^R\}$ is 1- but not 2-dimensionally substitutable [14, Section 3.1, Example 5]. Thus f is part of the bias. For every $\mathsf{Acc} \subseteq \mathsf{S}^\sharp$, $\widehat{s}^{-1}(\mathsf{Acc}) \in \mathcal{C}_{f,s}^{\mathrm{mcf}}$: it is regular, and equal componentwise types give equal total types in every context.

Proposition 2.10 (Intersection and abelian filtering)**.** *Fix f and $\widehat{s}$.*

(i) If L_1, L_2 are (f,s)-tuple-substitutable, so is $L_1 \cap L_2$.

(ii) If L is (f,s)-tuple-substitutable, $\mathbb{A}$ is an abelian group, $\varphi : \Sigma^* \to \mathbb{A}$ a monoid morphism, $\mathcal{H} \leq \mathbb{A}$ and $\gamma \in \mathbb{A}$, then $L \cap \varphi^{-1}(\gamma + \mathcal{H})$ is (f,s)-tuple-substitutable.

Whenever the intersection is an f-MCFL, it belongs to $\mathcal{C}^{\mathrm{mcf}}_{f,s}$.

Proof. (i) A context accepting equal-type $\vec{x}, \vec{y}$ in $L_1 \cap L_2$ accepts both in each L_i, so their distributions agree in each L_i, hence in the intersection. (ii) If a context accepts $\vec{x}, \vec{y}$ in $\varphi^{-1}(\gamma + \mathcal{H})$, then $\delta := \sum_i (\varphi(y_i) - \varphi(x_i)) \in \mathcal{H}$; by commutativity $\varphi(E'[\vec{y}]) - \varphi(E'[\vec{x}]) = \delta$ for every E', so coset membership transfers in both directions. Thus $\varphi^{-1}(\gamma + \mathcal{H})$ is (f, s_{triv})-, hence (f,s)-tuple-substitutable; apply (i). □

Consequently every f-MCFL of the form $\widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathcal{H})$ lies in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$; such factorizations are not supplied to the learner.

All interaction between templates and contexts below rests on one observation.

Lemma 2.11 (Induced child contexts)**.** *Let $\rho = (\alpha_1, \ldots, \alpha_e)$ be a good template of rank r with child arities $d_1, \ldots, d_r$, let $E \in \mathcal{E}_e$, let $1 \leq j \leq r$, and fix $\vec{u}_i \in (\Sigma^+)^{d_i}$ for $i \neq j$. Let E_j be the word obtained from $E[(\alpha_1, \ldots, \alpha_e)]$ by replacing ξ^κ_j by $\square_\kappa$ and every other variable ξ^ι_i by $u_{i,\iota}$. Then $E_j \in \mathcal{E}_{d_j}$ and $E_j[\vec{x}] = E[\rho(\vec{u}_1, \ldots, \vec{u}_{j-1}, \vec{x}, \vec{u}_{j+1}, \ldots, \vec{u}_r)]$ for every $\vec{x} \in (\Sigma^*)^{d_j}$.*

Proof. The identity holds because substitutions commute. By linearity and nondeletion each $\square_\kappa$ occurs exactly once in E_j, and by nonpermutation in the order $\kappa = 1, \ldots, d_j$, since E places $\alpha_1, \ldots, \alpha_e$ from left to right. Consider ξ^κ_j and $\xi^{\kappa+1}_j$. If both lie in one α_i, nonmerging puts a terminal or a variable ξ^ι_i with $i \neq j$ between them, which becomes a nonempty word. Otherwise they lie in $\alpha_i, \alpha_{i'}$ with $i < i'$, and the material between them contains the separator $v_i \in \Sigma^+$ of E. So every internal separator of E_j is nonempty. □

3 Exact Reconstruction from Positive Data

The reconstruction operator is Yoshinaka's observed-tuple grammar [14, Section 4.1] with type-guarded substitution links; its soundness and completeness proofs (Sections 3.2 and 3.4) transcribe [14, Lemmas 3–5]. They are given in full because they also cover unbounded rank and because the explicit witness set of Definition 3.10 is what Section 4 measures. Section 3.6 places the bounded-rank case in Yoshinaka's signature framework.

3.1 Set-Driven Reconstruction

Definition 3.1 (Composition witness)**.** Let K be finite and $r \geq 1$. A *composition witness of rank r* in K is a triple $(\vec{z}, \rho, (\vec{u}_1, \ldots, \vec{u}_r))$ with $\vec{z} \in \mathrm{Obs}(K)$, ρ a good template of rank r whose output and child arities are at most f, $\vec{u}_j \in (\Sigma^+)^{d_j}$, and $\vec{z} = \rho(\vec{u}_1, \ldots, \vec{u}_r)$.

Since each $\alpha_i(\vec{u}_1, \ldots, \vec{u}_r) = z_i$ fixes the positions of all child components inside $\vec{z}$, a composition witness is the same as a segmentation of an observed occurrence of $\vec{z}$ into terminal gaps and labelled child intervals. If $E[\vec{z}] \in K$ with $E \in \mathcal{E}_e$, Lemma 2.11 shows that every $\vec{u}_j \in \mathrm{Obs}(K)$. Linearity and nondeletion make the children contribute disjointly to $\vec{z}$, and each contributes at least one letter, so

$$r \leq \textstyle\sum_j \|\vec{u}_j\|_1 \leq \|\vec{z}\|_1 \leq \max_{w \in K} |w| \leq \|K\|_+. \tag{3.1}$$

Definition 3.2 (Capped and uncapped reconstruction operators)**.** Fix f and a finite-semigroup typing s with canonical completion $\widehat{s}$. For $r_{\max} \geq 1$ let $\mathcal{B}_{f,s}^{\leq r_{\max}}(K) = \widehat{G}_{\leq r_{\max}}(K)$ be the grammar with start symbol $\widehat{S}$, a nonterminal $[\vec{x}]$ of arity d for each $\vec{x} \in \mathrm{Obs}_d(K)$, $d \leq f$, and the rules

(Start) $\widehat{S} \to [(w)]$ for $w \in K \cap \Sigma^+$, and $\widehat{S} \to \varepsilon$ if $\varepsilon \in K$;

(Const) $[\vec{x}] \to \vec{x}$ for every observed $\vec{x}$;

(Comp) $[\vec{z}] \to \rho([\vec{u}_1], \ldots, [\vec{u}_r])$ for every composition witness of rank $1 \leq r \leq r_{\max}$;

(Link) $[\vec{x}] \to [\vec{y}]$ for observed $\vec{x}, \vec{y}$ of the same arity d with $\widehat{s}^{(d)}(\vec{x}) = \widehat{s}^{(d)}(\vec{y})$ and $E[\vec{x}], E[\vec{y}] \in K$ for some $E \in \mathcal{E}_d$.

For $K \neq \emptyset$ put $\mathcal{B}_{f,s}(K) := \widehat{G}(K) := \widehat{G}_{\leq \|K\|_+}(K)$; $\mathcal{B}_{f,s}(\emptyset)$ has no rules. Fixed encoding orders make both operators deterministic and set-driven.

By (3.1), $\mathcal{B}_{f,s}(K)$ contains a Comp rule for every composition witness in K. Both operators are sample-consistent, by $\widehat{S} \to [(w)] \to (w)$ and $\widehat{S} \to \varepsilon$. Every $[\vec{x}]$ derives only tuples with nonempty components, since Comp templates are λ-free; so $\widehat{S} \Rightarrow^* \varepsilon$ only by the rule $\widehat{S} \to \varepsilon$. For trivial typing and $\varepsilon \notin K$, Comp/Const, Link and Start are exactly the rules of Types I, II and III of [14, Section 4.1], and $\mathcal{B}_{f,s}$ is the operator of Yoshinaka's unbounded-rank learner $\mathcal{A}(p, *)$ [14, Corollary 2].

3.2 Soundness

Lemma 3.3 (Composition preserves equivalence)**.** *Let $L \subseteq \Sigma^*$, let ρ be a good template of rank r, and let $\vec{x}_j, \vec{u}_j \in (\Sigma^+)^{d_j}$ satisfy $\vec{u}_j \equiv_{L,s} \vec{x}_j$ for every j. Then $\rho(\vec{u}_1, \ldots, \vec{u}_r) \equiv_{L,s} \rho(\vec{x}_1, \ldots, \vec{x}_r)$.*

Proof. Each output component is the concatenation of the terminals of ρ and of child components in a fixed order, and is nonempty by λ-freeness; since $\widehat{s}$ is a monoid morphism, the parent types agree. Fix $E \in \mathcal{E}_e$ and put $\vec{p}_j := \rho(\vec{x}_1, \ldots, \vec{x}_j, \vec{u}_{j+1}, \ldots, \vec{u}_r)$ for $0 \leq j \leq r$. Lemma 2.11 with siblings $\vec{x}_1, \ldots, \vec{x}_{j-1}, \vec{u}_{j+1}, \ldots, \vec{u}_r$ gives $E_j \in \mathcal{E}_{d_j}$ with $E[\vec{p}_{j-1}] = E_j[\vec{u}_j]$ and $E[\vec{p}_j] = E_j[\vec{x}_j]$; as $\mathcal{D}_L(\vec{u}_j) = \mathcal{D}_L(\vec{x}_j)$, $E[\vec{p}_{j-1}] \in L$ iff $E[\vec{p}_j] \in L$. Chaining over j gives $\mathcal{D}_L(\vec{p}_0) = \mathcal{D}_L(\vec{p}_r)$. □

Proposition 3.4 (Soundness of every hypothesis)**.** *Let L be (f, s)-tuple-substitutable and $K \subseteq L$ finite. For every $r_{\max} \geq 1$, if $[\vec{x}] \Rightarrow^* \vec{u}$ in $\widehat{G}_{\leq r_{\max}}(K)$, then $\vec{u} \equiv_{L,s} \vec{x}$. Consequently $L(\mathcal{B}_{f,s}^{\leq r_{\max}}(K)) \subseteq L$ and $L(\mathcal{B}_{f,s}(K)) \subseteq L$.*

Proof. This is [14, Lemmas 3–4] with typed links. Induct on the derivation, according to its first rule. Const: $\vec{u} = \vec{x}$. Link $[\vec{x}] \to [\vec{y}]$: $\vec{x}, \vec{y}$ have equal types and share a context in $K \subseteq L$, so substitutability gives $\vec{x} \equiv_{L,s} \vec{y}$, and $\vec{u} \equiv_{L,s} \vec{y}$ by induction. Comp $[\vec{z}] \to \rho([\vec{u}_1], \ldots, [\vec{u}_r])$ with $[\vec{u}_j] \Rightarrow^* \vec{v}_j$: by induction $\vec{v}_j \equiv_{L,s} \vec{u}_j$, and Lemma 3.3 gives $\rho(\vec{v}_1, \ldots, \vec{v}_r) \equiv_{L,s} \vec{z}$. A nonempty word w is derived from $\widehat{S}$ via some $[(w_0)]$ with $w_0 \in K$; then $(w) \equiv_{L,s} (w_0)$ and $\square_1 \in \mathcal{D}_L((w_0))$, so $w \in L$; and ε is derived only if $\varepsilon \in K \subseteq L$. The uncapped operator is the capped one with $r_{\max} = \|K\|_+$. □

3.3 Yield-Typed Refinement and Witnesses

Definition 3.5 (Yield-type refinement)**.** Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a finite nondeleting MCFG and $s : \Sigma^+ \to \mathsf{S}$ a finite-semigroup typing, with canonical extension $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$. For a template ρ and $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{d_j}$, let $\mathrm{out}_\rho(\mathbf{t}_1, \ldots, \mathbf{t}_r) \in (\mathsf{S}^\sharp)^{d_0}$ have ith component the image of α_i under the monoid morphism $\Gamma_\rho^* \to \mathsf{S}^\sharp$ extending $a \mapsto \widehat{s}(a)$ and $\xi_j^\kappa \mapsto (\mathbf{t}_j)_\kappa$. The *full yield-type refinement* $G^{\widehat{s}}$ has a fresh start symbol $\widetilde{S}$, nonterminals $A_{\mathbf{t}}$ for $A \in V$ and $\mathbf{t} \in (\mathsf{S}^\sharp)^{\mathrm{fan}(A)}$, start rules $\widetilde{S} \to S_{(\eta)}$ for

$\eta \in \mathsf{S}^\sharp$, and for every rule $A \to \rho(B_1, \ldots, B_r)$ (including $r = 0$) and all $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{\mathrm{fan}(B_j)}$ the rule $A_{\mathrm{out}_\rho(\mathbf{t}_1,\ldots,\mathbf{t}_r)} \to \rho(B_{1,\mathbf{t}_1}, \ldots, B_{r,\mathbf{t}_r})$. The *trimmed refinement* $\widetilde{G}_0$ keeps exactly the typed nonterminals and rules occurring in some successful derivation from $\widetilde{S}$. Unless $G^{\widehat{s}}$ is named, *typed nonterminal* and *typed rule* refer to the non-start nonterminals and rules of $\widetilde{G}_0$, and $L_X := \{\vec{u} : X \Rightarrow^*_{\widetilde{G}_0} \vec{u}\}$.

Typed rules have the templates of their sources, so $\widetilde{G}_0$ is good whenever G is; each L_X is nonempty, with nonempty components if G is λ-free.

Proposition 3.6 (Yield-type invariants)**.** *(i) If $A_{\mathbf{t}} \Rightarrow^* \vec{u}$ in $G^{\widehat{s}}$, then $\widehat{s}^{(\mathrm{fan}(A))}(\vec{u}) = \mathbf{t}$, and erasing indices gives a G-derivation of $\vec{u}$.*

(ii) Every terminal derivation $A \Rightarrow^*_G \vec{u}$ lifts uniquely to $A_{\widehat{s}^{(\mathrm{fan}(A))}(\vec{u})} \Rightarrow^*_{G^{\widehat{s}}} \vec{u}$; a successful derivation of w lifts uniquely through $\widetilde{S} \to S_{(\widehat{s}(w))}$.

(iii) $L(\widetilde{G}_0) = L(G^{\widehat{s}}) = L(G)$.

Proof. At every rule application $\widehat{s}^{(\mathrm{fan}(A))}(\rho(\vec{u}_1, \ldots, \vec{u}_r)) = \mathrm{out}_\rho(\widehat{s}^{(d_1)}(\vec{u}_1), \ldots, \widehat{s}^{(d_r)}(\vec{u}_r))$, because $\widehat{s}$ is a monoid morphism. Induction on derivation height gives (i) and the unique lifting (ii). Erasing indices gives $L(G^{\widehat{s}}) \subseteq L(G)$, (ii) the converse, and trimming keeps every rule of every successful derivation. □

Lemma 3.7 (Monoid implementations)**.** *Let $h_0 : \Sigma^* \to M_0$ be a finite-monoid morphism, $S_0 := h_0(\Sigma^+)$ and $s_0 := h_0|_{\Sigma^+}$.*

(i) *For $1 \le d \le f$ and $\vec{x}, \vec{y} \in (\Sigma^+)^d$, the equalities $h_0^{(d)}(\vec{x}) = h_0^{(d)}(\vec{y})$, $s_0^{(d)}(\vec{x}) = s_0^{(d)}(\vec{y})$ and $\widehat{s_0}^{(d)}(\vec{x}) = \widehat{s_0}^{(d)}(\vec{y})$ are equivalent. Hence $\mathcal{C}^{\mathrm{mcf}}_{f,h_0} = \mathcal{C}^{\mathrm{mcf}}_{f,s_0}$, and the reconstruction operators built with h_0 and with $\widehat{s_0}$ return the same grammar on every sample.*

(ii) *For every good MCFG G, the trimmed refinements of G under h_0 and under $\widehat{s_0}$ coincide (the untrimmed ones need not); hence the same canonical tuples, contexts and witness words may be chosen for both.*

Proof. (i) The three maps agree on Σ^+ as maps into S_0, and the class and the Link guard compare types of nonempty components only. (ii) S_0 is a subsemigroup of M_0 and of $S_0^\sharp$ with the same multiplication. By λ-freeness and Proposition 3.6(i), every typed nonterminal occurring in a successful derivation under either morphism is some $A_{\mathbf{t}}$ with $\mathbf{t} \in S_0^{\mathrm{fan}(A)}$, and the start links used are $\widetilde{S} \to S_{(\eta)}$ with $\eta \in S_0$. For child indices in S_0, every component α_i of a good template is a nonempty word whose letters and variables are evaluated in S_0 identically under both morphisms, so both give the same out_ρ. Thus the two full refinements contain the same rules among S_0-indexed symbols, which carry all successful derivations. □

Lemma 3.8 (Canonical contexts for typed nonterminals)**.** *Let G be good. Every typed nonterminal X of arity d has a context $\chi(X) \in \mathcal{E}_d$ with $\chi(X)[\vec{u}] \in L(G)$ for all $\vec{u} \in L_X$; if $\widetilde{S} \to X$ is in $\widetilde{G}_0$, one may take $\chi(X) = \Box_1$.*

Proof. Take a successful typed derivation containing X, and follow the path from the root to an occurrence of X, starting with $\Box_1$. At each rule on the path apply Lemma 2.11 with the (nonempty) yields of the sibling subderivations fixed; this gives $\chi(X) \in \mathcal{E}_d$. The rules on the path depend only on the label X, so replacing the subtree of X by any X-derivation yields a successful derivation of $\chi(X)[\vec{u}]$. If $\widetilde{S} \to X$ is a rule, use the path of length zero. □

Lemma 3.9 (Typed nonterminals are distributionally homogeneous)**.** *Let G be good with $L := L(G) \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. For every typed nonterminal X of arity d, L_X lies in a single $\equiv^d_{L,s}$-class, denoted $[X]$.*

Proof. For $\vec{u}, \vec{v} \in L_X$, Proposition 3.6(i) gives equal types and $\chi(X)$ accepts both; substitutability gives $\mathcal{D}_L(\vec{u}) = \mathcal{D}_L(\vec{v})$. □

This parallels, but is weaker than, the homogeneity of congruential MCFGs [17], whose contexts may have empty separators.

Definition 3.10 (Canonical tuples, contexts and the witness set)**.** Let G be good. For every typed nonterminal X fix a *canonical tuple* $\omega(X) \in L_X$ and a *canonical context* $\chi(X)$ as in Lemma 3.8, with $\chi(X) = \Box_1$ whenever $\widetilde{S} \to X$ is in $\widetilde{G}_0$. Let $\mathsf{W}(\widetilde{G}_0)$ consist of the words

$$\chi(X)[\omega(X)] \ \ (X \text{ typed}), \qquad \chi(X)[\rho(\omega(Y_1), \ldots, \omega(Y_q))] \ \ (X \to \rho(Y_1, \ldots, Y_q) \text{ typed}, \ q \geq 0).$$

Since $\omega(X)$ and $\rho(\omega(Y_1), \ldots, \omega(Y_q))$ belong to L_X, $\mathsf{W}(\widetilde{G}_0)$ is a finite subset of $L(G)$, nonempty if $L(G) \neq \emptyset$. Components of canonical tuples are nonempty, so $\|\omega(X)\|_1 \geq \mathrm{fan}(X)$. Canonical objects are target witnesses, not inputs of the operator; qualitatively any choice works, while Section 4 chooses short ones.

Proposition 3.11 (Cardinality of the witness set)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be good of fan-out $\leq f$ and rule rank $\leq r_{\max}$, and let n_{typ} and n_{rule} be the numbers of typed nonterminals and typed rules of $\widetilde{G}_0$. Then*

$$|\mathsf{W}(\widetilde{G}_0)| \leq n_{\mathrm{typ}} + n_{\mathrm{rule}} \leq |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}}.$$

Proof. One word is selected per typed nonterminal and per typed rule. A typed nonterminal is some $A_{\mathbf{t}}$ with $\mathbf{t} \in (\mathsf{S}^\sharp)^{\mathrm{fan}(A)}$, and a typed rule is determined by its source rule and the child indices $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{\mathrm{fan}(B_j)}$. □

3.4 Completeness and Exact Reconstruction

The completeness argument [14, Lemma 5] needs only that the sample exposes, for every typed rule, one same-type link and one decomposition. We state it for arbitrary typed representatives, so that it also applies to the relaxed witnesses of Section 4. No assumption that $\iota(X) \in L_X$ is made: the two required witnessed sentence words are instead explicit hypotheses.

Proposition 3.12 (Completeness from a witness assignment)**.** *Let G be a finite good MCFG of fan-out $\leq f$, let r_G be the maximum rank of a typed rule of $\widetilde{G}_0$ ($r_G := 0$ if none), and assign to every typed nonterminal $X = A_{\mathbf{t}}$ a tuple $\iota(X) \in (\Sigma^+)^{\mathrm{fan}(A)}$ with $\widehat{s}^{(\mathrm{fan}(A))}(\iota(X)) = \mathbf{t}$ and a context $\chi(X) \in \mathcal{E}_{\mathrm{fan}(A)}$, where $\chi(X) = \Box_1$ whenever $\widetilde{S} \to X$ is in $\widetilde{G}_0$. If a finite sample K contains $\chi(X)[\iota(X)]$ for every X and $\chi(X)[\rho(\iota(Y_1), \ldots, \iota(Y_q))]$ for every typed rule $X \to \rho(Y_1, \ldots, Y_q)$, $q \geq 0$, then*

$$L(G) \subseteq L(\mathcal{B}^{\leq r_{\max}}_{f,s}(K)) \ \ (r_{\max} \geq \max\{1, r_G\}), \qquad L(G) \subseteq L(\mathcal{B}_{f,s}(K)).$$

Substitutability is not assumed. The canonical choice $\iota = \omega$ gives the hypothesis $\mathsf{W}(\widetilde{G}_0) \subseteq K$.

Proof. We show that $X \Rightarrow^*_{\widetilde{G}_0} \vec{u}$ implies $[\iota(X)] \Rightarrow^* \vec{u}$ in both grammars, by induction on the derivation. Each $\iota(X)$ is observed, since $\chi(X)[\iota(X)] \in K$. Let the derivation start with $X \to \rho(Y_1, \ldots, Y_q)$, followed by $Y_j \Rightarrow^* \vec{u}_j$, and put $\vec{z} := \rho(\iota(Y_1), \ldots, \iota(Y_q))$. Then $\vec{z}$ is observed, since $\chi(X)[\vec{z}] \in K$, and $\widehat{s}(\vec{z}) = \mathrm{out}_\rho(\text{indices of the } Y_j)$ is the index of X, which is $\widehat{s}(\iota(X))$; as $\chi(X)$ accepts both in K, the rule $[\iota(X)] \to [\vec{z}]$ is a Link. If $q = 0$, Const gives $[\vec{z}] \to \vec{z} = \vec{u}$. If $q \geq 1$, $(\vec{z}, \rho, (\iota(Y_j))_j)$ is a composition

witness of rank $q \le r_G \le r_{\max}$, and $q \le \|K\|_+$ by (3.1), so Comp gives $[\vec{z}] \to \rho([\iota(Y_1)], \dots, [\iota(Y_q)])$ in both operators; the induction hypothesis completes $[\iota(X)] \Rightarrow^* \rho(\vec{u}_1, \dots, \vec{u}_q)$. Finally a successful derivation begins with $\widetilde{S} \to X$, where $\chi(X) = \square_1$, so $\iota(X) = (w_0)$ with $w_0 \in K \cap \Sigma^+$ and $\widehat{S} \to [\iota(X)]$ is a Start rule; by Proposition 3.6(iii) this covers $L(G)$. □

Theorem 3.13 (Exact reconstruction for a fixed typing)**.** *Fix $f \ge 1$ and $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$, and put the completion $\widehat{s}$. Let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. If $L \cap \Sigma^+ = \emptyset$ put $\mathcal{F}_L := L$ and $r_G := 0$. Otherwise let G be a good presentation of $L \cap \Sigma^+$ (Proposition 2.6), and put $\mathcal{F}_L := \mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$. Then $\mathcal{F}_L \subseteq L$ is finite, and every finite K with $\mathcal{F}_L \subseteq K \subseteq L$ satisfies*

$$L(\mathcal{B}_{f,s}(K)) = L, \qquad L(\mathcal{B}^{\le r_{\max}}_{f,s}(K)) = L \quad \text{for every } r_{\max} \ge \max\{1, r_G\}.$$

If $L \in \mathsf{MCFL}(f, r)$, G can be chosen of rank $\le r$, so $r_{\max} = \max\{1, r\}$ suffices (also when $r = 0$).

Proof. If $L \cap \Sigma^+ = \emptyset$ then $K = L \in \{\emptyset, \{\varepsilon\}\}$ and the claim is immediate from Definition 3.2. Otherwise $\mathcal{F}_L \subseteq L$ is finite by the remark after Definition 3.10. Proposition 3.12 with $\iota = \omega$ (admissible by Proposition 3.6(i)) gives $L \cap \Sigma^+ \subseteq L(\mathcal{B}(K))$, and Proposition 3.4 gives $L(\mathcal{B}(K)) \subseteq L$, for both operators. The hypothesis contains ε iff $\varepsilon \in K$, iff $\varepsilon \in L$, because $\mathcal{F}_L \subseteq K \subseteq L$. Rank-preserving normalization gives the last sentence. □

3.5 Effective Reconstruction and Conservative Learning

Theorem 3.14 (Polynomial construction under a rule-rank cap)**.** *Fix $f, r_{\max} \ge 1$ and s, and implement type tests by the completion $\widehat{s}$. The encoding of $\widehat{G}_{\le r_{\max}}(K)$ is computable in time $(1 + \|K\|_+)^{O(f r_{\max})}$; its fan-out is at most f and its rank at most $r_{\max}$.*

Proof. Put $n_K = \max\{1, \|K\|_+\}$. An occurrence is a sample word with at most $2f$ cut positions, so there are $n_K^{O(f)}$ occurrences, typed by fixed-monoid products; grouping occurrences by context and type lists the Link rules. A composition witness of rank $r \le r_{\max}$ is determined by a parent occurrence and at most $2f r_{\max}$ further cut positions inside it, with a labelling whose number depends only on $f, r_{\max}$; goodness is checked directly. This gives $n_K^{O(f r_{\max})}$ candidate Comp rules; the remaining rules have rank at most one, and deduplication leaves an encoding of the same size. □

Definition 3.15 (Conservative sequential learner)**.** The learner $\mathcal{A}_{f,s}$ (respectively $\mathcal{A}^{\le r_{\max}}_{f,s}$) starts with H_0, keeps its hypothesis on $\perp$ and on every datum the hypothesis generates, and otherwise replaces it by $\mathcal{B}_{f,s}(K)$ (respectively $\mathcal{B}^{\le r_{\max}}_{f,s}(K)$), where K is the content read so far. This is the policy of [14, Algorithm 1].

Corollary 3.16 (Gold identification)**.** *Fix $f \ge 1$ and s.*

(i) *$\mathcal{A}_{f,s}$ TxtEx-identifies $\mathcal{C}^{\mathrm{mcf}}_{f,s}$.*

(ii) *For fixed $r_{\max}$, $\mathcal{A}^{\le r_{\max}}_{f,s}$ TxtEx-identifies $\mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathsf{MCFL}(f, r_{\max})$, in time per datum polynomial in the size of the content read so far.*

Proof. The learners are computable: hypotheses are computable and MCFG membership is decidable [30]. Let L be in the class and fix a text; let $\mathcal{F}_L$ be as in Theorem 3.13, with $r_{\max} \ge r_G$ in case (ii), and let n_0 be a stage whose content contains $\mathcal{F}_L$. Every hypothesis H satisfies $L(H) \subseteq L$, by Proposition 3.4 or because $H = H_0$. A rebuild at a stage $\ge n_0$ yields an exact hypothesis, which is never changed afterwards since all later data lie in L. If no rebuild occurs after n_0, let H be the

hypothesis current at n_0. If $H = H_0$, no positive datum was ever read, so $L = \emptyset = L(H)$. Otherwise H was built from the content of some stage m; it generates that content by sample-consistency and every later datum by conservativeness, hence all of L, so $L(H) = L$. In both cases the outputs stabilize on a correct grammar. For (ii), Theorem 3.14 bounds rebuilding, and membership of a datum w in a hypothesis $\widehat{G}$ of fan-out f and rank $r_{\max}$ takes $O(\|\widehat{G}\|^2|w|^{f(r_{\max}+1)})$ time [14, Proposition 4]. □

Update complexity versus convergence stage. The polynomial-time guarantee of Theorem 3.14 and Corollary 3.16 concerns hypothesis reconstruction and membership testing at each finite stage, measured against the data seen so far. TxtEx convergence in Corollary 3.16 does not bound the first stage of stabilization or the number of mind changes polynomially: a text may postpone the first presentation of a particular positive example arbitrarily long. Short characteristic data by itself gives no presentation-independent convergence-stage bound.

Every CFG has a rank-two presentation, so $\mathcal{A}_{1,s}^{\leq 2}$ covers $\mathcal{C}_{1,s}^{\mathrm{mcf}}$; this does not imply polynomial characteristic-data length at rank two (Section 4).

3.6 Specialization of Yoshinaka's 2015 Signature

We recall from [15, Sections 4 and 6.1] the notions used. A *signature* consists of pairwise disjoint sorts (sets of objects), a distinguished sort O_*, and typed function names; a grammar over it has rules $X_0 \leftarrow g\langle X_1, \ldots, X_n\rangle$ respecting sorts and initial symbols of sort O_*. $\mathcal{S}_O$ and $\mathcal{C}_O$ are the substructures and contexts (functions $O \to O_*$) of sort O realized in some grammar over the signature, and a language is *substitutable* if $c_1 \odot s_1, c_1 \odot s_2, c_2 \odot s_1 \in L$ imply $c_2 \odot s_2 \in L$ for all sorts O, $c_i \in \mathcal{C}_O$, $s_i \in \mathcal{S}_O$.

Fix $f, r \geq 1$ and $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$, and write $\widehat{s}$ for its canonical completion. The signature $\mathfrak{S}_{f,r,s}$ has the sort O_*, a tagged copy $\{\overline{w} : w \in \Sigma^*\}$ of Σ^*, and the tagged sorts $T_{d,\boldsymbol{\mu}} = \{\vec{w} \in (\Sigma^+)^d : s^{(d)}(\vec{w}) = \boldsymbol{\mu}\}$ for $1 \leq d \leq f$ and $\boldsymbol{\mu} \in \mathsf{S}^d$. Its functions are: every good template ρ of rank $q \leq r$ with arities at most f, typed as $T_{d_1,\boldsymbol{\mu}_1} \times \cdots \times T_{d_q,\boldsymbol{\mu}_q} \to T_{d_0,\mathrm{out}_\rho(\boldsymbol{\mu}_1,\ldots,\boldsymbol{\mu}_q)}$ for all $\boldsymbol{\mu}_j$ (the index lies in S^{d_0} by λ-freeness); promotions $\pi_\mu : T_{1,\mu} \to O_*$, $(w) \mapsto \overline{w}$; and a constant $\overline{\varepsilon}$ of sort O_*. Forgetting types, this is a finite sort refinement in the sense of [15, Section 8] of the signature of good MCFG operations on nonempty tuples, a nonempty-component variant of [15, Example 3]; at fan-out one the prefix–suffix case corresponds to [15, Example 4].

Proposition 3.17 (Bounded-rank coverage by Yoshinaka's 2015 framework)**.** (i) *$\mathcal{S}_{T_{d,\boldsymbol{\mu}}} = T_{d,\boldsymbol{\mu}}$, $\mathcal{C}_{T_{d,\boldsymbol{\mu}}} = \{\vec{x} \mapsto \overline{E[\vec{x}]} : E \in \mathcal{E}_d\}$ as sets of functions on $T_{d,\boldsymbol{\mu}}$, and $\mathcal{C}_{O_*}$ contains only the identity.*

(ii) $\mathsf{Subst}(\mathfrak{S}_{f,r,s}) = \mathcal{C}_{f,s}^{\mathrm{mcf}} \cap \mathsf{MCFL}(f, r)$.

(iii) *For fixed f, r, s, $\mathfrak{S}_{f,r,s}$ satisfies Assumptions 1–3 of [15]. Hence, by [15, Theorem 1] (proved there for* $\mathrm{SUBST}_{\mathrm{D}}$ *via its Lemmas 2–3), $\mathcal{C}_{f,s}^{\mathrm{mcf}} \cap \mathsf{MCFL}(f, r)$ is identifiable in the limit from positive data with polynomial-time updates, and for a target grammar without useless symbols, with N nonterminals and P rules,* $\mathrm{SUBST}_{\mathrm{D}}$ *has a characteristic set of cardinality at most $|N| + |P|$ (the bound stated after [15, Lemma 3]).*

Proof. (i) Constants (good rank-zero templates, i.e. tuples of nonempty words) give $\mathcal{S}_T = T$. No function takes an argument of sort O_* and only promotions and $\overline{\varepsilon}$ have codomain O_*; so $\mathcal{C}_{O_*}$ is the identity, and the context of a non-start nonterminal is obtained from the argument context $\square_1$ of a promotion by descending through good templates whose sibling arguments are nonempty tuples. By Lemma 2.11 each step yields an admissible context, so $\mathcal{C}_T \subseteq \{E[\cdot] : E \in \mathcal{E}_d\}$. Conversely $E = v_0\square_1 v_1 \cdots \square_d v_d \in \mathcal{E}_d$ is the context of Z in the grammar $X_* \leftarrow \pi_\nu\langle Y\rangle$, $Y \leftarrow (v_0\xi_1^1 v_1 \cdots \xi_1^d v_d)\langle Z\rangle$,

$Z \leftarrow \vec{x}_0$ for any $\vec{x}_0 \in T_{d,\boldsymbol{\mu}}$; the unary template is good because $v_1, \ldots, v_{d-1} \neq \varepsilon$ and its output is nonempty, and ν is its output type.

(ii) By (i), Yoshinaka's condition for the sort $T_{d,\boldsymbol{\mu}}$ says: for $E_1, E_2 \in \mathcal{E}_d$ and $\vec{x}, \vec{y}$ of type $\boldsymbol{\mu}$, $E_1[\vec{x}], E_1[\vec{y}], E_2[\vec{x}] \in L$ imply $E_2[\vec{y}] \in L$; for O_* it is vacuous. Over all $d \leq f$ and $\boldsymbol{\mu}$ this is (f,s)-tuple substitutability. It remains to see that $\mathfrak{S}_{f,r,s}$-grammars generate exactly $\mathsf{MCFL}(f,r)$. Erasing tags turns such a grammar into an MCFG of fan-out $\leq f$ and rank $\leq r$ (promotions become identity rules into a fresh start symbol, $\overline{\varepsilon}$ a constant rule). Conversely, for $L \in \mathsf{MCFL}(f,r)$ with $L \cap \Sigma^+ \neq \emptyset$, take a good presentation G of $L \cap \Sigma^+$ of rank $\leq r$ (Proposition 2.6). In $\widetilde{G}_0$ every non-start index lies in $\mathsf{S}^{\mathrm{fan}(A)}$ and every typed rule is an instance of a function of $\mathfrak{S}_{f,r,s}$; replace each start rule $\widetilde{S} \to S_{(\eta)}$ by a promotion π_η and add $\overline{\varepsilon}$ if $\varepsilon \in L$. The result generates L by Proposition 3.6(iii). The case $L \subseteq \{\varepsilon\}$ is trivial.

(iii) Assumption 2 holds since arities are at most r. For Assumption 1, sort membership, function evaluation, recognizing admissible contexts (by (i)), plugging, and membership in the language of a grammar over $\mathfrak{S}_{f,r,s}$ are polynomial; the last is bounded fan-out and rank MCFG membership after erasing tags [14, Proposition 4]. For Assumption 3, let $d \in \Sigma^*$. By (i) the substructures of non-start sorts occurring in d are the tuples of occurrences of arity $\leq f$ in d, typed by s, and their contexts are the corresponding admissible contexts; at O_* they are $\overline{d}$ and the identity. A typed template ρ with $c \odot \rho(\vec{u}_1, \ldots, \vec{u}_q) = \overline{d}$ yields an occurrence of $\vec{z} = \rho(\vec{u}_1, \ldots, \vec{u}_q)$ in d, i.e. a composition witness on it (a constant if $q = 0$); promotions and $\overline{\varepsilon}$ occur only for $d \neq \varepsilon$ and $d = \varepsilon$, respectively. Crucially, a function *name* includes its output sort and the ordered list of input sorts: for every candidate factorization the componentwise s-types of the observed child tuples determine these sorts, and the output sort is computed by out_ρ. Thus the enumeration records each typed instance, not only the untyped template, including the well-typed unary identity instances. As in Theorem 3.14, all these sets are computable in time $(|d|+1)^{O(f(r+1))}$. □

Lemma 3.18 (Substructure-indexed and guarded reconstruction)**.** *Let* $\mathrm{SUBST_P}$ *be Yoshinaka's substructure-indexed construction [15, Section 6.1] for* $\mathfrak{S}_{f,r,s}$ *with the sets of the preceding proof. For every finite* $K \subseteq \Sigma^*$, $L(\mathrm{SUBST_P}(K)) = L(\mathcal{B}^{\leq r}_{f,s}(K))$.

Proof. $\mathrm{SUBST_P}(K)$ has a nonterminal $[\![x]\!]$ for each substructure x occurring in K, initial symbols $[\![\overline{w}]\!]$ for $w \in K$, and a rule $[\![x_0]\!] \leftarrow g\langle [\![x_1]\!], \ldots, [\![x_n]\!] \rangle$ for each function g occurring in K and occurring substructures x_i of matching sorts whenever some context c occurring in K has $c \odot x_0, c \odot g(x_1, \ldots, x_n) \in K$. On non-start sorts identify $[\![\vec{x}]\!]$ with $[\vec{x}]$; these range over $\mathrm{Obs}(K)$.

Such a rule with $g = \rho$ of rank n is simulated in $\mathcal{B}^{\leq r}_{f,s}(K)$: put $\vec{z} := \rho(\vec{x}_1, \ldots, \vec{x}_n)$; by (i), $c = E[\cdot]$ with $E \in \mathcal{E}_d$ and $E[\vec{x}_0], E[\vec{z}] \in K$, so $\vec{z}$ is observed and has the sort, hence the type, of $\vec{x}_0$; thus $[\vec{x}_0] \to [\vec{z}]$ is a Link, followed by Const $[\vec{z}] \to \vec{z}$ if $n = 0$ or by the Comp rule of the composition witness $(\vec{z}, \rho, (\vec{x}_j)_j)$ if $n \geq 1$. Conversely, each non-start rule of $\mathcal{B}^{\leq r}_{f,s}(K)$ is a rule of $\mathrm{SUBST_P}(K)$: Const $[\vec{x}] \to \vec{x}$ and Comp $[\vec{z}] \to \rho(\ldots)$ use $x_0 = \vec{x}$, resp. $x_0 = \vec{z}$, and the context of an occurrence of x_0; Link $[\vec{x}] \to [\vec{y}]$ uses the identity template of the common sort, which occurs in K, and the shared context. For the Link case, the unary identity has a genuinely observed function instance: if $E[\vec{y}] \in K$, then $E[\mathrm{id}(\vec{y})] \in K$; hence its typed function name belongs to $\mathrm{FUN}(K)$. Likewise, each Comp uses its witnessed observed parent and child factorization, and each Const its rank-zero instance. By induction on derivations, $[\![\vec{x}]\!]$ and $[\vec{x}]$ derive the same tuples. At O_* the only context is the identity and only promotions and $\overline{\varepsilon}$ have codomain O_*; so for $K \neq \emptyset$ the rules for initial symbols are $[\![\overline{w}]\!] \leftarrow \pi\langle [\![(w')]\!] \rangle$ $(w, w' \in K,\ w' \neq \varepsilon)$ and $[\![\overline{w}]\!] \leftarrow \overline{\varepsilon}$ $(w \in K,\ \varepsilon \in K)$, and both grammars generate $\bigcup_{w' \in K \cap \Sigma^+} L([(w')]) \cup (\{\varepsilon\} \cap K)$; for $K = \emptyset$ both languages are empty. □

Remark 3.19 (What is and is not inherited). (a) Yoshinaka proves [15, Theorem 1] for $\mathrm{SUBST_D}$ and states without separate proof that his Algorithm 1 also works with $\mathrm{SUBST_P}$ [15, Sections 6.1 and 7].

For $\mathfrak{S}_{f,r,s}$, Lemma 3.18 together with Proposition 3.4 and Theorem 3.13 supplies that verification; Lemma 3.18 asserts equality of generated languages, not of grammars. (b) Cardinality bounds, whether $|N| + |P|$ or Proposition 3.11, do not bound total length; this is the subject of Section 4. (c) Assumption 2 bounds arity in each signature, so Proposition 3.17 covers $\mathsf{MCFL}(f, r)$ only for fixed r; unbounded rank is covered by Theorem 3.13, extending [14, Corollary 2] to typed links.

4 Characteristic-Data Bounds

Under a rule-rank cap, Proposition 3.11 already gives polynomially many characteristic words; this section controls their total length.

The bounds by typing regime are summarized in the table in the Introduction. For arbitrary fixed typings, the refined-thickness baseline is given by Lemma 4.4.

The first row is the unconditional baseline. Full yield-type refinement can make its thickness exponentially larger than the original presentation parameters (Proposition 4.5); rank one removes the thickness dependence, and the remaining rows transfer the *original* thickness. For window-only typings the window-skeleton lemma (Lemma 4.7) gives the bound; for window–group typings the marked heavy-node lemma (Lemma 4.12) combines boundary marks and simultaneous group-product preservation in one heavy-node count, invoking the block-replacement lemma (Appendix B). At fan-out one the companion CFG paper [25] excludes a universal original-thickness polynomial for arbitrary finite typing; its semilattice typing is not determined by any fixed prefix–suffix window. Polynomial update time in the observed sample is independent of these data-length refinements (Theorem 3.14 and Corollary 3.16).

4.1 Thickness and the witness-length template

For the quantitative results, presentations are reduced and *good*, as in Yoshinaka's fixed-dimension, fixed-rank good form [14, Section 2.3, after Example 2]. For a rule R of G let $\theta_G(R)$ be the minimum length of a successful sentence derivation using R, put $\theta_G := \max_{R \in P} \theta_G(R)$, $\bar{\theta}_G := \max\{1, \theta_G\}$, and $t_G := \max_A \min\{\|\vec{w}\|_1 : \vec{w} \in L_A(G)\}$. Thus θ_G is rule thickness and t_G nonterminal thickness; Yoshinaka writes them τ_G and t_G [14, Section 4.3]. If a target contains ε, every statement below is applied to its λ-free core and the single datum ε is added, exactly as in Theorem 3.13; this adds one to the encoded length.

Lemma 4.1 (Thickness lemma)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG.*

(i) *Every rule writes at most θ_G terminal symbols explicitly.*

(ii) *$t_G \le \theta_G$; that is, every $A \in V$ has a tuple $\vec{u} \in L_A(G)$ with $\|\vec{u}\|_1 \le \theta_G$.*

(iii) *$\theta_G \le 2\|G\| \max\{1, t_G\}$.*

(iv) *Every $A \in V$ of fan-out d has a context $E_A \in \mathcal{E}_d$ with $|E_A|_{\mathrm{fix}} \le \theta_G$ and $E_A[\vec{u}] \in L(G)$ for every $\vec{u} \in L_A(G)$, and $E_S = \square_1$. For every finite typing $\widehat{s}$ and every non-start typed nonterminal $X = A_{\mathbf{t}}$ of $\widetilde{G}_0$, the choice $\chi(X) := E_A$ meets Definition 3.10, and E_A arises from a successful typed derivation by deleting an occurrence of an X-subtree.*

Proof. Fix a rule R with left-hand side A and a successful derivation D_R using R of length $\theta_G(R)$; it exists because G is reduced. By nondeletion, every terminal written by R and every component of the tuple produced at that occurrence of A survives exactly once in the final sentence. This gives (i), and also (ii) because every productive A has a rule. For (iv) with $A \neq S$, delete the

subtree rooted at that occurrence of A and put ordered holes in place of its components; iterating Lemma 2.11 along the path to the root shows that the result E_A lies in $\mathcal{E}_d$, and its fixed material is part of the selected sentence. Substituting any $\vec{u} \in L_A(G)$ yields another successful derivation. Since $L_X \subseteq L_A(G)$ by Proposition 3.6(i), E_A accepts L_X; a typed nonterminal with a start link has the form $S_{\mathbf{t}}$, for which $E_S = \Box_1$. Filling E_A's derivation with an X-derivation and lifting by Proposition 3.6(ii) places X at the hole, so E_A is derivational for X.

For (iii), follow a shortest dependency path from S to A; it repeats no nonterminal, so its rules are distinct and differ from R. Keep the path child open, complete every sibling and every child of R by a shortest terminal tuple (of length at most t_G), and apply R. The distinct rules used write at most $\|G\|$ terminals and have at most $\|G\|$ child occurrences, so the resulting sentence has length at most $\|G\| + \|G\| t_G \le 2\|G\| \max\{1, t_G\}$. □

Hence the two thickness parameters are polynomially equivalent on reduced good presentations. Yoshinaka's efficiency criterion asks for polynomial hypothesis construction, characteristic-set cardinality, and total positive data in target size and rule thickness [14, Definition 1]; see also [10]. Proposition A.2 transfers the size–thickness scale to reduced λ-free inputs; genuine empty components may instead cause the exponential growth of Proposition A.3.

All higher-rank bounds below are assembled by the following two lemmas.

Lemma 4.2 (Witness-length template)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \ge 1$, let $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$ be a finite typing, and let $\widetilde{G}_0$ be the trimmed $\widehat{s}$-refinement. If every non-start typed nonterminal has a canonical tuple with $\|\omega(X)\|_1 \le b_{\mathrm{can}}$ and $\chi(X) := E_A$ is taken from Lemma 4.1(iv), then every word of $\mathsf{W}(\widetilde{G}_0)$ has length at most $2\bar{\theta}_G + r_{\max} b_{\mathrm{can}}$, and*

$$|\mathsf{W}(\widetilde{G}_0)| \le |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}}, \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ \le \big(|V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}}\big)\big(1 + 2\bar{\theta}_G + r_{\max} b_{\mathrm{can}}\big).$$

If moreover $L(G) = L \cap \Sigma^+$ for some $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, then $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}^{\le r_{\max}}_{f,s}$.

Proof. A witness is $\chi(X)[\omega(X)]$, $\chi(X)[\vec{c}]$ or $\chi(X)[\rho(\omega(Y_1), \dots, \omega(Y_q))]$. The context contributes at most $\bar{\theta}_G$ terminals, the source template at most $\bar{\theta}_G$ by Lemma 4.1(i), and there are at most $r_{\max}$ child tuples. The count is Proposition 3.11; multiply. The characteristic property is Proposition 3.12 (the cap $r_{\max}$ bounds every typed rank) together with Proposition 3.4 and the Start rule for ε. □

Lemma 4.3 (Type-preserving replacement)**.** *Let G be reduced and good, let $s : \Sigma^+ \to \mathsf{S}$ be a finite-semigroup typing and $\lambda : \Sigma^* \to N$ an auxiliary morphism with $s \preceq \lambda|_{\Sigma^+}$, and let $X = B_{\mathbf{t}}$ be a non-start typed nonterminal of the trimmed $\widehat{s}$-refinement $\widetilde{G}_0$. If $\vec{u} \in L_X$, $\vec{u}' \in L_B(G)$ and $\lambda^{(d)}(\vec{u}') = \lambda^{(d)}(\vec{u})$, then $\vec{u}' \in L_X$; in particular $\vec{u}'$ may be chosen as $\omega(X)$.*

Proof. Goodness makes all components nonempty, so $\widehat{s} \preceq \lambda$ gives $\widehat{s}^{(d)}(\vec{u}') = \mathbf{t}$. Proposition 3.6(ii) lifts a G-derivation of $\vec{u}'$ to a $G^{\widehat{s}}$-derivation rooted at X. Inserting it at an occurrence of X in a successful derivation of $\widetilde{G}_0$ gives a successful derivation, because typed rules depend only on typed symbols; hence all its rules survive trimming. □

4.2 Baselines and Rule Rank One

For a reduced good presentation G of an ε-free target, let $\widetilde{G}_0$ be its trimmed full yield-type refinement and put $\bar{\theta}_{\mathrm{typ}} := \max\{1, \theta_{\widetilde{G}_0}\}$, where $\theta_{\widetilde{G}_0}$ is rule thickness computed in the refined grammar. No comparison with the original untyped thickness is assumed.

Lemma 4.4 (Refined-thickness baseline)**.** *Fix $f \geq 1$, $r_{\max} \geq 1$, a finite alphabet Σ, and a finite typing $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$. Let G be a reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ presentation of an ε-free language in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, and let $\widetilde{G}_0$ be its trimmed yield-type refinement. Then the target witness set has cardinality polynomial in $\|G\|$ and total positive length polynomial in $\|G\|$ and $\bar{\theta}_{\mathrm{typ}} := \max\{1, \theta_{\widetilde{G}_0}\}$.*

Proof. $\widetilde{G}_0$ is reduced and good, so Lemma 4.1(i),(ii),(iv) applies to it with $\bar{\theta}_{\mathrm{typ}}$ in place of $\bar{\theta}_G$: canonical tuples and (derivational) canonical contexts of length at most $\bar{\theta}_{\mathrm{typ}}$ exist, with $\chi(X) = \square_1$ for start links. Each typed rule writes at most $\bar{\theta}_{\mathrm{typ}}$ terminals and has at most $r_{\max}$ children, so every witness has length at most $(r_{\max}+2)\bar{\theta}_{\mathrm{typ}}$, and

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \leq (|V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}})(1 + (r_{\max}+2)\bar{\theta}_{\mathrm{typ}})$$

by Proposition 3.11. □

For trivial typing no productive nonterminal has more than one yield type, so the original thickness scale is retained; this recovers Yoshinaka's untyped estimates $|K_G| \leq |P|$ and $\|K_G\| \leq |P|\tau_G$ [14, Lemma 7 and Theorem 1]. The quantitative claims below therefore concern typed, multi-component settings. Without uniqueness the refinement may be much thicker:

Proposition 4.5 (Exponential blow-up of full yield-type refinement)**.** *There are a fixed two-element typing h_{c}, reduced good fan-out-one rank-two grammars G_n of size $O(n)$, and a constant bound $\theta_{G_n} \leq 2$, such that $L(G_n) = \{\mathsf{a}, \mathsf{c}\}^+$ and the trimmed full yield-type refinement $\widetilde{G}_{n,0}$ satisfies $\theta_{\widetilde{G}_{n,0}} \geq 2^n$.*

Proof. Let $M_{\mathsf{c}} = \{1, \zeta\}$ with $\zeta^2 = \zeta$, $h_{\mathsf{c}}(\mathsf{a}) = 1$, and $h_{\mathsf{c}}(\mathsf{c}) = \zeta$. For $n \geq 1$, view

$$S \to SS \mid \mathsf{a} \mid \mathsf{c} \mid A_0 \mid \cdots \mid A_n, \qquad A_0 \to \mathsf{a}, \qquad A_i \to A_{i-1}A_{i-1} \mid \mathsf{c} \quad (1 \leq i \leq n)$$

as a fan-out-one MCFG. It is reduced and good, has size $O(n)$ and rule rank two, generates $\{\mathsf{a}, \mathsf{c}\}^+$, and every rule has a successful use of length at most two, so $\theta_{G_n} \leq 2$. In the full refinement, the retained copy $(A_n)_1$ can derive only words without c; recursively the unique such yield is a^{2^n}. Hence the typed rule $S_{(1)} \to (A_n)_1$ has successful thickness at least 2^n. □

This is a limitation of the full-refinement route, not of the language $\{\mathsf{a}, \mathsf{c}\}^+$. Separately, the companion CFG paper [25] (whose frozen TCS R1 source is linked in the bibliography) proves an operator-independent exponential characteristic-data lower bound for a fixed two-element semilattice typing at fan-out one and rule rank two. We do not reproduce it; it does not automatically extend to higher fan-out, whose tuple substitutability tests more arities.

Theorem 4.6 (Rank-one polynomial construction and characteristic data)**.** *Fix $f \geq 1$, finite Σ, and a finite-semigroup typing $s : \Sigma^+ \to \mathsf{S}$. Write $\widehat{s}$ for the canonical completion and $\mathsf{S}^\sharp$ for its codomain. For any finite standard fan-out-$\leq f$, rank-≤ 1 grammar G_{in} of $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, $\mathcal{B}^{\leq 1}_{f,s}$ is polynomial-time in $\|K\|_+$ and has a characteristic sample of total length $\mathrm{poly}_{f,s}(\|G_{\mathrm{in}}\|)$. The direct-start datum ε is added when $\varepsilon \in L$ without affecting this bound. The conservative learner $\mathcal{A}^{\leq 1}_{f,s}$ consequently TxtEx-identifies $\mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathsf{MCFL}(f, 1)$ with polynomial time per datum.*

Proof. Normalize without increasing rank (Proposition 2.6) and trim, obtaining a reduced good rank-$\leq$ 1 grammar G with $\|G\| = O_f(\|G_{\mathrm{in}}\|)$, and trim its typed refinement. Let $n_{\mathrm{typ}} \leq |V||\mathsf{S}^\sharp|^f$ be the number of non-start typed nonterminals and $b_G := \max\{1, \text{largest explicit terminal count of a rule}\} \leq \max\{1, \|G\|\}$. A terminal derivation from X is

a chain of unary rules ending in a constant rule. Choose one minimizing first the yield length and then the number of rules. No typed nonterminal repeats: splicing in the later suffix would not increase the yield (nondeletion) and would shorten the chain. Hence $\|\omega(X)\|_1 \le n_{\mathrm{typ}} b_G$. For the context, take $\Box_1$ if $\widetilde{S} \to X$ is present; otherwise follow a shortest path $\widetilde{S} \to X_1 \to \cdots \to X_p = X$ in the typed dependency graph, so $p \le n_{\mathrm{typ}}$. Composing its unary templates and putting ordered holes for the variables of X gives, by Lemma 2.11 iterated with no siblings, a derivational context $\chi(X) \in \mathcal{E}_{\mathrm{fan}(X)}$ with $|\chi(X)|_{\mathrm{fix}} \le n_{\mathrm{typ}} b_G$ accepting all of L_X. Every witness word therefore has length at most $n_{\mathrm{typ}} b_G + b_G + n_{\mathrm{typ}} b_G$, and Proposition 3.11 with $r_{\max} = 1$ gives

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \le (n_{\mathrm{typ}} + |P||\mathsf{S}^\sharp|^f)(1 + (2n_{\mathrm{typ}} + 1)b_G) = O_{f,s}(\|G_{\mathrm{in}}\|^3).$$

This set, with ε when $\varepsilon \in L$, is characteristic for $\mathcal{B}^{\le 1}_{f,s}$ by Propositions 3.12 and 3.4. The time and convergence claims are Theorem 3.14 and Corollary 3.16. □

This extends the rank-one untyped scale [14, Lemma 7] and the fan-out-one argument of [25] to arbitrary fixed finite typing at every fixed fan-out. Rank two cannot generally yield a grammar-size-only bound: the size-$O(n)$ grammar $A_0 \to \mathsf{c}$, $A_i \to A_{i-1}A_{i-1}$, $S \to A_n$ generates $\{\mathsf{c}^{2^n}\}$, which is (f,s)-tuple-substitutable for every fixed f, s (over a one-letter alphabet, sharing a sentence context forces equal total component lengths), and every characteristic sample for it contains a word of length at least 2^n.

4.3 Boundary Windows and Auxiliary Invariants

The refined thickness need not be polynomially controlled by the original good presentation (Proposition 4.5). A transfer is nevertheless possible whenever the typing is determined by a fixed prefix–suffix window together with an auxiliary finite invariant that can itself be preserved by short replacements. We use the thickness-sensitive representation model of Yoshinaka [14, Section 4.3].

Lemma 4.7 (Window skeletons)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \ge 1$. Fix $k, \ell \ge 0$, a morphism $\kappa : \Sigma^* \to N$ into a finite monoid, and a finite typing $\widehat{s} \preceq (\widehat{s}_{k,\ell}, \kappa)$. Let $\widetilde{G}_0$ be the trimmed $\widehat{s}$-refinement and n_{typ} its number of non-start typed nonterminals. Suppose that $\beta \ge 0$ is such that for every $C \in V$ and $\vec{x} \in L_C(G)$ there is $\vec{x}' \in L_C(G)$ with $\kappa^{(\mathrm{fan}(C))}(\vec{x}') = \kappa^{(\mathrm{fan}(C))}(\vec{x})$ and $\|\vec{x}'\|_1 \le \beta$. Then every non-start typed nonterminal X has a canonical tuple with*

$$\|\omega(X)\|_1 \le f(k+\ell+1)\, n_{\mathrm{typ}}\, (\bar{\theta}_G + r_{\max}\beta).$$

For trivial κ, $\beta = \bar{\theta}_G$ is admissible by Lemma 4.1(ii), so $\|\omega(X)\|_1 \le f(k+\ell+1) n_{\mathrm{typ}} (r_{\max}+1)\bar{\theta}_G$ whenever $\widehat{s} \preceq \widehat{s}_{k,\ell}$.

Proof. Choose a terminal derivation D from X in $\widetilde{G}_0$ with the minimum number of grammar-rule nodes, and let $\vec{u}$ be its root tuple. No typed nonterminal occurs twice on a root-to-leaf path: replacing the upper subtree by the lower one would give a derivation from X with fewer nodes. Hence every path has at most n_{typ} grammar nodes.

Mark in each component of $\vec{u}$ the terminals that determine its $\widehat{s}_{k,\ell}$-type (Definition 2.4): all terminals of a component shorter than $m_0 = \max\{1, k+\ell\}$, and otherwise its first k and last ℓ terminals (one arbitrary terminal when $k = \ell = 0$). There are at most $f(k+\ell+1)$ marks. Each terminal is written explicitly by exactly one node, so the union $\mathcal{K}$ of the root paths to the marked terminals has at most $f(k+\ell+1)n_{\mathrm{typ}}$ nodes. Call a child of a node of $\mathcal{K}$ that is not itself in $\mathcal{K}$

omitted. By linearity and nondeletion, each component of an omitted child is one contiguous block of exactly one root component, and this block contains no mark.

Replace each omitted subtree, with underlying nonterminal C and yield $\vec{x}$, by a G-derivation of a tuple $\vec{x}'$ as in the hypothesis, and forget all typed labels; this gives a G-derivation from the underlying nonterminal B of X with root tuple $\vec{u}'$. A short root component consists only of marked terminals, so it receives no replaced block and is unchanged. In a long component every replaced block lies strictly between the marked first k and last ℓ terminals, and goodness keeps it nonempty; the marked prefix and suffix remain literal, and the component stays on the long side of the threshold. Hence $\widehat{s}_{k,\ell}$-vectors agree. Each root component is the ordered product of explicit terminals and child components, and every replaced block keeps its κ-value; so κ-vectors agree, with no commutativity used. Thus $(\widehat{s}_{k,\ell},\kappa)^{(d)}(\vec{u}') = (\widehat{s}_{k,\ell},\kappa)^{(d)}(\vec{u})$, and Lemma 4.3 gives $\vec{u}' \in L_X$. Each node of $\mathcal{K}$ writes at most $\bar{\theta}_G$ terminals (Lemma 4.1(i)) and has at most $r_{\max}$ omitted children, each of length at most β; the bound follows. □

Theorem 4.8 (Polynomial data for prefix–suffix-determined typings)**.** *Fix $f \geq 1$, $r_{\max} \geq 1$, k, ℓ, a finite alphabet Σ, and a finite-semigroup typing $s : \Sigma^+ \to \mathsf{S}$ satisfying $s \preceq s_{k,\ell}$; write $\widehat{s}$ for its canonical completion. Let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and let G be a reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar with $L(G) = L \cap \Sigma^+$. Then $\mathcal{B}^{\leq r_{\max}}_{f,s}$ is polynomial-time in the positive sample. A characteristic sample for L is $K = \mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$, with*

$$|\mathsf{W}(\widetilde{G}_0)| = O_{f,r_{\max},s}(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},k,\ell,s,\Sigma}(\|G\|^2 \bar{\theta}_G),$$

where $\bar{\theta}_G = \max\{1, \theta_G\}$. For a reduced λ-free input presentation G_{in} (deleting rules allowed), the structural good-form conversion of Proposition A.2 yields the explicit presentation-relative bound

$$O_{f,r_{\max},k,\ell,s,\Sigma}\Big(\|G_{\mathrm{in}}\|^3 \max\{1, t_{G_{\mathrm{in}}}\}\Big).$$

Proof. By Lemma 3.7 we may use $\widehat{s}$. Lemma 4.7 with trivial κ gives $b_{\mathrm{can}} = f(k{+}\ell{+}1)n_{\mathrm{typ}}(r_{\max}{+}1)\bar{\theta}_G$, and $n_{\mathrm{typ}} \leq |V||\mathsf{S}^\sharp|^f = O(\|G\|)$. Lemma 4.2 bounds every witness by $O(\|G\|\bar{\theta}_G)$ and their number by $O(\|G\|)$, and gives the characteristic property; the datum ε adds at most one. Theorem 3.14 gives the construction time. For the input clause, Proposition A.2 gives $\|G\| = O(\|G_{\mathrm{in}}\|)$ and $\bar{\theta}_G = O(\|G_{\mathrm{in}}\| \max\{1, t_{G_{\mathrm{in}}}\})$; nonterminal thickness is used because deleting source rules may erase material before good-form conversion. □

Corollary 3.16 gives polynomial time per datum for the capped conservative learner. At fan-out one, polynomial time-and-thick-data learning for fixed windows is already due to Yoshinaka [12, Theorem 1]. As compared in the Introduction, our displayed quadratic total-length bound uses the original rule thickness of a reduced good presentation and has degree independent of the fixed window parameters. The contribution of Lemma 4.7 is to preserve boundary information simultaneously across *multiple* tuple components at every fixed $f \geq 2$. For fan-out-one window–group sufficient conditions, we credit the companion CFG manuscript [25]. Theorem 5.3 shows that this higher-rank regime is needed at the language level: it gives a non-context-free $(1,1)$-typed target of minimum fan-out two and minimum rule rank two over all finite fan-outs, whereas the rank-one examples are covered by Theorem 4.6.

4.4 Finite-Group Invariants

Window–group typings require short tuples with *prescribed* componentwise group values. The length bound is obtained from a single count of heavy nodes, which we state for an arbitrary finite-monoid invariant because Appendix C reuses it.

Lemma 4.9 (Heavy-node count)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \geq 1$, let $\lambda : \Sigma^* \to N$ be a morphism into a finite monoid, and let $b \geq \bar{\theta}_G$ and $c \geq 0$. Let T be a terminal derivation from $B \in V$ that minimizes first the total yield length and then the number of nodes among derivations from B with the same root λ-vector. Call a node* heavy *if its subtree yield has total length greater than b. If every set of pairwise incomparable heavy nodes of T has at most c elements, then*

$$\|\textit{root tuple of } T\|_1 \leq (r_{\max} + 1)\, b\, \max\{1, c\,|V||N|^f\}.$$

Proof. Label each node by its nonterminal and the λ-vector of its yield; there are at most $|V||N|^f$ labels. If an ancestor and a proper descendant had the same label, replacing the former subtree by the latter would keep the root λ-vector (templates evaluate homomorphically, as in Proposition 3.6), would not increase the yield (nondeletion) and would remove nodes. So every root-to-leaf path has at most $|V||N|^f$ nodes. If the root is light, its yield has length at most b. Otherwise the heavy nodes form an ancestor-closed subtree, whose leaves are pairwise incomparable; hence there are at most $c|V||N|^f$ heavy nodes. Every terminal of the root yield is written either by a heavy node (at most $\bar{\theta}_G \leq b$ per node, Lemma 4.1(i)) or inside a light child of a heavy node (at most $r_{\max}$ children of length at most b each). □

For a finite abelian group H write $\mathsf{d}(H)$ for the maximal length of a zero-sum-free sequence over H, so that $\mathsf{D}(H) = \mathsf{d}(H) + 1$ is its Davenport constant. For a finite group A and $f \geq 1$ put

$$N_{A,f} = \max\{3, |A|^f\}, \quad U_f = f^{f+1} f!, \quad M_{A,f} = U_f\, N_{A,f}^{2^{f^2}}, \qquad c_{A,f} := \begin{cases} \mathsf{d}(A^f), & A \text{ abelian}, \\ M_{A,f} - 1, & \text{otherwise}. \end{cases}$$

Proposition 4.10 (Simultaneous block replacement: the exact interface)**.** *Let A be a finite group and $f \geq 1$. There are d ordered output rows ($1 \leq d \leq f$). Row c is a fixed-order product of q_c blocks, interspersed with unaltered constants $a_{c,0}, \ldots, a_{c,q_c} \in A$. A block $\beta_{c,j}$ has an owner $o(c,j) \in [m]$ and an original and a replacement value $x_{c,j}, x'_{c,j} \in A$. Every owner has at most f blocks over all rows. Different owners' blocks may interleave arbitrarily; all material not replaced is absorbed into the intervening constants. For $J \subseteq [m]$ put*

$$\mathrm{val}_c(J) = a_{c,0} y_{c,1} a_{c,1} \cdots y_{c,q_c} a_{c,q_c}, \qquad y_{c,j} = \begin{cases} x'_{c,j}, & o(c,j) \in J, \\ x_{c,j}, & o(c,j) \notin J. \end{cases}$$

With $c_{A,f}$ as defined above, if $m > c_{A,f}$ there exists a nonempty $J \subseteq [m]$ such that $\mathrm{val}_c(J) = \mathrm{val}_c(\varnothing)$ $(c = 1, \ldots, d)$. For abelian A, one may take $c_{A,f} = \mathsf{d}(A^f)$; for an arbitrary finite group, $c_{A,f} = M_{A,f} - 1$ suffices. Individual block or owner values need not be preserved.

Proof. This is Lemma B.2 of Appendix B, written with root-component notation. For abelian A, associate with each owner its difference vector in A^d and select a nonempty zero-sum subsequence. In a nonabelian group one must not commute factors. Instead fix a common owner pattern and homogenize pairwise order comparisons by Lemma B.6; Lemma B.4 then decomposes the output rows into consecutive rank chunks, and Lemma B.5 selects a nonempty interval of owners preserving all ordered row products. The threshold is at most $M_{A,f}$. The detailed extraction and constant bounds are in the appendix. □

Theorem 4.11 (Short tuples with prescribed finite-group values)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and maximum rule rank $r_{\max} \geq 1$. Let $g : \Sigma^* \to A$ be a*

homomorphism into any *finite group A. For every $B \in V$ and every realizable $\mathbf{s} \in A^{\mathrm{fan}(B)}$ there is $\vec{u} \in L_B(G)$ with $g^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{s}$ and*

$$\|\vec{u}\|_1 \le \beta_g := (r_{\max}+1)\bar{\theta}_G \max\{1, c_{A,f}|V||A|^f\}.$$

In particular $\beta_g = O_{A,f,r_{\max}}(\|G\|\bar{\theta}_G)$; when A is abelian the antichain constant is the zero-sum constant $\mathsf{d}(A^f)$.

Proof. Let T be minimal as in Lemma 4.9 with $\lambda = g$ and $b = \bar{\theta}_G$, and let $v_1, \dots, v_m$ be pairwise incomparable heavy nodes. For each v_i, labelled C_i, fix a C_i-derivation of length at most $\bar{\theta}_G$ (Lemma 4.1(ii)); for $J \subseteq [m]$ let T_J replace the subtrees at v_i, $i \in J$. By linearity and nondeletion each component of v_i is a contiguous block of one root component, and blocks of different v_i are disjoint. The g-value of a root component of T_J is the ordered product of the g-values of its fixed segments and of its blocks, each block contributing its original or its replacement value according to whether its owner lies outside or inside J. This is a block-replacement family over A in the sense of Definition B.1. If $m > c_{A,f}$, Proposition 4.10 yields a nonempty J for which T_J has the root g-vector of T; but T_J is strictly shorter, since every replaced subtree has length greater than $\bar{\theta}_G$. This contradicts minimality, so Lemma 4.9 applies with $c = c_{A,f}$ and $N = A$. □

Lemma 4.12 (Marked heavy nodes for simultaneous window–group shortening)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \ge 1$. Fix $k, \ell \ge 0$, a finite group A, and a homomorphism $g : \Sigma^* \to A$. Put $\lambda = (\widehat{s}_{k,\ell}, g) : \Sigma^* \to N := M_{k,\ell} \times A$, $h := f(k+\ell+1)$, and $n_\lambda := |V||N|^f$. For every $B \in V$ and realizable vector $\mathbf{t} \in N^{\mathrm{fan}(B)}$, there exists $\vec{u} \in L_B(G)$ with $\lambda^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{t}$ and*

$$\|\vec{u}\|_1 \le b_{\mathrm{mark}} := (r_{\max}+1)\bar{\theta}_G \max\{1, (c_{A,f}+h)n_\lambda\}.$$

Hence $b_{\mathrm{mark}} = O(\|G\|\bar{\theta}_G)$ with $f, r_{\max}, k, \ell, A, g, \Sigma$ fixed. If A is abelian, $c_{A,f}$ can be taken as $\mathsf{d}(A^f)$.

Proof. Among B-derivations with root componentwise λ-value $\mathbf{t}$, choose T minimizing total yield length, then node count. As in Lemma 4.9, a root–leaf path repeats no pair of a nonterminal and its componentwise λ-value; its number of nodes is at most n_λ.

Mark the first k and last ℓ terminal positions in each root component of length at least $m_0 := \max\{1, k+\ell\}$, and all positions in a shorter component. If $k = \ell = 0$, mark one position per component to ensure that nonemptiness is explicitly preserved. There are at most $h = f(k+\ell+1)$ marks in total.

Call a node heavy if its subtree yields more than $\bar{\theta}_G$ terminals. Fix an arbitrary antichain Q of heavy nodes. Since incomparable subtrees occupy disjoint sets of terminal positions, at most h nodes of Q contain marked terminals. Let Q_0 consist of all other nodes. For each $v \in Q_0$ with nonterminal C_v, fix a derivation producing a tuple of total length at most $\bar{\theta}_G$, which exists by Lemma 4.1(ii). Each replacement strictly shortens the original heavy subtree.

Suppose $|Q_0| > c_{A,f}$. By linearity and nondeletion, every node in Q_0 contributes at most f nonoverlapping blocks distributed among at most f ordered root components. Everything outside these subtrees remains fixed. Since g is a homomorphism, each root group value is the ordered product of the unchanged pieces and the original or replacement block values. Proposition 4.10 thus selects a nonempty set of owners whose simultaneous replacement preserves all root group values, without assuming commutativity.

No selected subtree contains a mark. Short root components are untouched, while the marked prefixes and suffixes of long components retain their letters and order. Every replacement component is nonempty by goodness. Consequently all root $\widehat{s}_{k,\ell}$-values are preserved. The root λ-vector is

unchanged, while the total yield length strictly decreases, a contradiction to the choice of T. Every heavy antichain therefore has at most $c_{A,f} + h$ members. Lemma 4.9, with $b = \bar{\theta}_G$, $c = c_{A,f} + h$ and $N = M_{k,\ell} \times A$, gives the displayed bound. For abelian A use Lemma B.2(a). □

Example 4.13 (Two components with crossed owner order). Two owners with two blocks each may appear in the root tuple in the order $(\beta_{1,1}\beta_{2,1}, \beta_{2,2}\beta_{1,2})$. Take $A = S_3$, put $\sigma = (12)$ and $\tau = (23)$, and assign value σ to the blocks of owner 1 and τ to those of owner 2. The root group vector is $(\sigma\tau, \tau\sigma)$, with $\sigma\tau \neq \tau\sigma$. Thus commuting owner contributions would be invalid. This illustrates a block-replacement arrangement, not a proof that any particular target language requires a nonabelian typing. Proposition 4.10 preserves the two ordered root products simultaneously.

Theorem 4.14 (Polynomial characteristic data for window–finite-group typing)**.** *Fix $f, r_{\max}, k, \ell, \Sigma$, a finite group A, a homomorphism $g : \Sigma^* \to A$, and a fixed positive-word typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ whose fresh-identity completion satisfies $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$. For every target $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar G with $L(G) = L \cap \Sigma^+$, the capped reconstruction operator $\mathcal{B}^{\leq r_{\max}}_{f,s}$ has a characteristic sample with*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},k,\ell,\widehat{s},g,\Sigma}(\|G\|^2 \bar{\theta}_G).$$

Proof. Apply Lemma 4.12 to $\lambda = (\widehat{s}_{k,\ell}, g)$. Since $\widehat{s} \preceq \lambda$, every realizable $\widehat{s}$-typed non-start nonterminal $X = B_{\mathbf{s}}$ has an untyped B-tuple of the same root λ-value and total length at most $b_{\mathrm{mark}} = O(\|G\|\bar{\theta}_G)$. Lemma 4.3 lifts its derivation to X, so $b_{\mathrm{can}} \leq b_{\mathrm{mark}}$. The witness template (Lemma 4.2) gives $O(\|G\|)$ witnesses, each of length $O(\|G\|\bar{\theta}_G)$; the total encoded length is $O(\|G\|^2\bar{\theta}_G)$. Handle ε by its direct start rule. □

Theorem 4.15 (Polynomial data for window–abelian typing)**.** *Fix $f, r_{\max}, k, \ell, \Sigma$, a finite abelian group A, a homomorphism $g : \Sigma^* \to A$, and a fixed positive-word typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ whose fresh-identity completion satisfies $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$. For every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar G with $L(G) = L \cap \Sigma^+$, the capped operator $\mathcal{B}^{\leq r_{\max}}_{f,s}$ has a characteristic sample satisfying*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},k,\ell,\widehat{s},g,\Sigma}(\|G\|^2 \bar{\theta}_G).$$

In particular the bound uses the marked-node constant

$$b_{\mathrm{mark}} = (r_{\max} + 1)\bar{\theta}_G \max\{1, (\mathsf{d}(A^f) + f(k + \ell + 1))|V|(|M_{k,\ell}||A|)^f\}.$$

Proof. Apply the proof of Theorem 4.14 using the abelian case of Lemma 4.12. Lemma B.2(a) permits $c_{A,f} = \mathsf{d}(A^f)$ rather than $M_{A,f} - 1$, with $\mathsf{d}(A^f) = O(\exp(A) f \log |A|)$ (Remark B.3). □

Commutativity is used only in the zero-sum case of Lemma B.2. For nonabelian A, the order-homogenization proof of Appendix B.2 controls blocks interleaving across multiple root components.

4.5 An Algebraic Criterion for Polynomial Characteristic Data

Write **LI** for the pseudovariety of finite locally trivial semigroups and **Grp** for the pseudovariety of finite groups. Their join $\mathbf{LI} \vee \mathbf{Grp}$ consists of divisors of finite products $S \times A$, with $S \in \mathbf{LI}$ and A a finite group [28, 29]. The following criterion makes the window–group hypothesis intrinsic to the positive image of the fixed typing; its finite-semigroup ingredient also occurs in the companion CFG manuscript [25].

Lemma 4.16 (Locally trivial–group typings admit window–group refinements)**.** *Let Σ be a nonempty finite alphabet and let $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ be a finite-semigroup typing; write $\widehat{s}$ for its canonical completion. If $\mathsf{S} \in \mathbf{LI} \vee \mathbf{Grp}$, then there exist $k, \ell \geq 0$, a finite group A, and a morphism $g : \Sigma^* \to A$ such that $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$ on all of Σ^*. If $\mathsf{S} \in \mathbf{LI}$, the group factor can be taken trivial, so $\widehat{s} \preceq \widehat{s}_{k,\ell}$.*

Proof. Choose a finite locally trivial semigroup S, a finite group A, a subsemigroup $T \subseteq S \times A$, and an onto semigroup morphism $\pi : T \to \mathsf{S}$. Lifting each $s(a)$ to some $(s_a, g_a) \in T$ defines a morphism $\rho : \Sigma^+ \to T$ with $\pi\rho = s$; let $\sigma : \Sigma^+ \to S$ and $g : \Sigma^* \to A$ (with $g(\varepsilon) = 1_A$) be its coordinates. The positive image $S_0 = \sigma(\Sigma^+)$ is locally trivial. Put $q = |S_0| + 1$. The finite locally trivial identities $S_0^q = S_0 E(S_0) S_0$ and $euf = ef$ for idempotents e, f and $u \in S_0$ [25, 27] show that $\sigma(w)$ is determined by the length-q prefix and suffix of w when $|w| \geq 2q$: writing $w = pxz$ with $|p| = |z| = q$, $\sigma(p) = aeb$, $\sigma(z) = cfd$ and $\sigma^1(\varepsilon) = 1$, $\sigma(pxz) = ae\,(b\,\sigma^1(x)\,c)\,fd = aefd = \sigma(pz)$, where the parenthesized factor lies in S_0 because $b, c \in S_0$. Words shorter than $2q$ are recorded exactly by $\widehat{s}_{q,q}$. Hence equal $(\widehat{s}_{q,q}, g)$-values on Σ^+ give equal ρ- and $\widehat{s}$-values; the window typing separates ε from Σ^+. Take $k = \ell = q$. If $\mathsf{S} \in \mathbf{LI}$, use $S = \mathsf{S}$, $\pi = \mathrm{id}$ and the trivial group. □

Corollary 4.17 (Algebraic sufficient conditions for polynomial data)**.** *Fix a nonempty finite Σ, a fan-out cap $f \geq 1$, a rule-rank cap $r_{\max} \geq 1$, and a finite-semigroup typing $s : \Sigma^+ \to \mathsf{S}$; write $\widehat{s}$ for its canonical completion. Let G be any reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ presentation of a target in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, and put $\bar{\theta}_G = \max\{1, \theta_G\}$. The capped reconstruction operator has a characteristic sample of cardinality $O_{f,r_{\max},s}(\|G\|)$ with:*

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},s,\Sigma}(\|G\|^2 \bar{\theta}_G), \quad \text{if } \mathsf{S} \in \mathbf{LI},$$
$$\|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},s,\Sigma}(\|G\|^2 \bar{\theta}_G), \quad \text{if } \mathsf{S} \in \mathbf{LI} \vee \mathbf{Grp}.$$

Proof. Combine Lemma 4.16 with Theorem 4.8 or Theorem 4.14; the refining parameters depend only on s and Σ, not on G. □

The full algebraic sufficient condition is stated here; its structural representation and shortening proofs are in Appendix C. A finite semigroup belongs to **LG** (the *local groups*) if eSe is a group for every idempotent $e \in S$. This condition concerns the *positive image*, never the fresh-identity completion.

Theorem 4.18 (Polynomial characteristic data for local-group typing)**.** *Fix $f, r_{\max}, \Sigma$ and a finite-semigroup typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ with $\mathsf{S} \in \mathbf{LG}$; write $\widehat{s}$ for its canonical fresh-identity completion. For every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar G with $L(G) = L \cap \Sigma^+$, the capped operator $\mathcal{B}^{\leq r_{\max}}_{f,s}$ has a characteristic sample with*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},s,\Sigma}(\|G\|^3 \bar{\theta}_G).$$

Proof outline. The local-group representation of Theorem C.4 refines $\widehat{s}$ by a fixed prefix–suffix window and an ordered product of finite-group values attached to sliding factors. The window skeleton gives short replacement candidates of length $O(\|G\|\bar{\theta}_G)$. A simultaneous noncommutative block-replacement argument chooses which candidates to use while preserving every root component's ordered group product, yielding typed canonical tuples of length $O(\|G\|^2 \bar{\theta}_G)$. The witness-length template counts $O(\|G\|)$ typed rules. Appendix C gives the complete argument. □

The inclusion $\mathbf{LI} \vee \mathbf{Grp} \subsetneq \mathbf{LG}$ is strict (Proposition C.5). Thus the local-group bound already strictly strengthens the companion CFG sufficient regime when $f = 1$, while the algebraic characterization of local groups itself is classical. At fan-out one the two-element semilattice supplies

the pseudovariety-level lower-bound obstruction (Corollary C.7); that corollary's negative direction relies on the companion CFG paper [25]. No matching lower bound at fan-out $f \geq 2$ is asserted; see Open Problem 6.4. The empty-alphabet case is trivial.

4.6 Computing Short Characteristic Data

The bounds above exhibit some short tuple and some short context. Call a sentence context of a typed nonterminal X *derivational* if it arises from a successful typed derivation by deleting one occurrence of an X-subtree; such contexts are admissible. All canonical contexts used in Sections 4.1–4.5 and Appendix C are derivational (Lemma 4.1(iv) and the proofs of Lemma 4.4 and Theorem 4.6).

Proposition 4.19 (Effective short characteristic data)**.** *Fix $f, r_{\max}$ and a finite typing $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$. Given a reduced good MCFG G of fan-out at most f and rule rank at most $r_{\max}$, the following can be computed in time polynomial in $\|G\|$, with lengths written in binary:*

(i) $\theta_G(R)$ for every rule R, and hence θ_G and t_G;

(ii) the trimmed refinement $\widetilde{G}_0$ and, for each non-start typed nonterminal X, a canonical tuple $\omega(X) \in L_X$ of minimum total length and a derivational context $\chi(X)$ of minimum fixed length, with $\chi(X) = \square_1$ whenever $\widetilde{S} \to X$ belongs to $\widetilde{G}_0$;

(iii) derivation DAGs of size polynomial in $\|G\|$ representing these tuples and contexts, and hence every word of the witness set $\mathsf{W}(\widetilde{G}_0)$ of Definition 3.10 built from them.

Every word-length and total-length bound proved for $\mathsf{W}(\widetilde{G}_0)$ in Sections 4.1–4.5 and in Appendix C holds for this computed witness set. This does not apply to the relaxed samples $K^\dagger$ of Section 4.7, whose representatives need not be generated by their nonterminals. The words can be written out explicitly in time polynomial in $\|G\| + \|\mathsf{W}(\widetilde{G}_0)\|_+$.

Proof. For fixed $f, r_{\max}, s$ the full refinement $G^{\widehat{s}}$ has at most $|V||\mathsf{S}^\sharp|^f$ typed nonterminals and $|P||\mathsf{S}^\sharp|^{fr_{\max}}$ typed rules, and trimming is a linear-time productivity and reachability computation. View the typed rules as weighted deduction steps. A rule $X \to \rho(Y_1, \ldots, Y_r)$ maps child weights to $(c_\rho, 1) + \sum_j (x_j, n_j)$, where c_ρ is the number of terminals written by ρ and pairs are compared lexicographically (terminal length, derivation-node count). This function is strictly superior to each argument even when no terminal is written, so the shortest-derivation algorithm in the weighted deductive form of Nederhof [36], building on Knuth's generalization of Dijkstra's method (see also Ramalingam and Reps [35]), computes $\lambda(X) = \min\{\|\vec{u}\|_1 : \vec{u} \in L_X\}$ together with an acyclic choice of minimizing rules, that is, a derivation DAG for $\omega(X)$. A second Dijkstra pass computes outside lengths, breaking equal-length ties by the number of edges: it starts with $\mu(X) = 0$ when $\widetilde{S} \to X$ is present, and for every typed rule $X \to \rho(Y_1, \ldots, Y_r)$ and every j adds an edge from X to Y_j of weight $c_\rho + \sum_{i \neq j} \lambda(Y_i)$. Then $\mu(X)$ is the minimum fixed length of a derivational context of X, and the shortest-path tree supplies its DAG. The same two passes on G give $t_G = \max_A \lambda(A)$ and $\theta_G(R) = \mu(A) + c_\rho + \sum_j \lambda(B_j)$ for $R : A \to \rho(B_1, \ldots, B_r)$.

Each of the cited proofs bounds a witness word through the length of its context, of its rule template and of its canonical child tuples, and does so by exhibiting *some* tuple of L_X and *some* derivational context. The computed minima are no longer, and the number of witnesses is unchanged, so every stated bound persists. Expanding a DAG takes time linear in the length of the word it represents. □

Proposition 4.20 (Rule-only witnesses and a computable data budget)**.** *In the setting of Proposition 4.19, choose all canonical tuples by the consistent acyclic minimum-weight derivations described there, and let $\mathcal{R}$ be the non-start typed rules of $\widetilde{G}_0$ (rank-zero rules included).*

(i) For each non-start typed nonterminal X, let $R_X : X \to \rho_X(Y_1, \dots, Y_q)$ be the first rule of its chosen minimizing derivation (including $q = 0$). Then $\omega(X) = \rho_X(\omega(Y_1), \dots, \omega(Y_q))$. Consequently

$$\mathsf{W}(\widetilde{G}_0) = \{\chi(\mathrm{lhs}\, R)[\rho_R(\omega(Y_1), \dots, \omega(Y_q))] : R \in \mathcal{R}\}, \qquad |\mathsf{W}(\widetilde{G}_0)| \le |\mathcal{R}| \le |P||\mathsf{S}^\sharp|^{fr_{\max}},$$

and this set remains characteristic for $\mathcal{B}_{f,s}^{\le r_{\max}}$ whenever the target is in $\mathcal{C}_{f,s}^{\mathrm{mcf}}$.

(ii) Let $\lambda(X)$ be the minimum total yield length from X and $\mu(X)$ the minimum fixed length of a derivational context of X (with $\mu(X) = 0$ for a direct start link). For $R : X \to \rho(Y_1, \dots, Y_q)$ put

$$b(R) := c_\rho + \sum_{j=1}^{q} \lambda(Y_j), \qquad \mathsf{B}_{\mathrm{rule}}(G, \widehat{s}) := \sum_{R \in \mathcal{R}} \big(1 + \mu(\mathrm{lhs}\, R) + b(R)\big).$$

Then $\|\mathsf{W}(\widetilde{G}_0)\|_+ \le \mathsf{B}_{\mathrm{rule}}(G, \widehat{s})$, and the integer $\mathsf{B}_{\mathrm{rule}}(G, \widehat{s})$ is computable in time polynomial in $\|G\|$ for fixed $f, r_{\max}, s$, its value being represented in binary. If $L(G)$ also contains ε and its nonempty core is treated separately, the ε datum adds at most one to the budget.

Proof. (i) An optimal derivation of X starts with R_X, and each of its child derivations is optimal, since replacing a nonoptimal child by the chosen minimum would decrease the lexicographic weight. The node-count tie-breaker makes the simultaneous choice acyclic, which gives the identity, also for rank zero. Hence $\chi(X)[\omega(X)]$ is *identical* to the rule word of R_X, and every other term of Definition 3.10 is a rule word; the two sets coincide as sets of terminal words, so completeness and the characteristic-sample property are unchanged. There is at most one word per typed rule. (ii) The rule word of $R : X \to \rho(Y_1, \dots, Y_q)$ has length exactly $\mu(X) + b(R)$. Summing $|w| + 1$ over the indexed family gives $\mathsf{B}_{\mathrm{rule}}(G, \widehat{s})$; passing to distinct words can only lower it. The values λ, μ and the number of typed rules are computed, with polynomial binary length, by Proposition 4.19. □

This rule-only budget is presentation-sensitive and can be much smaller than the worst-case polynomial bounds; it is an efficiently computable *upper certificate*, not a claim of a minimum-size characteristic sample or an improved worst-case exponent. For the singleton $\{\mathsf{c}^{2^n}\}$ after Theorem 4.6, explicit words are exponentially longer than $\|G\|$, but the derivation DAGs remain of polynomial size.

4.7 Distributional Representatives and Abelian-Filtered Targets

For an existential characteristic-sample bound one may instead choose short representatives of the same *typed distribution classes*, even when they have no derivation from the corresponding nonterminal.

Proposition 4.21 (Relaxed distributional witnesses)**.** *Let G be a reduced good fan-out-$\le f$, rank-$\le r$ presentation of $L = L(G) \in \mathcal{C}_{f,s}^{\mathrm{mcf}}$, and let $\widetilde{G}_0$ be its trimmed yield-type refinement. For every non-start typed nonterminal X, let $[X]$ be the $\equiv_{L,s}$-class containing L_X, and put $\kappa_G(X) = \min_{\vec{v} \in [X]} \|\vec{v}\|_1$, $\kappa_G = \max_X \kappa_G(X)$. There exists a characteristic positive sample $K^\dagger$ for both the capped operator $\mathcal{B}_{f,s}^{\le r}$ (with cap at least the rank of G) and the uncapped $\mathcal{B}_{f,s}$ satisfying*

$$|K^\dagger| \le |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{fr},$$
$$\|K^\dagger\|_+ \le (|V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{fr})(1 + \theta_G + \|G\| + (r+1)\kappa_G).$$

The bound concerns the nonempty core; if $\varepsilon \in L$, add the sample word ε at encoded cost one.

Proof. For each $X = A_{\mathbf{t}}$ choose $\omega^\dagger(X) \in [X]$ of total length $\kappa_G(X)$, and let $\chi^\dagger(X) := E_A$ from Lemma 4.1(iv), of fixed length at most θ_G ($\square_1$ for a start link). It accepts every tuple of L_X, hence also $\omega^\dagger(X) \equiv_{L,s} \vec{u} \in L_X$. Let $K^\dagger$ consist of the words $\chi^\dagger(X)[\omega^\dagger(X)]$, the rule words $\chi^\dagger(X)[\rho(\omega^\dagger(Y_1), \dots, \omega^\dagger(Y_q))]$ and the rank-zero rule words. By Lemma 3.3 the rule tuple is equivalent to a tuple of L_X, so $K^\dagger \subseteq L$; it also has the $\widehat{s}$-type of X. Apply the generalized Proposition 3.12 with $\iota(X) = \omega^\dagger(X)$ and $\chi(X) = \chi^\dagger(X)$: the first family of sample words is its representative evidence and the rule-word family is precisely its rule evidence. Componentwise types agree by the typed rules and Lemma 3.3; all contexts are admissible and start-linked contexts are identities. The proposition supplies the capped and uncapped derivation simulations. Proposition 3.4 gives the reverse inclusion for all larger positive samples. A nonterminal word has length at most $\theta_G + \kappa_G$ and a rule word at most $\theta_G + \|G\| + r\kappa_G$; sum the encoded costs. Unlike Proposition 4.20, the nonterminal words cannot in general be omitted. □

For fixed $f, r, \widehat{s}$, a family G_n with $\|G_n\| \max\{1, \theta_{G_n}\} = n^{O(1)}$ and superpolynomial characteristic-data requirements must therefore have $\kappa_{G_n} = n^{\omega(1)}$; this is a necessary condition, not a proof that such a family exists.

Theorem 4.22 (Polynomial data for abelian-filtered targets)**.** *Fix a finite nonempty alphabet Σ, a finite semigroup typing $s : \Sigma^+ \to \mathsf{S}$ and bounds $f, r \geq 1$. Write $\widehat{s}$ for the canonical completion and $\mathsf{S}^\sharp$ for its codomain. There are constants C, e depending only on Σ, s, f, r with the following property. Let $\mathbb{A}$ be any abelian group, $\varphi : \Sigma^* \to \mathbb{A}$ a morphism, $\mathcal{H} \leq \mathbb{A}$, $\gamma \in \mathbb{A}$, and $\mathsf{Acc} \subseteq \mathsf{S}^\sharp$. If $L = \widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathcal{H})$ has a reduced good fan-out-$\leq f$, rank-$\leq r$ MCFG presentation G of $L \cap \Sigma^+$, then $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and admits a characteristic sample for $\mathcal{B}^{\leq r}_{f,s}$ with*

$$\|K^\dagger\|_+ \leq C\big(\|G\| \max\{1, \theta_G\}\big)^e.$$

The constants do not depend on $\mathbb{A}, \varphi, \mathcal{H}, \gamma$. The datum ε is included when $\varepsilon \in L$; the degenerate cases $L \subseteq \{\varepsilon\}$ are handled directly.

Proof. Membership in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ follows from Proposition 2.10 and the remark preceding it; the positive part has the assumed presentation, and ε is admitted by Start. By Proposition 4.21 it suffices to bound κ_G by a polynomial in $\|G\|\bar{\theta}_G$ with exponent depending only on $\Sigma, \widehat{s}, f, r$.

Let $\psi : \Sigma^* \to \mathbb{N}^\Sigma$ be the Parikh map, $\Psi(\vec{u}) = \sum_i \psi(u_i)$, and let $\overline{\varphi} : \mathbb{Z}^\Sigma \to \mathbb{A}$ be the group homomorphism through which φ factors; put $\Lambda = \{z \in \mathbb{Z}^\Sigma : \overline{\varphi}(z) \in \mathcal{H}\}$. Fix for each $A \in V$ a tuple $\omega_A \in L_A(G)$ with $\|\omega_A\|_1 \leq t_G \leq \theta_G$. The context E_A accepts every A-tuple, and all accepted words lie in one coset of $\mathcal{H}$, so

$$\Psi(\vec{u}) - \Psi(\omega_A) \in \Lambda \qquad (\vec{u} \in L_A(G)). \tag{$*$}$$

For $R : A \to \rho(B_1, \dots, B_q)$ let $c_R \in \mathbb{N}^\Sigma$ count its explicit letters and $\delta_R = \Psi(\omega_A) - c_R - \sum_j \Psi(\omega_{B_j})$. By ($*$) applied to $\rho(\omega_{B_1}, \dots, \omega_{B_q}) \in L_A(G)$, $\delta_R \in \Lambda$, and $\|\delta_R\|_1 \leq T := (r+1)t_G + \|G\|$. Let $\Gamma_T = \{z \in \Lambda : \|z\|_1 \leq T\}$ and $\Lambda_T = \langle \Gamma_T \rangle$. Induction on derivations, using $\Psi(\vec{u}) - \Psi(\omega_A) = \sum_j (\Psi(\vec{u}_j) - \Psi(\omega_{B_j})) - \delta_R$ for $\vec{u} = \rho(\vec{u}_1, \dots, \vec{u}_q)$, shows $\Psi(\vec{u}) - \Psi(\omega_A) \in \Lambda_T$ for every $\vec{u} \in L_A(G)$.

Now let $X = A_{\mathbf{t}}$ have arity $d \leq f$ and fix $\vec{u} \in L_X$. For each i choose, from a fixed semilinear representation of the Parikh image of the regular language $\widehat{s}^{-1}(t_i) \cap \Sigma^+$, a linear set $c_i + \sum_{p \in P_i} \mathbb{N}p$ containing $\psi(u_i)$; the number and size of all such generators depend only on Σ and $\widehat{s}$. The nonnegative integer system

$$\sum_{i=1}^{d} \Big(c_i + \sum_{p \in P_i} \nu_{i,p} p \Big) - \sum_{z \in \Gamma_T} (\mu_z^+ - \mu_z^-) z = \Psi(\omega_A)$$

is feasible (use the Parikh vectors of the u_i and the Λ_T-membership of $\Psi(\vec{u}) - \Psi(\omega_A)$). It has $|\Sigma|$ equations, $O_{f,s,\Sigma}((T+1)^{|\Sigma|})$ variables, and coefficients and right-hand side of absolute value $O_{f,s,\Sigma}(T+1)$; in particular, each coordinate of $\Psi(\omega_A)$ is at most $t_G \le \theta_G \le T$ for the nonempty core. Papadimitriou's bound for small nonnegative integral solutions of equality systems [41] (a solution with all entries at most $n(ma)^{2m+1}$ for m equations, n variables and entries bounded by a) gives a solution whose entries are bounded by a polynomial in $T+1$ with exponent depending only on $|\Sigma|$. Realize each resulting Parikh vector by a nonempty word $v_i \in \widehat{s}^{-1}(t_i)$ and put $\vec{v} = (v_1, \dots, v_d)$. Then $\|\vec{v}\|_1$ is polynomial in $T+1$, $\widehat{s}^{(d)}(\vec{v}) = \mathbf{t}$, and $\Psi(\vec{v}) - \Psi(\vec{u}) \in \Lambda_T \subseteq \Lambda$. For every sentence context E, the words $E[\vec{v}]$ and $E[\vec{u}]$ have the same $\widehat{s}$-value and $\varphi(E[\vec{v}]) - \varphi(E[\vec{u}]) = \overline{\varphi}(\Psi(\vec{v}) - \Psi(\vec{u})) \in \mathcal{H}$; hence $\vec{v} \equiv_{L,s} \vec{u}$ and $\kappa_G(X) \le \|\vec{v}\|_1$. Since $T \le (r+1)\theta_G + \|G\|$, the claim follows, with parameters independent of the auxiliary abelian group. □

Theorem 4.22 applies even when the positive image of $\widehat{s}$ lies *outside* **LG**: it is a sufficient condition on the target language, not on the finite control, and it claims no polynomial-time construction for ineffectively presented auxiliary groups.

5 Comparison with Multidimensional Substitutability

Yoshinaka's multidimensional substitutability is the direct positive-data baseline [13, 14]. Under trivial typing the present framework reproduces that semantic class; finite typing refines which observed tuples may be identified while preserving compositional rule witnesses, and $\widehat{s}_{k,\ell}$ lifts the boundary information of Yoshinaka's CFG (k,ℓ) hierarchy [12] componentwise. Since every (f,s) condition includes arity one, $\mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathrm{CFL} \subseteq \mathcal{C}^{\mathrm{mcf}}_{1,s}$.

5.1 Yoshinaka's Baseline and Prefix–Suffix Specialization

Definition 5.1 (Yoshinaka multidimensional substitutability)**.** Using the sentence-context sets $\mathcal{E}_d$ defined in Section 2, a language L is *fD-substitutable* if, for every $1 \le d \le f$, all $E, F \in \mathcal{E}_d$, and all $\vec{x}, \vec{y} \in (\Sigma^+)^d$, $E[\vec{x}], \quad E[\vec{y}], \quad F[\vec{x}] \in L \Longrightarrow F[\vec{y}] \in L$. Write $\mathsf{YSub}(f)$ for this semantic class and $\mathsf{YSL}(f,r) := \mathsf{YSub}(f) \cap \mathsf{MCFL}(f,r)$, where r bounds rule rank. We refer to the displayed four-instance implication as the *substitution implication*. Yoshinaka's original definition quantifies over every arity $d \le f$, uses nonempty tuple components, and requires nonempty internal separators in d-contexts; footnote 3 separately discusses the stronger alternative that permits empty components and separators [14, Section 3.1 and footnote 3].

The trivial typing removes the componentwise type premise, so $\mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{triv}}} = \mathsf{YSub}(f) \cap f\text{-MCFL} \subseteq \mathcal{C}^{\mathrm{mcf}}_{f,s}$. In Yoshinaka's word-level (k,ℓ) condition [12], with $u \in \Sigma^k$, $v \in \Sigma^\ell$ and $uy_1v, uy_2v \ne \varepsilon$,

$$w_1uy_1vw_1',\ w_1uy_2vw_1' \in L \quad \Longrightarrow \quad (w_2uy_1vw_2' \in L \iff w_2uy_2vw_2' \in L).$$

The compared fragments uy_iv are nonempty although the y_i may be empty, and equality of their $\widehat{s}_{k,\ell}$-values is exactly the boundary information used here, including exact identity of words shorter than $m_0 = \max\{1, k+\ell\}$. Hence $(1, \widehat{s}_{k,\ell})$-substitutability is precisely Yoshinaka's (k,ℓ) condition, and the same correspondence applies componentwise at larger fan-out.

A language is *strictly 2-local* if membership is determined by permitted adjacent pairs in $\vdash w \dashv$, for fixed endmarkers.

Lemma 5.2 (Strictly 2-local constraints are boundary typed)**.** *If $L_{\mathrm{sl}} \subseteq \Sigma^*$ is strictly 2-local, then for every $f \ge 1$, $L_{\mathrm{sl}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}}$.*

Proof. Let equal-$\widehat{s}_{1,1}$-type nonempty tuples $\vec{x}, \vec{y}$ share an accepting context E_0. Internal adjacent pairs of each y_i are permitted because they occur in $E_0[\vec{y}]$. If $E[\vec{x}] \in L_{\mathrm{sl}}$, every adjacent pair of $E[\vec{y}]$ is either in fixed material, internal to a y_i, or crosses a component boundary. These are permitted respectively by $E[\vec{x}]$, by $E_0[\vec{y}]$, or by equality of the first/last letters encoded by the boundary type; end-marker pairs are unchanged. Thus $E[\vec{y}] \in L_{\mathrm{sl}}$, and symmetry finishes the proof. □

5.2 An Intrinsically Rank-Two Boundary-Typed Target

In this subsection, $s_{1,1}$ denotes the positive-word restriction $\widehat{s}_{1,1}|_{\Sigma_{\mathrm{intr}}^+}$ of the prefix–suffix typing from Definition 2.4, with alphabet Σ_{intr} defined below.

Let

$$\Sigma_{\mathrm{intr}} := \{\mathtt{a}, \mathtt{b}, \mathtt{l}, \mathtt{r}, \mathtt{p}, \mathtt{q}, \mathtt{c}, \mathtt{d}, \#, \$, \mathtt{u}, \mathtt{v}, \mathtt{w}\}, \qquad P_{\mathrm{intr}} := \{\mathtt{lpc}, \mathtt{lpd}, \mathtt{lqc}, \mathtt{rpc}, \mathtt{rpd}\},$$

$$\mathrm{MIX}_2^+ := \{x \in \{\mathtt{a}, \mathtt{b}\}^+ : |x|_{\mathtt{a}} = |x|_{\mathtt{b}}\}, \qquad T_3^+ := \{\mathtt{u}^n\mathtt{v}^n\mathtt{w}^n : n \geq 1\}, \qquad L_{\mathrm{intr}} := P_{\mathrm{intr}} \# \mathrm{MIX}_2^+ \$ T_3^+.$$

The MIX_2^+ factor carries the rule-rank obstruction, which is not created by finite typing: $\mathrm{MIX}_2^+ = \{\mathtt{a}, \mathtt{b}\}^+ \cap \varphi_{\mathrm{mix}}^{-1}(0)$ for $\varphi_{\mathrm{mix}}(\mathtt{a}) = 1$, $\varphi_{\mathrm{mix}}(\mathtt{b}) = -1$ in $\mathbb{Z}$, so Proposition 2.10(ii) already places it in $\mathsf{YSub}(f)$ for every f. The T_3^+ factor forces minimum fan-out two, and the wrapper P_{intr} separates the target from $\mathsf{YSub}(f)$.

Theorem 5.3 (A non-context-free intrinsic-rank-two typed witness)**.** *The language L_{intr} satisfies $L_{\mathrm{intr}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}$, and, for every $f \geq 2$, $L_{\mathrm{intr}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}} \setminus \mathsf{YSub}(f)$. Moreover, for every finite $f' \geq 1$, $L_{\mathrm{intr}} \notin \mathsf{MCFL}(f', 1)$. Consequently the minimum fan-out of L_{intr} is exactly two, and its minimum rule rank among all finite-fan-out MCFG presentations is exactly two. In particular, Theorem 4.8 applies to a non-context-free language that cannot be covered by Theorem 4.6 at any finite fan-out.*

Proof. Grammar. Let B have the rules

$$B \to \mathtt{a}B\mathtt{b}B \mid \mathtt{b}B\mathtt{a}B \mid \mathtt{a}B\mathtt{b} \mid \mathtt{ab}B \mid \mathtt{ab} \mid \mathtt{b}B\mathtt{a} \mid \mathtt{ba}B \mid \mathtt{ba};$$

the first-return induction on $|x|_{\mathtt{a}} - |x|_{\mathtt{b}}$ gives $L_B = \mathrm{MIX}_2^+$. Let the fan-out-two nonterminal A have rules $A \to (\mathtt{u}, \mathtt{w}) \mid (\mathtt{u}\xi_1^1\mathtt{v}, \mathtt{w}\xi_1^2)(A)$, so $L_A = \{(\mathtt{u}^n\mathtt{v}^{n-1}, \mathtt{w}^n) : n \geq 1\}$, and for every $\alpha \in P_{\mathrm{intr}}$ add $S \to \alpha\#\xi_1^1\$\xi_2^1\mathtt{v}\xi_2^2(B, A)$. The grammar generates L_{intr}, is reduced, λ-free, linear, nondeleting and nonpermuting, and is nonmerging because the two components of A are separated by the literal $\mathtt{v}$. It is good, of fan-out two and rule rank two.

Non-context-freeness. With $R := \mathtt{lpc}\#\mathtt{ab}\$\,\mathtt{u}^+\mathtt{v}^+\mathtt{w}^+$, $L_{\mathrm{intr}} \cap R = \{\mathtt{lpc}\#\mathtt{ab}\$\mathtt{u}^n\mathtt{v}^n\mathtt{w}^n : n \geq 1\}$, which a homomorphism erasing the fixed prefix maps onto the non-context-free $\{\mathtt{u}^n\mathtt{v}^n\mathtt{w}^n : n \geq 1\}$. By closure of CFL under intersection with regular languages and homomorphism, L_{intr} is not context-free; fan-out-one MCFGs are exactly CFGs here, so the minimum fan-out is two.

Typed membership. Let $Q_{\mathrm{intr}} := P_{\mathrm{intr}} \# \{\mathtt{a}, \mathtt{b}\}^+ \$\,\mathtt{u}^+\mathtt{v}^+\mathtt{w}^+$. It is strictly 2-local, with the allowed adjacent pairs

$$\vdash \mathtt{l},\ \vdash \mathtt{r},\ \mathtt{lp},\ \mathtt{lq},\ \mathtt{rp},\ \mathtt{pc},\ \mathtt{pd},\ \mathtt{qc},\ \mathtt{c\#},\ \mathtt{d\#},\ \mathtt{\#a},\ \mathtt{\#b},\ \mathtt{aa},\ \mathtt{ab},\ \mathtt{ba},\ \mathtt{bb},$$
$$\mathtt{a\$},\ \mathtt{b\$},\ \mathtt{\$u},\ \mathtt{uu},\ \mathtt{uv},\ \mathtt{vv},\ \mathtt{vw},\ \mathtt{ww},\ \mathtt{w} \dashv .$$

Let $\Phi_{\mathrm{intr}} : \Sigma_{\mathrm{intr}}^* \to \mathbb{Z}^3$ send $\mathtt{a}, \mathtt{b}, \mathtt{u}, \mathtt{v}, \mathtt{w}$ to $(1,0,0)$, $(-1,0,0)$, $(0,1,0)$, $(0,-1,1)$, $(0,0,-1)$ and all other letters to 0. Its three coordinates impose $|x|_{\mathtt{a}} = |x|_{\mathtt{b}}$, $|x|_{\mathtt{u}} = |x|_{\mathtt{v}}$ and $|x|_{\mathtt{v}} = |x|_{\mathtt{w}}$, so $L_{\mathrm{intr}} = Q_{\mathrm{intr}} \cap \Phi_{\mathrm{intr}}^{-1}(0)$. Lemma 5.2 and Proposition 2.10(ii) make this intersection $(f, s_{1,1})$-tuple-substitutable for every f; with the grammar above, $L_{\mathrm{intr}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}}$ for every $f \geq 2$.

Separation. For $x = \mathtt{p}$, $y = \mathtt{q}$, $E = \mathtt{l}\Box_1\mathtt{c}\#\mathtt{ab}\$\mathtt{uvw}$ and $F = \mathtt{r}\Box_1\mathtt{d}\#\mathtt{ab}\$\mathtt{uvw}$, we have $E[x], E[y], F[x] \in L_{\mathrm{intr}}$ but $F[y] \notin L_{\mathrm{intr}}$, so the arity-one substitution implication fails and $L_{\mathrm{intr}} \notin \mathsf{YSub}(f)$ for every $f \geq 1$; the boundary typing separates the offending factors, $\widehat{s}_{1,1}(\mathtt{p}) \neq \widehat{s}_{1,1}(\mathtt{q})$.

Intrinsic rule rank. Gallot calls an MCFG *non-branching* when every positive-rank rule has a single child nonterminal, with no global fan-out bound [21, Definition 38]; his class is therefore $\bigcup_{f'<\infty} \mathsf{MCFL}(f', 1)$. Definition 38 permits a child variable to occur at most once, and Theorem 28 shows that non-permuting non-erasing presentations generate the same class. Theorem 29 identifies the class with finite-index EDT0L, and Theorem 31 (via Propositions 5–6) shows that $\mathrm{MIX}_2 := \{x \in \{\mathtt{a}, \mathtt{b}\}^* : |x|_{\mathtt{a}} = |x|_{\mathtt{b}}\}$ is not finite-index EDT0L [21, Theorems 28, 29, 31 and Propositions 5–6]; Bishop et al. give a peer-reviewed corroborating account in the R-MCFG formalism [22, Definition 6.1, Lemma 6.8, Section 7, Theorem C]. Let π fix $\mathtt{a}, \mathtt{b}$ and erase all other letters, so $\pi(L_{\mathrm{intr}}) = \mathrm{MIX}_2^+$. Applying π to the terminal material of a fan-out-f' rank-one presentation of L_{intr} would give one of MIX_2^+ (empty terminal pieces are allowed in the standard formalism), and a fresh start symbol with a unary identity link and a rank-zero ε rule would give one of MIX_2, a presentation in Gallot's class (Definition 38, Theorem 28), contradicting his lower bound. Hence $L_{\mathrm{intr}} \notin \mathsf{MCFL}(f', 1)$ for every finite f', and the displayed grammar proves the claimed intrinsic rule rank. □

The two lower bounds arise from separate factors rather than from an irreducible interaction between the resources, and the non-branching lower bound itself is prior work. What is specific here is that a strictly local wrapper and the $(1, 1)$ boundary typing place the target inside the typed class while the substitution implication keeps it outside Yoshinaka's trivial-typing baseline.

The quantitative gap between one- and higher-arity substitutability has a distinct source: in a two-hole context, a fixed separator can slide across different occurrences in one sentence. Repeated type pairs then force *distributional* equality, a constraint invisible at arity one. The next section makes this obstruction precise for the companion's doubling family without claiming a higher-arity characteristic-data lower bound.

6 A Higher-Arity Obstruction to Transferring the CFG Lower Bound

For a word m and positions $0 \leq p < q \leq |m|$, we write $m[0 : p]$ and $m[q : |m|]$ for its prefix and suffix. Only nonempty tuple components are used in the following lemma.

Lemma 6.1 (Sliding separators)**.** *Let $s : \Sigma^+ \to \mathsf{S}$ be a finite-semigroup typing and let L be $(2, s)$-tuple-substitutable. Suppose $v_0 m v_2 \in L$ and a nonempty word z occurs in m at $N > |\mathsf{S}|^2$ distinct positions p, each satisfying $0 < p$ and $p + |z| < |m|$. Then there are positions $p < q$ for which $\vec{x} = (m[0 : p], m[p + |z| : |m|]), \qquad \vec{y} = (m[0 : q], m[q + |z| : |m|])$ have equal componentwise s-types and $\mathcal{D}_L(\vec{x}) = \mathcal{D}_L(\vec{y})$.*

Proof. The two components of each tuple are nonempty, and all the tuples share the admissible context $E = v_0\Box_1 z\Box_2 v_2$ with $E[\vec{x}] = E[\vec{y}] = v_0 m v_2 \in L$. There are at most $|\mathsf{S}|^2$ componentwise type pairs, so two distinct positions give equal types. By $(2, s)$-tuple substitutability and the common accepting context, their full sentence-context distributions agree. □

We apply the sliding effect to the bracketed doubling family of the companion CFG work [25]. To make the scope of the claim explicit, recall its construction. Let $\Gamma_{\mathrm{br}} = \{\mathtt{a}, \lfloor, \rfloor, \lceil, \rceil\}$ and put $E_0 \to \mathtt{a}$ and

$$E_i \to \lfloor E_{i-1} E_{i-1} \rfloor \mid \lceil E_{i-1} \rceil \quad (1 \leq i \leq n), \qquad S_0 \to E_n.$$

Let L_n be the generated finite language, $w_0^* = \mathtt{a}$, $w_i^* = \lfloor w_{i-1}^* w_{i-1}^* \rfloor$, and $L_n^- = L_n \setminus \{w_n^*\}$. The following nonmembership result does *not* depend on the two-element semilattice typing used in the CFG lower bound.

Proposition 6.2 (No fixed finite typing lifts the doubling family)**.** *For every finite-semigroup typing $s : \Gamma_{\mathrm{br}}^+ \to \mathsf{S}$ and every $n \geq 2|\mathsf{S}| + 3$, neither L_n nor L_n^- is $(2, s)$-tuple-substitutable. Consequently, for every $f \geq 2$, $L_n, L_n^- \notin \mathcal{C}_{f,s}^{\mathrm{mcf}}$.*

Proof. Put $q = |\mathsf{S}|$, and let $a := s(\lfloor)$. Among $a, a^2, \ldots, a^{q+1}$ choose equal powers $a^i = a^j$ with $1 \leq i < j \leq q + 1$. Writing $\pi = j - i$, the identity $a^{i+\pi} = a^i$ implies $a^{t+\pi} = a^t$ for every $t \geq i$. Our lower bound on n ensures $n - j - 1 \geq q + 1 \geq i$ and $j \leq n - 2$. It follows that $s(\lfloor^i) = s(\lfloor^j)$, $\quad s(\lfloor^{n-i-1}) = s(\lfloor^{n-j-1})$. Thus the nonempty pairs $\vec{x} = (\lfloor^i, \lfloor^{n-i-1})$, $\quad \vec{y} = (\lfloor^j, \lfloor^{n-j-1})$ have equal componentwise types.

For L_n take $w = w_n^*$. For L_n^- take w obtained from w_n^* by replacing a deepest binary rule on the rightmost branch by its unary $\lceil E_0 \rceil$ alternative. This replacement lies outside the leftmost branch, so in both cases $w = \lfloor^n v$ for a suffix v. The admissible context $E = \Box_1 \lfloor \Box_2 v$ fills with both $\vec{x}$ and $\vec{y}$ to produce the same word w in the respective target.

Now use the binary node on the leftmost branch at depth i, whose two children are instances of E_{n-i-1}. Replace that rule in w by its unary $\lceil E_{n-i-1} \rceil$ alternative, retaining the original left child and dropping the right child. Call the resulting word w'. The replacement is grammatical, and $w' \in L_n^- \subseteq L_n$. Since w' begins with $\lfloor^i \lceil \lfloor^{n-i-1}$, it has a factorization $w' = \lfloor^i \lceil \lfloor^{n-i-1} v'$ for some v'. The context $F = \Box_1 \lceil \Box_2 v'$ therefore accepts $\vec{x}$ in both target languages.

By contrast, $F[\vec{y}] = \lfloor^j \lceil \lfloor^{n-j-1} v'$. Relative to w', the $\lceil$ opening symbol has moved $j - i > 0$ positions into the previously contiguous initial sequence of $\lfloor$ openings, while the corresponding closing symbols in v' have not moved. The first $\rceil$ that matches this unary opening in w' is now preceded by too many $\rfloor$ closings for the $n - j - 1$ $\lfloor$ openings following $\lceil$; a $\rfloor$ must close $\lceil$. Hence the bracket types do not match, so $F[\vec{y}] \notin L_n$ and, a fortiori, $F[\vec{y}] \notin L_n^-$. We have an accepted common context, equal types, and different distributions, contrary to $(2, s)$-tuple substitutability. □

Remark 6.3 (What the obstruction does not prove). Proposition 6.2 strengthens the non-transfer observation of Remark C.8: no *fixed* finite typing repairs two-dimensional substitutability for this particular doubling family when n grows. It is an instance of the sliding-separator phenomenon in Lemma 6.1. This rules out a direct reuse of the companion's examples; it does not rule out other higher-fan-out families with superpolynomial characteristic data. The family is not itself a higher-fan-out lower bound.

Open Problem 6.4 (Original-thickness bounds beyond fan-out one). Fix $f \geq 2$ and a rule-rank bound $r \geq 2$. Does there exist a fixed finite-semigroup typing s and a family of reduced good fan-out-$\leq f$, rank-$\leq r$ grammars G_n of targets in $\mathcal{C}_{f,s}^{\mathrm{mcf}}$ with $\|G_n\| \max\{1, \theta_{G_n}\} = n^{O(1)}$ but superpolynomial total positive length for every characteristic sample of $\mathcal{B}_{f,s}^{\leq r}$? The family above cannot provide such examples, and Proposition 4.21 shows that any such family must have superpolynomial minimum typed-distributional representative length κ_{G_n}.

6.1 Distinguishing Pairs and Operator-Independent Data Lower Bounds

The following distinguishing-pair argument is elementary for set-driven learning and introduces no new learning algorithm. The companion CFG paper uses a related argument [25, Proposition 6.4(iv)]. For MCFGs, we make explicit that the relevant filters preserve fan-out and rule rank.

Lemma 6.5 (Regular intersections preserve fan-out and rule rank)**.** *Let $f, r \geq 1$. If $L \in \mathsf{MCFL}(f, r)$ and $R \subseteq \Sigma^*$ is regular, then $L \cap R \in \mathsf{MCFL}(f, r)$.*

Proof. Let $\mathcal{D} = (Q, \Sigma, \delta, q_0, F)$ be a DFA for R. Refine each nonterminal A of fan-out $d \le f$ in a linear nondeleting MCFG for L by d pairs of DFA states, one entry and exit state for each output component. For each production retain exactly the finitely many assignments of states for which all terminal constants and occurrences of child components are compatible with the left-to-right DFA transitions. Each child component occurs exactly once, so the constraints are local to the production; neither the number of output components nor the number of children increases. At the start allow only pairs (q_0, q) with $q \in F$. The resulting finite grammar generates $L \cap R$ without increasing fan-out or rule rank. This is the standard regular-intersection product construction for MCFGs [30], with its rank bound made explicit. □

Lemma 6.6 (Forcing by a distinguishing pair)**.** *Let $\mathcal{C}$ be a language class and let* Op *be a set-driven operator having a characteristic sample for every language in $\mathcal{C}$. If $L', L \in \mathcal{C}$ and $L' \subsetneq L$, then every characteristic sample K of L for* Op *meets $L \setminus L'$.*

Proof. Otherwise $K \subseteq L'$. Choose a characteristic sample K' for L'. The finite sample $K \cup K' \subseteq L'$ contains both characteristic samples, so $\mathrm{Op}(K \cup K')$ would generate both L and L', a contradiction. □

Corollary 6.7 (Forcing by fibres and finite-index cosets)**.** *Fix $f, r \ge 1$ and a finite-semigroup typing s, and put $\mathcal{C} := \mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathsf{MCFL}(f, r)$. Let $L \in \mathcal{C}$ and let R be one of the following:*

(i) $R = \widehat{s}^{-1}(F)$ *with* $F \subseteq \mathsf{S}^\sharp$*;*

(ii) $R = \varphi^{-1}(\gamma + \mathcal{H})$*, where* $\varphi : \Sigma^* \to \mathbb{A}$ *is a morphism into an abelian group,* $\mathcal{H} \le \mathbb{A}$ *has* finite index*, and* $\gamma \in \mathbb{A}$*.*

If $L \setminus R \neq \emptyset$, every characteristic sample K of L for every set-driven operator having characteristic samples throughout $\mathcal{C}$ satisfies $K \cap (L \setminus R) \neq \emptyset$. In particular, K meets every nonempty $\widehat{s}$-type fibre of L.

Proof. In (i), R is regular as an inverse image under a finite monoid morphism. In (ii), R is regular since it is recognized by the finite quotient $\mathbb{A}/\mathcal{H}$. For the finite type-fibre filter, equality of componentwise s-values ensures that every admissible sentence context gives the same filter value, so intersection preserves (f, s)-tuple substitutability. For the finite-index abelian-coset filter, the same preservation follows from Proposition 2.10(ii). By Lemma 6.5, it is also in $\mathsf{MCFL}(f, r)$. Therefore $L' = L \cap R \in \mathcal{C}$ is a proper sublanguage, and Lemma 6.6 applies. To force a word in a nonempty type fibre of type μ, use $R = \widehat{s}^{-1}(\mathsf{S}^\sharp \setminus \{\mu\})$. □

Let $\mathcal{F}_s$ be the union of the two families of filters above and define their *forcing depth* for $L \in \mathcal{C}$ by

$$\Delta_s(L) := \sup_{\substack{R \in \mathcal{F}_s \\ L \setminus R \neq \emptyset}} \min_{w \in L \setminus R} (|w| + 1),$$

where the supremum of the empty family is defined to be zero. The corollary gives $\|K\|_+ \ge \Delta_s(L)$ for every characteristic sample of the operators in question. Theorem 3.13 provides such an operator under a fixed rank cap, so $\Delta_s(L)$ is finite for every fixed target $L \in \mathcal{C}$. Superpolynomial forcing depth along a target family would provide an operator-independent lower bound; no converse is claimed.

Open Problem 6.8 (Operator-independent higher-fan-out lower bounds). Fix $f \geq 2$ and $r \geq 2$. Does there exist a fixed finite-semigroup typing s and a family of reduced good fan-out-$\leq f$, rank-$\leq r$ grammars G_n generating targets in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, with $\|G_n\|\bar{\theta}_{G_n} = n^{O(1)}$, such that for *every* set-driven operator having characteristic samples on all of $\mathcal{C}$, every characteristic sample of $L(G_n)$ has superpolynomial total encoded length? The condition $\Delta_s(L(G_n)) = n^{\omega(1)}$ would suffice; it is not known to be necessary.

6.2 Short coset certificates and obstructions for finite targets

We record two limitations on lower-bound constructions. The first restricts *which competitors* can force large positive samples; the second restricts *which finite grammars* can hide very long examples behind short rule witnesses. Neither settles Open Problems 6.4 and 6.8.

For a tuple $\vec{u}$ let $\Psi(\vec{u})$ denote the sum of the Parikh vectors of its components, and write $\psi(w)$ for the Parikh vector of a word. The following identity makes explicit the rule-difference mechanism used in Theorem 4.22.

Proposition 6.9 (Short Parikh-difference certificates)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of rule rank at most r, with $\emptyset \neq L(G) \subseteq \Sigma^+$, and put $\theta = \theta_G$. For each $A \in V$ choose a minimum-length tuple $\omega_A \in L_A(G)$. For a rule $R : A \to \rho(B_1, \ldots, B_q)$ let c_R count the explicit terminals in ρ and put $\delta_R = c_R + \sum_{j=1}^{q} \Psi(\omega_{B_j}) - \Psi(\omega_A) \in \mathbb{Z}^{\Sigma}$. Then*

$$\langle \psi(u) - \psi(v) : u, v \in L(G) \rangle_{\mathbb{Z}} = \langle \delta_R : R \in P \rangle_{\mathbb{Z}}.$$

Each δ_R is the difference of the Parikh vectors of two positive words of length at most $(r+2)\theta$. Moreover, $\|\delta_R\|_1 \leq (r+2)\theta$.

Proof. For a derivation tree T rooted at A and yielding $\vec{u}$, expand $\sum_{v \in T} \delta_{R(v)}$. Each non-root reference vector $\Psi(\omega_B)$ cancels between its own node and its parent's rule; linearity and nondeletion give $\Psi(\vec{u}) - \Psi(\omega_A) = \sum_{v \in T} \delta_{R(v)}$. For the start symbol this proves one inclusion.

Conversely, Lemma 4.1(iv) supplies for each A an admissible context E_A with at most θ fixed terminals that accepts every tuple in $L_A(G)$. Both $u_R = E_A[\omega_A]$ and $v_R = E_A[\rho(\omega_{B_1}, \ldots, \omega_{B_q})]$ are in $L(G)$, and $\psi(v_R) - \psi(u_R) = \delta_R$. Lemma 4.1(i)–(ii) bounds each reference tuple and each rule's explicit terminal material by θ. Thus $|u_R| \leq 2\theta$, $|v_R| \leq (q+2)\theta$, and the triangle inequality applied directly to the definition of δ_R gives $\|\delta_R\|_1 \leq (q+2)\theta$. □

Corollary 6.10 (Uniform certificates for abelian-coset competitors)**.** *Under the assumptions of Proposition 6.9, there is a single set $T_G \subseteq L(G)$ of at most $|V| + |P|$ words with $\|T_G\|_+ \leq (|V| + |P|)(1 + (r+2)\theta_G)$, such that for* every *morphism $\varphi : \Sigma^* \to \mathbb{A}$ into an abelian group, subgroup $H \leq \mathbb{A}$ and coset $\gamma + H$,*

$$T_G \subseteq \varphi^{-1}(\gamma + H) \quad \Longrightarrow \quad L(G) \subseteq \varphi^{-1}(\gamma + H).$$

For a fixed finite typing s put $M_s(L) = \max_{t \in \widehat{s}(L)} \min\{|w| + 1 : w \in L, \widehat{s}(w) = t\}$. For the filter family $\mathcal{F}_s$ defined above, when $L = L(G) \subseteq \Sigma^+$ belongs to the corresponding typed class,

$$M_s(L) \leq \Delta_s(L) \leq \max\{M_s(L), (r+2)\theta_G + 1\}.$$

All proper competitors $L \cap R$ with $R \in \mathcal{F}_s$ have their difference sets simultaneously met by some positive set of weight at most $(|V| + |P|)(1 + (r+2)\theta_G) + |\mathsf{S}| M_s(L)$.

Proof. Take T_G to contain $E_A[\omega_A]$ for each A, and $E_A[\rho(\omega_{B_1}, \ldots, \omega_{B_q})]$ for each rule R. If all these words lie in a coset, every δ_R maps into H; the preceding lattice identity then puts every word of $L(G)$ into that coset. Hence each coset filter which excludes a target word is violated by a member of T_G, of length at most $(r+2)\theta_G$. For a type filter which excludes a realized type, a shortest word of that type has encoded length at most $M_s(L)$; conversely, excluding a type achieving $M_s(L)$ gives the lower bound. Adjoining one shortest word for each realized type to T_G proves the simultaneous bound. □

The uniform set above need not be a characteristic sample of any learner. Thus the corollary excludes superpolynomial *hitting-set* arguments based only on these abelian-coset competitors when the original grammar thickness is polynomial; it does not exclude other competitors or general operator-independent lower bounds.

We turn to finite languages. A *componentwise factor embedding* $\vec{x} \preceq_{\mathrm{fac}} \vec{y}$ means that each x_i is a contiguous factor of y_i.

Lemma 6.11 (Finite-target tuple-nesting obstruction)**.** *Let L be finite and (f, s)-tuple-substitutable. If nonempty d-tuples $\vec{x}, \vec{y}$, $d \le f$, have equal componentwise types and a common accepting admissible context, then $\vec{x} \preceq_{\mathrm{fac}} \vec{y}$ implies $\|\vec{x}\|_1 = \|\vec{y}\|_1$. In particular, if G is a reduced good grammar for L, each componentwise-type fibre of $L_A(G)$ is an antichain under proper componentwise factor embedding.*

Proof. Suppose $y_i = p_i x_i q_i$ and $h = \sum_i (|p_i| + |q_i|) > 0$. Equality of the full tuple distributions follows from substitutability. Write $T(\vec{z}) = (p_i z_i q_i)_i$ and let E be a common accepting context. Every $E[T^k(\,\cdot\,)]$ is still an admissible d-hole context: the original nonempty separators remain in place. Inductively, the equality of distributions forces $E[T^k(\vec{x})] \in L$ for all $k \ge 0$, producing words of strictly increasing lengths, a contradiction. For grammar yields use the common successful context E_A from Lemma 4.1(iv). □

For the next result fix $\Sigma = \{\mathtt{a}, \#\}$ and the non-local-group two-element semilattice typing $s(\mathtt{a}) = 0$, $s(\#) = 1$, with multiplication max. The positive type-zero fibre is $L \cap \mathtt{a}^+$.

Lemma 6.12 (Transport of short unary mass)**.** *Suppose L is $(2, s)$-tuple-substitutable and $N = \min\{k \ge 1 : \mathtt{a}^k \in L\}$ exists. For $2 \le m < N$, the tuples $(\mathtt{a}^i, \mathtt{a}^{m-i})$, $1 \le i < m$, have the same complete two-hole distribution. Consequently, if $\mathtt{a}^p z \mathtt{a}^{m-p} \in L$ for any nonempty z and $1 \le p < m$, then $\mathtt{a}^i z \mathtt{a}^{m-i} \in L$ for all $1 \le i < m$.*

Proof. All the displayed tuples have type $(0, 0)$, and the admissible context $\Box_1 \mathtt{a} \Box_2 \mathtt{a}^{N-m-1}$ accepts every one, producing $\mathtt{a}^N$. Apply two-dimensional substitutability, then use $\Box_1 z \Box_2$. □

Theorem 6.13 (Linear short-fibre bound with an interior marker)**.** *Let G be a reduced good MCFG of rule rank at most r with finite nonempty $L = L(G) \subseteq \mathtt{a}^+ \cup \mathtt{a}^+ \# \mathtt{a}^+$. If L is $(2, s)$-tuple-substitutable, then*

$$M_s(L) \le (r+2)\theta_G + 2.$$

This holds for arbitrary fan-out, provided the relevant two-hole substitutability condition is imposed.

Proof. For a finite $(1, s)$-substitutable language, two accepted words of the same type cannot properly contain one another as contiguous factors: apply Lemma 6.11 to the identity context. In particular, $L \cap \mathtt{a}^+$ has at most one word. If only one type is realized, its shortest word has length at most θ_G; suppose henceforth that both are realized, with type-zero word $\mathtt{a}^N$. Some rule writes $\#$, and its shortest successful use gives a type-one word of length at most θ_G. If $N \le \theta_G$ the conclusion follows. Otherwise write this short word as $\mathtt{a}^p \# \mathtt{a}^q$, $p, q \ge 1$, and put $m = p + q < N$. If

$m \geq 3$, Lemma 6.12 produces every interior split of total unary mass m. Comparing $\#\mathtt{a}^{m-1}$ and $\mathtt{a}\#\mathtt{a}^{m-2}$ under the common context $\mathtt{a}\square_1$, and then using $\square_1\mathtt{a}$, forces $\#\mathtt{a}^m \in L$, contrary to the interior-marker hypothesis. Thus $m = 2$ and the short word is $\mathtt{a}\#\mathtt{a}$. Every other word of type one properly contains this word as a factor, which contradicts Lemma 6.11. Hence $L = \{\mathtt{a}^N, \mathtt{a}\#\mathtt{a}\}$.

The Parikh-difference group is generated by $(N-2,-1)$ and is nontrivial. Proposition 6.9 therefore supplies a nonzero rule defect. As its defining positive words belong to the two-word language, the defect is $\pm(N-2,-1)$, of norm $N-1$. Consequently $N-1 \leq (r+2)\theta_G$, giving the claimed bound including encoded word cost. □

The finite-language hypothesis above is essential to the factor argument; the interior-marker hypothesis excludes boundary markers. Neither assumption follows from general two-dimensional substitutability.

Theorem 6.14 (Length spectrum of a finite good MCFG)**.** *Let G be a reduced good MCFG of rule rank at most r with nonempty finite language $L(G)$. Put $\theta = \theta_G$ and $t = |\{|w| : w \in L(G)\}|$. Then*

$$\max_{w \in L(G)} |w| \leq (r+1)|V|\,\theta\,\max\{1, t-1\}.$$

This bound does not assume any semigroup typing or substitutability.

Proof. Choose, for each accepted word, a successful derivation with the fewest nodes. No root-to-leaf path of such a derivation repeats a nonterminal. Indeed, a return path from A to A either introduces some terminal or a productive side-child, in which case repeating it makes $L(G)$ infinite by linearity and nondeletion; or it introduces nothing, in which case every step is a unary, terminal-free variable redistribution. Since each component remains nonempty, and the endpoints have the same fan-out, the composite redistributes the components bijectively. The nonpermuting condition makes that bijection the identity, so the zero-growth return can be removed, contradicting minimality. Thus every such derivation has depth at most $|V|$.

Call a node heavy if the total length of its tuple yield exceeds θ. Every nonterminal has a yield of total length at most θ by Lemma 4.1(ii). If there are k pairwise incomparable heavy nodes, replace these subtrees, one at a time, by short yields. Each replacement strictly decreases the final length, so the resulting $k+1$ accepted lengths are distinct. Hence $k \leq t-1$. The heavy nodes form an ancestor-closed subtree of depth at most $|V|$, and therefore number at most $|V|(t-1)$. If the root is light, $|w| \leq \theta$. Otherwise every terminal is written either by a heavy node itself (at most θ terminals per node) or by one of its at most r light children (at most θ terminals in each). It follows that $|w| \leq (r+1)\theta|V|(t-1)$, as required. □

Thus, if a family of finite good grammars has polynomial original size and thickness but accepts a superpolynomial-length word, it must have superpolynomially many distinct accepted lengths. This necessary condition alone does not settle either higher-fan-out open problem.

7 Conclusion

Qualitative positive-data identification of typed MCFGs follows Yoshinaka's distributional framework. We obtain quantitative bounds on the total encoded length of characteristic samples relative to the original size and rule thickness of reduced good MCFG presentations. Prefix–suffix window typing, and window typing combined with an arbitrary finite group, admit $O(\|G\|^2\bar{\theta}_G)$ bounds. The group proof simultaneously preserves ordered products across output components without assuming commutativity. For positive-image local-group typings we establish $O(\|G\|^3\bar{\theta}_G)$. Even at

fan-out one, this strictly extends the **LI** $\vee$ **Grp** sufficient regime from the companion CFG work. We additionally establish bounds for rank-one grammars and abelian-filtered targets and give a computable rule-only data budget.

These algebraic hypotheses are sufficient conditions, and the displayed degrees have not been shown optimal. The companion CFG lower-bound obstruction is established at fan-out one; its exponential example family cannot be transferred to the two-dimensional substitutability setting. Distinguishing-pair arguments give an operator-independent forcing-depth criterion. Short Parikh-difference certificates show that abelian-coset competitors alone cannot yield a superpolynomial hitting-set obstruction under polynomial original rule thickness. For finite targets we give a tuple-nesting obstruction, a linear short-fibre bound in an interior-marker case, and a general finite-MCFG length-spectrum bound. These partial results do not settle whether higher fan-out can require superpolynomial characteristic-data length or forcing depth.

Acknowledgments

The author thanks Makoto Kanazawa and Ryo Yoshinaka for valuable discussions and encouragement.

Declaration of competing interest

The author declares that he has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data availability

No empirical datasets were used. The source manuscript and scripts for reproducing finite illustrative checks are available in the public repository `https://github.com/growupkuriyama-hub/Papers/tree/main/02_fixed-h-mcfg`. The main branch contains the canonical working source for this manuscript. The general theorems are supported by mathematical proofs in this article, not by exhaustive computational testing.

Use of AI tools in the research process

ChatGPT (OpenAI) and Claude (Anthropic) were used as aids in exploring proof approaches, critically examining proof drafts and potential counterexamples, cross-checking literature, and examining finite cases. Such output was treated as a proposal or sanity check, never as a substitute for mathematical proof. The general statements are justified by the arguments in the article and its appendices; finite computations alone do not establish the general theorems.

A Details of the Good Normalization

We record the four conversions used in Proposition 2.6 once, together with the tuple-language identity that each satisfies. Both the fixed-parameter size bound of that proposition and the thickness transfer of Proposition A.2 are read off from these identities. Semantic correctness of the nondeleting and nonmerging conversions is standard [14,30,32]; Kanazawa's existence statement [32, Section 2.1] does not by itself give the size or thickness bounds, which are therefore derived here. For $\vec{w} = (w_1, \dots, w_d)$ and nonempty $I \subseteq [d]$, $\pi_I(\vec{w})$ retains the components indexed by I in increasing order, $\sigma\vec{w}$ permutes components, and $\mu(\vec{w})$ concatenates consecutive blocks of components; thus $\|\pi_I(\vec{w})\|_1 \le \|\vec{w}\|_1$ and $\|\sigma\vec{w}\|_1 = \|\mu(\vec{w})\|_1 = \|\vec{w}\|_1$.

Lemma A.1 (Stagewise tuple languages)**.** *Fix f and $r_{\max}$, and let G_0 be a reduced f-MCFG of rule rank at most $r_{\max}$. Each of the following conversions preserves linearity, does not increase fan-out or rule rank, and multiplies the number of nonterminal and rule copies by a factor depending only on f and $r_{\max}$, without enlarging any template.*

(a) Nondeletion *(the construction (f3) in the proof of [30, Lemma 2.2]). Introduce A^I for nonempty $I \subseteq [\mathrm{fan}(A)]$. For $R : A \to \rho(B_1, \dots, B_q)$ let J_j be the set of components of B_j used by the I-indexed outputs of ρ; keep these outputs, rename their variables, use $B_j^{J_j}$ when $J_j \neq \emptyset$ and omit B_j otherwise. The resulting rules are nondeleting, and $L_{A^I} = \pi_I(L_A(G_0))$.*

(b) Nonpermutation. *Decorate a fan-out-d nonterminal by $\sigma \in \mathfrak{S}_d$; the order in which a permuted parent template uses a child's variables determines the child's permutation. The decorated nonterminal generates σL_A.*

(c) Empty components. *Let $\mathsf{Pat}(A)$ be the set of nonempty-component patterns of tuples in L_A, computed by finite saturation. For realizable child patterns, replace variables of absent components by ε, delete empty parent components and omit all-empty children; introduce A^I for each nonempty $I \in \mathsf{Pat}(A)$. Then L_{A^I} is the set of π_I-projections of the tuples of L_A with pattern I, the result is λ-free, and $S^{\{1\}}$ generates the nonempty part of the language. If every derived component is already nonempty, this stage is the identity.*

(d) Nonmerging *(the construction in the proof of [14, Section 2.3, Lemma 1]). Decorate a nonterminal by a composition μ of its fan-out; in a merged parent template, maximal adjacent runs of consecutive variables of a child determine that child's merge. The merge-decorated nonterminal generates the* image $L_{A^\mu} = \mu(L_A)$.

Stages (b)–(d) preserve the conditions established by earlier stages.

Proof. (a) For arbitrary child tuples, $\pi_I(\rho(\vec{v}_1, \dots, \vec{v}_q)) = \rho^I\big((\pi_{J_j}\vec{v}_j)_{J_j \neq \emptyset}\big)$, so induction on derivations gives $\subseteq$. For $\supseteq$, lift every retained child tuple by the induction hypothesis and supply a terminal derivation of each omitted child; this uses that G_0 is reduced, hence every omitted child is productive. Without prior trimming only $\pi_I(L_A) \subseteq L_{A^I}$ would be automatic. There are at most 2^f copies of each nonterminal, and I determines the J_j, so at most 2^f copies of each rule; omitting children cannot increase rank.

(b) Linearity and nondeletion make the induced child permutations well defined; induction in both directions identifies the language with σL_A. There are at most $f!$ copies, no child is duplicated, and total length, λ-freeness, fan-out and rank are preserved.

(c) Retained variables still occur once and in order, so nondeletion, nonpermutation, fan-out and rank are preserved. Induction on derivations gives the projection identity; omitted all-empty

children are restored by terminal derivations realizing their all-empty pattern. There are at most 2^f patterns per nonterminal and $2^{f r_{\max}}$ choices per rule, and projection only shortens templates.

(d) Nonpermutation makes the maximal runs compatible with each child's component order, and nondeletion makes them cover all its components. The merged template satisfies $\rho^{\mu}(\mu_1(\vec{v}_1), \ldots, \mu_q(\vec{v}_q)) = \mu(\rho(\vec{v}_1, \ldots, \vec{v}_q))$, and induction gives the image equality. No inverse-image claim is made, because the merge map need not be injective. Concatenation preserves total length, λ-freeness, nonpermutation, fan-out and rank; for fixed dimension and rank the conversion has linear size [14, Section 2.3, after Example 2]. □

Proof of the size bound in Proposition 2.6. Apply (a)–(d) to the trimmed input and trim again. Each stage multiplies the encoding by a factor depending only on $f, r_{\max}$ and trimming can only decrease it, so $\|G_{\text{good}}\| = O_{f,r_{\max}}(\|G_{\text{in}}\|)$, and $L(G_{\text{good}}) = L(G_{\text{in}}) \setminus \{\varepsilon\}$ by (c).

Thickness transfer for λ-free presentations

Proposition A.2 (Thickness under good-form conversion)**.** *Fix f and $r_{\max}$. Let G_{in} be a reduced λ-free f-MCFG of maximum rule rank at most $r_{\max}$, possibly with deleting rules, and suppose $L(G_{\text{in}}) \cap \Sigma^+ \neq \emptyset$. The good normalization of Proposition 2.6 yields an equivalent reduced good grammar G with the same fan-out and rule-rank bounds such that*

$$\|G\| = O_{f,r_{\max}}(\|G_{\text{in}}\|), \qquad t_G \leq t_{G_{\text{in}}}, \qquad \theta_G = O_{f,r_{\max}}(\|G_{\text{in}}\| \max\{1, t_{G_{\text{in}}}\}).$$

Proof. Because G_{in} is λ-free, every component of every derived tuple is nonempty, by induction on derivations; this persists through (a) and (b), so stage (c) is the identity. By Lemma A.1, every nonterminal X of the final grammar is associated with an original nonterminal A such that

$$L_X(G) = \mu(\sigma\pi_I(L_A(G_{\text{in}}))), \qquad \min_{\vec{u} \in L_X(G)} \|\vec{u}\|_1 \leq \min_{\vec{w} \in L_A(G_{\text{in}})} \|\vec{w}\|_1.$$

Trimming leaves the tuple languages of surviving nonterminals unchanged, so $t_G \leq t_{G_{\text{in}}}$, and the size bound was shown above. Finally Lemma 4.1(iii) gives $\theta_G \leq 2\|G\| \max\{1, t_G\} = O_{f,r_{\max}}(\|G_{\text{in}}\| \max\{1, t_{G_{\text{in}}}\})$. □

Thus deleting rules do not by themselves obstruct polynomial thickness transfer from a reduced λ-free presentation; genuine empty components are the problematic step.

Proposition A.3 (Empty-component elimination can blow up thickness)**.** *For every $n \geq 1$ there is a reduced fan-out-two, rule-rank-two MCFG E_n of size $O(n)$ with $t_{E_n} \leq 2$ and $\theta_{E_n} = O(n)$ such that the empty-component elimination stage of Proposition 2.6 contains a retained nonterminal and a retained rule of thickness at least 2^n.*

Proof. Let E_n have fan-out-one S, fan-out-two $A_0, \ldots, A_n$, and

$$A_0 \to (\mathtt{a}, \varepsilon), \qquad A_i \to (\varepsilon, \mathtt{c}) \mid (\xi_1^1 \xi_2^1, \xi_1^2 \xi_2^2)(A_{i-1}, A_{i-1}) \quad (1 \leq i \leq n), \qquad S \to (\xi_1^1 \mathtt{d}\, \xi_1^2)(A_n).$$

It is reduced, nondeleting, nonpermuting, and nonmerging. Short side derivations give $t_{E_n} \leq 2$, and placing any chosen rule on one branch of the binary spine while completing siblings minimally gives $\theta_{E_n} = O(n)$. If an A_i-tuple has nonempty first and empty second component, the rule $(\varepsilon, \mathtt{c})$ cannot occur in it; inductively the shortest such tuple is $(\mathtt{a}^{2^i}, \varepsilon)$. Thus pattern $\{1\}$ for A_n survives, and the projected start rule $S \to \xi_1^1 \mathtt{d}\, (A_n^{\{1\}})$ succeeds only with a child of length at least 2^n. □

This is the obstruction of Proposition 4.5: finite state refinement can force exponentially long shortest witnesses inside one refined state.

B Simultaneous block replacement in finite groups

Semigroup convention. The finite group A and the morphisms into it are auxiliary to the shortening proofs. The semantic typing is always the positive-image restriction of the window–group (or local-group) product, with the fresh identity used only to evaluate empty template factors; by Lemma 3.7 this changes neither nonempty tuple comparisons nor the selected witnesses.

This appendix proves the purely group-theoretic statement used in Theorem 4.11 and Lemma C.6. The commutative case is a zero-sum argument; the general case needs an order-homogenization argument because, for $f \geq 2$, blocks of different indices can interleave differently on different lines.

Definition B.1 (Block-replacement family)**.** Let A be a finite group and $f \geq 1$. A *block-replacement family* of width f over A with index set $[m]$ consists of $d \leq f$ *lines*; on line c, a left-to-right sequence of blocks $\beta_{c,1}, \dots, \beta_{c,q_c}$ and constants $a_{c,0}, \dots, a_{c,q_c} \in A$; for every block β an *owner* $o(\beta) \in [m]$ and two values $x_\beta, x'_\beta \in A$; every index owns at most f blocks. For $J \subseteq [m]$ put

$$\mathrm{val}_c(J) := a_{c,0}\, y_{c,1}\, a_{c,1} \cdots y_{c,q_c}\, a_{c,q_c}, \qquad y_{c,j} := \begin{cases} x'_{\beta_{c,j}}, & o(\beta_{c,j}) \in J, \\ x_{\beta_{c,j}}, & o(\beta_{c,j}) \notin J. \end{cases}$$

Lemma B.2 (Block replacement)**.** *In a block-replacement family of width f over A there is a nonempty $J \subseteq [m]$ with $\mathrm{val}_c(J) = \mathrm{val}_c(\emptyset)$ for every line c, provided that either* (a) *A is abelian and $m > \mathsf{d}(A^f)$, or* (b) *A is arbitrary and $m \geq M_{A,f}$, where d and $M_{A,f}$ are as in Section 4.4.*

Part (a) is proved in Appendix B.1 and part (b) in Appendix B.2.

B.1 Commutative groups: zero sums

Proof of Lemma B.2(a). Write A additively. By commutativity $\mathrm{val}_c(J) - \mathrm{val}_c(\emptyset) = \sum_{i \in J} \delta_{i,c}$, where $\delta_{i,c}$ is the sum of $x'_\beta - x_\beta$ over the blocks β of line c owned by i. Put $\delta_i = (\delta_{i,c})_c \in A^d$. A sequence over A^d of length $m > \mathsf{d}(A^f) \geq \mathsf{d}(A^d)$ has a nonempty zero-sum subsequence $(\delta_i)_{i \in J}$; the inequality holds because A^d embeds in A^f and a zero-sum-free sequence in a subgroup is zero-sum free in the group. □

Remark B.3 (Size of the zero-sum constant)*.* The prefix-sum pigeonhole argument gives $\mathsf{d}(A^f) \leq |A|^f - 1$; the zero-sum bound is usually much stronger. Write $A \cong C_{n_1} \oplus \cdots \oplus C_{n_s}$ and $n = \exp(A) \geq 2$. Then

$$f \sum_{i=1}^{s} (n_i - 1) \;\leq\; \mathsf{d}(A^f) \;\leq\; (n-1) + n \log \frac{|A|^f}{n},$$

where the lower bound is attained by a basis sequence and the upper bound is due to van Emde Boas and Kruyswijk [34, Proposition 5.1.8 and Theorem 5.5.5]; for a p-group the lower bound is an equality [34, Theorem 5.5.9] (Olson). Hence $\mathsf{d}(A^f) = O(\exp(A)\, f \log |A|)$ is linear in f, whereas $|A|^f - 1$ is exponential; for $A = C_2$ the antichain bound drops from $2^f - 1$ to f. The value $\mathsf{d}(A^d)$ is also the limit of the argument as stated, since a zero-sum-free sequence of that length exists; we do not claim that a grammar realizes such a configuration.

B.2 Arbitrary finite groups: order homogenization

Lemma B.4 (A homogeneous block-order lemma)**.** *Fix positive integers e, n with $n \geq 3$. For each $t \in [e]$ and $i \in [n]$, let (t, i) be a distinct point on a line, with a total left-to-right order. Suppose that (i) the order of (t, i) and (u, i) is independent of i; and (ii) for each ordered pair (t, u), including*

$t = u$, the order of (t, i) and (u, j) for $i < j$ is independent of the choice of $i < j$. Then the types at the points have a partition into classes, each occupying a consecutive region of the line. Every class has one direction (increasing or decreasing in i); inside its region the points occur in n successive chunks*, one per rank in that direction. Each chunk consists of exactly one point of each type in the class (in a fixed within-chunk order).*

Proof. By (ii) with $t = u$, each type's points are monotone. For distinct types t, u write X for $(t, i) < (u, j)$ and Y for $(u, i) < (t, j)$, $i < j$; both are uniform by (ii). Suppose first that both types increase. $X \wedge Y$ puts all points of rank i before all points of rank $j > i$. Under $X \wedge \neg Y$ all off-diagonal comparisons put t left of u, and so do equal ranks: $(u, i) < (t, i)$ would give $(t, 1) < (t, 2) < (u, 1) < (t, 1)$; thus the types are separated. The case $\neg X \wedge Y$ is symmetric, and $\neg X \wedge \neg Y$ gives $(t, 1) < (t, 2) < (u, 1) < (u, 2) < (t, 1)$, a contradiction. If t increases and u decreases, $X \wedge Y$ gives $(t, 2) < (u, 3) < (u, 1) < (t, 2)$ and $\neg X \wedge \neg Y$ gives $(t, 3) < (u, 2) < (t, 1) < (t, 3)$; under $X \wedge \neg Y$ all off-diagonal comparisons put t left of u, and a contrary equal-rank comparison would give $(t, 1) < (u, 2) < (u, 1) < (t, 1)$, contrary to (i), so the types are separated; $\neg X \wedge Y$ is symmetric. Reversing indices covers two decreasing types. Thus any two types are either *separated*, or have the same direction and are *interleaved by rank*.

Interleaving is transitive: if t interleaves with u and u with v, all three have one direction, say increasing; if t, v were separated with t to the left we would have $(v, 1) < (u, 2) < (t, 3) < (v, 1)$, and with v to the left $(t, 1) < (u, 2) < (v, 3) < (t, 1)$. Hence interleaving, with equality, partitions the types into classes. A type outside a class is separated from each type inside it and lies on the same side of all of them, since otherwise two interleaved types would be separated. So class regions are consecutive, the rank comparisons give successive chunks, and the within-chunk order is independent of rank by (i). □

Lemma B.5 (Simultaneous finite-group interval shortening)**.** *Let A be any finite group, and consider a block-replacement family of width f over A with $n \geq \max\{3, |A|^f\}$ indices in which all indices have the same number of blocks, the same lines for them and the same within-index order, and in which every cross-index comparison between the t-th block of index i and the u-th block of index j on a common line has the same truth value for all $i < j$. Then a nonempty interval J of indices satisfies $\mathrm{val}_c(J) = \mathrm{val}_c(\emptyset)$ for every line c.*

Proof. Apply Lemma B.4 on each line to the block types going to that line. Across all lines there are $q \leq f$ class regions. For such a region K, enumerate its chunks $1, \ldots, n$ from left to right, including the constants between the blocks of the region. Let P_p be the original value of its prefix ending after chunk p and Q_p the value of that prefix when chunks $1, \ldots, p$ are replaced; set $P_0 = Q_0 = 1_A$ and $\kappa_p^K = P_p Q_p^{-1}$. Replacing exactly chunks $p+1, \ldots, p'$ gives the ordered region value $P_p Q_p^{-1} Q_{p'} P_{p'}^{-1} P_n$, because $Q_p^{-1} Q_{p'}$ is the replaced value of the segment between the two prefix boundaries and $P_{p'}^{-1} P_n$ is the unchanged suffix. Hence the region keeps its value P_n whenever $\kappa_p^K = \kappa_{p'}^K$; inverses are used, but factors are never commuted.

If K increases with the index order, the chunks of indices $(a, b]$ occupy $(a, b]$; if it decreases, they occupy $(n - b, n - a]$. For $j = 0, \ldots, n$ define a state in A^q by taking κ_j^K from each increasing and κ_{n-j}^K from each decreasing class. There are $n + 1 > |A|^q$ states, so two coincide, at $a < b$. Replacing the indices $a + 1, \ldots, b$ then preserves every region value; everything else on every line is unchanged. □

Independent Ramsey proof of the qualitative assertion in Lemma B.2(b). We first establish existence of a simultaneous replacement set without using the quantitative order-homogenization bound below. Let $N = \max\{3, |A|^f\}$, and let $R(N, c)$ denote a finite two-uniform Ramsey number for c

colors. An owner with no blocks can be selected alone. Otherwise, among sufficiently many owners, select $R(N, 2^{f^2})$ of the same unary pattern (number of blocks, assigned lines, and within-owner order); there are at most U_f patterns. Absorb all other owners into the unchanged constants.

The common pattern has $e \le f$ block types. Color a pair of selected owners $i < j$ by the truth values of $p_t(i) < p_u(j)$ for all ordered pairs of types (t, u) on the same line. There are at most $2^{e^2} \le 2^{f^2}$ colors. By finite Ramsey, there is an ordered set of N owners on which every comparison has a fixed truth value, independently of $i < j$. Within-owner comparisons are already fixed by the common pattern. On each line, the argument of Lemma B.4 decomposes this homogeneous configuration into consecutive rank-ordered regions, all increasing or decreasing. Lemma B.5 now supplies a nonempty interval of owners whose simultaneous replacement leaves the ordered product on every line unchanged.

This proves the *existence* assertion for an arbitrary finite, possibly nonabelian group, without assuming any commutation of factors. Ramsey's bound does not supply the asserted polynomial dependence on $|A|$ for fixed f; that stronger, explicit assertion is established next by the monotone-comparison extraction. □

Lemma B.6 (Polynomial order homogenization)**.** *Let $N \ge 3$, and define $m_0 = N$ and $m_{s+1} = (m_s - 1)^2 + 1$, so that $m_k \le N^{2^k}$. Let T be a finite set of point types, each assigned to one of several lines. For every index $i \in [n]$ and $t \in T$ let $p_t(i)$ be a point on the line of t, the points on each line being distinct. Let Π be the set of ordered pairs (t, u), including $t = u$, of types on a common line, and put $k = |\Pi|$. If $n \ge m_k$, there is a set $I \subseteq [n]$ of N indices such that, for every $(t, u) \in \Pi$, the truth value of $p_t(i) < p_u(j)$ is the same for all $i < j$ in I.*

Proof. Treat the pairs one at a time, the pairs (t, t) first. Each step passes from a current index set of size at least $(m-1)^2 + 1$ to a subset of size at least m; homogeneity already obtained, and monotonicity of each type, persist in subsets. For (t, t), homogeneity means that p_t is monotone, and the Erdős–Szekeres theorem gives a monotone subsequence of length m.

For $t \ne u$, both p_t and p_u are already monotone. Relabel the current indices $1, \dots, m'$ in order, with $m' \ge (m-1)^2 + 1$, and put $X(i, j) = [p_t(i) < p_u(j)]$ for $i < j$. Each row $X(i, \cdot)$ and each column $X(\cdot, j)$ is monotone; replacing X by $1 - X$ if necessary, the rows go from 0 to 1.

If the columns go from 1 to 0, put $J(i) = \min\{j > i : X(i, j) = 1\}$ (or ∞); J is nondecreasing. Set $i_1 = 1$ and $i_{s+1} = J(i_s)$ while this is at most m'. For $a < b$, $i_b \ge J(i_a)$ and so $X(i_a, i_b) = 1$; if the chain has m terms we are done. Otherwise its at most $m - 1$ terms cut $[m']$ into at most $m - 1$ intervals $[i_s, i_{s+1})$ and $[i_L, m']$, one of which has more than $m'/(m-1) > m - 1$ indices; inside it every $i < j$ satisfies $j < J(i_s) \le J(i)$, so $X(i, j) = 0$ (and $J \equiv \infty$ in the last interval).

If the columns go from 0 to 1, the set $\{X = 1\}$ is closed upward in both coordinates. Let κ be the least index with $X(\kappa, \kappa + 1) = 1$, or $\kappa = m'$. For $\kappa \le i < j$, $X(i, j) \ge X(\kappa, \kappa + 1) = 1$; for $i < j \le \kappa$, $X(i, j) \le X(j - 1, j) = 0$. One of $[1, \kappa]$ and $[\kappa, m']$ has at least $(m' + 1)/2 \ge m$ indices, because $(m-1)^2 + 1 \ge 2m - 1$.

After all k pairs at least $m_0 = N$ indices remain, and $m_{s+1} \le m_s^2$ gives $m_k \le N^{2^k}$. □

Proof of Lemma B.2(b). If an index owns no blocks, choosing that index alone preserves every line, so assume each index owns at least one block. The *unary pattern* of an index records the number of its blocks, the line of each, and their order within any shared line; there are at most $f \cdot f^f \cdot f! = U_f$ patterns. Hence at least $N_{A,f}^{2^{f^2}}$ indices share one pattern; the blocks of all other indices are never replaced and are absorbed into the constants. Number the blocks of the common pattern as types $1, \dots, f'$ ($f' \le f$) and use their positions as points on their lines. There are at most f^2 ordered pairs of types on a common line, including $t = u$, and $m_k \le N^{2^k} \le N^{2^{f^2}}$ for $k \le f^2$, so Lemma B.6 with

$N = N_{A,f}$ yields $N_{A,f}$ indices with constant cross-index comparisons for all $i < j$. These satisfy the hypotheses of Lemma B.5, which supplies the required nonempty interval. □

For fixed f the constant $M_{A,f} = U_f \max\{3, |A|^f\}^{2^{f^2}}$ is polynomial in $|A|$; it replaces the multicolor Ramsey bound $U_f\, c_f^{\,c_f(N_{A,f}-1)+1}$, $c_f = 2^{f^2}$, of earlier versions, which is exponential in $|A|^f$. The dependence on f remains doubly exponential.

C Local-group typings and the semilattice boundary

Semigroup convention. Throughout this appendix **LG** is tested on the *positive image* $S = \widehat{s}(\Sigma^+)$ of the typing, not on the canonical completion $S^\sharp$. Auxiliary group-valued maps and the displayed monoid morphisms are implementations whose restrictions to nonempty words define the semantic bias (Lemma 3.7).

LG is the pseudovariety of finite semigroups S in which every local monoid eSe, with e idempotent, is a group; classically $\mathbf{LG} = \mathbf{G} * \mathbf{D}$ [38, 39]. It is exactly the class of finite semigroups containing no copy of the two-element semilattice $U_1 = \{1, \zeta\}$ ($\zeta^2 = \zeta$) [37, Proposition 2.3]; that semilattice is the positive image of the typing behind the exponential lower bound of [25]. Theorem C.4 adds to this *known* algebra the sliding-factor group typing used for the MCFG quantitative bound.

Throughout, Σ is nonempty, $k \geq 1$, and A is a finite group. For a word w with $|w| \geq k+1$ let $\sigma_k(w)$ be the sequence of its factors of length $k+1$, ordered by their starting positions.

Definition C.1 (Local-group typing)**.** Let $g : \Sigma^{k+1} \to A$ be any map. For $|w| \geq 2k$ put

$$\gamma_{k,g}(w) := g(F_1)g(F_2)\cdots g(F_{|w|-k}), \qquad \sigma_k(w) = (F_1, \ldots, F_{|w|-k}),$$

the ordered product in A, and put $\gamma_{k,g}(w) := 1_A$ for $|w| < 2k$. The *local-group typing* is $\lambda_{k,g}(w) := \big(\widehat{s}_{k,k}(w), \gamma_{k,g}(w)\big)$. A finite typing $\widehat{s}$ is *local-group-determined* if $\widehat{s} \preceq \lambda_{k,g}$ for some k, A and g.

For $\ell = k$, words of length $< 2k$ are *short* and recorded exactly, while a *long* w is recorded by $\langle \mathrm{pref}_k(w), \mathrm{suf}_k(w)\rangle$ (Definition 2.4). For $s, p \in \Sigma^k$ write $J_g(s,p) := \prod_{i=1}^{k} g((sp)[i..i+k])$, the ordered product over the k factors of length $k+1$ of sp that start in s.

Lemma C.2 (Local-group typings are morphisms)**.** *The kernel of $\lambda_{k,g}$ is a congruence of finite index on Σ^*. Hence $\lambda_{k,g}$ is a homomorphism onto a finite monoid $M_{k,g}$. For long u, v,*

$$\lambda_{k,g}(uv) = \big(\langle \mathrm{pref}_k\, u, \mathrm{suf}_k\, v\rangle,\ \gamma_{k,g}(u)\, J_g(\mathrm{suf}_k\, u, \mathrm{pref}_k\, v)\, \gamma_{k,g}(v)\big).$$

Proof. The window component is a morphism by Definition 2.4. For long u, v, the factors of length $k+1$ of uv are, in order, those of u, the k factors starting in the last k positions of u (inside $\mathrm{suf}_k(u)\,\mathrm{pref}_k(v)$), and those of v. If u is short and v long, the factors starting in u lie inside $u\,\mathrm{pref}_k(v)$ and the rest are those of v, so $\gamma_{k,g}(uv)$ is determined by u, $\mathrm{pref}_k(v)$ and $\gamma_{k,g}(v)$; the case of long u and short v is symmetric, and if both are short uv is determined exactly. Thus $\lambda_{k,g}(uv)$ is a function of $\lambda_{k,g}(u)$ and $\lambda_{k,g}(v)$, and the image has at most $|\Sigma|^{<2k} + |\Sigma|^{2k}|A|$ elements. □

Lemma C.3 (Block decomposition)**.** *Let $w = x_0\beta_1 x_1 \cdots \beta_q x_q$ be a factorization of a word into designated blocks β_j and fixed words x_j. Let $w' = x_0\beta'_1 x_1 \cdots \beta'_q x_q$, where $\widehat{s}_{k,k}(\beta'_j) = \widehat{s}_{k,k}(\beta_j)$ for every j. Then $\widehat{s}_{k,k}(w') = \widehat{s}_{k,k}(w)$. If w is long, there are $c_0, \ldots, c_q \in A$, depending only on the x_j and the window types of the β_j, such that*

$$\gamma_{k,g}(w) = c_0\, \gamma_{k,g}(\beta_1)\, c_1 \cdots \gamma_{k,g}(\beta_q)\, c_q, \qquad \gamma_{k,g}(w') = c_0\, \gamma_{k,g}(\beta'_1)\, c_1 \cdots \gamma_{k,g}(\beta'_q)\, c_q.$$

Proof. A short block is recorded exactly, so $\beta'_j = \beta_j$, and we absorb it into the fixed words. Every remaining block is long and its replacement is long with the same first and last k letters; hence w' is long exactly when w is, with the same first and last k letters, which gives the window claim. A factor F of length $k+1$ of w either lies inside a long block β_j or meets each long block it touches only in a prefix or a suffix of length at most k (meeting both ends would force $|\beta_j| \le k-1$). Factors of the second kind are spelled by the fixed words and the windows of the blocks, and occur in the same order in w'. Ordering all factors by starting position, those inside β_j form one consecutive run contributing $\gamma_{k,g}(\beta_j)$; let c_j be the ordered product of the remaining factors between consecutive runs. □

Theorem C.4 (Characterizations of local-group positive images)**.** *For a finite typing $s : \Sigma^+ \to \mathsf{S}$ put $S := \widehat{s}(\Sigma^+)$. The following are equivalent.*

(i) *$S \in \mathbf{LG}$; that is, eSe is a group for every idempotent e of S.*

(ii) *S has no subsemigroup isomorphic to U_1.*

(iii) *Every idempotent of S lies in the minimal ideal of S.*

(iv) *$\widehat{s}$ is local-group-determined. One may take $k = |S|$.*

Proof. We include elementary arguments for the standard equivalence of (i)–(iii) (cf. [37, Proposition 2.3]), independently of the new sliding-factor statement (iv). Let K be the minimal ideal of S.

(i)⇒(ii). If $\{e, z\} \cong U_1$ with $ez = ze = z \neq e$, then z is a nonidentity idempotent in the group eSe, which is impossible.

(ii)⇒(iii). For an idempotent $e \in S$ choose $a \in K$ and put $z = (eae)^\omega$. Then z is an idempotent in $K \cap eSe$ with $ez = ze = z$. If $z \neq e$, the pair $\{e, z\}$ is a copy of U_1, a contradiction. Hence $e = z \in K$.

There is also a direct proof of (i)⇒(iii) that does not invoke the two-element semilattice: for any idempotent e and $a \in K$, the element eae lies in $K \cap eSe$; since eSe is a group with identity e, it has an inverse $y \in eSe$. Consequently $e = (eae)y \in K$, because K is an ideal.

(iii)⇒(i). The minimal ideal of a finite semigroup is completely simple [40]. For an idempotent $e \in K$, the ideal property gives $eSe = eKe$; this is the maximal subgroup (the $\mathcal{H}$-class) containing e, so it is a group.

(iii)⇒(iv). By the Rees theorem [40] write $K = \mathcal{M}(H; I, \Lambda; P)$, with product $(i, a, \lambda)(j, b, \mu) = (i, ap_{\lambda j}b, \mu)$, and put $n := |S|$. *Long products lie in K.* The Rees quotient S/K has no idempotent other than 0. For $s_1, \dots, s_m \in S$ with $m := |S/K| \le n$, if no prefix product in S/K were 0, two would coincide, $p_i = p_j$ with $i < j$, and then $p_i = p_i x^\omega = 0$ for $x = s_{i+1} \cdots s_j$. Hence $\widehat{s}(w) \in K$ whenever $|w| \ge n$. *Right action.* Fix $i_0 \in I$ and for $\lambda \in \Lambda$ let $e_\lambda := (i_0, p^{-1}_{\lambda i_0}, \lambda)$, an idempotent with $xe_\lambda = x$ for every $x = (i, a, \lambda) \in K$. Writing $e_\lambda s = (j, c, \mu) \in K$ for $s \in S$ gives

$$(i, a, \lambda)\, s = (i,\ a\,\rho(\lambda, s),\ \mu(\lambda, s)), \qquad \rho(\lambda, s) := p_{\lambda j} c,$$

with $\rho(\lambda, s) \in H$ and $\mu(\lambda, s) \in \Lambda$ depending only on (λ, s). Right multiplication by s preserves the I-coordinate, as is explicit in this formula. Dually, for $z = (i, a, \lambda) \in K$ choose an idempotent $f_i \in K$ with $f_i z = z$. Then $sz = (sf_i)z$ and $sf_i \in K$, so the Rees product shows that left multiplication preserves the Λ-coordinate. In particular, for $|w| \ge n$ the Λ-coordinate of $\widehat{s}(w)$ is that of $\widehat{s}(\mathrm{suf}_n\, w)$,

and the I-coordinate that of $\widehat{s}(\mathrm{pref}_n\, w)$. *Window factorization.* Let $w = a_1 \cdots a_m$ with $m \geq 2n$, and let λ_j be the Λ-coordinate of $\widehat{s}(a_{j-n+1} \cdots a_j)$ for $j \geq n$. Iterating the right action gives

$$\widehat{s}(w) = \Big(i,\ a \prod_{j=n}^{m-1} \rho(\lambda_j, \widehat{s}(a_{j+1})),\ \lambda_m\Big), \qquad \widehat{s}(a_1 \cdots a_n) = (i, a, \lambda_n),$$

and the j-th factor depends only on the factor $a_{j-n+1} \cdots a_{j+1}$ of length $n+1$. With $g(F) := \rho(\lambda(\widehat{s}(F[1..n])), \widehat{s}(F[n+1]))$ for $F \in \Sigma^{n+1}$, the product is $\gamma_{n,g}(w)$. So $\widehat{s}(w)$ is determined by $\mathrm{pref}_n\, w$, $\mathrm{suf}_n\, w$ and $\gamma_{n,g}(w)$; words of length $< 2n$ are recorded exactly by $\widehat{s}_{n,n}$, and ε is alone in its $\lambda_{n,g}$-class. Thus $\widehat{s} \preceq \lambda_{n,g}$ with $A = H$.

(iv)⇒(i). **LG** is closed under homomorphic images, so it suffices to show $T := \lambda_{k,g}(\Sigma^+) \in \mathbf{LG}$. A short element is never idempotent, since the square of a short word is longer. Let $e = (\langle p, s\rangle, \gamma_e)$ be idempotent. By Lemma C.2, $\gamma_e J_g(s,p) \gamma_e = \gamma_e$, and for every $x \in T$, exe has window $\langle p, s\rangle$. Hence eTe lies in $H_e := \{(\langle p, s\rangle, \alpha) : \alpha \in A\}$, whose product $\alpha * \beta = \alpha J_g(s,p) \beta$ makes $\alpha \mapsto \alpha J_g(s,p)$ an isomorphism onto A. Thus eTe is a finite subsemigroup of a group, hence a group. □

Proposition C.5 (Window–group typings are local-group typings, strictly)**.** (i) *If $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$ for a homomorphism g into a finite group, then $\widehat{s} \preceq \lambda_{K,g'}$ with $K = \max\{k, \ell, 1\}$ and $g'(F) := g(\text{last letter of } F)$. Hence $\mathbf{LI} \vee \mathbf{G} \subseteq \mathbf{LG}$.*

(ii) *Let $\Sigma = \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}$. Let $h_{\mathtt{ab}} := \lambda_{1,g}$ with $g(\mathtt{ab}) = 1 \in \mathbb{Z}_2$ and $g(F) = 0$ otherwise; it records the parity of the number of occurrences of $\mathtt{ab}$, together with the window. Then $h_{\mathtt{ab}}$ is not window–group-determined. Hence $h_{\mathtt{ab}}(\Sigma^+) \in \mathbf{LG} \setminus (\mathbf{LI} \vee \mathbf{G})$.*

Proof. (i) For $|w| \geq 2K$, $g(w) = g(\mathrm{pref}_K\, w)\, \gamma_{K,g'}(w)$, and shorter words are recorded exactly. For the inclusion, let $S \in \mathbf{LI} \vee \mathbf{G}$ and take $\Sigma = S$ with the evaluation typing $s : \Sigma^+ \twoheadrightarrow S$. Lemma 4.16 makes $\widehat{s}$ window–group determined, hence local-group-determined by the first part, so $S \in \mathbf{LG}$ by Theorem C.4.

(ii) Suppose $h_{\mathtt{ab}} \preceq (\widehat{s}_{k,\ell}, g')$ with $g' : \Sigma^* \to A'$. Let $K = \max\{k, \ell, 1\}$ and let q be an even multiple of the exponent of A'. Put $u = \mathtt{a}^K(\mathtt{abc})^q \mathtt{a}^K$ and $v = \mathtt{a}^K(\mathtt{bac})^q \mathtt{a}^K$. Both are long, with prefix $\mathtt{a}^k$ and suffix $\mathtt{a}^\ell$, and $g'(u) = g'(\mathtt{a})^{2K} = g'(v)$. But u contains q occurrences of $\mathtt{ab}$, while v contains exactly one, at $\mathtt{a}^K\mathtt{b}$, so $h_{\mathtt{ab}}(u) \neq h_{\mathtt{ab}}(v)$. Membership in **LG** follows from Theorem C.4; if the positive image were in $\mathbf{LI} \vee \mathbf{G}$, Lemma 4.16 would make $h_{\mathtt{ab}}$ window–group-determined. □

Lemma C.6 (Short tuples with prescribed local-group type)**.** *Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a reduced good MCFG of fan-out at most f and rule rank at most $r_{\max} \geq 1$, and let $\lambda = \lambda_{k,g}$. Let $b_{\mathrm{w}} := f(2k+1) n_{\mathrm{w}} (r_{\max}+1) \bar{\theta}_G$, where $n_{\mathrm{w}} \leq |V| |M_{k,k}|^f$ is the number of non-start typed nonterminals of the trimmed $\widehat{s}_{k,k}$-refinement. Put $n_\lambda := |V| |M_{k,g}|^f$, and let $M_{A,f}$ be the constant of Section 4.4. For every $B \in V$ and every realizable $\mathbf{t} \in M_{k,g}^{\mathrm{fan}(B)}$ there is $\vec{u} \in L_B(G)$ with $\lambda^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{t}$ and*

$$\|\vec{u}\|_1 \leq \beta_\lambda := (r_{\max}+1)\, b_{\mathrm{w}} \max\{1, (M_{A,f} - 1)\, n_\lambda\}.$$

If A is abelian, $M_{A,f} - 1$ may be replaced by $\mathtt{d}(A^f)$.

Proof. Let T be minimal as in Lemma 4.9 with root λ-vector $\mathbf{t}$ and threshold $b = b_{\mathrm{w}} \geq \bar{\theta}_G$. For a heavy node v labelled C, with yield $\vec{x}$ and window vector $\boldsymbol{\omega} = \widehat{s}_{k,k}(\vec{x})$, the typed nonterminal $C_{\boldsymbol{\omega}}$ is retained in the trimmed $\widehat{s}_{k,k}$-refinement, because T can be inserted into a successful derivation (G is reduced). Lemma 4.7 (trivial κ) gives a replacement tuple of C with the same window vector and length at most b_{w}. For pairwise incomparable heavy nodes $v_1, \ldots, v_m$ and $J \subseteq [m]$, let T_J replace the subtrees at v_i, $i \in J$. The components of the v_i are disjoint contiguous blocks of the root components.

By Lemma C.3, every T_J has the root window vector of T; a short root component is unchanged, and on each long root component $\gamma_{k,g}$ is an ordered product $c_0 y_1 c_1 \cdots y_q c_q$ of block values between constants that do not depend on J. The long root components thus form a block-replacement family over A (Definition B.1). If $m \geq M_{A,f}$, or $m > \mathtt{d}(A^f)$ for abelian A, Lemma B.2 gives a nonempty J with the root λ-vector of T, while T_J is strictly shorter. This contradicts minimality, and Lemma 4.9 with $N = M_{k,g}$ gives the bound. □

Detailed proof of Theorem 4.18. The theorem is stated in Section 4.5; we verify here the bound using the sliding-factor characterization and simultaneous ordered group-block replacement.

Proof of Theorem 4.18. By Theorem C.4, $\widehat{s} \preceq \lambda_{k,g}$ for some k, A, g depending only on $\widehat{s}$. For a non-start typed nonterminal $X = B_{\mathbf{s}}$ and $\vec{u} \in L_X$, Lemma C.6 gives $\vec{u}' \in L_B(G)$ with the same λ-vector and $\|\vec{u}'\|_1 \leq \beta_\lambda$, and Lemma 4.3 gives $\vec{u}' \in L_X$. The role of the two shortening stages is distinct: the window stage provides replacements of length b_{w} without prescribing their group values, while the block-replacement lemma selects a nonempty family of such replacements whose *combined* ordered group products are unchanged. No per-subtree group-value preservation is assumed. Since $b_{\mathrm{w}} = O(\|G\|\bar{\theta}_G)$ and $n_\lambda = O(\|G\|)$, the heavy-node bound gives $\beta_\lambda = O(\|G\|^2\bar{\theta}_G)$ for fixed $f, r_{\max}, s, \Sigma$. There are $O(\|G\|)$ typed rules, and Lemma 4.2 then yields the displayed cubic bound on the *total* encoded positive-data length. □

By Proposition C.5(i), this theorem also applies to window and window–group typings, but its general cubic bound is weaker than the quadratic bounds proved for those subclasses in Theorems 4.8 and 4.14.

Corollary C.7 (Pseudovariety dichotomy)**.** *Let* $\mathbf{V}$ *be a pseudovariety of finite semigroups.*

(i) *If* $\mathbf{V} \subseteq \mathbf{LG}$*, then for every alphabet, every typing* $\widehat{s}$ *with* $\widehat{s}(\Sigma^+) \in \mathbf{V}$*, and all fixed* $f, r_{\max}$*, characteristic data of total length* $O(\|G\|^3\bar{\theta}_G)$ *exist as in Theorem 4.18.*

(ii) *If* $\mathbf{V} \not\subseteq \mathbf{LG}$*, then* $U_1 \in \mathbf{V}$*. At fan-out one there is a typing* h_{U} *with* $h_{\mathrm{U}}(\Sigma^+) = U_1$ *and targets with* $\|G\|, \theta_G = O(n)$*, every characteristic sample of which for* $\mathcal{B}^{\leq 2}_{1,h_{\mathrm{U}}}$ *contains a word of length* $3 \cdot 2^n - 2$*. The same holds for every set-driven operator admitting characteristic samples throughout* $\mathcal{C}^{\mathrm{mcf}}_{1,h_{\mathrm{U}}}$*.*

Proof. (i) is Theorem 4.18. For (ii), choose $S \in \mathbf{V} \setminus \mathbf{LG}$; by Theorem C.4 it has a subsemigroup isomorphic to U_1, and pseudovarieties are closed under subsemigroups. Take the doubling family L_n, the typing h_{U} and the long type-1 word w^*_n of [25, Proposition 6.4], where L_n is shown to be Clark–Eyraud substitutable, hence $(1, h_{\mathrm{U}})$-tuple-substitutable. The type filter $h_{\mathrm{U}}^{-1}(\zeta)$ is $(1, h_{\mathrm{U}})$-tuple-substitutable by the remark preceding Proposition 2.10, so by Proposition 2.10(i) so is $L^-_n := L_n \cap h_{\mathrm{U}}^{-1}(\zeta) = L_n \setminus \{w^*_n\}$. If $K \subseteq L_n$ omits w^*_n, then $K \subseteq L^-_n$ and Proposition 3.4 gives $L(\mathcal{B}^{\leq 2}_{1,h_{\mathrm{U}}}(K)) \subseteq L^-_n$; thus every characteristic sample contains w^*_n. The set-driven statement is the distinguishing-pair argument of [25, Proposition 6.4(iv)] for the pair L_n, L^-_n. □

Remark C.8 (Scope of the dichotomy). (i) Part (i) holds at every fixed fan-out; part (ii) is established only at fan-out one. The doubling family does not lift directly: L_2 is not $(2, h_{\mathrm{U}})$-tuple-substitutable, since the type-1 pairs $([, \mathtt{a}])$ and $([\mathtt{a},])$ share the context $[[\mathtt{aa}]\,\square_1\,\mathtt{a}\,\square_2\,]$ but have different distributions (checked exhaustively for $n \leq 3$; under the finer typing recording each component's set of letters, still with semilattice image, the failure appears at $n = 3$). Whether part (ii) holds at fan-out $f \geq 2$ is the pseudovariety form of Open Problem 6.4.

(ii) The dichotomy concerns pseudovarieties. For an individual typing $\widehat{s}$ with $\widehat{s}(\Sigma^+) \notin \mathbf{LG}$, the class $\mathcal{C}^{\mathrm{mcf}}_{1,s}$ need not contain a coded copy of the doubling family, and whether every such typing forces superpolynomial characteristic data remains open.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During preparation of this work, the author used ChatGPT (OpenAI) and Claude (Anthropic) to assist with organization of the manuscript, expression and translation of text, cross-checks of references and notation, and editing of LaTeX source. After using these tools, the author reviewed and edited the material as needed and takes full responsibility for the content of the publication.

References

[1] E. M. Gold. Language identification in the limit. *Information and Control*, 10(5):447–474, 1967.

[2] D. Angluin. Inductive inference of formal languages from positive data. *Information and Control*, 45(2):117–135, 1980.

[3] A. Clark and R. Eyraud. Polynomial identification in the limit of substitutable context-free languages. *Journal of Machine Learning Research*, 8:1725–1745, 2007.

[4] C. de la Higuera. Characteristic sets for polynomial grammatical inference. *Machine Learning*, 27:125–138, 1997.

[5] F. Coste, D. Fredouille, C. Kermorvant, and C. de la Higuera. Introducing domain and typing bias in automata inference. In *Grammatical Inference: Algorithms and Applications – ICGI 2004*, LNCS 3264, pp. 115–126, Springer, 2004.

[6] Y. Takada. Grammatical inference for even linear languages based on control sets. *Information Processing Letters*, 28(4):193–199, 1988. doi:10.1016/0020-0190(88)90208-6.

[7] Y. Takada. Learning formal languages based on control sets. In K. P. Jantke and S. Lange (eds.), *Algorithmic Learning for Knowledge-Based Systems*, LNCS 961, pp. 316–339, Springer, 1995.

[8] A. Clark. Learning trees from strings: A strong learning algorithm for some context-free grammars. *Journal of Machine Learning Research*, 14:3537–3559, 2013.

[9] F. Coste, G. Garet, and J. Nicolas. Local substitutability for sequence generalization. In *Proceedings of the 11th International Conference on Grammatical Inference*, JMLR Workshop and Conference Proceedings 21, pp. 97–111, 2012.

[10] R. Eyraud, J. Heinz, and R. Yoshinaka. Efficiency in the identification in the limit learning paradigm. In J. Heinz and J. M. Sempere (eds.), *Topics in Grammatical Inference*, pp. 25–46. Springer, 2016. doi:10.1007/978-3-662-48395-4_2.

[11] F. Coste and J. Nicolas. Learning local substitutable context-free languages from positive examples in polynomial time and data by reduction. In *Proceedings of the 14th International*

Conference on Grammatical Inference, Proceedings of Machine Learning Research 93, pp. 155–168, 2019.

[12] R. Yoshinaka. Identification in the limit of k, l-substitutable context-free languages. In *Grammatical Inference: Algorithms and Applications*, LNCS 5278, pp. 266–279. Springer, 2008.

[13] R. Yoshinaka. Learning mildly context-sensitive languages with multidimensional substitutability from positive data. In *Algorithmic Learning Theory*, LNCS 5809, pp. 278–292. Springer, 2009. doi:10.1007/978-3-642-04414-4_24.

[14] R. Yoshinaka. Efficient learning of multiple context-free languages with multidimensional substitutability from positive data. *Theoretical Computer Science*, 412(19):1821–1831, 2011. doi:10.1016/j.tcs.2010.12.058.

[15] R. Yoshinaka. General perspective on distributionally learnable classes. In *Proceedings of the 14th Meeting on the Mathematics of Language (MoL 2015)*, pp. 87–98. Association for Computational Linguistics, 2015. doi:10.3115/v1/W15-2308.

[16] R. Yoshinaka. Polynomial-time identification of multiple context-free languages from positive data and membership queries. In *Grammatical Inference: Theoretical Results and Applications*, LNCS 6339, pp. 230–244. Springer, 2010. doi:10.1007/978-3-642-15488-1_19.

[17] A. Clark and R. Yoshinaka. Distributional learning of context-free and multiple context-free grammars. In J. Heinz and J. M. Sempere (eds.), *Topics in Grammatical Inference*, pp. 143–172. Springer, 2016. doi:10.1007/978-3-662-48395-4_6.

[18] A. Clark and R. Yoshinaka. Distributional learning of parallel multiple context-free grammars. *Machine Learning*, 96(1–2):5–31, 2014. doi:10.1007/s10994-013-5403-2.

[19] T. Oates, T. Armstrong, L. Becerra-Bonache, and M. Atamas. Inferring grammars for mildly context-sensitive languages in polynomial-time. In *Grammatical Inference: Algorithms and Applications – ICGI 2006*, LNCS 4201, pp. 137–147. Springer, 2006.

[20] M. Kanazawa and R. Yoshinaka. Extending distributional learning from positive data and membership queries. In *Proceedings of the 16th International Conference on Grammatical Inference*, Proceedings of Machine Learning Research 217, pp. 8–22, 2023.

[21] P. D. Gallot. *Safety of transformations of data trees: Tree transducer theory applied to a verification problem on shell scripts*. PhD thesis, Université de Lille, 2021. `https://pepite-depot.univ-lille.fr/LIBRE/EDMADIS/2021/2021LILUB030.pdf`.

[22] A. Bishop, M. Elder, A. Evetts, P. Gallot, and A. Levine. On groups with EDT0L word problem. *International Journal of Algebra and Computation*, 36(4):425–486, 2026. doi:10.1142/S0218196726500165.

[23] A. Duncan, M. Elder, L. Frenkel, and M. Lyu. A substitution lemma for multiple context-free languages. *International Journal of Algebra and Computation*, 2026 (online publication). doi:10.1142/S0218196726500463.

[24] O. Rambow and G. Satta. Independent parallelism in finite copying parallel rewriting systems. *Theoretical Computer Science*, 223(1–2):87–120, 1999. doi:10.1016/S0304-3975(97)00190-4.

[25] T. Kuriyama. Distributional learning of context-free languages under fixed finite-semigroup typing. arXiv:1409.6247v5 [cs.FL], 2026. `https://arxiv.org/abs/1409.6247v5`. The companion revision submitted to *Theoretical Computer Science* as TCS-D-26-00494R1 is preserved in the frozen Round-1 manuscript (TeX source). For precise version-locked locations, the cited R1 semilattice lower bound is Proposition 6.4 (Appendix C gives the normalization details); the positive-typing convention is defined in the R1 manuscript. The public arXiv record and frozen R1 submission are separate version references.

[26] R. McNaughton and S. Papert. *Counter-Free Automata*. MIT Press, Cambridge, MA, 1971.

[27] J.-E. Pin. *Varieties of Formal Languages*. North Oxford Academic, London, and Plenum, New York, 1986.

[28] J. C. Costa. Some pseudovariety joins involving locally trivial semigroups. *Semigroup Forum*, 64(1):12–28, 2002. doi:10.1007/s002330010060.

[29] J. C. Costa. Some pseudovariety joins involving groups and locally trivial semigroups. In *Semigroups, Algorithms, Automata and Languages*, pp. 341–348. World Scientific, 2002.

[30] H. Seki, T. Matsumura, M. Fujii, and T. Kasami. On multiple context-free grammars. *Theoretical Computer Science*, 88(2):191–229, 1991. doi:10.1016/0304-3975(91)90374-B.

[31] L. Kallmeyer. *Parsing Beyond Context-Free Grammars*. Springer, Berlin, Heidelberg, 2010. doi:10.1007/978-3-642-14846-0.

[32] M. Kanazawa. Ogden's lemma, multiple context-free grammars, and the control language hierarchy. *Information and Computation*, 269:104449, 2019. doi:10.1016/j.ic.2019.104449.

[33] C. Gómez-Rodríguez, M. Kuhlmann, G. Satta, and D. Weir. Optimal reduction of rule length in linear context-free rewriting systems. In *Proceedings of Human Language Technologies: The 2009 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 539–547. Association for Computational Linguistics, 2009.

[34] A. Geroldinger and F. Halter-Koch. *Non-Unique Factorizations: Algebraic, Combinatorial and Analytic Theory*, volume 278 of *Pure and Applied Mathematics*. Chapman & Hall/CRC, Boca Raton, 2006.

[35] G. Ramalingam and T. Reps. An incremental algorithm for a generalization of the shortest-path problem. *Journal of Algorithms*, 21(2):267–305, 1996. doi:10.1006/jagm.1996.0046.

[36] M.-J. Nederhof. Weighted deductive parsing and Knuth's algorithm. *Computational Linguistics*, 29(1):135–143, 2003.

[37] J. Almeida, S. Margolis, B. Steinberg, and M. Volkov. Representation theory of finite semigroups, semigroup radicals and formal language theory. *Transactions of the American Mathematical Society*, 361(3):1429–1461, 2009. doi:10.1090/S0002-9947-08-04712-0.

[38] H. Straubing. Finite semigroup varieties of the form $\mathbf{V} * \mathbf{D}$. *Journal of Pure and Applied Algebra*, 36:53–94, 1985.

[39] B. Tilson. Categories as algebra: an essential ingredient in the theory of monoids. *Journal of Pure and Applied Algebra*, 48:83–198, 1987.

[40] J. M. Howie. *Fundamentals of Semigroup Theory*, volume 12 of *London Mathematical Society Monographs, New Series*. Clarendon Press, Oxford, 1995.

[41] C. H. Papadimitriou. On the complexity of integer programming. *Journal of the ACM*, 28(4):765–768, 1981. doi:10.1145/322276.322287.