Polynomial Upper Bounds on Characteristic-Data Length for Multiple Context-Free Languages under Fixed Finite-Semigroup Typing

Takayuki Kuriyama
Independent Researcher, Tokyo, Japan
growup.kuriyama@gmail.com

Abstract

We study positive-data learning of multiple context-free languages that are tuple-substitutable under a fixed finite-semigroup typing and bound the total encoded length of characteristic samples relative to their MCFG presentations. Qualitative identification follows Yoshinaka's distributional framework; our contribution is quantitative. With the alphabet, fan-out, rule rank, and typing fixed, a reduced good grammar G of rule thickness $\bar{\theta}_G$ admits characteristic data of length $O(\|G\|^2\bar{\theta}_G)$ for prefix–suffix-window typings combined with any finite, possibly noncommutative, group, and of length $O(\|G\|^3\bar{\theta}_G)$ when the positive typing image is a local group. The proofs shorten interleaved tuple components simultaneously while preserving ordered group products. A further refinement-thickness dichotomy characterizes local groups as precisely the finite-semigroup regime admitting polynomially bounded full yield-type refinement thickness at every fixed fan-out; this is distinct from a lower bound on characteristic-data length. General input MCFGs are covered by parameterized bounds using component-pattern thickness, which may be exponential in grammar size. A separate conservative learner gives Gold identification. We exhibit a class-wide finite-feature bias for cross-serial languages and explicit characteristic samples for that family. Whether the local-group condition is necessary for polynomial characteristic-data length at fan-out two or more remains open.

Keywords: grammatical inference; multiple context-free grammars; positive data; finite-semigroup typing; multidimensional substitutability; Gold identification; rule rank.

1 Introduction

Positive-data grammatical inference asks when a language can be identified in the limit from positive examples alone [1, 2]. Distributional learning answers this question by treating words, or tuples of words, that occur in common contexts as candidate constituents. The approach began with substitutable CFGs [3, 12] and was extended to multiple context-free grammars (MCFGs), whose tuple-valued nonterminals express discontinuous and cross-serial dependencies [13, 14, 29, 30]. Yoshinaka's finite-signature generalization [15] provides positive-data identification and, in the corresponding bounded-rank specialization, characteristic samples consisting of polynomially many words.

The question: how long must characteristic data be? A sample of polynomial *cardinality* may still contain exponentially long words. We therefore measure a finite sample K by its total encoded length $\|K\|_+ = \sum_{w\in K}(|w| + 1)$, not merely by $|K|$. This follows the polynomial *thick-data*

perspective reviewed by Eyraud, Heinz and Yoshinaka [10, Section 2.5.2]: the size of a grammar and the length needed to exercise its productive rules are distinct resources. Our aim is to bound the total encoded length of characteristic data relative to an MCFG presentation of the target. Qualitative identification is inherited from Yoshinaka's framework; the contribution of this paper is quantitative.

The setting: a fixed finite-semigroup typing. We study this question for learners equipped with a finite bias. Before any data are seen, we fix a fan-out bound f and a surjective finite-semigroup homomorphism $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$, called a *typing*. The learner compares observed *nonempty* tuple components only when their s-values agree; the target class $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ consists of the f-MCFLs that are tuple-substitutable under this restriction (Definition 2.8). The canonical fresh-identity completion $\widehat{s} : \Sigma^* \to \mathsf{S}^\sharp$ is used to type grammar yields and evaluate rule templates, including empty terminal factors. The substitutability condition itself compares nonempty tuple components using s. Fixing the typing in advance separates the choice of a finite bias from learning under that bias [5]. Prefix–suffix windows, values in a finite group, their products, and DFA transition morphisms are typical instances. Other biases restrict derivations by regular control sets [6, 7] or restrict the evidence used to compare factors [9, 11].

The typing is a *class-level* bias: it is supplied once, before the target is known, and is not chosen to recognize each target separately. Finite features, such as morphological compatibility or case licensing, can constrain substitution between discontinuous constituents while the MCFG accounts for unbounded dependencies. For example, fix one window–group typing $s = (s_{1,1}, g)|_{\Sigma^+_{\mathrm{feat}}}$ for *every* $R \subseteq \{0,1\}^2$, and consider

$$L_R = \{\mathtt{a}^m \mathtt{p}_i \mathtt{b}^n \mathtt{c}^m \mathtt{q}_j \mathtt{d}^n : m, n \geq 1,\ (i,j) \in R\}.$$

Every L_R belongs to $\mathcal{C}^{\mathrm{mcf}}_{2,s}$, and for nonempty R these are non-context-free cross-serial languages (Corollary 6.4). The group component g records the two feature *values*, not the compatibility relation R; neither R nor a target grammar is an input to the learner, which must recover both the licensing information and the unbounded crossing dependencies from positive examples. Explicit characteristic samples of encoded length at most $36|R| + 22$ exist for $R \neq \emptyset$, while every sample characteristic for a single set-driven operator on the whole family must exhibit every licensed pair (Proposition 6.5). The nonrectangular instance $R = \{(0,0), (0,1), (1,0)\}$ shows that untyped and boundary-only substitutability both fail here (Proposition 6.3).

Why bound the original presentation? A natural route is to refine the grammar by types and apply untyped estimates to the refinement. This route can be expensive. The full typed construction may contain as many as $|V||\mathsf{T}^\sharp|^f$ nonterminals and $|P||\mathsf{T}^\sharp|^{fr}$ rules, and its thickness can explode. This is a *potential* cost, not a monotonicity theorem for minimal characteristic-data length, but it can be exponential: for a fixed two-element typing $s_{\mathtt{c}}$ there are rank-two CFG presentations G_n of $\{\mathtt{a}, \mathtt{c}\}^+$ with $\|G_n\| = O(n)$ and rule thickness $\theta_{G_n} \leq 2$, whose trimmed full yield-type refinement has thickness $\theta_{\widetilde{G}_{n,0}} \geq 2^n$ (Proposition 4.4). We therefore bound characteristic data directly in terms of the original presentation, and ask when the algebra of the fixed typing permits short witnesses *before* this cost is paid. For the same reason we treat the typing as given rather than something to infer or optimize: a finer typing $s \preceq t$ (that is, $\ker t \subseteq \ker s$ on Σ^+) enlarges the class, $\mathcal{C}^{\mathrm{mcf}}_{f,s} \subseteq \mathcal{C}^{\mathrm{mcf}}_{f,t}$, but it also changes the reconstruction operator and may leave the algebraic regimes covered by our bounds. We ask when a fixed, externally given typing has class-uniform short witnesses.

The technical obstacle. To obtain short representatives of prescribed root types, we replace selected long subderivations by shorter derivations that need not preserve the componentwise group values of the replaced subtrees. In a multi-component MCFG derivation, one subtree may contribute several blocks to different root components, interleaved with blocks from other subtrees. For noncommutative groups, the effect of simultaneous replacements depends on the order of these blocks. The key quantitative step is an *owner-wise simultaneous selection* of replacements that preserves the ordered group product of every root component, even though individual replacement blocks may change their group values (Proposition 4.6).

Main results. Two input-sensitive thickness scales appear. For a reduced good MCFG G, the primary scale is its *rule thickness* $\bar{\theta}_G = \max\{1, \theta_G\}$: the maximum, over rules, of the shortest accepted sentence using that rule (Section 4.1). For an arbitrary standard MCFG H, the normalization bounds use instead the *component-pattern thickness* $\bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})$, which measures the shortest way to realize each feasible pattern of nonempty tuple components (Definition 4.18). The two scales are not interchangeable for inputs with empty components; Table 1 compares them with the auxiliary quantities used in the proofs. With the alphabet, fan-out f, rule-rank bound $r_{\max}$ and typing fixed, write $\bar{\theta} = \bar{\theta}_G$ for a reduced good grammar G, and $n = \max\{1, \|H\|\}$, $t = \bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})$ for a standard input grammar H. The principal bounds on total encoded characteristic-data length for the set-driven Const/Link/Comp reconstruction operator are:

Fixed typing / rank	Reduced good $(\|G\|, \bar{\theta})$	Any standard input (n, t)	Source
Rule rank ≤ 1; any typing	—	$O(n^3)$	Thm. 5.1
Prefix–suffix window	$O(\|G\|^2\bar{\theta})$	$O(n^3 t)$	Thm. 4.9
Window with finite group	$O(\|G\|^2\bar{\theta})$	$O(n^3 t)$	Thm. 4.10
Positive image in **LG**	$O(\|G\|^3\bar{\theta})$	$O(n^4 t)$	Thm. 4.15

The finite group may be noncommutative. The arbitrary-input column follows from Corollary 4.20 via normalization controlled by component-pattern thickness (Proposition 4.19), and the rank-one row is a separate normalization argument. The bounds are *class-uniform*: they apply to every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ admitting a presentation of the stated form, with constants independent of the particular target. The hidden constants depend on f, $r_{\max}$, $|\Sigma|$ and the typing, and may grow rapidly with them; for noncommutative-group shortening they contain $M_{A,f} = U_f \max\{3, |A|^f\}^{2^{f^2}}$, which may grow doubly exponentially with f. The abelian case has sharper constants (Theorem 4.11). A separate conservative learner gives Gold identification.

The obstruction of Proposition 4.4 has an exact algebraic boundary. Corollary 4.16 shows that the typings whose positive image is a local group, and only those, give uniformly polynomial full-refinement thickness at every fan-out; unlike the characteristic-data theorems, this needs no target-specific substitutability hypothesis. It concerns refinement thickness, and is *not* a lower bound on characteristic-data length.

Further results include a local-group sufficient condition strictly beyond $\mathbf{LI} \vee \mathbf{Grp}$ (Proposition B.5), a rule-only computable witness budget (Proposition 5.2), and an intrinsically rank-two example (Theorem 6.11). Theorem 5.4 complements the class-uniform bounds: it gives a polynomial bound for targets that *have* a representation $L = \widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathsf{H})$ for an abelian-group morphism φ and a subgroup H, even when the auxiliary abelian group is infinite. This is not a bound for every member of $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, and neither the window–group nor the local-group bound assumes such a representation. A refinement s' of s may put an individual example into this form, but it changes the supplied typing and hence the reconstruction operator. The typing-class separations in Section 6 do not assert that either shortening technique is necessary for any single example.

Relation to prior work. The following comparison concerns the guarantees used in this paper; it does not attempt an exhaustive classification of prior learning models.

Source	Established scope	Present additional question
Yoshinaka (2008) [12]	(k,ℓ)-substitutable CFGs; characteristic-data size with a grammar-thickness parameter	Higher-fan-out, fixed-group, and LG-typed shortening
Yoshinaka (2009, 2011) [13, 14]	Multidimensionally substitutable MCFGs and positive-data learning; polynomial identification in the corresponding grammar regime	Original-thickness bounds under supplied finite-semigroup typing
Yoshinaka (2015) [15]	Finite-signature and disjoint-sort learning principles, including finite witness cardinality under fixed bounded-rank signatures	Encoded-length upper bounds that survive the typing refinement

The Const/Link/Comp construction specializes Yoshinaka's qualitative learning principle rather than claiming it as new. The finite-semigroup structure theory and local-group/delay representations used here are classical [36–38]. The qualitative learning principles and finite witness-cardinality results are due to Yoshinaka [14, 15]. Classical sliding-factor coordinates serve as input to our new shortening step rather than a separate novelty claim; the rank-one obstruction uses [21, 22]. Our independent contribution is the noncommutative owner-wise shortening, the resulting window–group and local-group bounds (Theorems 4.10 and 4.15), and the normalization controlled by component-pattern thickness. A related, separately submitted manuscript addresses finite-semigroup-typed CFG learning [25]; no lower-bound construction or sharpness theorem from that manuscript is reproduced or used as a premise here. Positive-data learning of other mildly context-sensitive formalisms is studied in [19], and MCFG learning with membership queries or dual context sets in [16, 18, 20]. A recent substitution lemma for MCFLs [23] is a necessary condition for membership, not a learning criterion.

Scope and organization. The characteristic-data guarantees apply to the specified reconstruction operator and the stated grammar presentations; no presentation-independent stage of Gold convergence is claimed. The bounds measure total encoded length, not an intrinsic minimum grammar size; for general standard inputs, component-pattern thickness may be exponential in grammar size. The polynomial exponents are not established as optimal. The local-group condition is sufficient at every fixed fan-out; whether it is necessary for polynomial characteristic-data length at fan-out two or more is open (Open Problem 7.1). Section 2 fixes the setting, Section 3 defines the reconstruction operator and its qualitative guarantees, Sections 4 and 5 prove the quantitative bounds, and Section 6 compares the typed classes with multidimensional substitutability and develops the examples. The appendices prove the technical shortening, normalization and reconstruction statements.

2 Preliminaries

Throughout, Σ is a finite alphabet, ε the empty word and $\Sigma^+ := \Sigma^* \setminus \{\varepsilon\}$. If $\Sigma = \emptyset$, the only languages are $\emptyset$ and $\{\varepsilon\}$ and are handled directly. For a finite $K \subseteq \Sigma^*$ put $\|K\|_+ := \sum_{w\in K}(|w|+1)$, and for a tuple $\vec{u} = (u_1, \ldots, u_d)$ put $\|\vec{u}\|_1 := \sum_i |u_i|$. Every polynomial bound below states whether it concerns cardinality or $\|\cdot\|_+$. The table fixes symbols used across sections; all others (sliding-factor labellings, auxiliary morphisms, padding lengths) are defined where they occur.

Symbol	Meaning
$G,\ H$	a (reduced good) MCFG; an arbitrary standard input MCFG
B, C, X, Y	nonterminals; X, Y also denote typed nonterminals where specified
$\|G\|$, θ_G, t_G	encoding size (Definition 2.5), rule thickness, nonterminal thickness
s, $\widehat{s}$, S	positive typing, its fresh-identity completion, its positive image
$\preceq$	kernel refinement (Definition 2.3), also between monoid morphisms on Σ^*
A	a finite group in group-theoretic lemmas (otherwise locally defined); A an abelian group in Section 5.3
g	a monoid morphism $\Sigma^* \to A$ into a finite group
T	a derivation tree

2.1 Positive-Data Learning

Definition 2.1 (Text, TxtEx identification, conservativeness)**.** Fix a pause symbol $\perp \notin \Sigma$. A *text* for $L \subseteq \Sigma^*$ is an infinite sequence over $L \cup \{\perp\}$ in which every element of L occurs; the *content* of a finite prefix is the set of elements of Σ^* occurring in it. A *learner* is a computable map from finite prefixes to grammars. It *TxtEx-identifies* L [1] if on every text for L its outputs are, from some stage on, one fixed grammar H with $L(H) = L$; it TxtEx-identifies a class if it does so for each member. It is *conservative* if it changes its hypothesis only on a datum $w \in \Sigma^*$ not generated by the current hypothesis. We fix a hypothesis H_0 with $L(H_0) = \emptyset$.

Definition 2.2 (Set-driven operator and characteristic sample)**.** A *set-driven reconstruction operator* $\mathcal{B}$ maps each finite $K \subseteq \Sigma^*$ to a grammar $\mathcal{B}(K)$ depending only on the set K; it is *sample-consistent* if $K \subseteq L(\mathcal{B}(K))$ for all K. A finite $\mathcal{F} \subseteq L$ is a *characteristic sample for $\mathcal{B}$ and L* if $L(\mathcal{B}(K)) = L$ for every finite K with $\mathcal{F} \subseteq K \subseteq L$ (cf. [4]). This constrains generated languages only: $\mathcal{B}(K)$ need not stabilize syntactically as K grows.

2.2 Fixed Finite-Semigroup Typing

Definition 2.3 (Typing, fresh completion, refinement)**.** A *typing* is a surjective semigroup morphism $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ onto a finite semigroup, given by the multiplication table of S and the images of the letters, so that $s(w)$ is computable in time $O(|w|)$; S is its *positive image*. (A morphism into a larger finite semigroup is replaced by its corestriction to the image.) Let $\mathsf{S}^\sharp := \mathsf{S} \uplus \{\mathbf{1}_\sharp\}$ be obtained by adjoining a *fresh* identity, even if S already has one, and let $\widehat{s} : \Sigma^* \twoheadrightarrow \mathsf{S}^\sharp$ be the monoid morphism with $\widehat{s}(\varepsilon) = \mathbf{1}_\sharp$ and $\widehat{s}|_{\Sigma^+} = s$. For $\vec{w} \in (\Sigma^*)^d$ put $\widehat{s}^{(d)}(\vec{w}) := (\widehat{s}(w_1), \dots, \widehat{s}(w_d))$; on $(\Sigma^+)^d$ this is $s^{(d)}(\vec{w})$. For typings s, t write $s \preceq t$ if $\ker t \subseteq \ker s$ on Σ^+, and let (s, t) be the typing $w \mapsto (s(w), t(w))$ onto its image in $\mathsf{S} \times \mathsf{T}$; thus $s, t \preceq (s, t)$. Write s_{triv} for the one-element typing.

The fresh completion evaluates grammar templates (including empty terminal gaps) and types tuple yields in the refinement of Section 3.2; its new identity is not a positive-word type. Algebraic conditions such as **LG** concern S, not $\mathsf{S}^\sharp$. A finite-monoid morphism $q : \Sigma^* \to N$ determines the positive typing $q|_{\Sigma^+}$: by Lemma E.1(i), it gives the same target class and set-driven reconstruction rules as the fresh completion of that typing. Lemma E.1(ii) additionally identifies their trimmed yield-type refinements for *good* grammars. This last identification is not asserted for arbitrary standard grammars with empty-component yields.

Definition 2.4 (Prefix–suffix typing)**.** For $k, \ell \geq 0$ put $m_0 := \max\{1, k + \ell\}$, and for $|w| \geq m_0$ let $\mathrm{pref}_k(w)$, $\mathrm{suf}_\ell(w)$ be the prefix and suffix of the indicated lengths (length-zero windows are ε). Define

$$\widehat{s}_{k,\ell}(w) = \begin{cases} w, & |w| < m_0, \\ \langle \mathrm{pref}_k(w), \mathrm{suf}_\ell(w) \rangle, & |w| \geq m_0, \end{cases}$$

with disjoint tags for short words and pairs, and put $M_{k,\ell} := \widehat{s}_{k,\ell}(\Sigma^*)$. For short x, y and pairs $\langle u, v\rangle, \langle u', v'\rangle$ define

$$x \cdot y := \widehat{s}_{k,\ell}(xy), \qquad x \cdot \langle u, v\rangle := \langle \mathrm{pref}_k(xu), v\rangle,$$
$$\langle u, v\rangle \cdot y := \langle u, \mathrm{suf}_\ell(vy)\rangle, \quad \langle u, v\rangle \cdot \langle u', v'\rangle := \langle u, v'\rangle.$$

The *boundary typing* is $s_{k,\ell} := \widehat{s}_{k,\ell}|_{\Sigma^+}$; a typing s with $s \preceq s_{k,\ell}$ is (k,ℓ)-*boundary-determined*, or briefly a *boundary typing*.

A case check on the lengths of $x, y \in \Sigma^*$ gives $\widehat{s}_{k,\ell}(xy) = \widehat{s}_{k,\ell}(x) \cdot \widehat{s}_{k,\ell}(y)$. Since $\widehat{s}_{k,\ell}$ is onto $M_{k,\ell}$, the product is associative and $\widehat{s}_{k,\ell}(\varepsilon)$ is an identity, so $\widehat{s}_{k,\ell}$ is a finite-monoid morphism. The short-word tag separates ε from $\widehat{s}_{k,\ell}(\Sigma^+)$, so $\widehat{s}_{k,\ell}$ is isomorphic to the completion $\widehat{s_{k,\ell}}$.

2.3 Multiple Context-Free Grammars

Definition 2.5 (Standard MCFG presentation and size)**.** An MCFG [29, 30] is $G = (V, \Sigma, P, S, \mathrm{fan})$ with finite V, $S \in V$, $\mathrm{fan}(S) = 1$ and $\mathrm{fan}(A) \geq 1$. A rule of rank $r \geq 0$ is $A \to \rho(B_1, \ldots, B_r)$ with template $\rho = (\alpha_1, \ldots, \alpha_{d_0})$, $d_0 = \mathrm{fan}(A)$, $d_j = \mathrm{fan}(B_j)$, where each α_i is a word over $\Gamma_\rho := \Sigma \uplus \Xi_1 \uplus \cdots \uplus \Xi_r$, $\Xi_j := \{\xi_j^1, \ldots, \xi_j^{d_j}\}$, and every variable occurs at most once in $\alpha_1 \cdots \alpha_{d_0}$ (*linearity*). For $\vec{u}_j \in (\Sigma^*)^{d_j}$, $\rho(\vec{u}_1, \ldots, \vec{u}_r)$ substitutes $u_{j,\kappa}$ for ξ_j^κ; for $r = 0$ the rule is a *constant rule*. The tuple languages $L_A(G) \subseteq (\Sigma^*)^{\mathrm{fan}(A)}$ form the least family closed under the rules, and $L(G) := \{w : (w) \in L_S(G)\}$. We write $A \Rightarrow_G^* \vec{u}$ for a *terminal derivation* (a finite derivation tree rooted at A, evaluated bottom-up); one from S is *successful*. G is an f-MCFG if $\mathrm{fan}(A) \leq f$ for all A, and has rule-rank cap $r_{\max}$ if all ranks are at most $r_{\max}$. Such a presentation, with no further restriction, is called *standard*. Its *size* is

$$\|G\| := |V| + \sum_{(A \to \rho(B_1,\ldots,B_r)) \in P} \Big(1 + r + \textstyle\sum_i |\alpha_i|\Big),$$

where $|\alpha_i|$ counts terminal and variable occurrences. All bounds use this symbol-level size; a bit-level encoding with binary nonterminal names changes it by at most a factor $O(\log \|G\|)$.

$\mathsf{MCFL}(f, r)$ is the class of languages of MCFGs of fan-out $\leq f$ and rule rank $\leq r$; f-MCFL $= \bigcup_{r<\infty} \mathsf{MCFL}(f, r)$, and the minimum rule rank of a language is taken over all its fan-out-$\leq f$ presentations. For nonterminals of equal fan-out d, $A \to B$ abbreviates $A \to (\xi_1^1, \ldots, \xi_1^d)(B)$, and $A \to \vec{c}$ a constant rule. Rule variables are always ξ_j^κ; terminals in examples are in `typewriter` font. Fan-out and rank are distinct resources. At fixed fan-out at least two the rule-rank hierarchy does not collapse in general, although adjacent levels can coincide; in particular $\mathsf{MCFL}(2, 2) = \mathsf{MCFL}(2, 3)$ [24, 32]. By contrast, fan-out-one CFGs can be binarized. The reconstruction below therefore assumes neither a binary normal form nor a known rank cap.

Definition 2.6 (Grammar properties)**.** A rule is *nondeleting* if every variable occurs exactly once in its template; a nondeleting rule is *nonpermuting* if for each j the variables $\xi_j^1, \ldots, \xi_j^{d_j}$ occur in this order in $\alpha_1 \cdots \alpha_{d_0}$, and *nonmerging* if moreover no α_i contains a factor $\xi_j^\kappa \xi_j^{\kappa+1}$. Following [14, Section 2.3], a rule is λ-*free* if no α_i is empty. A rule, template or grammar is *good* if it is λ-free, nondeleting, nonpermuting and nonmerging. A nonterminal A is *reachable* if there is a chain $S = A_0, \ldots, A_m = A$ in which each A_{i+1} is a child in a rule for A_i, and *productive* if $L_A(G) \neq \emptyset$; G is *reduced* if every nonterminal is both.

In a λ-free grammar every component of every derived tuple is nonempty (induction on derivations). λ-freeness concerns template components and is distinct from $\varepsilon \notin L(G)$.

2.4 Tuple Contexts and Substitutability

A *sentence context of arity d* is a word $E = v_0\Box_1 v_1 \cdots v_{d-1}\Box_d v_d$ with each hole occurring once, $v_0, v_d \in \Sigma^*$ and $v_1, \ldots, v_{d-1} \in \Sigma^+$; $\mathcal{E}_d$ is the set of these *admissible* contexts [14, Section 3.1]. $E[\vec{w}]$ fills $\Box_i$ by w_i, $|E|_{\text{fix}} := \sum_{i=0}^{d} |v_i|$, and E *accepts* $\vec{w}$ in L if $E[\vec{w}] \in L$. The *d-tuple distribution* of $\vec{x} \in (\Sigma^*)^d$ is $\mathcal{D}_L^{(d)}(\vec{x}) := \{E \in \mathcal{E}_d : E[\vec{x}] \in L\}$; at fan-out one it is the two-sided distribution $\{(u, v) : uxv \in L\}$.

Definition 2.7 (Tuple occurrence)**.** For $w = a_1 \cdots a_n$ number the cut positions $⓪a_1① \cdots a_n⓷$ and write $w[i : i'] := a_{i+1} \cdots a_{i'}$. An *arity-d tuple occurrence* in w is a sequence $0 \le i_1 < i'_1 < i_2 < i'_2 < \cdots < i_d < i'_d \le n$; it determines $\vec{x} = (w[i_1 : i'_1], \ldots, w[i_d : i'_d]) \in (\Sigma^+)^d$ and the context $E \in \mathcal{E}_d$ obtained by replacing these intervals by $\Box_1, \ldots, \Box_d$, so that $E[\vec{x}] = w$. For finite K, $\text{Obs}_d(K) := \{\vec{x} \in (\Sigma^+)^d : E[\vec{x}] \in K \text{ for some } E \in \mathcal{E}_d\}$ is the set of tuples arising from such occurrences, and $\text{Obs}(K) := \bigcup_{1 \le d \le f} \text{Obs}_d(K)$ for the fixed fan-out bound f.

Definition 2.8 ((f, s)-tuple substitutability and the target class)**.** Let $f \ge 1$ and let s be a typing. A language $L \subseteq \Sigma^*$ is (f, s)*-tuple-substitutable* if for every $1 \le d \le f$ and all $\vec{x}, \vec{y} \in (\Sigma^+)^d$,

$$s^{(d)}(\vec{x}) = s^{(d)}(\vec{y}) \;\text{ and }\; \mathcal{D}_L^{(d)}(\vec{x}) \cap \mathcal{D}_L^{(d)}(\vec{y}) \neq \emptyset \implies \mathcal{D}_L^{(d)}(\vec{x}) = \mathcal{D}_L^{(d)}(\vec{y}).$$

$\mathcal{C}_{f,s}^{\text{mcf}}$ is the class of (f, s)-tuple-substitutable f-MCFLs; a member of $\mathcal{C}_{f,s}^{\text{mcf}}$ for a boundary typing s is *boundary-typed*. For $\vec{x}, \vec{y} \in (\Sigma^+)^d$ write $\vec{x} \equiv_{L,s}^{d} \vec{y}$ (or $\vec{x} \equiv_{L,s} \vec{y}$) if $s^{(d)}(\vec{x}) = s^{(d)}(\vec{y})$ and $\mathcal{D}_L^{(d)}(\vec{x}) = \mathcal{D}_L^{(d)}(\vec{y})$; this is an equivalence relation.

Empty components are never compared, following [14, Section 3.1 and footnote 3] (see [8] for CFGs), and $\varepsilon \in L$ is handled at the start symbol; otherwise already $\mathtt{a}^+$ would fail, since $\mathtt{a}$ and ε share a context but not their distributions. The bias consists of both s and f:

(a) If $s \preceq t$, equal t-types are equal s-types, so $\mathcal{C}_{f,s}^{\text{mcf}} \subseteq \mathcal{C}_{f,t}^{\text{mcf}}$; in particular $\mathcal{C}_{f,s_{\text{triv}}}^{\text{mcf}} \subseteq \mathcal{C}_{f,s}^{\text{mcf}}$. This compares classes only: the operators $\mathcal{B}_{f,s}$ and $\mathcal{B}_{f,t}$ of Section 3 differ in general (they coincide when $\ker s = \ker t$), and no monotonicity of characteristic-data length in the typing is claimed.

(b) $\mathcal{C}_{f+1,s}^{\text{mcf}} \cap f\text{-MCFL} \subseteq \mathcal{C}_{f,s}^{\text{mcf}}$, but the class is not monotone in f: for trivial typing $\{w\#w^R\}$ is 1- but not 2-dimensionally substitutable [14, Section 3.1, Example 5].

(c) For every $\mathsf{Acc} \subseteq \mathsf{S}^\sharp$, $\widehat{s}^{-1}(\mathsf{Acc}) \in \mathcal{C}_{f,s}^{\text{mcf}}$: it is regular, and componentwise equal types give equal total types in every context.

Proposition 2.9 (Intersection and abelian filtering)**.** *Fix f and s.*

(i) If L_1, L_2 are (f, s)-tuple-substitutable, so is $L_1 \cap L_2$.

(ii) If L is (f, s)-tuple-substitutable, A is an abelian group, $\varphi : \Sigma^* \to \mathsf{A}$ a monoid morphism, $\mathsf{H} \le \mathsf{A}$ and $\gamma \in \mathsf{A}$, then $L \cap \varphi^{-1}(\gamma + \mathsf{H})$ is (f, s)-tuple-substitutable.

Whenever such an intersection is an f-MCFL, it belongs to $\mathcal{C}_{f,s}^{\text{mcf}}$; in particular so does every f-MCFL $\widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathsf{H})$. Such factorizations are not supplied to the learner.

Proof. (i) A context accepting equal-type $\vec{x}, \vec{y}$ in $L_1 \cap L_2$ accepts both in each L_i, so their distributions agree in each L_i, hence in the intersection. (ii) If some context E accepts $\vec{x}, \vec{y}$ in $\varphi^{-1}(\gamma + \mathsf{H})$, then $\delta := \sum_i (\varphi(y_i) - \varphi(x_i)) \in \mathsf{H}$, and by commutativity $\varphi(E'[\vec{y}]) - \varphi(E'[\vec{x}]) = \delta$ for every $E' \in \mathcal{E}_d$, so coset membership transfers in both directions. Thus $\varphi^{-1}(\gamma + \mathsf{H})$ is (f, s_{triv})-, hence (f, s)-tuple-substitutable; apply (i). The last two claims follow from (i), (ii) and item (c) above. □

3 Reconstruction and Qualitative Learning

This section fixes the reconstruction operator analysed later and records its qualitative guarantees, none of which is claimed as new. For fixed f, rank cap and typing, our operator follows the distributional reconstruction pattern for substitutable languages over finite signatures in [15]. We use the MCFG signature of [15, Example 3] with tuple sorts refined by s-types (cf. [15, Section 8]). The related qualitative results are

Here	Related prior result	Treatment here
Prop. 3.10 (soundness)	[15, Lemma 2]	—
Prop. 3.11 (completeness)	[15, Lemma 3]	arbitrary same-type representatives
Thm. 3.13 (exact reconstruction)	[15, Lemmas 2–3]	as for Prop. 3.11
Cor. 3.16 (identification)	[15, Theorem 1]	unbounded rank over the class

The table records qualitative analogues rather than literal identities of reconstruction operators. The concrete grammar statements and their adaptations are verified directly in Appendices D and E; no embedding into the signature framework is used as a premise. For a fixed sample the uncapped operator is the finite cap $\|K\|_+$, so soundness needs no addition; but identifying all of $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, whose members have unbounded minimum rank, uses a sample-dependent cap, outside the bounded-arity assumption of [15]. Representatives that are not derivable from their nonterminal are used in Proposition 5.3 and hence Theorem 5.4. Finally, [15] bounds the *number* of characteristic words by the number of target rules and nonterminals; their *length*, which this paper bounds, is not addressed there.

3.1 Set-Driven Reconstruction

All interaction between templates and contexts rests on one observation.

Lemma 3.1 (Induced child contexts)**.** *Let $\rho = (\alpha_1, \dots, \alpha_e)$ be a good template of rank r with child arities $d_1, \dots, d_r$, let $E \in \mathcal{E}_e$, let $1 \le j \le r$, and fix $\vec{u}_i \in (\Sigma^+)^{d_i}$ for $i \neq j$. Let E_j be the word obtained from $E[(\alpha_1, \dots, \alpha_e)]$ by replacing ξ^κ_j by $\square_\kappa$ and every other variable ξ^ι_i by $u_{i,\iota}$. Then $E_j \in \mathcal{E}_{d_j}$ and $E_j[\vec{x}] = E[\rho(\vec{u}_1, \dots, \vec{u}_{j-1}, \vec{x}, \vec{u}_{j+1}, \dots, \vec{u}_r)]$ for every $\vec{x} \in (\Sigma^*)^{d_j}$.*

Proof. The identity holds because substitutions commute. By linearity and nondeletion each $\square_\kappa$ occurs exactly once in E_j, and by nonpermutation in the order $\kappa = 1, \dots, d_j$, since E places $\alpha_1, \dots, \alpha_e$ from left to right. Consider ξ^κ_j and $\xi^{\kappa+1}_j$. If both lie in one α_i, nonmerging puts a terminal or a variable ξ^ι_i with $i \neq j$ between them, which becomes a nonempty word. Otherwise they lie in $\alpha_i, \alpha_{i'}$ with $i < i'$, and the material between them contains the separator $v_i \in \Sigma^+$ of E. So every internal separator of E_j is nonempty. □

Lemma 3.2 (Composition preserves equivalence)**.** *Let $L \subseteq \Sigma^*$, let ρ be a good template of rank r, and let $\vec{x}_j, \vec{u}_j \in (\Sigma^+)^{d_j}$ satisfy $\vec{u}_j \equiv_{L,s} \vec{x}_j$ for every j. Then $\rho(\vec{u}_1, \dots, \vec{u}_r) \equiv_{L,s} \rho(\vec{x}_1, \dots, \vec{x}_r)$.*

Definition 3.3 (Composition witness)**.** Let K be finite and $r \ge 1$. A *composition witness of rank r* in K is a triple $(\vec{z}, \rho, (\vec{u}_1, \dots, \vec{u}_r))$ in which ρ is a good template of rank r with output arity $e \le f$ and child arities $d_1, \dots, d_r \le f$, $\vec{z} \in \mathrm{Obs}_e(K)$, $\vec{u}_j \in (\Sigma^+)^{d_j}$, and $\vec{z} = \rho(\vec{u}_1, \dots, \vec{u}_r)$.

Equivalently, a composition witness is a segmentation of the components of an observed $\vec{z}$ into terminal gaps and nonempty intervals labelled by variables, subject to the good-template conditions. If $E[\vec{z}] \in K$, Lemma 3.1 gives $E_j[\vec{u}_j] = E[\vec{z}] \in K$ with E_j admissible, so every $\vec{u}_j \in \mathrm{Obs}(K)$. The children occupy disjoint, nonempty parts of $\vec{z}$, so

$$r \le \textstyle\sum_j \|\vec{u}_j\|_1 \le \|\vec{z}\|_1 \le \max_{w \in K} |w| \le \|K\|_+. \tag{3.1}$$

Definition 3.4 (Capped and uncapped reconstruction operators)**.** Fix f and a typing s. For $r_{\max} \geq 1$ let $\mathcal{B}_{f,s}^{\leq r_{\max}}(K) = \widehat{G}_{\leq r_{\max}}(K)$ be the grammar with start symbol $\widehat{S}$, a nonterminal $[\vec{x}]$ of arity d for each $\vec{x} \in \mathrm{Obs}_d(K)$, $1 \leq d \leq f$, and the rules

(Start) $\widehat{S} \to [(w)]$ for $w \in K \cap \Sigma^+$, and $\widehat{S} \to \varepsilon$ if $\varepsilon \in K$;

(Const) $[\vec{x}] \to \vec{x}$ for every $\vec{x} \in \mathrm{Obs}(K)$;

(Comp) $[\vec{z}] \to \rho([\vec{u}_1], \ldots, [\vec{u}_r])$ for every composition witness of rank $1 \leq r \leq r_{\max}$;

(Link) $[\vec{x}] \to [\vec{y}]$ for $\vec{x}, \vec{y} \in \mathrm{Obs}_d(K)$ with $s^{(d)}(\vec{x}) = s^{(d)}(\vec{y})$ and $E[\vec{x}], E[\vec{y}] \in K$ for some $E \in \mathcal{E}_d$.

For $K \neq \emptyset$ put $\mathcal{B}_{f,s}(K) := \widehat{G}(K) := \widehat{G}_{\leq \|K\|_+}(K)$; $\mathcal{B}_{f,s}(\emptyset)$ has no rules. Fixed encoding orders make both operators deterministic and set-driven.

By (3.1), $\mathcal{B}_{f,s}(K)$ contains a Comp rule for every composition witness in K. Both operators are sample-consistent, via $\widehat{S} \to [(w)] \to (w)$ and $\widehat{S} \to \varepsilon$. All non-start templates are λ-free, so $\widehat{S} \Rightarrow^* \varepsilon$ only by the rule $\widehat{S} \to \varepsilon$. For trivial typing and $\varepsilon \notin K$, Comp/Const, Link and Start correspond to the rule Types I–III of [14, Section 4.1], and $\mathcal{B}_{f,s}$ corresponds to the operator of the unbounded-rank learner $\mathcal{A}(p, *)$ of [14, Corollary 2].

3.2 Yield-Typed Refinement and Witnesses

Definition 3.5 (Yield-type refinement)**.** Let $G = (V, \Sigma, P, S, \mathrm{fan})$ be a finite nondeleting MCFG and s a typing. For a template ρ and $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{d_j}$, let $\mathrm{out}_\rho(\mathbf{t}_1, \ldots, \mathbf{t}_r) \in (\mathsf{S}^\sharp)^{d_0}$ have ith component the image of α_i under the monoid morphism $\Gamma_\rho^* \to \mathsf{S}^\sharp$ extending $a \mapsto \widehat{s}(a)$ and $\xi_j^\kappa \mapsto (\mathbf{t}_j)_\kappa$. The *full yield-type refinement* $G^{\widehat{s}}$ has a fresh start symbol $\widetilde{S}$, nonterminals $A_{\mathbf{t}}$ for $A \in V$ and $\mathbf{t} \in (\mathsf{S}^\sharp)^{\mathrm{fan}(A)}$, start rules $\widetilde{S} \to S_{(\eta)}$ for $\eta \in \mathsf{S}^\sharp$, and for every rule $A \to \rho(B_1, \ldots, B_r)$ (including $r = 0$) and all $\mathbf{t}_j \in (\mathsf{S}^\sharp)^{\mathrm{fan}(B_j)}$ the rule $A_{\mathrm{out}_\rho(\mathbf{t}_1,\ldots,\mathbf{t}_r)} \to \rho(B_{1,\mathbf{t}_1}, \ldots, B_{r,\mathbf{t}_r})$. The *trimmed refinement* $\widetilde{G}_0$ keeps exactly the nonterminals and rules occurring in some successful derivation from $\widetilde{S}$. Unless $G^{\widehat{s}}$ is named, *typed nonterminals* and *typed rules* are the non-start nonterminals and rules of $\widetilde{G}_0$, and $L_X := \{\vec{u} : X \Rightarrow^*_{\widetilde{G}_0} \vec{u}\}$.

Typed rules have the templates of their sources, so $\widetilde{G}_0$ is good whenever G is; each L_X is nonempty, and its tuples have nonempty components if G is λ-free.

Proposition 3.6 (Yield-type invariants)**.** *(i) If $A_{\mathbf{t}} \Rightarrow^* \vec{u}$ in $G^{\widehat{s}}$, then $\widehat{s}^{(\mathrm{fan}(A))}(\vec{u}) = \mathbf{t}$, and erasing indices gives a G-derivation of $\vec{u}$.*

(ii) Every terminal derivation $A \Rightarrow^*_G \vec{u}$ lifts uniquely to $A_{\widehat{s}^{(\mathrm{fan}(A))}(\vec{u})} \Rightarrow^*_{G^{\widehat{s}}} \vec{u}$; a successful derivation of w lifts uniquely through $\widetilde{S} \to S_{(\widehat{s}(w))}$.

(iii) $L(\widetilde{G}_0) = L(G^{\widehat{s}}) = L(G)$.

Lemma 3.7 (Canonical contexts for typed nonterminals)**.** *Let G be good. Every typed nonterminal X of arity d has a context $\chi(X) \in \mathcal{E}_d$ with $\chi(X)[\vec{u}] \in L(G)$ for all $\vec{u} \in L_X$; if $\widetilde{S} \to X$ is in $\widetilde{G}_0$, one may take $\chi(X) = \Box_1$.*

Lemma 3.8 (Typed nonterminals are distributionally homogeneous)**.** *Let G be good of fan-out $\leq f$ with $L := L(G) \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. For every typed nonterminal X of arity d, L_X lies in a single $\equiv^d_{L,s}$-class, denoted $[X]$.*

The proofs of these three statements are in Appendix E. The homogeneity of Lemma 3.8 is weaker than that of congruential MCFGs [17], whose contexts may have empty separators.

Definition 3.9 (Canonical tuples, contexts and the witness set)**.** Let G be good. For every typed nonterminal X fix a *canonical tuple* $\omega(X) \in L_X$ and a *canonical context* $\chi(X)$ as in Lemma 3.7, with $\chi(X) = \square_1$ whenever $\widetilde{S} \to X$ is in $\widetilde{G}_0$. Let $\mathsf{W}(\widetilde{G}_0)$ consist of the words

$$\chi(X)[\omega(X)] \ \ (X \text{ typed}), \qquad \chi(X)[\rho(\omega(Y_1), \ldots, \omega(Y_q))] \ \ (X \to \rho(Y_1, \ldots, Y_q) \text{ typed}, \ q \geq 0).$$

Since $\omega(X)$ and $\rho(\omega(Y_1), \ldots, \omega(Y_q))$ belong to L_X, $\mathsf{W}(\widetilde{G}_0)$ is a finite subset of $L(G)$, nonempty if $L(G) \neq \emptyset$, and $|\mathsf{W}(\widetilde{G}_0)|$ is at most the number of typed nonterminals and rules. Canonical objects are properties of the target, not inputs of the operator: qualitatively any choice works, while Section 4 chooses short ones.

3.3 Qualitative Guarantees

Proposition 3.10 (Soundness of every hypothesis)**.** *Let L be (f,s)-tuple-substitutable and $K \subseteq L$ finite. For every $r_{\max} \geq 1$, if $[\vec{x}] \Rightarrow^* \vec{u}$ in $\widehat{G}_{\leq r_{\max}}(K)$, then $\vec{u} \equiv_{L,s} \vec{x}$. Consequently $L(\mathcal{B}^{\leq r_{\max}}_{f,s}(K)) \subseteq L$ and $L(\mathcal{B}_{f,s}(K)) \subseteq L$.*

Proposition 3.11 (Completeness from a witness assignment)**.** *Let G be a finite good MCFG of fan-out $\leq f$, let r_G be the maximum rank of a typed rule of $\widetilde{G}_0$ ($r_G := 0$ if none), and assign to each typed nonterminal $X = A_{\mathbf{t}}$ a tuple $\iota(X) \in (\Sigma^+)^{\mathrm{fan}(A)}$ with $\widehat{s}^{(\mathrm{fan}(A))}(\iota(X)) = \mathbf{t}$ and a context $\chi(X) \in \mathcal{E}_{\mathrm{fan}(A)}$, with $\chi(X) = \square_1$ if $\widetilde{S} \to X$ is in $\widetilde{G}_0$. Suppose K contains $\chi(X)[\iota(X)]$ for each X and $\chi(X)[\rho(\iota(Y_1), \ldots, \iota(Y_q))]$ for each typed rule $X \to \rho(Y_1, \ldots, Y_q)$, $q \geq 0$. Then*

$$L(G) \subseteq L(\mathcal{B}^{\leq r_{\max}}_{f,s}(K)) \quad (r_{\max} \geq \max\{1, r_G\}), \qquad L(G) \subseteq L(\mathcal{B}_{f,s}(K)).$$

Neither substitutability nor $\iota(X) \in L_X$ is assumed; $\iota = \omega$ is the case $\mathsf{W}(\widetilde{G}_0) \subseteq K$.

Exact reconstruction needs a good presentation of the target; the following normalization controls size, fan-out and rank, but not thickness (for reduced λ-free inputs, see Proposition C.1).

Proposition 3.12 (Good normalization with rank preservation)**.** *Let G_{in} be a finite standard f-MCFG of rule rank at most $r_{\max}$. There is a finite good f-MCFG G_{good} of rule rank at most $r_{\max}$ with $L(G_{\mathrm{good}}) = L(G_{\mathrm{in}}) \setminus \{\varepsilon\}$ and, for fixed $f, r_{\max}$, $\|G_{\mathrm{good}}\| = O_{f,r_{\max}}(\|G_{\mathrm{in}}\|)$. If $L(G_{\mathrm{in}}) \cap \Sigma^+ \neq \emptyset$, G_{good} is reduced.*

Theorem 3.13 (Exact reconstruction for a fixed typing)**.** *Fix $f \geq 1$ and a typing s, and let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. If $L \cap \Sigma^+ = \emptyset$ put $\mathcal{F}_L := L$ and $r_G := 0$. Otherwise let G be a reduced good fan-out-$\leq f$ presentation of $L \cap \Sigma^+$ (Proposition 3.12), let r_G be as in Proposition 3.11, and put $\mathcal{F}_L := \mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$. Then $\mathcal{F}_L \subseteq L$ is finite, and every finite K with $\mathcal{F}_L \subseteq K \subseteq L$ satisfies*

$$L(\mathcal{B}_{f,s}(K)) = L, \qquad L(\mathcal{B}^{\leq r_{\max}}_{f,s}(K)) = L \quad \text{for every } r_{\max} \geq \max\{1, r_G\}.$$

If $L \in \mathsf{MCFL}(f, r)$, G can be chosen of rank $\leq r$, so $r_{\max} = \max\{1, r\}$ suffices (also when $r = 0$).

Theorem 3.14 (Polynomial construction under a rule-rank cap)**.** *Fix $f, r_{\max} \geq 1$ and s. The encoding of $\widehat{G}_{\leq r_{\max}}(K)$ is computable in time $(1 + \|K\|_+)^{O(f r_{\max})}$; its fan-out is at most f and its rank at most $r_{\max}$.*

Definition 3.15 (Conservative sequential learner)**.** The learner $\mathcal{A}_{f,s}$ (respectively $\mathcal{A}_{f,s}^{\leq r_{\max}}$) starts with H_0, keeps its hypothesis on $\perp$ and on every datum the hypothesis generates, and otherwise replaces it by $\mathcal{B}_{f,s}(K)$ (respectively $\mathcal{B}_{f,s}^{\leq r_{\max}}(K)$), where K is the content read so far. This is the policy of [14, Algorithm 1].

Corollary 3.16 (Gold identification)**.** *Fix $f \geq 1$ and s.*

(i) *$\mathcal{A}_{f,s}$ TxtEx-identifies $\mathcal{C}_{f,s}^{\mathrm{mcf}}$.*

(ii) *For fixed $r_{\max}$, $\mathcal{A}_{f,s}^{\leq r_{\max}}$ TxtEx-identifies $\mathcal{C}_{f,s}^{\mathrm{mcf}} \cap \mathsf{MCFL}(f, r_{\max})$, and processes each datum in time polynomial in $\|K\|_+$ for the content K read so far.*

Proofs are in Appendix D (Proposition 3.12: Appendix F). None of these statements bounds the length of characteristic data or the stage of convergence.

4 Characteristic-Data Bounds

Throughout this section f, $r_{\max} \geq 1$, Σ and the typing $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ are fixed; $G = (V, \Sigma, P, S, \mathrm{fan})$ is a reduced good MCFG of fan-out $\leq f$ and rule rank $\leq r_{\max}$, and $\widetilde{G}_0$ is its trimmed $\widehat{s}$-refinement. Under a rank cap, Theorem 3.13 already supplies $O(\|G\|)$ characteristic words; the task is to bound their total length *in terms of G*. All bounds follow one scheme.

(1) *Witness length* (Lemma 4.2). If, for an auxiliary finite-monoid morphism λ refining s, every tuple of every nonterminal can be exchanged for a tuple of the same nonterminal and the same λ-vector of total length at most β, then $\mathsf{W}(\widetilde{G}_0)$ is characteristic and has total length $O(\|G\|(\bar{\theta}_G + \beta))$.

(2) *Minimal counterexamples* (Section 4.3). In a shortest derivation with a prescribed root λ-vector, a large antichain of heavy subtrees would admit a simultaneous replacement preserving that vector; so antichains are small, and Lemma 4.5 turns this into $\beta = O(\|G\|\bar{\theta}_G)$ or $O(\|G\|^2\bar{\theta}_G)$.

(3) *Regimes.* Only λ and the replacement step change: window typings keep marked boundary letters; window–group typings in addition preserve ordered group products (Proposition 4.6); local-group typings reduce to these through sliding-factor coordinates (Section 4.5).

The full yield-type refinement is not used as the primary thickness measure, since it may be exponentially thicker than G (Proposition 4.4). Corollary 4.16 characterizes the semigroup pseudovarieties admitting a *uniform polynomial bound* for this refinement; it does not decide whether a particular grammar exhibits exponential blow-up. Section 4.6 transfers all bounds to arbitrary standard input grammars.

4.1 Thickness Scales and Witness Length

For a rule R of G let

$$\theta_G(R) := \min\{\, |w| : w \in L(G) \text{ has a successful derivation using } R \,\},$$

the length of the derived *sentence*, not of the derivation. Put $\theta_G := \max_{R \in P} \theta_G(R)$, $\bar{\theta}_G := \max\{1, \theta_G\}$ and $t_G := \max_{A \in V} \min\{\|\vec{w}\|_1 : \vec{w} \in L_A(G)\}$ (rule and nonterminal thickness; Yoshinaka writes τ_G for θ_G [14, Section 4.3]). Table 1 lists these scales together with the pattern scale of Section 4.6

Table 1: Thickness scales: two theorem parameters and three auxiliary quantities.

Scale	What it measures / role	Comparison or obstruction
θ_G	Shortest accepted sentence using each rule, maximized over rules; *primary for reduced good G.*	$t_G \le \theta_G \le 2\|G\| \max\{1, t_G\}$ (Lemma 4.1).
t_G	Largest minimum total yield length of a nonterminal; proof-scale for good forms.	Polynomially equivalent to θ_G on reduced good presentations.
$t_{\mathrm{pat}}(H)$	Largest minimum yield length over *realized nonempty-component support patterns*; *primary for arbitrary H.*	$t_H \le t_{\mathrm{pat}}(H)$, with equality for reduced λ-free H; the family E_n has $\theta_{E_n} = O(n)$ but $t_{\mathrm{pat}}(E_n) \ge 2^n$ (Prop. F.2).
θ_{typ}	Rule thickness of the trimmed full yield-type refinement $\widetilde{G}_0$; baseline only.	$\theta_{\mathrm{typ}} \ge \theta_G$; it can be $\ge 2^n$ with $\theta_G \le 2$ (Prop. 4.4).
κ_G	Shortest representative of each typed distribution class, maximized over classes; auxiliary for abelian-filtered targets.	$\kappa_G \le \theta_{\mathrm{typ}}$; representatives need not be derivable from the typed nonterminal (Prop. 5.3).

and two auxiliary quantities; θ_{typ} and κ_G compare G with its own trimmed refinement and are not parameters of the main theorems.

If a target L contains ε, every statement below is applied to $L \cap \Sigma^+$ and the datum ε is added, as in Theorem 3.13; this adds one to the encoded length.

Lemma 4.1 (Thickness lemma)**.** *Let G be reduced and good.*

(i) *Every rule writes at most θ_G terminal symbols explicitly.*

(ii) *$t_G \le \theta_G$: every $A \in V$ has some $\vec{u} \in L_A(G)$ with $\|\vec{u}\|_1 \le \theta_G$.*

(iii) *$\theta_G \le 2\|G\| \max\{1, t_G\}$.*

(iv) *Every $A \in V$ of fan-out d has a context $E_A \in \mathcal{E}_d$ with $|E_A|_{\mathrm{fix}} \le \theta_G$ that results from a successful derivation by deleting one occurrence of an A-subtree; $E_S = \square_1$, and $E_A[\vec{u}] \in L(G)$ for all $\vec{u} \in L_A(G)$. For every typed nonterminal $X = A_{\mathbf{t}}$, the choice $\chi(X) := E_A$ meets Definition 3.9 and is* derivational *for X: it results from a successful derivation of $\widetilde{G}_0$ by deleting an occurrence of an X-subtree.*

Proof. For a rule R with left-hand side A fix a successful derivation D_R of length $\theta_G(R)$ that uses R; it exists because G is reduced. By nondeletion every terminal written by R, and every component of the tuple produced at that occurrence of A, survives exactly once in the sentence. This gives (i), and (ii) because every productive A has a rule.

(iv) For $A = S$ take $\square_1$. Otherwise delete the A-subtree of D_R and put ordered holes in place of its components; iterating Lemma 3.1 along the path to the root (siblings have nonempty yields) shows $E_A \in \mathcal{E}_d$, and its fixed letters belong to the sentence of D_R. Inserting any A-derivation gives a successful derivation, so E_A accepts $L_A(G) \supseteq L_X$ (Proposition 3.6(i)). A typed nonterminal with a start link has the form $S_{\mathbf{t}}$, and $E_S = \square_1$. For derivationality, insert the erasure of any $\widetilde{G}_0$-derivation D_X from X into the hole; the unique lift of the resulting successful derivation (Proposition 3.6(ii)) labels every node by its nonterminal and the $\widehat{s}$-vector of its yield, as D_X does, so it contains D_X at the hole, and deleting that X-subtree leaves E_A.

(iii) Follow a shortest dependency path from S to A; it repeats no nonterminal, so its rules are distinct and differ from R. Complete every sibling on the path and every child of R by a shortest terminal tuple (length $\le t_G$) and apply R. The distinct rules used write at most $\|G\|$ terminals and have at most $\|G\|$ child occurrences, so the sentence has length at most $\|G\| + \|G\| t_G \le 2\|G\| \max\{1, t_G\}$. □

Thus θ_G and t_G are polynomially equivalent on reduced good presentations; Yoshinaka's efficiency criterion measures data in target size and rule thickness [14, Definition 1]; see also [10].

For a finite-monoid morphism $\lambda : \Sigma^* \to N$ we write $s \preceq \lambda$ if $\ker(\lambda|_{\Sigma^+}) \subseteq \ker s$; since $\widehat{s}$ separates ε, this is the relation of Definition 2.3 whenever λ also separates ε from Σ^+. For $\beta \geq 0$ consider

($\mathrm{R}^{\lambda}_{\beta}$) for every $B \in V$ and $\vec{u} \in L_B(G)$ there is $\vec{u}' \in L_B(G)$ with $\lambda^{(\mathrm{fan}(B))}(\vec{u}') = \lambda^{(\mathrm{fan}(B))}(\vec{u})$ and $\|\vec{u}'\|_1 \leq \beta$.

Put

$$C_s(G) := |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}} \;\leq\; |\mathsf{S}^\sharp|^{f r_{\max}} \|G\|,$$

using $|V| + |P| \leq \|G\|$; it bounds the numbers of typed nonterminals and typed rules together, since a typed rule is determined by its source rule and the type vectors of its at most $r_{\max}$ children.

Lemma 4.2 (Witness length)**.** *Let $\lambda : \Sigma^* \to N$ be a finite-monoid morphism with $s \preceq \lambda$.*

(i) (Type-preserving exchange.) *If $X = B_{\mathbf{t}}$ is a typed nonterminal, $\vec{u} \in L_X$, $\vec{u}' \in L_B(G)$ and $\lambda^{(\mathrm{fan}(B))}(\vec{u}') = \lambda^{(\mathrm{fan}(B))}(\vec{u})$, then $\vec{u}' \in L_X$. Hence under ($\mathrm{R}^{\lambda}_{\beta}$) there are canonical tuples with $\|\omega(X)\|_1 \leq \beta$ for every X.*

(ii) (Length.) *If $\|\omega(X)\|_1 \leq b$ for every X and $\chi(X) := E_B$ as in Lemma 4.1(iv), every word of $\mathsf{W}(\widetilde{G}_0)$ has length at most $2\bar{\theta}_G + r_{\max} b$, and*

$$|\mathsf{W}(\widetilde{G}_0)| \leq C_s(G), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ \leq C_s(G)\,(1 + 2\bar{\theta}_G + r_{\max} b).$$

(iii) (Refined thickness.) *In (ii), each rule word $\chi(X)[\rho(\omega(Y_1), \ldots, \omega(Y_q))]$ is the yield of a successful derivation of $\widetilde{G}_0$ using the typed rule $X \to \rho(Y_1, \ldots, Y_q)$; hence $\theta_{\widetilde{G}_0} \leq 2\bar{\theta}_G + r_{\max} b$.*

(iv) (Characteristic sample.) *If $L(G) = L \cap \Sigma^+$ for some $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, then $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}^{\leq r_{\max}}_{f,s}$ and for $\mathcal{B}_{f,s}$.*

Items (i)–(iii) do not use substitutability.

Proof. (i) All components are nonempty because G is λ-free, so $s \preceq \lambda$ and Proposition 3.6(i) give $\widehat{s}^{(\mathrm{fan}(B))}(\vec{u}') = \widehat{s}^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{t}$. By Proposition 3.6(ii) a G-derivation of $\vec{u}'$ lifts to a $G^{\widehat{s}}$-derivation from X. Inserted at an occurrence of X in a successful derivation of $\widetilde{G}_0$ it gives a successful derivation, because typed rules depend only on typed symbols; so all its rules are retained and $\vec{u}' \in L_X$. For the second claim apply ($\mathrm{R}^{\lambda}_{\beta}$) to some $\vec{u} \in L_X \neq \emptyset$.

(ii) A witness word is $\chi(X)[\omega(X)]$ or a rule word. The context has at most $\bar{\theta}_G$ fixed letters, a typed rule has the template of its source and so writes at most $\bar{\theta}_G$ letters (Lemma 4.1(i)), and there are at most $r_{\max}$ child tuples. There is one word per typed nonterminal and per typed rule.

(iii) Apply the typed rule to $\widetilde{G}_0$-derivations of the $\omega(Y_j) \in L_{Y_j}$, erase the labels, and insert the result into the successful G-derivation from which E_B was obtained. As in the proof of Lemma 4.1(iv), the unique lift contains the typed derivation at the hole, and every rule of a successful derivation of $G^{\widehat{s}}$ is retained. A start link $\widetilde{S} \to S_{\mathbf{t}}$ is used by the successful derivation of $\omega(S_{\mathbf{t}})$, of length at most b.

(iv) Proposition 3.11 with $\iota = \omega$ ($r_G \leq r_{\max}$, and $\chi(S_{\mathbf{t}}) = E_S = \Box_1$) and Proposition 3.10 give $L(\mathcal{B}(K)) \cap \Sigma^+ = L \cap \Sigma^+$ for every finite K with $\mathsf{W}(\widetilde{G}_0) \subseteq K \subseteq L$; the Start rule handles ε. □

4.2 Refined-Thickness Baseline

Put $\bar{\theta}_{\mathrm{typ}} := \max\{1, \theta_{\widetilde{G}_0}\}$, the rule thickness of the trimmed refinement itself. Every rule R of G has a retained typed copy (lift a successful derivation using R), and every successful derivation using a copy erases to one using R; hence $\theta_G \le \theta_{\widetilde{G}_0}$.

Lemma 4.3 (Refined-thickness baseline)**.** *Canonical tuples can be chosen so that* $\|\mathsf{W}(\widetilde{G}_0)\|_+ \le C_s(G)\,(1 + (r_{\max}+2)\bar{\theta}_{\mathrm{typ}})$.

Proof. $\widetilde{G}_0$ is reduced and good, so Lemma 4.1(ii) applied to it gives $\omega(X) \in L_X$ with $\|\omega(X)\|_1 \le \theta_{\widetilde{G}_0}$. Apply Lemma 4.2(ii) with $b = \theta_{\widetilde{G}_0}$ and use $\bar{\theta}_G \le \bar{\theta}_{\mathrm{typ}}$. □

For trivial typing a nonterminal of a λ-free grammar has a single realized type, so $\widetilde{G}_0$ is G with start links. The resulting baseline is of the same polynomial scale as Yoshinaka's sharper *untyped* witness estimates $|K_G| \le |P|$ and $\|K_G\| \le |P|\theta_G$ [14, Lemma 7 and Theorem 1]. The displayed bound here concerns our particular typed witness construction and does not itself assert those sharper cardinality or length constants. Under a nontrivial typing the refinement can be much thicker:

Proposition 4.4 (Exponential blow-up of full yield-type refinement)**.** (a) *Let $e \ne z$ be idempotents of S with $ez = ze = z$, and let $u, v \in \Sigma^+$ with $s(u) = e$ and $s(v) = z$. For $n \ge 1$ let G_n be the fan-out-one, rank-two grammar*

$$S \to SS \mid a\ (a \in \Sigma) \mid A_0 \mid \cdots \mid A_n, \qquad A_0 \to u, \qquad A_i \to A_{i-1}A_{i-1} \mid v \quad (1 \le i \le n),$$

where u, v are constant templates. Then G_n is reduced and good, $L(G_n) = \Sigma^+$, $\|G_n\| = O_{\Sigma,u,v}(n)$, $\theta_{G_n} \le \max\{2, 2|u|, 2|v|\}$, and $\theta_{\widetilde{G}_{n,0}} \ge |u|2^n$.

(b) *In particular, for $\Sigma = \{\mathtt{a}, \mathtt{c}\}$ and the two-element typing $s_{\mathtt{c}}$ onto $\{1, \zeta\}$ ($\zeta^2 = \zeta$) with $s_{\mathtt{c}}(\mathtt{a}) = 1$, $s_{\mathtt{c}}(\mathtt{c}) = \zeta$, taking $u = \mathtt{a}$, $v = \mathtt{c}$ gives $\|G_n\| = O(n)$, $\theta_{G_n} \le 2$ and $\theta_{\widetilde{G}_{n,0}} \ge 2^n$.*

Proof. All templates are nonempty and use each variable once in order, so G_n is good; it is reduced and generates Σ^+ via $S \to SS \mid a$. Each rule has a successful use of length at most $\max\{2, 2|u|, 2|v|\}$: complete every A_i by v ($i \ge 1$) or u ($i = 0$). A yield of A_n is a product of leaf values $s(u) = e$ and $s(v) = z$; since $e^2 = e$ and z is absorbing in $\{e, z\}$, its type is e iff no leaf uses v. Hence the unique yield of $(A_n)_{(e)}$ is u^{2^n}; the typed rule $S_{(e)} \to (A_n)_{(e)}$ is retained (via $S \to A_n$), and every successful derivation using it has length at least $|u|2^n$. In (b), $e = 1$ and $z = \zeta$. □

This obstruction concerns the refinement, not characteristic-data length for Σ^+ or for any operator. The proofs below therefore work on G itself.

4.3 Short Representatives by Minimal Counterexamples

Let $\lambda : \Sigma^* \to N$ be a finite-monoid morphism. A vector $\mathbf{t} \in N^d$ is *realizable at B* (of fan-out d) if $\mathbf{t} = \lambda^{(d)}(\vec{u})$ for some $\vec{u} \in L_B(G)$. A *λ-minimal* derivation for $(B, \mathbf{t})$ is a terminal derivation from B with root λ-vector $\mathbf{t}$ that minimizes first the total yield length and then the number of nodes. For a threshold b, a node is *heavy* if its yield has total length greater than b.

Lemma 4.5 (Heavy-node count)**.** *Let $b \ge \bar{\theta}_G$ and $c \ge 0$, and let T be λ-minimal for $(B, \mathbf{t})$. If every antichain of heavy nodes of T has at most c elements, then*

$$\|\textit{root tuple of } T\|_1 \le (r_{\max}+1)\, b\, \max\{1, c\,|V|\,|N|^f\}.$$

Proof. Label each node by its nonterminal and the λ-vector of its yield; there are at most $|V||N|^f$ labels. If an ancestor and a proper descendant had the same label, replacing the ancestor's subtree by the descendant's would keep the root λ-vector (templates evaluate λ homomorphically), would not lengthen the root yield (by linearity the descendant's yield is part of the ancestor's), and would remove nodes, contradicting minimality. So every root-to-leaf path has at most $|V||N|^f$ nodes. If the root is light its yield has length at most b. Otherwise the heavy nodes form an ancestor-closed subtree whose leaves form an antichain, so there are at most $c|V||N|^f$ heavy nodes. Every terminal of the root yield is written by a heavy node (at most $\bar{\theta}_G \le b$ each, Lemma 4.1(i)) or lies in a light child of a heavy node (at most $r_{\max}$ children of length at most b). □

Group constants. For a finite abelian group H, $\mathsf{d}(H)$ is the maximal length of a zero-sum-free sequence over H (so $\mathsf{d}(H)+1$ is the Davenport constant; $\mathsf{d}(\{1\}) = 0$). For a finite group A and $f \ge 1$ put

$$N_{A,f} = \max\{3, |A|^f\}, \quad U_f = f^{f+1} f!, \quad M_{A,f} = U_f\, N_{A,f}^{2^{f^2}}, \qquad c_{A,f} := \begin{cases} \mathsf{d}(A^f), & A \text{ abelian}, \\ M_{A,f} - 1, & \text{otherwise}. \end{cases}$$

U_f bounds the number of block patterns of one owner, and $N_{A,f} \ge 3$ is needed by the order classification in Appendix A.1.

Simultaneous block replacement. A *block-replacement family* of width f over A (Definition A.1) consists of $d \le f$ rows, each an ordered product in A of fixed constants and *blocks*; each block has an *owner* in $[m]$, an original value and a replacement value, and each owner owns at most f blocks, which may interleave with other owners' blocks arbitrarily. $\mathrm{val}_c(J)$ is the value of row c when exactly the blocks of owners in $J \subseteq [m]$ take their replacement values.

Proposition 4.6 (Simultaneous block replacement)**.** *In a block-replacement family of width f over a finite group A with $m > c_{A,f}$ owners there is a nonempty $J \subseteq [m]$ with $\mathrm{val}_c(J) = \mathrm{val}_c(\varnothing)$ for every row c. Individual block or owner values need not be preserved.*

Proof. This is Lemma A.2. For abelian A one selects a nonempty zero-sum subsequence of the owners' difference vectors in $A^d \le A^f$. For arbitrary A no factors are commuted: among owners with a common block pattern (at most U_f patterns) one homogenizes all pairwise order comparisons (Lemma A.7), which decomposes every row into consecutive rank chunks (Lemma A.4); a pigeonhole on prefix-product boundary states then yields an interval of owners preserving all ordered row products (Lemma A.5). □

Owner contributions cannot be commuted: if the two rows are $\beta_{1,1}\beta_{2,1}$ and $\beta_{2,2}\beta_{1,2}$, with value $\sigma = (12)$ on the blocks of owner 1 and $\tau = (23)$ on those of owner 2 in S_3, the row values are $\sigma\tau \ne \tau\sigma$. Example A.6 shows a replacement preserving rows of opposite owner orientation.

Lemma 4.7 (Short representatives for window–group invariants)**.** *Fix $k, \ell \ge 0$, a finite group A (possibly trivial) and a morphism $g : \Sigma^* \to A$. Put $\lambda := (\widehat{s}_{k,\ell}, g) : \Sigma^* \to N := M_{k,\ell} \times A$ and $h := f(k+\ell+1)$. Then $(\mathrm{R}^{\lambda}_{\beta})$ holds with*

$$\beta = b_{k,\ell,A} := (r_{\max}+1)\,\bar{\theta}_G\,(c_{A,f}+h)\,|V||N|^f \;\le\; (r_{\max}+1)(c_{A,f}+h)|N|^f\,\|G\|\,\bar{\theta}_G.$$

Proof. Let $\vec{u} \in L_B(G)$, $\mathbf{t} := \lambda^{(\mathrm{fan}(B))}(\vec{u})$, and let T be λ-minimal for $(B, \mathbf{t})$, with heavy threshold $b = \bar{\theta}_G$. In each root component of length at least $m_0 = \max\{1, k+\ell\}$ mark the first k and the last

ℓ positions (one arbitrary position if $k = \ell = 0$), and in each shorter component mark all positions: at most h marks. Let Q be an antichain of heavy nodes. Incomparable subtrees have disjoint yields, so at most h nodes of Q contain a mark; let Q_0 be the others.

Suppose $|Q_0| > c_{A,f}$. For $v \in Q_0$ with nonterminal C_v fix a tuple of $L_{C_v}(G)$ of total length at most $\bar{\theta}_G$ (Lemma 4.1(ii)). By linearity and nondeletion each $v \in Q_0$ contributes at most f disjoint contiguous blocks to the root components, and the g-value of each root component is the ordered product of the g-values of fixed material and of these blocks. This is a block-replacement family over A with owners Q_0, so Proposition 4.6 gives a nonempty $J \subseteq Q_0$ whose simultaneous replacement preserves every root g-value. The replaced blocks contain no mark: short root components are untouched, and in a long component they lie strictly between the marked prefix and suffix, which keep their letters and positions relative to the ends, while the component keeps its $k + \ell$ marked letters (at least one if $k = \ell = 0$) and stays long. So the root $\widehat{s}_{k,\ell}$-vector is preserved as well, and the result is a derivation from B with root λ-vector $\mathbf{t}$ that is strictly shorter, since each replaced yield exceeded $\bar{\theta}_G$. This contradicts minimality.

Hence every heavy antichain has at most $c_{A,f} + h$ elements, and Lemma 4.5 with $c = c_{A,f} + h \geq 1$ gives the bound; the second inequality uses $|V| \leq \|G\|$. □

Corollary 4.8 (Window representatives)**.** *For $k, \ell \geq 0$, $(\mathrm{R}_\beta^{\widehat{s}_{k,\ell}})$ holds with $\beta = b_{k,\ell} := (r_{\max} + 1)\bar{\theta}_G\, f(k + \ell + 1)\, |V||M_{k,\ell}|^f$.*

Proof. Take A trivial in Lemma 4.7: then $c_{A,f} = 0$ and $N \cong M_{k,\ell}$. □

4.4 Window and Window–Group Typings

Theorem 4.9 (Polynomial data for prefix–suffix-determined typings)**.** *Let $s \preceq s_{k,\ell}$, let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and $L(G) = L \cap \Sigma^+$. Then $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}^{\leq r_{\max}}_{f,s}$ with $|\mathsf{W}(\widetilde{G}_0)| \leq C_s(G) = O(\|G\|)$ and*

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \leq K_{k,\ell}\, \|G\|^2 \bar{\theta}_G, \qquad K_{k,\ell} := |\mathsf{S}^\sharp|^{f r_{\max}} \big(3 + r_{\max}(r_{\max} + 1) f(k + \ell + 1) |M_{k,\ell}|^f\big).$$

Proof. Corollary 4.8 and Lemma 4.2(i),(ii),(iv) with $\lambda = \widehat{s}_{k,\ell}$ and $b = b_{k,\ell}$ give the sample and $\|\mathsf{W}\|_+ \leq C_s(G)(1 + 2\bar{\theta}_G + r_{\max} b_{k,\ell})$; insert $C_s(G) \leq |\mathsf{S}^\sharp|^{f r_{\max}} \|G\|$, $1 + 2\bar{\theta}_G \leq 3\|G\|\bar{\theta}_G$ and $|V| \leq \|G\|$. □

Theorem 4.10 (Polynomial characteristic data for window–finite-group typing)**.** *Let A be a finite group, $g : \Sigma^* \to A$ a morphism, and $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$; let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and $L(G) = L \cap \Sigma^+$. Then $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}^{\leq r_{\max}}_{f,s}$ with $|\mathsf{W}(\widetilde{G}_0)| \leq C_s(G) = O(\|G\|)$ and*

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \leq K_{k,\ell,A}\, \|G\|^2 \bar{\theta}_G,$$
$$K_{k,\ell,A} := |\mathsf{S}^\sharp|^{f r_{\max}} \big(3 + r_{\max}(r_{\max} + 1)(c_{A,f} + f(k + \ell + 1))(|M_{k,\ell}||A|)^f\big).$$

The group may be noncommutative and its blocks may interleave across components.

Proof. As for Theorem 4.9, with Lemma 4.7 and $b = b_{k,\ell,A}$. □

Theorem 4.11 (Polynomial data for window–abelian typing)**.** *If in Theorem 4.10 the group A is abelian, then $c_{A,f} = \mathsf{d}(A^f) = O(\exp(A)\, f \log |A|)$ (Remark A.3), so $K_{k,\ell,A}$ is polynomial in f apart from the factors $|\mathsf{S}^\sharp|^{f r_{\max}}$ and $(|M_{k,\ell}||A|)^f$; for nonabelian A, $c_{A,f} = M_{A,f} - 1$ grows doubly exponentially in f.*

Proof. Lemma A.2(a). □

At fan-out one, polynomial time and thick-data learning for fixed windows is due to Yoshinaka [12, Theorem 1]; the bounds above are relative to a reduced good presentation and are not claimed to improve on that result after normalization. What is new is that window and ordered group values are preserved *simultaneously* across up to f interleaved components. This higher-rank regime is needed at the language level: Theorem 6.11 gives a non-context-free $(1,1)$-typed target of minimum fan-out two and minimum rule rank two, whereas rank one is covered by Theorem 5.1.

4.5 An Algebraic Criterion for Polynomial Characteristic Data

LI is the pseudovariety of finite locally trivial semigroups ($eSe = \{e\}$ for every idempotent e), **Grp** that of finite groups, and $\mathbf{LI} \vee \mathbf{Grp}$ consists of the divisors of products $S \times A$ with $S \in \mathbf{LI}$ and A a finite group [27, 28]. The following lemma makes the window–group hypothesis intrinsic to the positive image S; the semigroup facts are classical [26].

Lemma 4.12 (Locally trivial–group typings admit window–group refinements)**.** *Let $\Sigma \neq \emptyset$. If $\mathsf{S} \in \mathbf{LI} \vee \mathbf{Grp}$, there are $k \geq 1$, a finite group A and a morphism $g : \Sigma^* \to A$ with $\widehat{s} \preceq (\widehat{s}_{k,k}, g)$; if $\mathsf{S} \in \mathbf{LI}$, A can be trivial.*

Proof. Choose a finite $S' \in \mathbf{LI}$, a finite group A, a subsemigroup $T \leq S' \times A$ and a surjective morphism $\pi : T \to \mathsf{S}$ (for $\mathsf{S} \in \mathbf{LI}$: $S' = \mathsf{S}$, $A = \{1\}$). Lifting each $s(a)$ to $(\sigma_a, g_a) \in T$ defines a morphism $\rho = (\sigma, g) : \Sigma^+ \to T$ with $\pi\rho = s$; put $g(\varepsilon) = 1_A$. The subsemigroup $S_0 := \sigma(\Sigma^+) \leq S'$ is locally trivial; let $n := |S_0|$. Two facts:

(a) $S_0^n \subseteq S_0 E(S_0) S_0$. Let t_i be the prefix products of $s_1 \cdots s_n$. If $t_i = t_j$ with $i < j$, then $t_i = t_i u = t_i u^\omega$ for $u = s_{i+1} \cdots s_j$, so the product is $t_i u^\omega y$ with $y \in S_0^1$; otherwise the t_i exhaust S_0, one of them is an idempotent e, and the product is ey. In both cases it is xey with $x, y \in S_0^1$, and $xey = (xe)e(ey)$.

(b) $euf = ef$ for $e, f \in E(S_0)$ and $u \in S_0^1$: $euf = eu(fef) = (e(uf)e)f = ef$.

For $|w| \geq 2n$ write $w = pxz$ with $|p| = |z| = n$. By (a), $\sigma(p) = aeb$ and $\sigma(z) = cfd$ with $a, b, c, d \in S_0$ and e, f idempotent, and by (b) $\sigma(w) = ae\,(b\,\sigma(x)c)\,fd = aefd$ (with $\sigma(\varepsilon) := 1$), which depends only on $p = \mathrm{pref}_n(w)$ and $z = \mathrm{suf}_n(w)$. Shorter words are recorded exactly by $\widehat{s}_{n,n}$, which also separates ε. Hence equal $(\widehat{s}_{n,n}, g)$-values imply equal ρ-values and equal s-values; take $k = n$. □

Corollary 4.13 (Algebraic sufficient condition for polynomial data)**.** *If $\mathsf{S} \in \mathbf{LI} \vee \mathbf{Grp}$, then for every $L \in \mathcal{C}_{f,s}^{\mathrm{mcf}}$ and every G as above with $L(G) = L \cap \Sigma^+$, $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}_{f,s}^{\leq r_{\max}}$ of cardinality $O(\|G\|)$ and total length $O_{f,r_{\max},s,\Sigma}(\|G\|^2 \bar{\theta}_G)$.*

Proof. Lemma 4.12 and Theorem 4.10; k, A and g depend only on s. □

A finite semigroup is a *local group*, $S \in \mathbf{LG}$, if eSe is a group for every idempotent $e \in S$. This is always tested on the positive image S, never on $\mathsf{S}^\sharp$. By Proposition B.5, $\mathbf{LI} \vee \mathbf{Grp} \subsetneq \mathbf{LG}$.

Definition 4.14 (Local-group typing)**.** Fix $k \geq 1$, a finite group A and a map $\nu : \Sigma^{k+1} \to A$ (the *sliding-factor labelling*). For $|w| \geq 2k$ let $\sigma_k(w) = (F_1, \ldots, F_{|w|-k})$ be the factors of length $k+1$ of w in order of their starting positions, and put $\gamma_{k,\nu}(w) := \nu(F_1)\nu(F_2)\cdots\nu(F_{|w|-k})$ (ordered product in A); put $\gamma_{k,\nu}(w) := 1_A$ for $|w| < 2k$. The *local-group typing* is $\lambda_{k,\nu}(w) := (\widehat{s}_{k,k}(w), \gamma_{k,\nu}(w))$; by Lemma B.1 it is a morphism onto a finite monoid $M_{k,\nu}$. The typing s is *local-group-determined* if $\widehat{s} \preceq \lambda_{k,\nu}$ for some k, A and ν.

Theorem 4.15 (Polynomial characteristic data for local-group typing)**.** *Let* $\mathsf{S} \in \mathbf{LG}$*,* $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ *and* $L(G) = L \cap \Sigma^+$*. Then* $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ *is a characteristic sample for* $\mathcal{B}^{\leq r_{\max}}_{f,s}$ *with*

$$|\mathsf{W}(\widetilde{G}_0)| = O(\|G\|), \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ = O_{f,r_{\max},s,\Sigma}(\|G\|^3 \bar{\theta}_G).$$

Proof. By Theorem B.4, $\widehat{s} \preceq \lambda := \lambda_{k,\nu}$ for $k = |\mathsf{S}|$ and some A, ν depending only on s. Lemma B.6 gives $(\mathrm{R}^{\lambda}_{\beta})$ with $\beta = \beta_\lambda = O(\|G\|^2 \bar{\theta}_G)$: heavy subtrees are first exchanged for short tuples with the same window vector (Corollary 4.8), and then a nonempty family is selected whose ordered sliding-factor products remain unchanged (Lemma B.2 and Proposition 4.6). Lemma 4.2 gives $\|\mathsf{W}\|_+ \leq C_s(G)(1 + 2\bar{\theta}_G + r_{\max}\beta_\lambda) = O(\|G\|^3 \bar{\theta}_G)$. □

The Rees and sliding-factor coordinates are classical; their combination with owner-wise block shortening is the quantitative step. For window and window–group typings the quadratic bounds of Theorems 4.9 and 4.10 are sharper.

The algebraic boundary for full refinement thickness. Put $\theta_{\mathrm{typ}}(G, s) := \max\{1, \theta_{\widetilde{G}_0}\}$. For $f \geq 1$, a pseudovariety $\mathbf{V}$ has *polynomial refinement thickness at fan-out* f if for every nonempty finite Σ, every $r_{\max} \geq 1$ and every surjective $s : \Sigma^+ \twoheadrightarrow \mathsf{S}$ with $\mathsf{S} \in \mathbf{V}$ there is a polynomial $q = q_{f,r_{\max},s,\Sigma}$ such that every reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ grammar G over Σ satisfies $\theta_{\mathrm{typ}}(G, s) \leq q(\|G\|, \bar{\theta}_G)$.

Corollary 4.16 (Refinement-thickness dichotomy)**.** *For every* $f \geq 1$ *and every pseudovariety* $\mathbf{V}$ *of finite semigroups,* $\mathbf{V}$ *has polynomial refinement thickness at fan-out* f *if and only if* $\mathbf{V} \subseteq \mathbf{LG}$*; for* $\mathbf{V} \subseteq \mathbf{LG}$ *one may take* $\theta_{\mathrm{typ}}(G, s) = O_{f,r_{\max},s,\Sigma}(\|G\|^2 \bar{\theta}_G)$*. The equivalence persists if* G *ranges only over presentations of* $L \cap \Sigma^+$ *with* $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$*. For each single typing with* $\mathsf{S} \notin \mathbf{LG}$ *the polynomial bound already fails at fan-out one and rule rank two.*

Proof. If $\mathsf{S} \in \mathbf{LG}$, the proof of Theorem 4.15 gives $(\mathrm{R}^{\lambda}_{\beta})$ with $\beta_\lambda = O(\|G\|^2 \bar{\theta}_G)$, and Lemma 4.2(i),(iii), which do not use substitutability, give $\theta_{\widetilde{G}_0} \leq 2\bar{\theta}_G + r_{\max}\beta_\lambda$.

If $\mathsf{S} \notin \mathbf{LG}$, Theorem B.4(ii) gives a subsemigroup $\{e, z\} \cong U_1$: idempotents $e \neq z$ with $ez = ze = z$. By surjectivity choose $u, v \in \Sigma^+$ with $s(u) = e$, $s(v) = z$. Proposition 4.4(a) gives fan-out-one, rank-two grammars G_n with $\|G_n\| = O(n)$, $\bar{\theta}_{G_n} = O(1)$ and $\theta_{\mathrm{typ}}(G_n, s) \geq 2^n$; they present $\Sigma^+ = \widehat{s}^{-1}(\mathsf{S}) \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. Finally, if $\mathbf{V} \not\subseteq \mathbf{LG}$, pick $\mathsf{S} \in \mathbf{V} \setminus \mathbf{LG}$; it is the positive image of the evaluation typing $\mathsf{S}^+ \to \mathsf{S}$ over the alphabet $\Sigma := \mathsf{S}$. □

Remark 4.17 (Refinement thickness is not a data lower bound)*.* The corollary concerns one particular refinement, not the minimal encoded length of characteristic data and not arbitrary learners; its negative direction uses rank two. Polynomial refinement thickness yields polynomial data (Lemma 4.3), but not conversely: short distributional representatives may exist although the refinement is exponentially thick (Proposition 5.3). Whether $\mathbf{LG}$ is necessary for polynomial characteristic data at fan-out ≥ 2 is open (Open Problem 7.1).

4.6 Unrestricted Standard Input MCFGs

The input grammar need not be good: deleting, permuting and merging rules and empty components are all allowed. What then matters is the cost of realizing each nonempty-component pattern.

Definition 4.18 (Component-pattern thickness)**.** For a productive nonterminal A of a standard MCFG H let $\mathrm{supp}(\vec{u}) := \{i : u_i \neq \varepsilon\}$ and $\mathsf{Pat}_H(A) := \{\mathrm{supp}(\vec{u}) : \vec{u} \in L_A(H)\}$. Put

$$t_{\mathrm{pat}}(H) := \max_{\substack{A \text{ productive} \\ I \in \mathsf{Pat}_H(A)}} \min_{\substack{\vec{u} \in L_A(H) \\ \mathrm{supp}(\vec{u}) = I}} \|\vec{u}\|_1, \qquad \bar{t}_{\mathrm{pat}}(H) := \max\{1, t_{\mathrm{pat}}(H)\},$$

with $\max \emptyset := 0$. If H is reduced and λ-free, only the full pattern occurs and $t_{\mathrm{pat}}(H) = t_H$.

Proposition 4.19 (Good normalization controlled by pattern thickness)**.** *Let H be a standard fan-out-$\leq f$, rank-$\leq r$ MCFG with $L(H) \cap \Sigma^+ \neq \emptyset$, and H_{red} its trimmed form. There is a reduced good MCFG G for $L(H) \cap \Sigma^+$ with the same fan-out and rank bounds and*

$$\|G\| = O_{f,r}(\|H\|), \qquad t_G \leq t_{\mathrm{pat}}(H_{\mathrm{red}}), \qquad \bar{\theta}_G = O_{f,r}(\|H\|\, \bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})).$$

Proof. Take G from Proposition 3.12. By Lemma F.1, each nonterminal language of G is an image of a union of support-pattern fibres of H_{red} under component deletion, permutation and concatenation, which do not increase total length; so $t_G \leq t_{\mathrm{pat}}(H_{\mathrm{red}})$ (Appendix C). Lemma 4.1(iii) gives the bound on $\bar{\theta}_G$. □

Corollary 4.20 (Total characteristic-data length from any standard MCFG)**.** *Let H be any finite standard fan-out-$\leq f$, rank-$\leq r$ MCFG for $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$; put $n := \max\{1, \|H\|\}$ and $t := \bar{t}_{\mathrm{pat}}(H_{\mathrm{red}})$. The operator $\mathcal{B}^{\leq r}_{f,s}$ has a characteristic sample $K \subseteq L$ with $|K| = O(n)$ and*

$$\|K\|_+ = \begin{cases} O(n^3 t), & \text{if } \widehat{s} \preceq (\widehat{s}_{k,\ell}, g) \text{ for some } k, \ell \text{ and finite-group morphism } g \text{ (possibly trivial)}, \\ O(n^4 t), & \text{if } \mathsf{S} \in \mathbf{LG}, \end{cases}$$

with constants depending on f, r, s, Σ (and k, ℓ, g), not on H.

Proof. If $L \subseteq \{\varepsilon\}$ take $K = L$. Otherwise Proposition 4.19 gives G with $\|G\| = O(n)$ and $\bar{\theta}_G = O(nt)$; apply Theorem 4.10 or Theorem 4.15. □

5 Further Quantitative Results

This section gives three results that complement Section 4 and are not used there: a bound in grammar size alone for rank one under every typing (Section 5.1); polynomial-time computation of the short witness sets together with a computable data budget (Section 5.2); and characteristic samples from short *distributional* representatives, which yield polynomial data for abelian-filtered targets under every typing (Section 5.3). Notation is that of Section 4; in particular $C_s(G) = |V||\mathsf{S}^\sharp|^f + |P||\mathsf{S}^\sharp|^{f r_{\max}} \leq |\mathsf{S}^\sharp|^{f r_{\max}} \|G\|$.

5.1 Rule Rank One

Theorem 5.1 (Rank-one characteristic-data bound)**.** *Fix f and a typing s.*

(i) *If G is a reduced good fan-out-$\leq f$, rank-≤ 1 presentation of $L \cap \Sigma^+$ with $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, then $\mathsf{W}(\widetilde{G}_0) \cup (\{\varepsilon\} \cap L)$ is a characteristic sample for $\mathcal{B}^{\leq 1}_{f,s}$ and*

$$\|\mathsf{W}(\widetilde{G}_0)\|_+ \leq |\mathsf{S}^\sharp|^f (3 + |\mathsf{S}^\sharp|^f) \|G\|^2 \bar{\theta}_G \leq 2|\mathsf{S}^\sharp|^f (3 + |\mathsf{S}^\sharp|^f) \|G\|^3.$$

(ii) *For any standard fan-out-$\le f$, rank-≤ 1 grammar G_{in} for $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, the operator $\mathcal{B}^{\le 1}_{f,s}$ has a characteristic sample K with $\|K\|_+ = O_{f,s}(\|G_{\mathrm{in}}\|^3)$.*

Proof. (i) A terminal derivation in a rank-≤ 1 grammar is a chain of unary rules ending in a constant rule. In a shortest chain no nonterminal repeats (splicing out the segment between two occurrences does not lengthen the yield, by linearity), so its rules are distinct. Hence $t_G \le \|G\|$, and the proof of Lemma 4.1(iii), whose path rules together with R are distinct and whose child chain is distinct, gives $\theta_G \le 2\|G\|$, so $\bar{\theta}_G \le 2\|G\|$. In $\widetilde{G}_0$, which also has rank ≤ 1, a shortest chain from a typed nonterminal X repeats no typed nonterminal, so it has at most $|V||\mathsf{S}^\sharp|^f$ rules, each writing at most $\bar{\theta}_G$ letters (Lemma 4.1(i)). So $\|\omega(X)\|_1 \le b := |V||\mathsf{S}^\sharp|^f\bar{\theta}_G$ for suitable canonical tuples. Lemma 4.2(ii),(iv) with $r_{\max} = 1$ gives the sample and $\|\mathsf{W}\|_+ \le C_s(G)(1 + 2\bar{\theta}_G + b)$; use $C_s(G) \le |\mathsf{S}^\sharp|^f\|G\|$, $1 + 2\bar{\theta}_G \le 3\bar{\theta}_G$ and $|V| \le \|G\|$.

(ii) If $L \subseteq \{\varepsilon\}$ take $K = L$. Otherwise Proposition 3.12 gives a reduced good rank-≤ 1 presentation G of $L \cap \Sigma^+$ with $\|G\| = O_f(\|G_{\mathrm{in}}\|)$; apply (i). □

No thickness parameter is needed at rank one, and the typing affects only the constant; this extends the rank-one untyped scale [14, Lemma 7]. Rank two needs thickness: the size-$O(n)$ grammar $S \to A_n$, $A_i \to A_{i-1}A_{i-1}$, $A_0 \to \mathsf{c}$ generates $\{\mathsf{c}^{2^n}\}$, which lies in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ for every f and s (over one letter, a shared sentence context forces equal total component length, hence equal distributions), and every characteristic sample for $\mathcal{B}^{\le r_{\max}}_{f,s}$ contains c^{2^n}, because it is a nonempty subset of the target ($L(\mathcal{B}(\emptyset)) = \emptyset$).

5.2 Computing Short Characteristic Data

Let G be reduced and good of fan-out $\le f$ and rank $\le r_{\max}$. For a typed rule $R : X \to \rho(Y_1, \dots, Y_q)$ let c_R be the number of terminals written by ρ and $\operatorname{lhs} R := X$; let $\mathcal{R}$ be the set of typed rules of $\widetilde{G}_0$ (rank zero included). Put

$$\ell_{\mathrm{in}}(X) := \min_{\vec{u} \in L_X} \|\vec{u}\|_1, \qquad \ell_{\mathrm{out}}(X) := \min\{|\chi|_{\mathrm{fix}} : \chi \text{ derivational for } X\},$$

(derivational as in Lemma 4.1(iv); so $\ell_{\mathrm{out}}(X) = 0$ if $\widetilde{S} \to X$ is a rule), and

$$b(R) := c_R + \sum_{j=1}^{q} \ell_{\mathrm{in}}(Y_j), \qquad \mathsf{B}_{\mathrm{rule}}(G, s) := \sum_{R \in \mathcal{R}} \big(1 + \ell_{\mathrm{out}}(\operatorname{lhs} R) + b(R)\big).$$

Proposition 5.2 (Effective short characteristic data and a rule-only budget)**.** *Fix f, $r_{\max}$ and s. Given G, the following are computable in time polynomial in $\|G\|$, numbers being written in binary.*

(i) *$\widetilde{G}_0$, all values $\ell_{\mathrm{in}}(X)$, $\ell_{\mathrm{out}}(X)$ and $\mathsf{B}_{\mathrm{rule}}(G, s)$, and $\theta_G(R)$ for every rule R of G (hence θ_G and t_G).*

(ii) *For each typed nonterminal X, a rule $R_X : X \to \rho_X(Y_1, \dots, Y_q)$ such that the recursion $\omega(X) := \rho_X(\omega(Y_1), \dots, \omega(Y_q))$ is well-founded and gives $\omega(X) \in L_X$ with $\|\omega(X)\|_1 = \ell_{\mathrm{in}}(X)$; and a derivational context $\chi(X)$ with $|\chi(X)|_{\mathrm{fix}} = \ell_{\mathrm{out}}(X)$, equal to $\Box_1$ whenever $\widetilde{S} \to X$ is a rule. Both are given by derivation DAGs of size polynomial in $\|G\|$.*

For these choices,

$$\mathsf{W}(\widetilde{G}_0) = \{\chi(\operatorname{lhs} R)[\rho_R(\omega(Y_1), \dots, \omega(Y_q))] : R \in \mathcal{R}\},$$
$$|\mathsf{W}(\widetilde{G}_0)| \le |\mathcal{R}|, \qquad \|\mathsf{W}(\widetilde{G}_0)\|_+ \le \mathsf{B}_{\mathrm{rule}}(G, s),$$

the words can be written out in time polynomial in $\|G\| + \|\mathsf{W}(\widetilde{G}_0)\|_+$, and every bound on $\mathsf{W}(\widetilde{G}_0)$ proved in Sections 4.1–4.5, Section 5.1 and Appendix B holds for this set; in particular it is characteristic whenever $L(G) = L \cap \Sigma^+$ with $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$.

Proof. For fixed $f, r_{\max}, s$ the full refinement has $O(C_s(G))$ symbols and rules, and trimming is a linear-time productivity and reachability computation.

Inside values. Give a derivation the weight (terminal length, number of nodes), ordered lexicographically. A rule R maps child weights (x_j, n_j) to $(c_R, 1) + \sum_j (x_j, n_j)$, which is strictly larger than every argument. Knuth's generalization of Dijkstra's algorithm, in Nederhof's weighted-deduction form [35] (see also [34]), therefore computes the minimum weight of every typed nonterminal together with a minimizing rule R_X whose children are finalized before X. This order makes the recursion for ω well-founded, the chosen rules form a DAG, and the first coordinate of the weight is $\ell_{\mathrm{in}}(X) = \|\omega(X)\|_1$.

Outside values. Run Dijkstra's algorithm on the graph with an edge $X \to Y_j$ of weight $c_R + \sum_{i \neq j} \ell_{\mathrm{in}}(Y_i)$ for every typed rule $R : X \to \rho(Y_1, \ldots, Y_q)$, from the sources X with $\widetilde{S} \to X$ at distance 0. A path from a source to X yields, by composing its templates and filling the siblings with $\omega(Y_i)$, a context of X of that fixed length; it is admissible by Lemma 3.1 and derivational because it arises from a successful derivation by deleting the X-subtree. Conversely every derivational context of X comes from a root-to-X path whose siblings have yields at least ℓ_{in}. Hence the distance is $\ell_{\mathrm{out}}(X)$, and the shortest-path tree with the inside DAG represents $\chi(X)$. The same two passes on G give $t_G = \max_A \ell_{\mathrm{in}}(A)$ and $\theta_G(R) = \ell_{\mathrm{out}}(A) + c_R + \sum_j \ell_{\mathrm{in}}(B_j)$ for $R : A \to \rho(B_1, \ldots, B_q)$, since outside and inside parts of a derivation can be chosen independently. All values are at most exponential in $\|G\|$, so have polynomial binary length.

Witness set and budget. By definition of ω, $\chi(X)[\omega(X)]$ is the rule word of R_X, and every other word of Definition 3.9 is a rule word. The rule word of R has length exactly $\ell_{\mathrm{out}}(\mathrm{lhs}\, R) + b(R)$; summing $|w| + 1$ over $\mathcal{R}$ gives $\mathsf{B}_{\mathrm{rule}}$, and passing to distinct words can only lower it. Expanding a DAG takes time linear in the length of the word it represents.

Inherited bounds. Each cited bound is proved by exhibiting *some* $\omega(X) \in L_X$ and the derivational context E_B (Lemma 4.1(iv)); the computed minima are no longer and the number of words is unchanged. The characteristic property is Lemma 4.2(iv). □

The budget $\mathsf{B}_{\mathrm{rule}}$ is a presentation-sensitive *upper certificate*, often far below the worst-case bounds; it is not the size of a minimum characteristic sample. For $\{\mathsf{c}^{2^n}\}$ of Section 5.1 the words are exponentially longer than $\|G\|$, but their DAGs have polynomial size.

5.3 Distributional Representatives and Abelian-Filtered Targets

Proposition 3.11 does not require a representative to be derivable from its nonterminal. One may therefore use the shortest member of the whole typed distribution class. Let G be a reduced good fan-out-$\leq f$, rank-$\leq r$ presentation of $L \cap \Sigma^+$ with $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$. For a typed nonterminal X let $[X]$ be the $\equiv_{L,s}$-class containing L_X (Lemma 3.8) and put

$$\kappa_G(X) := \min_{\vec{v} \in [X]} \|\vec{v}\|_1 \leq \ell_{\mathrm{in}}(X), \qquad \kappa_G := \max_X \kappa_G(X).$$

Proposition 5.3 (Relaxed distributional witnesses)**.** *For each typed nonterminal $X = B_{\mathbf{t}}$ choose $\iota(X) \in [X]$ with $\|\iota(X)\|_1 = \kappa_G(X)$ and put $\chi(X) := E_B$ (Lemma 4.1(iv)). The words $\chi(X)[\iota(X)]$ and $\chi(X)[\rho(\iota(Y_1), \ldots, \iota(Y_q))]$, over all typed nonterminals and typed rules $X \to \rho(Y_1, \ldots, Y_q)$, form a set $K^\dagger \subseteq L$ such that $K^\dagger \cup (\{\varepsilon\} \cap L)$ is characteristic for $\mathcal{B}^{\leq r}_{f,s}$ and $\mathcal{B}_{f,s}$, and*

$$|K^\dagger| \leq C_s(G), \qquad \|K^\dagger\|_+ \leq C_s(G)(1 + 2\bar{\theta}_G + r\kappa_G).$$

Proof. Fix $\vec{u}_X \in L_X$ for every X. Since $\iota(X) \equiv_{L,s} \vec{u}_X$ and E_B accepts $\vec{u}_X$, it accepts $\iota(X)$; so $\chi(X)[\iota(X)] \in L$. By Lemma 3.2, $\rho(\iota(Y_1), \dots, \iota(Y_q)) \equiv_{L,s} \rho(\vec{u}_{Y_1}, \dots, \vec{u}_{Y_q}) \in L_X$, so the rule words also lie in L. Each $\iota(X)$ has $\widehat{s}$-vector $\mathbf{t}$, nonempty components, and $\chi(S_{\mathbf{t}}) = E_S = \square_1$, so Proposition 3.11 applies with this ι and gives $L \cap \Sigma^+ \subseteq L(\mathcal{B}(K))$ for $K \supseteq K^\dagger$; Proposition 3.10 gives the converse for $K \subseteq L$, and the Start rule handles ε. The length bound is Lemma 4.2(ii) with $b = \kappa_G$, whose proof uses only $\|\iota(X)\|_1 \le b$. □

Unlike Proposition 5.2, the words $\chi(X)[\iota(X)]$ cannot in general be omitted. For fixed f, r, s, a family G_n with $\|G_n\|\bar{\theta}_{G_n} = n^{O(1)}$ that requires superpolynomial characteristic data must therefore have $\kappa_{G_n} = n^{\omega(1)}$; this is a necessary condition, not a proof that such a family exists.

Theorem 5.4 (Polynomial data for abelian-filtered targets)**.** *Fix a nonempty alphabet Σ, a typing s and $f, r \ge 1$. There is a constant C depending only on Σ, s, f, r such that the following holds for every abelian group A, morphism $\varphi : \Sigma^* \to \mathsf{A}$, subgroup $\mathsf{H} \le \mathsf{A}$, $\gamma \in \mathsf{A}$ and $\mathsf{Acc} \subseteq \mathsf{S}^\sharp$. If $L := \widehat{s}^{-1}(\mathsf{Acc}) \cap \varphi^{-1}(\gamma + \mathsf{H})$ and $L \cap \Sigma^+$ has a reduced good fan-out-$\le f$, rank-$\le r$ presentation G, then $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and $\mathcal{B}^{\le r}_{f,s}$ has a characteristic sample K with*

$$\|K\|_+ \le C\,(\|G\|\bar{\theta}_G)^{3|\Sigma|+2}.$$

Proof. $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ by Proposition 2.9. By Proposition 5.3 and $C_s(G) \le |\mathsf{S}^\sharp|^{fr}\|G\|$ it suffices to show $\kappa_G \le C'(T+1)^{3|\Sigma|+1}$ with $T := (r+1)t_G + \|G\| \le (r+2)\|G\|\bar{\theta}_G$ and C' depending only on Σ, s, f.

A lattice of admissible Parikh shifts. Let ψ be the Parikh map, $\Psi(\vec{u}) := \sum_i \psi(u_i)$, and $\overline{\varphi} : \mathbb{Z}^\Sigma \to \mathsf{A}$ the homomorphism with $\varphi = \overline{\varphi} \circ \psi$; put $\Lambda := \overline{\varphi}^{-1}(\mathsf{H})$. For each $C \in V$ fix $\omega_C \in L_C(G)$ with $\|\omega_C\|_1 \le t_G$. For $\vec{u} \in L_C(G)$ the context E_C accepts $\vec{u}$ and ω_C, and $L \subseteq \varphi^{-1}(\gamma + \mathsf{H})$, so

$$\Psi(\vec{u}) - \Psi(\omega_C) \in \Lambda. \tag{$*$}$$

For a rule $R : C \to \rho(B_1, \dots, B_q)$ let $c_R \in \mathbb{N}^\Sigma$ count its letters and $\delta_R := \Psi(\omega_C) - c_R - \sum_j \Psi(\omega_{B_j})$. Applying $(*)$ to $\rho(\omega_{B_1}, \dots, \omega_{B_q})$ gives $\delta_R \in \Lambda$, and $\|\delta_R\|_1 \le T$. Let $\Gamma_T := \{z \in \Lambda : \|z\|_1 \le T\}$ and $\Lambda_T := \langle \Gamma_T \rangle$. For $\vec{u} = \rho(\vec{u}_1, \dots, \vec{u}_q)$ with $\vec{u}_j \in L_{B_j}(G)$, $\Psi(\vec{u}) - \Psi(\omega_C) = \sum_j (\Psi(\vec{u}_j) - \Psi(\omega_{B_j})) - \delta_R$, so induction on derivations gives $\Psi(\vec{u}) - \Psi(\omega_C) \in \Lambda_T$ for all $\vec{u} \in L_C(G)$.

A short representative. Let $X = C_{\mathbf{t}}$ have arity $d \le f$ and $\vec{u} \in L_X$. For each i fix, once and for all, a semilinear representation of the Parikh image of the regular language $\widehat{s}^{-1}(t_i) \cap \Sigma^+$, and choose a linear set $c_i + \sum_{p \in P_i} \mathbb{N}p$ of it containing $\psi(u_i)$; the number and size of these generators depend only on Σ and s. Write $\Psi(\vec{u}) - \Psi(\omega_C) = \sum_{z \in \Gamma_T} m_z z$ with $m_z \in \mathbb{Z}$. Then the system

$$\sum_{i=1}^{d} \sum_{p \in P_i} \eta_{i,p}\, p - \sum_{z \in \Gamma_T} (\mu_z^+ - \mu_z^-)\, z = \Psi(\omega_C) - \sum_{i=1}^{d} c_i$$

has a solution in nonnegative integers (the coefficients of $\psi(u_i)$ and $\mu_z^\pm := \max\{\pm m_z, 0\}$). It has $m = |\Sigma|$ equations, $n \le D + 2(2T+1)^{|\Sigma|}$ unknowns and entries of absolute value at most $a \le D(T+1)$, with D depending only on Σ, s, f. By Papadimitriou's bound [39] it has a nonnegative integer solution with all entries at most $n(ma)^{2m+1} = O\big((T+1)^{3|\Sigma|+1}\big)$. For each i the vector $c_i + \sum_p \eta_{i,p} p$ is the Parikh vector of some nonempty $v_i \in \widehat{s}^{-1}(t_i)$; put $\vec{v} := (v_1, \dots, v_d)$. Then $\|\vec{v}\|_1 = O((T+1)^{3|\Sigma|+1})$, $\widehat{s}^{(d)}(\vec{v}) = \mathbf{t}$, and $\Psi(\vec{v}) - \Psi(\vec{u}) \in \Lambda_T \subseteq \Lambda$. For every admissible E, $E[\vec{v}]$ and $E[\vec{u}]$ have the same $\widehat{s}$-value, and $\varphi(E[\vec{v}]) - \varphi(E[\vec{u}]) = \overline{\varphi}(\Psi(\vec{v}) - \Psi(\vec{u})) \in \mathsf{H}$; so $E[\vec{v}] \in L$ iff $E[\vec{u}] \in L$, i.e. $\vec{v} \equiv_{L,s} \vec{u}$ and $\kappa_G(X) \le \|\vec{v}\|_1$. □

The theorem holds even when $\mathsf{S} \notin \mathbf{LG}$: its hypothesis concerns the target, not the typing. It is not a bound for all of $\mathcal{C}^{\mathrm{mcf}}_{f,s}$, the representation of L is not supplied to the learner, and no polynomial-time construction is claimed for ineffectively presented groups A.

6 Comparison with Multidimensional Substitutability

Yoshinaka's multidimensional substitutability [13, 14] is the direct positive-data baseline. This section shows that (i) trivial typing recovers it and window typing recovers the CFG (k, ℓ) hierarchy of [12], while a fixed finite typing admits non-context-free targets outside the baseline (Section 6.1); (ii) the algebraic regimes of Section 4 are separated at the level of targets (Section 6.2); and (iii) a boundary-typed target can require rule rank two at every finite fan-out (Section 6.3). Since every (f, s) condition includes arity one, $\mathcal{C}^{\mathrm{mcf}}_{f,s} \cap \mathrm{CFL} \subseteq \mathcal{C}^{\mathrm{mcf}}_{1,s}$.

6.1 Yoshinaka's Baseline and Finite Feature Licensing

Definition 6.1 (Multidimensional substitutability)**.** A language L is *fD-substitutable* if for every $1 \le d \le f$, all $E, F \in \mathcal{E}_d$ and all $\vec{x}, \vec{y} \in (\Sigma^+)^d$, $E[\vec{x}], E[\vec{y}], F[\vec{x}] \in L \Rightarrow F[\vec{y}] \in L$ (the *substitution implication*); $\mathsf{YSub}(f)$ is the class of these languages. This is Yoshinaka's definition, with nonempty components and nonempty internal separators [14, Section 3.1 and footnote 3].

The substitution implication is (f, s_{triv})-tuple substitutability, so $\mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{triv}}} = \mathsf{YSub}(f) \cap f\text{-MCFL} \subseteq \mathcal{C}^{\mathrm{mcf}}_{f,s}$ for every s. Yoshinaka's word-level (k, ℓ) condition [12] requires, for $u \in \Sigma^k$, $v \in \Sigma^\ell$ and nonempty uy_1v, uy_2v, that $w_1uy_1vw_1', w_1uy_2vw_1' \in L$ imply $w_2uy_1vw_2' \in L \iff w_2uy_2vw_2' \in L$. Two nonempty words have equal $s_{k,\ell}$-type iff they are identical or both have length at least $\max\{1, k + \ell\}$ with the same first k and last ℓ letters, i.e. iff they are of the forms uy_1v, uy_2v above (or equal). Hence $(1, s_{k,\ell})$-tuple substitutability is exactly Yoshinaka's (k, ℓ) condition; the componentwise version at higher fan-out is a definition of this paper.

A language L is *strictly 2-local* if there is a set Π of pairs over $\Sigma \cup \{\vdash, \dashv\}$ such that a word w lies in L iff every pair of adjacent symbols of $\vdash w \dashv$ lies in Π.

Lemma 6.2 (Strictly 2-local languages are boundary typed)**.** *Every strictly 2-local L is $(f, s_{1,1})$-tuple-substitutable for every f; in particular $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{1,1}}$, and $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ whenever $s_{1,1} \preceq s$.*

Proof. Let $\vec{x}, \vec{y} \in (\Sigma^+)^d$ have equal $s_{1,1}$-types, so corresponding components are identical letters or share first and last letters, and let $E_0[\vec{x}], E_0[\vec{y}] \in L$. Suppose $E[\vec{x}] \in L$. Since the internal separators of E are nonempty, every adjacent pair of $\vdash E[\vec{y}] \dashv$ either lies inside some y_i, and then occurs in $E_0[\vec{y}] \in L$, or joins a first or last letter of some y_i to a symbol of $\vdash E \dashv$, or lies in $\vdash E \dashv$, and then occurs in $\vdash E[\vec{x}] \dashv$. So $E[\vec{y}] \in L$; symmetry and $s_{1,1} \preceq s$ finish the proof (regular languages are MCFLs). □

Proposition 6.3 (Cross-serial dependencies with finite feature licensing)**.** *Let* $\Sigma_{\mathrm{feat}} = \{\mathtt{p}_0, \mathtt{p}_1, \mathtt{q}_0, \mathtt{q}_1, \mathtt{a}, \mathtt{b}, \mathtt{c}, \mathtt{d}\}$, $\mathcal{R}_{\mathrm{feat}} = \{(0,0), (0,1), (1,0)\}$, *and*

$$L_{\mathrm{feat}} = \{\mathtt{a}^m \mathtt{p}_i \mathtt{b}^n \mathtt{c}^m \mathtt{q}_j \mathtt{d}^n : m, n \ge 1,\ (i,j) \in \mathcal{R}_{\mathrm{feat}}\}.$$

Let $g : \Sigma^*_{\mathrm{feat}} \to \mathbb{Z}_2^2$ map $\mathtt{p}_1 \mapsto (1,0)$, $\mathtt{q}_1 \mapsto (0,1)$ and all other letters to 0, and let $s := (s_{1,1}, g)|_{\Sigma^+_{\mathrm{feat}}}$ (corestricted to its image). Then $L_{\mathrm{feat}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}$ and $L_{\mathrm{feat}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s} \setminus \mathsf{YSub}(f)$ for every $f \ge 2$; moreover L_{feat} is not $(1, s_{1,1})$-tuple-substitutable.

Proof. Grammar. With fan-out-two A, B and rules $A \to (\mathtt{a}, \mathtt{c}) \mid (\mathtt{a}\xi^1_1, \mathtt{c}\xi^2_1)(A)$, $B \to (\mathtt{b}, \mathtt{d}) \mid (\mathtt{b}\xi^1_1, \mathtt{d}\xi^2_1)(B)$ and $S \to (\xi^1_1 \mathtt{p}_i \xi^1_2 \xi^2_1 \mathtt{q}_j \xi^2_2)(A, B)$ for $(i,j) \in \mathcal{R}_{\mathrm{feat}}$, we get $L_A = \{(\mathtt{a}^m, \mathtt{c}^m)\}$, $L_B = \{(\mathtt{b}^n, \mathtt{d}^n)\}$ and $L(G) = L_{\mathrm{feat}}$; the grammar is good, of fan-out and rank two. Intersecting with $\mathtt{a}^+\mathtt{p}_0\mathtt{b}^+\mathtt{c}^+\mathtt{q}_0\mathtt{d}^+$ and erasing the markers gives the classical non-context-free $\{\mathtt{a}^m\mathtt{b}^n\mathtt{c}^m\mathtt{d}^n\}$.

Typed substitutability. Let $Q = \mathtt{a}^+\{\mathtt{p}_0,\mathtt{p}_1\}\mathtt{b}^+\mathtt{c}^+\{\mathtt{q}_0,\mathtt{q}_1\}\mathtt{d}^+$ (strictly 2-local), $K_{\mathrm{feat}} := \{w \in \Sigma^+_{\mathrm{feat}} : g(w) \in \mathcal{R}_{\mathrm{feat}}\}$, and let φ map $\mathtt{a},\mathtt{c},\mathtt{b},\mathtt{d}$ to $(1,0),(-1,0),(0,1),(0,-1) \in \mathbb{Z}^2$ and the markers to 0. On Q, $g(w) = (i,j)$ for the markers $\mathtt{p}_i,\mathtt{q}_j$ of w, so $L_{\mathrm{feat}} = Q \cap K_{\mathrm{feat}} \cap \varphi^{-1}(0)$. The three languages are (f,s)-tuple-substitutable by Lemma 6.2, by item (c) after Definition 2.8 (K_{feat} is a union of s-classes), and by Proposition 2.9(ii); Proposition 2.9(i) gives the intersection.

Failure without g. For $x = \mathtt{ap_0b}$, $y = \mathtt{ap_1b}$, $E = \square_1\mathtt{cq_0d}$ and $F = \square_1\mathtt{cq_1d}$ we have $E[x], E[y], F[x] \in L_{\mathrm{feat}}$ but $F[y] \notin L_{\mathrm{feat}}$, since $(1,1) \notin \mathcal{R}_{\mathrm{feat}}$. Here $s_{1,1}(x) = s_{1,1}(y)$ but $g(x) \neq g(y)$. □

Corollary 6.4 (Uniform finite feature licensing)**.** *With the* same *typing s, for every $R \subseteq \{0,1\}^2$ the language $L_R := \{\mathtt{a}^m\mathtt{p}_i\mathtt{b}^n\mathtt{c}^m\mathtt{q}_j\mathtt{d}^n : m,n \geq 1,\ (i,j) \in R\}$ lies in $\mathcal{C}^{\mathrm{mcf}}_{f,s}$ for all $f \geq 2$, and for $R \neq \emptyset$ it is not context-free. The typing is fixed once for the whole family; the same holds for any fixed inventories $\{0,\dots,p-1\} \times \{0,\dots,q-1\}$ with g into $\mathbb{Z}_p \times \mathbb{Z}_q$.*

Proof. Keep the start rules indexed by R, replace $\mathcal{R}_{\mathrm{feat}}$ by R in K_{feat} (still a union of s-classes), and repeat the proof; for $(i,j) \in R$ use $\mathtt{a}^+\mathtt{p}_i\mathtt{b}^+\mathtt{c}^+\mathtt{q}_j\mathtt{d}^+$ for non-context-freeness. □

Proposition 6.5 (Explicit data budget for a common feature typing)**.** *Let $f \geq 2$ and $R \neq \emptyset$. The operator $\mathcal{B}^{\leq 2}_{f,s}$ has a characteristic sample $K_R \subseteq L_R$ with $\|K_R\|_+ = 36|R| + 22$ (for $R = \emptyset$ the empty sample is characteristic). Conversely, if a set-driven operator $\mathcal{D}$ has characteristic samples for* all *$L_{R'}$, $R' \subseteq \{0,1\}^2$, then every characteristic sample K for L_R under $\mathcal{D}$ contains a word with markers $\mathtt{p}_i,\mathtt{q}_j$ for each $(i,j) \in R$; hence $\|K\|_+ \geq 7|R|$.*

Proof. Write $w^{m,n}_{ij} := \mathtt{a}^m\mathtt{p}_i\mathtt{b}^n\mathtt{c}^m\mathtt{q}_j\mathtt{d}^n$, of length $2m + 2n + 2$, fix $(i_0,j_0) \in R$, and let

$$K_R := \{w^{m,n}_{ij} : (i,j) \in R,\ m,n \in \{1,2\}\} \cup \{w^{3,1}_{i_0j_0}, w^{1,3}_{i_0j_0}\}.$$

The encoded costs are $7 + 9 + 9 + 11 = 36$ per pair and $11 + 11$ for the two extra words.

Typed refinement. For the grammar G of Proposition 6.3 with start rules indexed by R, the s-type of $(\mathtt{a}^m,\mathtt{c}^m)$ depends only on whether $m = 1$ (short words are recorded exactly), so A has two typed copies A_1, A_+, and likewise B_1, B_+; the type of $w^{m,n}_{ij}$ is determined by (i,j), giving one typed start nonterminal S_{ij} per $(i,j) \in R$. With $\rho_A := (\mathtt{a}\xi^1_1, \mathtt{c}\xi^2_1)$, the typed rules are $A_1 \to (\mathtt{a},\mathtt{c})$, $A_+ \to \rho_A(A_1)$, $A_+ \to \rho_A(A_+)$, their B-analogues, and four start rules $S_{ij} \to \rho_{ij}(A_\bullet, B_\circ)$ per $(i,j) \in R$.

Witnesses. Take $\iota(A_1) = (\mathtt{a},\mathtt{c})$, $\iota(A_+) = (\mathtt{a}^2,\mathtt{c}^2)$, $\iota(B_1) = (\mathtt{b},\mathtt{d})$, $\iota(B_+) = (\mathtt{b}^2,\mathtt{d}^2)$, $\iota(S_{ij}) = w^{1,1}_{ij}$, the contexts $\chi(A_\bullet) = \square_1\mathtt{p}_{i_0}\mathtt{b}\square_2\mathtt{q}_{j_0}\mathtt{d}$, $\chi(B_\circ) = \mathtt{ap}_{i_0}\square_1\mathtt{cq}_{j_0}\square_2$ and $\chi(S_{ij}) = \square_1$. The required words of Proposition 3.11 are: $w^{1,1}_{i_0j_0}$ for A_1, B_1 and their constant rules; $w^{2,1}_{i_0j_0}$ and $w^{1,2}_{i_0j_0}$ for A_+, B_+ and the rules from A_1, B_1; $w^{3,1}_{i_0j_0}$ and $w^{1,3}_{i_0j_0}$ for the rules from A_+, B_+; and $w^{m,n}_{ij}$, $m,n \in \{1,2\}$, for S_{ij} and its four rules. All lie in K_R. Since $K_R \subseteq L_R \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$, Propositions 3.11 and 3.10 make K_R characteristic for $\mathcal{B}^{\leq 2}_{f,s}$.

Lower bound. If K contains no word with markers $\mathtt{p}_i,\mathtt{q}_j$, then $K \subseteq L_{R'}$ for $R' = R \setminus \{(i,j)\}$. For a characteristic sample F of $L_{R'}$, the set $K \cup F$ lies between K and L_R and between F and $L_{R'}$, so $L(\mathcal{D}(K \cup F))$ would equal both $L_R \neq L_{R'}$. Distinct pairs need distinct words, each of length at least 6. □

The example models a finite case-licensing or feature-compatibility constraint; no claim about a specific natural language is made. The relation R varies with the target, while the morphism g and the operator do not.

6.2 Separating the Typing Classes

The feature example is covered by a window and an abelian group. Two further fan-out-two targets separate the regimes of Section 4: L_{A} is admitted by a local-group typing but by no window–group-determined typing, and L_{B} by a window–S_3 typing but by no window–abelian one. Both have the form $\sigma w\$c$: a flag $\sigma = \mathtt{r}$ leaves w unconstrained, $\sigma = \mathtt{l}$ imposes a typing-dependent condition on w, and $c \in C_\times := \{\mathtt{1}^m\mathtt{2}^n\mathtt{3}^m\mathtt{4}^n : m, n \geq 1\}$ makes the target non-context-free. Let $\varphi_\times$ map $\mathtt{1}, \mathtt{3}, \mathtt{2}, \mathtt{4}$ to $(1,0), (-1,0), (0,1), (0,-1) \in \mathbb{Z}^2$ and all other letters to 0, and let $\#_{\mathtt{ab}}(w)$ be the number of occurrences of the factor $\mathtt{ab}$ in w.

Lemma 6.6 (Flag separation)**.** *Let t be a typing, $u, v \in \Sigma^+$ with $t(u) = t(v)$, and $z \in \Sigma^*$. If L contains $\mathtt{r}uz$, $\mathtt{r}vz$ and $\mathtt{l}uz$ but not $\mathtt{l}vz$, then $L \notin \mathcal{C}^{\mathrm{mcf}}_{f,t}$ for every $f \geq 1$.*

Proof. $\mathtt{r}\Box_1 z \in \mathcal{E}_1$ accepts u and v, while $\mathtt{l}\Box_1 z$ accepts u but not v; so $\mathcal{D}^{(1)}_L(u) \neq \mathcal{D}^{(1)}_L(v)$, and every (f,t) condition includes arity one. □

The targets. Let $\Sigma_{\mathrm{A}} = \{\mathtt{l}, \mathtt{r}, \mathtt{a}, \mathtt{b}, \mathtt{c}, \$, \mathtt{1}, \mathtt{2}, \mathtt{3}, \mathtt{4}\}$,

$$L_{\mathrm{A}} := \{\sigma w\$c : \sigma \in \{\mathtt{l}, \mathtt{r}\},\ w \in \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}^+,\ c \in C_\times,\ \sigma = \mathtt{r} \text{ or } \#_{\mathtt{ab}}(w) \text{ even}\},$$

and $s_{\mathrm{A}} := \lambda_{1,\nu}|_{\Sigma_{\mathrm{A}}^+}$ with $\nu : \Sigma_{\mathrm{A}}^2 \to \mathbb{Z}_2$, $\nu(\mathtt{ab}) = 1$ and $\nu = 0$ otherwise (Definition 4.14); it records the first and last letters and, for $|w| \geq 2$, the parity of $\#_{\mathtt{ab}}(w)$. Let $\Sigma_{\mathrm{B}} = \Sigma_{\mathrm{A}} \setminus \{\mathtt{c}\}$, $\alpha = (12)$, $\beta = (23) \in S_3$, $g_{\mathrm{B}} : \Sigma_{\mathrm{B}}^* \to S_3$ with $\mathtt{a} \mapsto \alpha$, $\mathtt{b} \mapsto \beta$ and all other letters $\mapsto 1$,

$$L_{\mathrm{B}} := \{\sigma w\$c : \sigma \in \{\mathtt{l}, \mathtt{r}\},\ w \in \{\mathtt{a}, \mathtt{b}\}^+,\ c \in C_\times,\ \sigma = \mathtt{r} \text{ or } g_{\mathrm{B}}(w) = 1\},$$

and $s_{\mathrm{B}} := (s_{1,1}, g_{\mathrm{B}})|_{\Sigma_{\mathrm{B}}^+}$, corestricted. **Ab** is the pseudovariety of finite abelian groups, and $\mathbf{LI} \vee \mathbf{Ab}$ consists of the divisors of $T \times A$ with $T \in \mathbf{LI}$ and A finite abelian.

Proposition 6.7 (A target separating local-group and window–group typings)**.** (i) $L_{\mathrm{A}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}$.

(ii) $L_{\mathrm{A}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{A}}}$ *for every $f \geq 2$, and $s_{\mathrm{A}}(\Sigma_{\mathrm{A}}^+) \in \mathbf{LG}$.*

(iii) *If $t \preceq (s_{k,\ell}, g'|_{\Sigma_{\mathrm{A}}^+})$ for some k, ℓ and a morphism g' into a finite group (in particular if $t(\Sigma_{\mathrm{A}}^+) \in \mathbf{LI} \vee \mathbf{Grp}$), then $L_{\mathrm{A}} \notin \mathcal{C}^{\mathrm{mcf}}_{f,t}$ for all $f \geq 1$.*

Hence $s_{\mathrm{A}}(\Sigma_{\mathrm{A}}^+) \in \mathbf{LG} \setminus (\mathbf{LI} \vee \mathbf{Grp})$, and L_{A} is covered by Theorem 4.15 but, for no typing, by Theorems 4.9, 4.10, 4.11 or Corollary 4.13.

Proposition 6.8 (A target separating nonabelian and window–abelian typings)**.** (i) $L_{\mathrm{B}} \in \mathsf{MCFL}(2,2) \setminus \mathrm{CFL}$.

(ii) $L_{\mathrm{B}} \in \mathcal{C}^{\mathrm{mcf}}_{f,s_{\mathrm{B}}}$ *for every $f \geq 2$, and s_{B} is window–group-determined with group S_3.*

(iii) *If $t \preceq (s_{k,\ell}, g'|_{\Sigma_{\mathrm{B}}^+})$ for some k, ℓ and a morphism g' into a finite* abelian *group (in particular if $t(\Sigma_{\mathrm{B}}^+) \in \mathbf{LI} \vee \mathbf{Ab}$), then $L_{\mathrm{B}} \notin \mathcal{C}^{\mathrm{mcf}}_{f,t}$ for all $f \geq 1$.*

Hence L_{B} is covered by Theorem 4.10 with the nonabelian constant $c_{S_3,f} = M_{S_3,f} - 1$, but for no typing by Theorem 4.9 or 4.11.

Proof of Propositions 6.7 and 6.8. Write $X \in \{\mathrm{A}, \mathrm{B}\}$, let Σ_X° be $\{\mathtt{a}, \mathtt{b}, \mathtt{c}\}$ or $\{\mathtt{a}, \mathtt{b}\}$, and let $P_X(w)$ be "$\#_{\mathtt{ab}}(w)$ even" or "$g_{\mathrm{B}}(w) = 1$".

(i) Let W be a right-linear (fan-out-one, rank-one, good) grammar for the regular language $\{\sigma w : \sigma \in \{\mathtt{1}, \mathtt{r}\}, w \in (\Sigma_X^\circ)^+, \sigma = \mathtt{r} \text{ or } P_X(w)\}$; it tracks the last letter and the parity, respectively the value in S_3. Add

$$C \to (\xi_1^1\$\mathtt{1}, \mathtt{3})(W) \mid (\xi_1^1\mathtt{1}, \xi_1^2\mathtt{3})(C), \quad D \to (\mathtt{2}, \mathtt{4}) \mid (\xi_1^1\mathtt{2}, \xi_1^2\mathtt{4})(D), \quad S \to (\xi_1^1\xi_2^1\xi_1^2\xi_2^2)(C, D).$$

Then $L_C = \{(\sigma w\$\mathtt{1}^m, \mathtt{3}^m)\}$, $L_D = \{(\mathtt{2}^n, \mathtt{4}^n)\}$, and S generates L_X; the grammar is good (adjacent variables in the start template belong to different children). Intersecting with $\mathtt{r}\mathtt{a}\$\mathtt{1}^+\mathtt{2}^+\mathtt{3}^+\mathtt{4}^+$ and erasing $\mathtt{r}\mathtt{a}\$$ gives the non-context-free $C_\times$.

(ii) $L_X = Q_X \cap \Pi_X \cap \varphi_\times^{-1}(0)$ with $Q_X = \{\mathtt{1}, \mathtt{r}\}(\Sigma_X^\circ)^+\$\mathtt{1}^+\mathtt{2}^+\mathtt{3}^+\mathtt{4}^+$ and Π_X the set of nonempty words that begin with $\mathtt{r}$, or begin with $\mathtt{1}$ and satisfy P_X for the whole word. On Q_X the condition P_X of the whole word is that of w: no factor $\mathtt{ab}$ meets σ, \$ or the digits, and g_{B} is trivial on all other letters. Q_X is strictly 2-local (its pairs are $\vdash\sigma$, σx, xy, $x\$$, $\$\mathtt{1}$, $\mathtt{11}, \mathtt{12}, \mathtt{22}, \mathtt{23}, \mathtt{33}, \mathtt{34}, \mathtt{44}$, $\mathtt{4}\dashv$ for $\sigma \in \{\mathtt{1}, \mathtt{r}\}$, $x, y \in \Sigma_X^\circ$) and $s_{1,1} \preceq s_X$, so Lemma 6.2 applies. Π_X is a union of s_X-classes: s_X records the first letter, and for words of length at least two the parity of $\#_{\mathtt{ab}}$, respectively the value of g_{B} (the one-letter words $\mathtt{1}, \mathtt{r}$ lie in Π_X); so item (c) after Definition 2.8 applies. Proposition 2.9 handles $\varphi_\times^{-1}(0)$ and the intersection; with (i), $L_X \in \mathcal{C}^{\mathrm{mcf}}_{f,s_X}$ for $f \geq 2$. s_{A} is local-group-determined, so its positive image lies in **LG** by Theorem B.4.

(iii) for A. Let $\hat{k} = \max\{k, \ell, 1\}$, $q \geq 2$ an even multiple of the exponent of the group of g', $u = \mathtt{a}^{\hat{k}}(\mathtt{abc})^q\mathtt{a}^{\hat{k}}$ and $v = \mathtt{a}^{\hat{k}}(\mathtt{bac})^q\mathtt{a}^{\hat{k}}$. Both are long for $s_{k,\ell}$ with prefix $\mathtt{a}^k$ and suffix $\mathtt{a}^\ell$, and $g'(u) = g'(\mathtt{a})^{2\hat{k}} = g'(v)$; so $t(u) = t(v)$. But $\#_{\mathtt{ab}}(u) = q$ is even and $\#_{\mathtt{ab}}(v) = 1$ (only at $\mathtt{a}^{\hat{k}}\mathtt{b}$). With $z = \$\mathtt{1234}$, Lemma 6.6 applies. A typing with positive image in $\mathbf{LI} \vee \mathbf{Grp}$ satisfies the hypothesis by Lemma 4.12.

(iii) for B. Let $\hat{k} \geq \max\{k, \ell, 1\}$ be even, $u = \mathtt{a}^{\hat{k}}(\mathtt{ab})^3\mathtt{a}^{\hat{k}}$ and $v = \mathtt{a}^{\hat{k}}(\mathtt{ab})^2\mathtt{ba}\,\mathtt{a}^{\hat{k}}$. They have the same prefix $\mathtt{a}^k$, suffix $\mathtt{a}^\ell$ and Parikh vector, so $t(u) = t(v)$ for abelian g'. Since $\alpha^2 = \beta^2 = 1$, $\alpha\beta$ has order 3 and $\beta\alpha = (\alpha\beta)^{-1}$, we get $g_{\mathrm{B}}(u) = (\alpha\beta)^3 = 1$ and $g_{\mathrm{B}}(v) = (\alpha\beta)^2(\alpha\beta)^{-1} = \alpha\beta \neq 1$. With $z = \$\mathtt{1234}$, Lemma 6.6 applies. For $\mathbf{LI} \vee \mathbf{Ab}$, the proof of Lemma 4.12 with an abelian group factor gives an abelian g'.

Consequences. If $s_{\mathrm{A}}(\Sigma_{\mathrm{A}}^+)$ were in $\mathbf{LI} \vee \mathbf{Grp}$, (iii) with $t = s_{\mathrm{A}}$ would contradict (ii). The window, window–group and window–abelian theorems and Corollary 4.13 assume a typing of the excluded kind. □

Remark 6.9 (Scope of the separations). These propositions separate *typing classes*; they do not show that local-group or nonabelian shortening is needed for short data for L_{A} or L_{B}. Indeed, refining s_{A} by the syntactic morphism of Q_{A} gives a typing s' with $L_{\mathrm{A}} = \widehat{s'}^{-1}(\mathsf{Acc}') \cap \varphi_\times^{-1}(0)$, so Theorem 5.4 gives polynomial data for $\mathcal{B}^{\leq 2}_{f,s'}$ (similarly for L_{B}); this concerns the refined operator, not $\mathcal{B}^{\leq 2}_{f,s_{\mathrm{A}}}$. The separation also happens at arity one (Lemma 6.6), so these examples isolate the expressiveness of the typing, not its interaction with interleaved components (Example A.6).

6.3 An Intrinsically Rank-Two Boundary-Typed Target

Let $\Sigma_{\mathrm{intr}} := \{\mathtt{a}, \mathtt{b}, \mathtt{1}, \mathtt{r}, \mathtt{p}, \mathtt{q}, \mathtt{c}, \mathtt{d}, \#, \$, \mathtt{u}, \mathtt{v}, \mathtt{w}\}$, $P_{\mathrm{intr}} := \{\mathtt{lpc}, \mathtt{lpd}, \mathtt{lqc}, \mathtt{rpc}, \mathtt{rpd}\}$,

$$\mathrm{MIX}_2^+ := \{x \in \{\mathtt{a}, \mathtt{b}\}^+ : |x|_{\mathtt{a}} = |x|_{\mathtt{b}}\}, \quad T_3^+ := \{\mathtt{u}^n\mathtt{v}^n\mathtt{w}^n : n \geq 1\}, \quad L_{\mathrm{intr}} := P_{\mathrm{intr}} \,\#\, \mathrm{MIX}_2^+ \,\$\, T_3^+.$$

The factor MIX_2^+ forces rule rank two, T_3^+ forces fan-out two, and the wrapper P_{intr} separates the target from $\mathsf{YSub}(f)$.

Lemma 6.10 (Rank-one MCFLs are finite-index EDT0L)**.** *Every language in* $\mathsf{MCFL}(f', 1)$ *has an EDT0L system of index at most* f'*.*

Proof. Let H be reduced, of fan-out $\leq f'$ and rank ≤ 1. Use a symbol $[B, j]$ for each nonterminal B and component j, axiom $[S, 1]$, and for each rule $A \to \rho(B)$ (or $A \to \rho$) a table mapping $[A, j]$ to the jth component of ρ with ξ_1^k replaced by $[B, k]$, and fixing every other symbol. Since each rule has at most one child and is linear, every reachable sentential form contains symbols of a single nonterminal, each at most once, so at most f' nonterminal occurrences. Table sequences that produce terminal words correspond to derivations of H (components of the child not used by a deleting rule are completed by any derivation, as H is reduced), with the same terminal word. □

Theorem 6.11 (A non-context-free intrinsic-rank-two typed witness)**.** $L_{\mathrm{intr}} \in \mathsf{MCFL}(2, 2) \setminus \mathrm{CFL}$*, and* $L_{\mathrm{intr}} \in \mathcal{C}^{\mathrm{mcf}}_{f, s_{1,1}} \setminus \mathsf{YSub}(f)$ *for every* $f \geq 2$*, but* $L_{\mathrm{intr}} \notin \mathsf{MCFL}(f', 1)$ *for every finite* f'*. Thus its minimum fan-out is two and its minimum rule rank over all finite-fan-out presentations is two: Theorem 4.9 applies to it, while Theorem 5.1 does not at any finite fan-out.*

Proof. Grammar. The fan-out-one rules $B \to \mathtt{a}B\mathtt{b}B \mid \mathtt{b}B\mathtt{a}B \mid \mathtt{a}B\mathtt{b} \mid \mathtt{ab}B \mid \mathtt{ab} \mid \mathtt{b}B\mathtt{a} \mid \mathtt{ba}B \mid \mathtt{ba}$ give $L_B = \mathrm{MIX}_2^+$: every word they generate is balanced, and conversely a balanced word decomposes at the first return of $|\cdot|_{\mathtt{a}} - |\cdot|_{\mathtt{b}}$ to zero as $\mathtt{a}y\mathtt{b}z$ or $\mathtt{b}y\mathtt{a}z$ with y, z balanced and possibly empty. With $A \to (\mathtt{u}, \mathtt{w}) \mid (\mathtt{u}\xi_1^1\mathtt{v}, \mathtt{w}\xi_1^2)(A)$ we get $L_A = \{(\mathtt{u}^n\mathtt{v}^{n-1}, \mathtt{w}^n)\}$, and the rules $S \to (\alpha\#\xi_1^1\$\xi_2^1\mathtt{v}\xi_2^2)(B, A)$, $\alpha \in P_{\mathrm{intr}}$, generate L_{intr}. The grammar is reduced and good (the components of A are separated by $\mathtt{v}$), of fan-out and rank two.

Not context-free. Intersecting with $\mathtt{lpc}\#\mathtt{ab}\$\mathtt{u}^+\mathtt{v}^+\mathtt{w}^+$ and erasing the prefix gives T_3^+.

Typed membership. $L_{\mathrm{intr}} = Q_{\mathrm{intr}} \cap \Phi^{-1}(0)$, where $Q_{\mathrm{intr}} := P_{\mathrm{intr}}\#\{\mathtt{a}, \mathtt{b}\}^+\$\,\mathtt{u}^+\mathtt{v}^+\mathtt{w}^+$ is strictly 2-local, with pairs $\vdash \mathtt{l}$, $\vdash \mathtt{r}$, $\mathtt{lp}, \mathtt{lq}, \mathtt{rp}, \mathtt{pc}, \mathtt{pd}, \mathtt{qc}, \mathtt{c}\#, \mathtt{d}\#$, $\#x$, xy, $x\$$ $(x, y \in \{\mathtt{a}, \mathtt{b}\})$, $\$\mathtt{u}, \mathtt{uu}, \mathtt{uv}, \mathtt{vv}, \mathtt{vw}, \mathtt{ww}, \mathtt{w} \dashv$ (the first six letter pairs admit exactly the words of P_{intr}), and $\Phi : \Sigma_{\mathrm{intr}}^* \to \mathbb{Z}^3$ maps $\mathtt{a}, \mathtt{b}, \mathtt{u}, \mathtt{v}, \mathtt{w}$ to $(1,0,0), (-1,0,0), (0,1,0), (0,-1,1), (0,0,-1)$ and other letters to 0. Lemma 6.2 and Proposition 2.9 give $L_{\mathrm{intr}} \in \mathcal{C}^{\mathrm{mcf}}_{f, s_{1,1}}$ for $f \geq 2$.

Outside $\mathsf{YSub}(f)$*.* For $x = \mathtt{p}$, $y = \mathtt{q}$, $E = \mathtt{l}\square_1\mathtt{c}\#\mathtt{ab}\$\mathtt{uvw}$, $F = \mathtt{r}\square_1\mathtt{d}\#\mathtt{ab}\$\mathtt{uvw}$: $E[x], E[y], F[x] \in L_{\mathrm{intr}}$ but $F[y] \notin L_{\mathrm{intr}}$; the boundary typing separates x and y.

Rank one is impossible. $\mathsf{MCFL}(f', 1)$ is closed under letter-to-word homomorphisms (apply them to the templates). If $L_{\mathrm{intr}} \in \mathsf{MCFL}(f', 1)$, erasing all letters except $\mathtt{a}, \mathtt{b}$ would put MIX_2^+ in $\mathsf{MCFL}(f', 1)$, hence make it EDT0L by Lemma 6.10, and then also $\mathrm{MIX}_2 = \mathrm{MIX}_2^+ \cup \{\varepsilon\}$ (EDT0L is closed under union). But MIX_2 is the word problem of $\mathbb{Z}$, which is not EDT0L [22]; Lemma 6.10 agrees with Gallot's characterization of non-branching MCFGs [21, Definition 38 and Theorems 28–29]. □

The two lower bounds come from separate factors, and the non-branching lower bound is prior work. What is specific here is that a strictly local wrapper and the $(1, 1)$ boundary typing place the target inside the typed class while the substitution implication keeps it outside Yoshinaka's baseline.

7 Conclusion

This paper turns the qualitative positive-data learning principle for typed MCFGs into upper bounds on the *total encoded length* of characteristic samples, measured against a given presentation. For a reduced good grammar G and fixed alphabet, fan-out, rank and typing:

- window and window–finite-group typings give $O(\|G\|^2\bar{\theta}_G)$, for noncommutative groups and interleaved tuple components alike (Theorems 4.9 and 4.10);

- local-group typings give $O(\|G\|^3\bar{\theta}_G)$, by combining classical sliding-factor coordinates with owner-wise block shortening (Theorem 4.15);
- rank one gives $O(\|G\|^3)$ under every typing (Theorem 5.1), and arbitrary standard input grammars are covered through component-pattern thickness (Corollary 4.20).

The refinement-thickness dichotomy (Corollary 4.16) places the exact boundary for polynomial thickness of the full yield-type refinement at **LG**, at every fan-out; it is not a statement about characteristic-data length. The feature-licensing family of Section 6.1 shows that a single fixed typing admits non-context-free cross-serial targets for which untyped and boundary-only substitution are unsound. No lower bound or separating family from the related CFG manuscript is used.

Open Problem 7.1 (Characteristic-data length at higher fan-out). Fix $f \geq 2$ and $r_{\max} \geq 2$. Is there a typing s with $\mathsf{S} \notin \mathbf{LG}$ and a polynomial q such that, for every $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ and every reduced good fan-out-$\leq f$, rank-$\leq r_{\max}$ presentation G of $L \cap \Sigma^+$, the *fixed* operator $\mathcal{B}^{\leq r_{\max}}_{f,s}$ has a characteristic sample of total length at most $q(\|G\|, \bar{\theta}_G)$?

For the opposite direction, take a fixed surjective typing s with positive image U_1. Do there exist fixed $f \geq 2$, $r_{\max} \geq 2$ and targets $L_n \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ with reduced good presentations G_n of fan-out $\leq f$ and rank $\leq r_{\max}$, such that $\|G_n\|\bar{\theta}_{G_n} = n^{O(1)}$, but *every* characteristic sample for $\mathcal{B}^{\leq r_{\max}}_{f,s}$ on L_n has total length $n^{\omega(1)}$? This would be a lower bound for the specified operator, not automatically for every set-driven learner. The refinement-thickness dichotomy (Corollary 4.16) does not answer either characteristic-data question.

Thus the outstanding issue concerns characteristic-data complexity for a fixed reconstruction operator across a typed target class, not whether an individual grammar has a thick full yield-type refinement.

Acknowledgments

The author thanks Makoto Kanazawa and Ryo Yoshinaka for valuable discussions and encouragement.

Declaration of competing interest

The author declares that he has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data availability

No empirical dataset is associated with this theoretical paper.

A Simultaneous Block Replacement in Finite Groups

This appendix proves Proposition 4.6, used in Lemmas 4.7 and B.6. The abelian case is a zero-sum argument. The general case never commutes factors; it has three parts: an order classification for homogeneous point configurations (Lemma A.4), an interval selection by boundary states (Lemma A.5), and a Ramsey-type extraction of homogeneous indices (Lemma A.7).

Definition A.1 (Block-replacement family)**.** Let A be a finite group and $f \geq 1$. A *block-replacement family* of width f over A with index set $[m]$ consists of $d \leq f$ *lines*; on line c, a left-to-right sequence of blocks $\beta_{c,1}, \ldots, \beta_{c,q_c}$ and constants $a_{c,0}, \ldots, a_{c,q_c} \in A$; for every block β an *owner* $o(\beta) \in [m]$ and two values $x_\beta, x'_\beta \in A$; every index owns at most f blocks. For $J \subseteq [m]$ put

$$\mathrm{val}_c(J) := a_{c,0}\, y_{c,1}\, a_{c,1} \cdots y_{c,q_c}\, a_{c,q_c}, \qquad y_{c,j} := \begin{cases} x'_{\beta_{c,j}}, & o(\beta_{c,j}) \in J, \\ x_{\beta_{c,j}}, & o(\beta_{c,j}) \notin J. \end{cases}$$

Lemma A.2 (Block replacement)**.** *In a block-replacement family of width f over A there is a nonempty $J \subseteq [m]$ with $\mathrm{val}_c(J) = \mathrm{val}_c(\emptyset)$ for every line c, provided that either* (a) *A is abelian and $m > \mathsf{d}(A^f)$, or* (b) *A is arbitrary and $m \geq M_{A,f}$, with d and $M_{A,f}$ as in Section 4.3.*

Proof of (a). Write A additively. By commutativity $\mathrm{val}_c(J) - \mathrm{val}_c(\emptyset) = \sum_{i \in J} \delta_{i,c}$, where $\delta_{i,c}$ is the sum of $x'_\beta - x_\beta$ over the blocks β of line c owned by i. The sequence $(\delta_i)_{i \in [m]}$ with $\delta_i := (\delta_{i,c})_c \in A^d$ has length $m > \mathsf{d}(A^f) \geq \mathsf{d}(A^d)$ (A^d embeds in A^f), so it has a nonempty zero-sum subsequence $(\delta_i)_{i \in J}$. □

Remark A.3 (Size of the zero-sum constant). The prefix-sum pigeonhole gives $\mathsf{d}(A^f) \leq |A|^f - 1$, but much more is true. Write $A \cong C_{n_1} \oplus \cdots \oplus C_{n_s}$ and $n = \exp(A) \geq 2$. Then

$$f \sum_{i=1}^{s} (n_i - 1) \;\leq\; \mathsf{d}(A^f) \;\leq\; (n-1) + n \log \frac{|A|^f}{n},$$

where the lower bound is attained by a basis sequence and the upper bound is due to van Emde Boas and Kruyswijk [33, Proposition 5.1.8 and Theorem 5.5.5]; for a p-group the lower bound is an equality [33, Theorem 5.5.9]. Hence $\mathsf{d}(A^f) = O(\exp(A)\, f \log |A|)$; for $A = C_2$ the antichain bound drops from $2^f - 1$ to f. Since zero-sum-free sequences of length $\mathsf{d}(A^d)$ exist, the argument cannot do better; we do not claim that a grammar realizes them.

A.1 Arbitrary Finite Groups: Order Homogenization

Lemma A.4 (A homogeneous block-order lemma)**.** *Let $e \geq 1$ and $n \geq 3$, and let (t, i), $t \in [e]$, $i \in [n]$, be distinct points on a line. Suppose that* (i) *the order of (t, i) and (u, i) does not depend on i, and* (ii) *for each ordered pair (t, u), including $t = u$, the order of (t, i) and (u, j) is the same for all $i < j$. Then $[e]$ is partitioned into classes such that: the points of each class occupy a consecutive region containing no other points; each class has a direction (increasing or decreasing in i); and its region consists of n successive* chunks*, one per index in that direction, each containing exactly one point of every type of the class, in an order independent of the index.*

Proof. By (ii) with $t = u$, the points of each type are monotone in i. For distinct t, u let X mean $(t, i) < (u, j)$ and Y mean $(u, i) < (t, j)$ for $i < j$; both are constant by (ii).

Both increasing. $X \wedge Y$ puts all points of index i of both types before those of index $j > i$ (*interleaved by rank*). Under $X \wedge \neg Y$ all comparisons of different indices put t left of u, and so do

equal indices, since $(u,i) < (t,i)$ would give $(t,1) < (t,2) < (u,1) < (t,1)$; so t and u are *separated*. $\neg X \wedge Y$ is symmetric, and $\neg X \wedge \neg Y$ gives $(t,1) < (t,2) < (u,1) < (u,2) < (t,1)$.

t increasing, u decreasing. $X \wedge Y$ gives $(t,2) < (u,3) < (u,1) < (t,2)$ and $\neg X \wedge \neg Y$ gives $(t,3) < (u,2) < (t,1) < (t,3)$. Under $X \wedge \neg Y$ the comparisons of different indices put t left of u, and $(u,i) < (t,i)$ would give $(t,1) < (u,2) < (u,1) < (t,1)$; so they are separated, and $\neg X \wedge Y$ is symmetric. Reversing indices covers two decreasing types.

Thus two types are separated, or have the same direction and are interleaved. Interleaving is transitive: if t,u and u,v are interleaved (all increasing, say) but t,v are separated, then $(v,1) < (u,2) < (t,3) < (v,1)$ if t is left of v, and $(t,1) < (u,2) < (v,3) < (t,1)$ otherwise. So interleaving is an equivalence relation. A type outside a class is separated from every type of the class and lies on the same side of all of them, since otherwise two interleaved types would be separated; hence class regions contain no other points. Inside a class, interleaving gives the chunks, and (i) gives the within-chunk order. □

Lemma A.5 (Simultaneous finite-group interval shortening)**.** *Let A be a finite group and consider a block-replacement family of width f over A with $n \geq \max\{3, |A|^f\}$ indices, in which all indices have the same number of blocks, on the same lines, in the same within-index order, and in which, for every two block positions t,u on a common line, the truth value of "block t of index i lies left of block u of index j" is the same for all $i < j$. Then a nonempty interval J of indices satisfies* $\mathrm{val}_c(J) = \mathrm{val}_c(\emptyset)$ *for every line c.*

Proof. Apply Lemma A.4 on each line; there are $q \leq f$ class regions in total. Fix a region $\mathcal{Q}$ and number its chunks $1,\dots,n$ from left to right. Assign constants between chunks to the chunk on their left, and absorb fixed constants at the ends of the region into its first or last chunk. Each chunk has original and replacement products with the *same* fixed constants in the same order. Let P_p be the original value of chunks $1,\dots,p$ and Q_p their value when all of them are replaced ($P_0 = Q_0 = 1_A$), and put $\kappa_p^{\mathcal{Q}} := P_p Q_p^{-1}$. Replacing exactly chunks $p+1,\dots,p'$ gives the ordered value $P_p\,(Q_p^{-1}Q_{p'})\,(P_{p'}^{-1}P_n)$, which equals P_n iff $\kappa_p^{\mathcal{Q}} = \kappa_{p'}^{\mathcal{Q}}$; no factors are commuted. If $\mathcal{Q}$ is increasing, the owners in $(a,b]$ occupy exactly chunks $(a,b]$; if it is decreasing, chunks $(n-b,n-a]$. For $j = 0,\dots,n$ let the state be the vector in A^q whose $\mathcal{Q}$-entry is $\kappa_j^{\mathcal{Q}}$ for increasing and $\kappa_{n-j}^{\mathcal{Q}}$ for decreasing regions. Since $n+1 > |A|^q$, two states coincide, at $a < b$. Replacing the owners $a+1,\dots,b$ then preserves every region value, and nothing outside the regions changes. □

Example A.6 (Opposite orientations and noncommuting values). In S_3 let $\alpha = (12)$, $\beta = (23)$. Three owners have blocks U_i, V_i on row 1 and W_i on row 2:

$$\text{row 1: } (U_1\beta V_1)\,\alpha\,(U_2\beta V_2)\,\alpha\,(U_3\beta V_3), \qquad \text{row 2: } W_3\,\beta\,W_2\,\beta\,W_1,$$

so row 1 is one increasing class and row 2 one decreasing class. Originally $U_i = \beta$, $V_i = W_i = 1$, and both rows have value 1. Replace exactly owners $J = \{2,3\}$ by $U_2' = 1$, $W_2' = \beta\alpha\beta$, $U_3' = \alpha\beta\alpha\beta$, $W_3' = \alpha$, $V_i' = 1$. Then row 1 becomes $\alpha \cdot \beta\alpha \cdot \alpha\beta\alpha = 1$ and row 2 becomes $\alpha\beta \cdot \beta\alpha\beta \cdot \beta = 1$, computed in order. In the notation of the proof, $(a,b] = (1,3]$ and $\kappa_1 = \kappa_3 = 1$ on row 1, while $\kappa_{3-b} = \kappa_0 = 1 = \kappa_2 = \kappa_{3-a}$ on row 2: one owner interval matches two different boundary-index orientations, although no owner preserves its own contribution.

Lemma A.7 (Polynomial order homogenization)**.** *Let $N \geq 3$, $m_0 := N$ and $m_{s+1} := (m_s - 1)^2 + 1$, so $m_k \leq N^{2^k}$. Let T be a finite set of types, each assigned to a line, and for $i \in [n]$ and $t \in T$ let $p_t(i)$ be a point on the line of t, the points on each line being distinct. Let Π be the set of ordered pairs (t,u), $t = u$ allowed, of types on a common line, and $k := |\Pi|$. If $n \geq m_k$, there is $I \subseteq [n]$ with $|I| = N$ such that for every $(t,u) \in \Pi$ the truth value of $p_t(i) < p_u(j)$ is the same for all $i < j$ in I.*

Proof. Treat the pairs one at a time, the pairs (t,t) first; each step passes from an index set of size at least $(m-1)^2+1$ to a subset of size at least m, and earlier homogeneity persists in subsets. For (t,t), Erdős–Szekeres gives a monotone subsequence of length m. For $t \neq u$ (both now monotone), relabel the current indices $1,\dots,m'$ and put $X(i,j) := [p_t(i) < p_u(j)]$ for $i<j$; rows and columns of X are monotone, with directions fixed by those of p_u and p_t. Replacing X by $1-X$ if necessary, rows are nondecreasing in j.

If columns are nonincreasing in i, let $J(i) := \min\{j > i : X(i,j) = 1\}$ (∞ if none); J is nondecreasing. The chain $i_1 = 1$, $i_{s+1} = J(i_s)$ satisfies $X(i_a, i_b) = 1$ for $a < b$; if it has m terms we are done. Otherwise its at most $m-1$ terms cut $[m']$ into at most $m-1$ intervals $[i_s, i_{s+1})$, $[i_L, m']$, one of which has at least m indices, and inside it $j < J(i_s) \le J(i)$ gives $X(i,j) = 0$.

If columns are nondecreasing, let κ be the least index with $X(\kappa, \kappa+1) = 1$ ($\kappa = m'$ if none). Then $X(i,j) = 1$ for $\kappa \le i < j$ and $X(i,j) \le X(j-1,j) = 0$ for $i < j \le \kappa$, and one of $[1,\kappa]$, $[\kappa, m']$ has at least $(m'+1)/2 \ge m$ indices because $(m-1)^2 + 1 \ge 2m-1$. □

Proof of Lemma A.2(b). An index owning no block can be chosen alone, so assume every index owns a block. The *pattern* of an index records the number of its blocks, the line of each, and their order on shared lines; there are at most $f \cdot f^f \cdot f! = U_f$ patterns, so at least $N_{A,f}^{2^{f^2}}$ indices share one. Absorb the blocks of all other indices into the constants. The blocks of the common pattern give at most f types and at most f^2 ordered pairs on common lines, so Lemma A.7 with $N = N_{A,f}$ yields $N_{A,f}$ indices satisfying the hypotheses of Lemma A.5, which gives the interval. □

For fixed f, $M_{A,f} = U_f \max\{3, |A|^f\}^{2^{f^2}}$ is polynomial in $|A|$; its degree is doubly exponential in f.

B Local-Group Coordinates and MCFG Block Shortening

LG is always tested on the positive image $S = \widehat{s}(\Sigma^+)$. The classical identity $\mathbf{LG} = \mathbf{Grp} * \mathbf{D}$ places it within the sliding-window and delay theory [37, 38], and the exclusion of $U_1 = \{1, \zeta\}$ ($\zeta^2 = \zeta$) is [36, Proposition 2.3]. The coordinates below are a self-contained implementation of this algebra, not a new characterization; what is specific to MCFG learning is the factorization of group values into fixed separators and replaceable blocks (Lemma B.2). Throughout, $k \ge 1$, A is a finite group, $\lambda_{k,\nu}$ is as in Definition 4.14, a word is *short* if $|w| < 2k$ (recorded exactly by $\widehat{s}_{k,k}$) and *long* otherwise, and for $s, p \in \Sigma^k$ we put $J_\nu(s,p) := \prod_{i=1}^{k} \nu\big((sp)[i..i+k]\big)$, the ordered product over the k factors of length $k+1$ of sp that start in s.

Lemma B.1 (Local-group typings are morphisms)**.** *$\lambda_{k,\nu}(uv)$ is a function of $\lambda_{k,\nu}(u)$ and $\lambda_{k,\nu}(v)$; hence $\lambda_{k,\nu}$ is a morphism onto a finite monoid $M_{k,\nu}$ with at most $|\Sigma|^{<2k} + |\Sigma|^{2k}|A|$ elements. For long u, v,*

$$\lambda_{k,\nu}(uv) = \big(\langle \mathrm{pref}_k\, u, \mathrm{suf}_k\, v\rangle,\ \gamma_{k,\nu}(u)\, J_\nu(\mathrm{suf}_k\, u, \mathrm{pref}_k\, v)\, \gamma_{k,\nu}(v)\big).$$

Proof. The window component is a morphism. For long u, v the factors of length $k+1$ of uv are, in order, those of u, the k factors starting in the last k positions of u (inside $\mathrm{suf}_k(u)\,\mathrm{pref}_k(v)$), and those of v. If u is short and v long, the factors starting in u lie inside $u\,\mathrm{pref}_k(v)$ and the others are those of v; the case of long u and short v is symmetric; and if both are short, uv is determined exactly. □

Lemma B.2 (Block decomposition)**.** *Let $w = x_0\beta_1 x_1 \cdots \beta_q x_q$ and $w' = x_0\beta'_1 x_1 \cdots \beta'_q x_q$ with $\widehat{s}_{k,k}(\beta'_j) = \widehat{s}_{k,k}(\beta_j)$ for all j. Then $\widehat{s}_{k,k}(w') = \widehat{s}_{k,k}(w)$, and if w is long there are $c_0, \dots, c_q \in A$, depending only on the x_j and the window types of the β_j, with*

$$\gamma_{k,\nu}(w) = c_0\,\gamma_{k,\nu}(\beta_1)\,c_1 \cdots \gamma_{k,\nu}(\beta_q)\,c_q, \qquad \gamma_{k,\nu}(w') = c_0\,\gamma_{k,\nu}(\beta'_1)\,c_1 \cdots \gamma_{k,\nu}(\beta'_q)\,c_q.$$

Proof. The window claim holds because $\widehat{s}_{k,k}$ is a morphism. A short block equals its replacement and is absorbed into the fixed words, so let all blocks be long. A factor of length $k+1$ either lies inside some block or meets every block it touches in a prefix or suffix of length at most k; factors of the second kind are spelled by the fixed words and the first and last k letters of the blocks, so they are the same in w and w' and occur in the same order. Ordered by starting position, the factors inside β_j form one run with product $\gamma_{k,\nu}(\beta_j)$; let c_j be the ordered product of the factors between consecutive runs. □

Lemma B.3 (Sliding-factor coordinates)**.** *Let S be a finite semigroup whose idempotents all lie in its minimal ideal $\mathsf{K} = \mathcal{M}(H; I, \Lambda; P)$ (a Rees matrix semigroup with P over H), let $n := |S|$, and let $h : \Sigma^+ \to S$ be a morphism. There is $\nu : \Sigma^{n+1} \to H$ such that every w with $|w| \geq 2n$ satisfies*

$$h(w) = \big(\mathrm{row}(h(\mathrm{pref}_n\, w)),\ \mathrm{grp}(h(\mathrm{pref}_n\, w))\,\gamma_{n,\nu}(w),\ \mathrm{col}(h(\mathrm{suf}_n\, w))\big), \tag{$\star$}$$

where row, grp, col *are the coordinates of $(i, a, \lambda) \in \mathsf{K}$. In particular $h(w)$ is determined by its first and last n letters and an ordered product in H.*

Proof. (1) Products of n letters lie in K. The Rees quotient S/K has no idempotent but 0, so each element has a power equal to 0. If a product of $m = |S/\mathsf{K}| \leq n$ elements were nonzero, two of its m nonzero prefix products would coincide, $p_i = p_j$ with $i < j$; then $p_i = p_i x^q$ for $x = s_{i+1} \cdots s_j$ and all q, impossible for $x^q = 0$.

(2) Coordinates. With $(i, a, \lambda)(j, b, \mu) = (i, ap_{\lambda j}b, \mu)$, fix $i_0 \in I$ and $e_\lambda := (i_0, p^{-1}_{\lambda i_0}, \lambda)$, a right identity for elements with Λ-coordinate λ. For $s \in S$ write $e_\lambda s = (j, c, \mu)$ and $\rho(\lambda, s) := p_{\lambda j}c$; then $(i, a, \lambda)s = (i, a\rho(\lambda, s), \mu)$, so right multiplication fixes the I-coordinate. Dually, left multiplication fixes the Λ-coordinate of an element of K (write $s(i, a, \lambda) = (sf)(i, a, \lambda)$ for an idempotent $f \in \mathsf{K}$ fixing (i, a, λ) on the left).

(3) Group coordinate. Let $w = a_1 \cdots a_m$, $m \geq 2n$, and $h(a_1 \cdots a_n) = (i, a, \lambda_n)$. For $j \geq n$ the Λ-coordinate λ_j of $h(a_1 \cdots a_j)$ equals that of $h(a_{j-n+1} \cdots a_j)$ by (1) and (2). Put $\nu(F) := \rho(\mathrm{col}(h(F[1..n])), h(F[n+1]))$. Applying the right-action identity for $j = n, \dots, m-1$ in order gives $h(w) = (i,\, a \prod_{j=n}^{m-1} \nu(a_{j-n+1} \cdots a_{j+1}),\, \lambda_m)$, and the product is $\gamma_{n,\nu}(w)$. □

Theorem B.4 (Local-group criteria in sliding-factor coordinates)**.** *For a typing s with positive image S the following are equivalent:* (i) *$S \in \mathbf{LG}$;* (ii) *S has no subsemigroup isomorphic to U_1;* (iii) *every idempotent of S lies in the minimal ideal of S;* (iv) *s is local-group-determined, and one may take $k = |S|$.*

Proof. By [36, Proposition 2.3], (i) and (ii) are equivalent to $S^m = \mathsf{K}$ for some m. If $S^m = \mathsf{K}$ then $e = e^m \in \mathsf{K}$ for each idempotent e; conversely, if $E(S) \subseteq \mathsf{K}$ then $S^n = SE(S)S \subseteq \mathsf{K}$ for $n = |S|$ [36, Lemma 2.2], and $S^n \supseteq \mathsf{K}$ since S^n is an ideal. So (i)–(iii) are equivalent.

(iii)⇒(iv): Lemma B.3 with $h = s$ and $n = |S|$ determines $s(w)$ for $|w| \geq 2n$ from $\lambda_{n,\nu}(w)$, and $\widehat{s}_{n,n}$ records shorter words and ε exactly; so $\widehat{s} \preceq \lambda_{n,\nu}$.

(iv)⇒(i): S is a homomorphic image of $T := \lambda_{k,\nu}(\Sigma^+)$ and $\mathbf{LG}$ is closed under homomorphic images, so it suffices that $T \in \mathbf{LG}$. No short element is idempotent (squaring changes the exactly recorded word or makes it long). Let $e = (\langle p, s\rangle, \gamma_e)$ be idempotent. By Lemma B.1, every element

of eTe has window $\langle p, s\rangle$, and long elements with this window multiply by $(\langle p, s\rangle, \alpha)(\langle p, s\rangle, \beta) = (\langle p, s\rangle, \alpha J \beta)$, $J := J_\nu(s, p)$. This operation makes $\{\langle p, s\rangle\} \times A$ a group isomorphic to A via $\alpha \mapsto \alpha J$, and eTe is a finite nonempty subsemigroup of it, hence a subgroup. □

Proposition B.5 (Window–group typings are local-group typings, strictly)**.** (i) *If $\widehat{s} \preceq (\widehat{s}_{k,\ell}, g)$ for a morphism g into a finite group, then $\widehat{s} \preceq \lambda_{\hat{k},\nu}$ with $\hat{k} = \max\{k, \ell, 1\}$ and $\nu(F) := g(\text{last letter of } F)$. Hence* $\mathbf{LI} \vee \mathbf{Grp} \subseteq \mathbf{LG}$.

(ii) *For $\Sigma = \{\mathtt{a}, \mathtt{b}, \mathtt{c}\}$ and $h_{\mathtt{ab}} := \lambda_{1,\nu}$ with $\nu(\mathtt{ab}) = 1 \in \mathbb{Z}_2$ and $\nu = 0$ otherwise, $h_{\mathtt{ab}}(\Sigma^+) \in \mathbf{LG} \setminus (\mathbf{LI} \vee \mathbf{Grp})$.*

Proof. (i) For $|w| \geq 2\hat{k}$, $g(w) = g(\mathrm{pref}_{\hat{k}}\, w)\, \gamma_{\hat{k},\nu}(w)$, and $\mathrm{pref}_{\hat{k}}, \mathrm{suf}_{\hat{k}}$ determine $\widehat{s}_{k,\ell}(w)$; shorter words are recorded exactly. For the inclusion, a semigroup $S \in \mathbf{LI} \vee \mathbf{Grp}$ is the positive image of the evaluation typing over the alphabet S, which is window–group-determined by Lemma 4.12, hence local-group-determined, so $S \in \mathbf{LG}$ by Theorem B.4. (ii) Membership in $\mathbf{LG}$ is Theorem B.4. The words u, v in the proof of Proposition 6.7(iii) lie in $\{\mathtt{a}, \mathtt{b}, \mathtt{c}\}^+$, have equal $(\widehat{s}_{k,\ell}, g')$-values for any given k, ℓ, g', and differ in the parity of $\#_{\mathtt{ab}}$; so $h_{\mathtt{ab}}$ is not window–group-determined, and Lemma 4.12 excludes $\mathbf{LI} \vee \mathbf{Grp}$. □

Lemma B.6 (Short tuples with prescribed local-group type)**.** *Let G be reduced and good of fan-out $\leq f$ and rank $\leq r_{\max}$, and $\lambda = \lambda_{k,\nu}$. Let $b_{\mathrm{w}} := b_{k,k} = (r_{\max} + 1)\bar{\theta}_G\, f(2k+1)\, |V||M_{k,k}|^f$ (Corollary 4.8) and $n_\lambda := |V||M_{k,\nu}|^f$. For every $B \in V$ and every realizable $\mathbf{t} \in M_{k,\nu}^{\mathrm{fan}(B)}$ there is $\vec{u} \in L_B(G)$ with $\lambda^{(\mathrm{fan}(B))}(\vec{u}) = \mathbf{t}$ and*

$$\|\vec{u}\|_1 \leq \beta_\lambda := (r_{\max} + 1)\, b_{\mathrm{w}} \max\{1, c_{A,f}\, n_\lambda\}.$$

Proof. Let T be λ-minimal for $(B, \mathbf{t})$ with heavy threshold $b = b_{\mathrm{w}} \geq \bar{\theta}_G$. For a heavy node v with nonterminal C and yield $\vec{x}$, Corollary 4.8 gives $\vec{x}' \in L_C(G)$ with the same $\widehat{s}_{k,k}$-vector and $\|\vec{x}'\|_1 \leq b_{\mathrm{w}}$. For an antichain $v_1, \ldots, v_m$ of heavy nodes and $J \subseteq [m]$, let T_J replace the subtrees at v_i, $i \in J$; their components are disjoint contiguous blocks of the root components. By Lemma B.2, every T_J has the root window vector of T, short root components are unchanged, and on each long root component $\gamma_{k,\nu}$ is an ordered product of block values between constants independent of J: a block-replacement family over A with owners $v_1, \ldots, v_m$. If $m > c_{A,f}$, Lemma A.2 gives a nonempty J with λ-vector $\mathbf{t}$, and T_J is strictly shorter, contradicting minimality. Lemma 4.5 with $c = c_{A,f}$ and $N = M_{k,\nu}$ gives the bound. □

C Pattern-Thickness Transfer for General MCFG Presentations

Proof of $t_G \leq t_{\mathrm{pat}}(H_{\mathrm{red}})$ in Proposition 4.19. For a nonterminal A of H_{red} and a realized pattern $I \in \mathsf{Pat}(A)$ let $F_{A,I} := \{\vec{u} \in L_A(H_{\mathrm{red}}) : \mathrm{supp}(\vec{u}) = I\}$; it contains a tuple of length at most $t_{\mathrm{pat}}(H_{\mathrm{red}})$, and $L_A = \bigcup_I F_{A,I}$. Call a map *admissible* if it is a composition of projections π_J, permutations σ and merges μ; it does not increase total length, and the support of $\phi(\vec{u})$ depends only on that of $\vec{u}$. Follow the stages of Lemma F.1 applied to H_{red}. After (a) and (b), every nonterminal X has $L_X = \phi_X(L_A)$ for an original A and an admissible ϕ_X. Stage (c) creates $X^{I'}$, for realized I', with $L_{X^{I'}} = \pi_{I'}\{\vec{u} \in L_X : \mathrm{supp}(\vec{u}) = I'\}$; as supports of images depend only on supports, this contains $\pi_{I'}\phi_X(F_{A,I})$ for some realized I. Stage (d) replaces languages by their images under merges, and final trimming does not change surviving languages. Hence every nonterminal Y of G satisfies $\psi(F_{A,I}) \subseteq L_Y$ for some A, realized I and admissible ψ, so $\min_{\vec{u} \in L_Y} \|\vec{u}\|_1 \leq t_{\mathrm{pat}}(H_{\mathrm{red}})$. □

Proposition C.1 (Thickness under good-form conversion of λ-free inputs)**.** *Let G_{in} be a reduced λ-free fan-out-$\leq f$, rank-$\leq r_{\max}$ MCFG, possibly with deleting rules, with $L(G_{\mathrm{in}}) \cap \Sigma^+ \neq \emptyset$. The good normalization yields a reduced good G with the same bounds and*

$$\|G\| = O_{f,r_{\max}}(\|G_{\mathrm{in}}\|), \qquad t_G \leq t_{G_{\mathrm{in}}}, \qquad \bar{\theta}_G = O_{f,r_{\max}}(\|G_{\mathrm{in}}\| \max\{1, t_{G_{\mathrm{in}}}\}).$$

Proof. For reduced λ-free G_{in} only full patterns occur, so $t_{\mathrm{pat}}(G_{\mathrm{in}}) = t_{G_{\mathrm{in}}}$; apply Proposition 4.19. □

Thus deleting rules alone do not obstruct thickness transfer; genuinely empty components can (Proposition F.2).

D Auxiliary Verification of Qualitative Reconstruction

Proof of Lemma 3.2. Put $\vec{z}_j := \rho(\vec{u}_1, \ldots, \vec{u}_j, \vec{x}_{j+1}, \ldots, \vec{x}_r)$. For $E \in \mathcal{E}_e$, Lemma 3.1 gives an admissible E_j with $E[\vec{z}_{j-1}] = E_j[\vec{x}_j]$ and $E[\vec{z}_j] = E_j[\vec{u}_j]$; since $\vec{u}_j \equiv_{L,s} \vec{x}_j$, E accepts $\vec{z}_{j-1}$ iff it accepts $\vec{z}_j$. So $\mathcal{D}_L(\vec{z}_0) = \mathcal{D}_L(\vec{z}_r)$. The types agree because $\widehat{s}^{(e)}(\rho(\cdots)) = \mathrm{out}_\rho$ of the children's types. □

Proof of Proposition 3.10. Induction on the height of a derivation from $[\vec{x}]$. Const derives $\vec{x}$. Link $[\vec{x}] \to [\vec{y}]$ has same-type $\vec{x}, \vec{y}$ with a common context in $K \subseteq L$, so $\vec{x} \equiv_{L,s} \vec{y}$ by substitutability; combine with the hypothesis for $[\vec{y}]$. For Comp $[\vec{z}] \to \rho([\vec{u}_1], \ldots, [\vec{u}_r])$ with $\vec{z} = \rho(\vec{u}_1, \ldots, \vec{u}_r)$, the children derive tuples equivalent to the $\vec{u}_j$, and Lemma 3.2 transfers this to $\vec{z}$. For $\widehat{S} \to [(w)]$ with $w \in K \subseteq L$, a derived (v) satisfies $(v) \equiv_{L,s} (w)$, and $\square_1$ accepts (w), so $v \in L$; $\widehat{S} \to \varepsilon$ occurs only if $\varepsilon \in K \subseteq L$. The uncapped grammar is the case of cap $\|K\|_+$. □

Proof of Proposition 3.11. This is the argument of [15, Lemma 3] for the concrete grammar. Each $\iota(X)$ and each $\vec{z}_R := \rho(\iota(Y_1), \ldots, \iota(Y_q))$ for a typed rule $R : X \to \rho(Y_1, \ldots, Y_q)$ is observed, via $\chi(X)$. They have the same type (out_ρ of the children's types is the type of X) and the common context $\chi(X)$, so Link gives $[\iota(X)] \to [\vec{z}_R]$. If $q = 0$, Const gives $[\vec{z}_R] \to \vec{z}_R$; otherwise $(\vec{z}_R, \rho, (\iota(Y_j))_j)$ is a composition witness of rank $q \leq r_G$ (and $q \leq \|K\|_+$ by (3.1)), so Comp gives $[\vec{z}_R] \to \rho([\iota(Y_1)], \ldots, [\iota(Y_q)])$. By induction, $X \Rightarrow^*_{\widetilde{G}_0} \vec{u}$ implies $[\iota(X)] \Rightarrow^* \vec{u}$. For $w \in L(G) = L(\widetilde{G}_0)$ (Proposition 3.6(iii)) a successful derivation starts with a start link $\widetilde{S} \to X$; then $\chi(X) = \square_1$, so $\iota(X) = (w')$ with $w' \in K$, and Start gives $\widehat{S} \to [\iota(X)] \Rightarrow^* (w)$. □

Proof of Theorem 3.13. If $L \cap \Sigma^+ = \emptyset$, then $K = L \subseteq \{\varepsilon\}$ and Definition 3.4 gives the claim. Otherwise $\mathcal{F}_L \subseteq L$ is finite, Proposition 3.11 with $\iota = \omega$ (canonical tuples have the right type by Proposition 3.6(i)) gives $L \cap \Sigma^+ \subseteq L(\mathcal{B}(K))$, Proposition 3.10 gives $L(\mathcal{B}(K)) \subseteq L$, and $\varepsilon \in L(\mathcal{B}(K))$ iff $\varepsilon \in K$ iff $\varepsilon \in L$. Rank-preserving normalization gives the last sentence. □

Proof of Theorem 3.14. Put $n_K = \max\{1, \|K\|_+\}$. An occurrence is a sample word with at most $2f$ cut positions, so there are $n_K^{O(f)}$ of them, each typed in linear time; grouping occurrences by context and type lists the Link rules. A composition witness of rank $\leq r_{\max}$ is a parent occurrence with at most $2fr_{\max}$ further cut positions and a labelling whose number depends only on $f, r_{\max}$; goodness is checked directly. This gives $n_K^{O(fr_{\max})}$ Comp rules, and all other rules are fewer. □

Proof of Corollary 3.16. Let $L \in \mathcal{C}^{\mathrm{mcf}}_{f,s}$ (for (ii), also $L \in \mathsf{MCFL}(f, r_{\max})$), let $\mathcal{F}_L$ be the characteristic sample of Theorem 3.13 for the operator used, and consider a text for L; membership in MCFLs is decidable [29], so the learner is computable. Let n_0 be a stage by which all of $\mathcal{F}_L$ has appeared. If the learner changes its hypothesis at some stage $n \geq n_0$, the new hypothesis $\mathcal{B}(K_n)$ generates exactly L, so it is never changed again. Otherwise the hypothesis H held at n_0 is final. If $H = H_0$,

every datum was generated by $\emptyset$, so $L = \emptyset$. Else $H = \mathcal{B}(K_m)$ for some $m \le n_0$, so $L(H) \subseteq L$ by Proposition 3.10; every word of L appears in the text, before m (then it lies in $K_m \subseteq L(H)$ by sample consistency) or after it (then it is generated by H, which is never changed). So $L(H) = L$. For (ii), each update is a construction (Theorem 3.14) plus a membership test, both polynomial in $\|K\|_+$ for fixed f and $r_{\max}$ [14, Proposition 4]. □

These bounds concern the work per datum. They do not bound the stage of convergence, since a text may postpone any example arbitrarily long; short characteristic data alone gives no presentation-independent convergence stage.

E Auxiliary Proofs for Typed Witness Construction

Proof of Proposition 3.6. At every rule application the morphism property gives

$$\widehat{s}^{(\mathrm{fan}(A))}(\rho(\vec{u}_1, \dots, \vec{u}_r)) = \mathrm{out}_\rho\big(\widehat{s}^{(d_1)}(\vec{u}_1), \dots, \widehat{s}^{(d_r)}(\vec{u}_r)\big).$$

Induction on height gives (i) and the existence and uniqueness of the lift in (ii), since each typed label is forced by its yield. Erasing labels gives $L(G^{\widehat{s}}) \subseteq L(G)$, (ii) gives the converse, and trimming keeps every rule of every successful derivation. □

Definitions 2.8, 3.4 and 3.5 apply verbatim with any finite-monoid morphism $h : \Sigma^* \to N$ in place of $\widehat{s}$ (tuple types being $h^{(d)}$); write $\mathcal{C}^{\mathrm{mcf}}_{f,h}$ for the class. The class and the operator use types of nonempty components only, so they depend only on $h|_{\Sigma^+}$.

Lemma E.1 (Monoid implementations)**.** *Let $h_0 : \Sigma^* \to M_0$ be a finite-monoid morphism, $S_0 := h_0(\Sigma^+)$ and $s_0 := h_0|_{\Sigma^+}$.*

(i) *On $(\Sigma^+)^d$, equality of $h_0^{(d)}$-, $s_0^{(d)}$- and $\widehat{s_0}^{(d)}$-values coincide. Hence $\mathcal{C}^{\mathrm{mcf}}_{f,h_0} = \mathcal{C}^{\mathrm{mcf}}_{f,s_0}$, and the operators built with h_0 and with $\widehat{s_0}$ return the same grammar on every sample.*

(ii) *For every good MCFG G, the trimmed refinements of G under h_0 and under $\widehat{s_0}$ coincide (the untrimmed ones need not); hence the same canonical tuples, contexts and witness words may be chosen.*

Proof. (i) The three maps agree on Σ^+. (ii) By λ-freeness and Proposition 3.6(i), every typed nonterminal of a successful derivation under either morphism is $A_{\mathbf{t}}$ with $\mathbf{t} \in S_0^{\mathrm{fan}(A)}$, and every start link used is $\widetilde{S} \to S_{(\eta)}$ with $\eta \in S_0$. For children indexed in S_0, each nonempty template component is evaluated in S_0 identically by both morphisms, so both refinements have the same rules among S_0-indexed symbols, which carry all successful derivations. □

Proof of Lemma 3.7. Take a successful typed derivation containing X and follow the path from the root to an occurrence of X, starting with $\square_1$; at each rule apply Lemma 3.1 with the (nonempty) sibling yields fixed. This gives $\chi(X) \in \mathcal{E}_d$, and replacing the X-subtree by any X-derivation gives a successful derivation of $\chi(X)[\vec{u}]$. If $\widetilde{S} \to X$ is a rule, take the path of length zero. □

Proof of Lemma 3.8. For $\vec{u}, \vec{v} \in L_X$, Proposition 3.6(i) gives equal types, $\chi(X)$ accepts both, and the arity is at most f; substitutability gives $\mathcal{D}_L(\vec{u}) = \mathcal{D}_L(\vec{v})$. □

F Details of the Good Normalization

The four conversions below give Proposition 3.12; each satisfies a tuple-language identity from which the size bound and the thickness transfer of Appendix C are read off. Correctness of the nondeleting and nonmerging conversions is standard [14, 29, 31]; the size and thickness bounds are not stated there and are derived here. For $\vec{w} = (w_1, \dots, w_d)$ and nonempty $I \subseteq [d]$, $\pi_I(\vec{w})$ keeps the components indexed by I in order, $\sigma\vec{w}$ permutes components and $\mu(\vec{w})$ concatenates consecutive blocks of components; so $\|\pi_I(\vec{w})\|_1 \le \|\vec{w}\|_1 = \|\sigma\vec{w}\|_1 = \|\mu(\vec{w})\|_1$.

Lemma F.1 (Stagewise tuple languages)**.** *Let G_0 be a reduced f-MCFG of rank at most $r_{\max}$. Each conversion below preserves linearity, does not increase fan-out, rank or any template, and multiplies the numbers of nonterminals and rules by a factor depending only on f and $r_{\max}$.*

(a) Nondeletion *(construction (f3) in the proof of [29, Lemma 2.2]). Introduce A^I for nonempty $I \subseteq [\mathrm{fan}(A)]$. For $R : A \to \rho(B_1, \dots, B_q)$ let J_j be the components of B_j used by the I-indexed outputs of ρ; keep these outputs, use $B_j^{J_j}$ if $J_j \neq \emptyset$ and omit B_j otherwise. The rules are nondeleting and $L_{A^I} = \pi_I(L_A(G_0))$.*

(b) Nonpermutation. *Decorate a fan-out-d nonterminal by $\sigma \in \mathfrak{S}_d$, the child's permutation being determined by the order in which the parent template uses its variables. The decorated nonterminal generates σL_A.*

(c) Empty components. *Let $\mathsf{Pat}(A)$ be the set of support patterns of L_A, computed by saturation. For realizable child patterns, replace variables of empty components by ε, delete empty parent components, omit all-empty children, and introduce A^I for each nonempty $I \in \mathsf{Pat}(A)$. Then $L_{A^I} = \pi_I\{\vec{u} \in L_A : \mathrm{supp}(\vec{u}) = I\}$, the result is λ-free, and $S^{\{1\}}$ generates the nonempty part of the language.*

(d) Nonmerging *(construction in the proof of [14, Section 2.3, Lemma 1]). Decorate a nonterminal by a composition μ of its fan-out, the child's composition being given by the maximal runs of its consecutive variables in the parent template. The decorated nonterminal generates the image $\mu(L_A)$.*

Stages (b)–(d) preserve the properties established by earlier stages.

Proof. (a) For any child tuples, $\pi_I(\rho(\vec{v}_1, \dots, \vec{v}_q)) = \rho^I((\pi_{J_j}\vec{v}_j)_{J_j \neq \emptyset})$; induction gives $\subseteq$, and for $\supseteq$ the omitted children are completed by terminal derivations, which exist because G_0 is reduced. There are at most 2^f copies of each nonterminal and rule. (b) Linearity and nondeletion make the child permutations well defined; induction gives σL_A. There are at most $f!$ copies. (c) Retained variables occur once and in order; induction gives the identity, and omitted all-empty children are completed by derivations of their empty pattern. There are at most 2^f patterns per nonterminal and $2^{f r_{\max}}$ choices per rule. (d) Nonpermutation makes the runs compatible with each child's order and nondeletion makes them cover its components; then $\rho^\mu(\mu_1(\vec{v}_1), \dots, \mu_q(\vec{v}_q)) = \mu(\rho(\vec{v}_1, \dots, \vec{v}_q))$, and induction gives the image identity (merging need not be injective, so no inverse-image claim is made). The size is linear for fixed $f, r_{\max}$ [14, Section 2.3, after Example 2]. □

Proof of Proposition 3.12. Apply (a)–(d) to the trimmed input and trim again. Each stage multiplies the size by a factor depending only on $f, r_{\max}$, so $\|G_{\mathrm{good}}\| = O_{f,r_{\max}}(\|G_{\mathrm{in}}\|)$, and $L(G_{\mathrm{good}}) = L(G_{\mathrm{in}}) \setminus \{\varepsilon\}$ by (c). If this language is nonempty, the final trimming leaves a reduced grammar, and trimming preserves goodness. □

Proposition F.2 (Empty-component elimination can blow up thickness)**.** *For every $n \geq 1$ there is a reduced fan-out-two, rank-two MCFG E_n of size $O(n)$ with $t_{E_n} \leq 2$ and $\theta_{E_n} = O(n)$ whose stage-(c) conversion contains a nonterminal and a rule of thickness at least 2^n; in particular $t_{\mathrm{pat}}(E_n) \geq 2^n$.*

Proof. Let E_n have rules $A_0 \to (\mathtt{a}, \varepsilon)$, $A_i \to (\varepsilon, \mathtt{c}) \mid (\xi_1^1\xi_2^1, \xi_1^2\xi_2^2)(A_{i-1}, A_{i-1})$ for $1 \leq i \leq n$, and $S \to (\xi_1^1 \mathtt{d}\, \xi_1^2)(A_n)$. It is reduced, nondeleting, nonpermuting and nonmerging. Completing every A_i by $(\varepsilon, \mathtt{c})$ (or A_0 by $(\mathtt{a}, \varepsilon)$) gives $t_{E_n} \leq 2$ and, along a path of length at most n, $\theta_{E_n} = O(n)$. An A_i-tuple of support $\{1\}$ cannot use $(\varepsilon, \mathtt{c})$, so it is $(\mathtt{a}^{2^i}, \varepsilon)$. Hence $A_n^{\{1\}}$ and the rule $S \to (\xi_1^1 \mathtt{d})(A_n^{\{1\}})$ have thickness at least 2^n. □

This differs from Proposition 4.4: here a support pattern of the source forces an exponentially long tuple, there the typed refinement does. Neither is a characteristic-data lower bound.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During preparation of this work, the author used ChatGPT (OpenAI) and Claude (Anthropic) to assist with organization of the manuscript, expression and translation of text, cross-checks of references and notation, and editing of LaTeX source. After using these tools, the author reviewed and edited the material as needed and takes full responsibility for the content of the publication.

References

[1] E. M. Gold. Language identification in the limit. *Information and Control*, 10(5):447–474, 1967.

[2] D. Angluin. Inductive inference of formal languages from positive data. *Information and Control*, 45(2):117–135, 1980.

[3] A. Clark and R. Eyraud. Polynomial identification in the limit of substitutable context-free languages. *Journal of Machine Learning Research*, 8:1725–1745, 2007.

[4] C. de la Higuera. Characteristic sets for polynomial grammatical inference. *Machine Learning*, 27:125–138, 1997.

[5] F. Coste, D. Fredouille, C. Kermorvant, and C. de la Higuera. Introducing domain and typing bias in automata inference. In *Grammatical Inference: Algorithms and Applications – ICGI 2004*, LNCS 3264, pp. 115–126, Springer, 2004.

[6] Y. Takada. Grammatical inference for even linear languages based on control sets. *Information Processing Letters*, 28(4):193–199, 1988. doi:10.1016/0020-0190(88)90208-6.

[7] Y. Takada. Learning formal languages based on control sets. In K. P. Jantke and S. Lange (eds.), *Algorithmic Learning for Knowledge-Based Systems*, LNCS 961, pp. 316–339, Springer, 1995.

[8] A. Clark. Learning trees from strings: A strong learning algorithm for some context-free grammars. *Journal of Machine Learning Research*, 14:3537–3559, 2013.

[9] F. Coste, G. Garet, and J. Nicolas. Local substitutability for sequence generalization. In *Proceedings of the 11th International Conference on Grammatical Inference*, JMLR Workshop and Conference Proceedings 21, pp. 97–111, 2012.

[10] R. Eyraud, J. Heinz, and R. Yoshinaka. Efficiency in the identification in the limit learning paradigm. In J. Heinz and J. M. Sempere (eds.), *Topics in Grammatical Inference*, pp. 25–46. Springer, 2016. doi:10.1007/978-3-662-48395-4_2.

[11] F. Coste and J. Nicolas. Learning local substitutable context-free languages from positive examples in polynomial time and data by reduction. In *Proceedings of the 14th International Conference on Grammatical Inference*, Proceedings of Machine Learning Research 93, pp. 155–168, 2019.

[12] R. Yoshinaka. Identification in the limit of k, l-substitutable context-free languages. In *Grammatical Inference: Algorithms and Applications*, LNCS 5278, pp. 266–279. Springer, 2008.

[13] R. Yoshinaka. Learning mildly context-sensitive languages with multidimensional substitutability from positive data. In *Algorithmic Learning Theory*, LNCS 5809, pp. 278–292. Springer, 2009. doi:10.1007/978-3-642-04414-4_24.

[14] R. Yoshinaka. Efficient learning of multiple context-free languages with multidimensional substitutability from positive data. *Theoretical Computer Science*, 412(19):1821–1831, 2011. doi:10.1016/j.tcs.2010.12.058.

[15] R. Yoshinaka. General perspective on distributionally learnable classes. In *Proceedings of the 14th Meeting on the Mathematics of Language (MoL 2015)*, pp. 87–98. Association for Computational Linguistics, 2015. doi:10.3115/v1/W15-2308.

[16] R. Yoshinaka. Polynomial-time identification of multiple context-free languages from positive data and membership queries. In *Grammatical Inference: Theoretical Results and Applications*, LNCS 6339, pp. 230–244. Springer, 2010. doi:10.1007/978-3-642-15488-1_19.

[17] A. Clark and R. Yoshinaka. Distributional learning of context-free and multiple context-free grammars. In J. Heinz and J. M. Sempere (eds.), *Topics in Grammatical Inference*, pp. 143–172. Springer, 2016. doi:10.1007/978-3-662-48395-4_6.

[18] A. Clark and R. Yoshinaka. Distributional learning of parallel multiple context-free grammars. *Machine Learning*, 96(1–2):5–31, 2014. doi:10.1007/s10994-013-5403-2.

[19] T. Oates, T. Armstrong, L. Becerra-Bonache, and M. Atamas. Inferring grammars for mildly context-sensitive languages in polynomial-time. In *Grammatical Inference: Algorithms and Applications – ICGI 2006*, LNCS 4201, pp. 137–147. Springer, 2006.

[20] M. Kanazawa and R. Yoshinaka. Extending distributional learning from positive data and membership queries. In *Proceedings of the 16th International Conference on Grammatical Inference*, Proceedings of Machine Learning Research 217, pp. 8–22, 2023.

[21] P. D. Gallot. *Safety of transformations of data trees: Tree transducer theory applied to a verification problem on shell scripts*. PhD thesis, Université de Lille, 2021. `https://pepite-depot.univ-lille.fr/LIBRE/EDMADIS/2021/2021LILUB030.pdf`.

[22] A. Bishop, M. Elder, A. Evetts, P. Gallot, and A. Levine. On groups with EDT0L word problem. *International Journal of Algebra and Computation*, 36(4):425–486, 2026. doi:10.1142/S0218196726500165.

[23] A. Duncan, M. Elder, L. Frenkel, and M. Lyu. A substitution lemma for multiple context-free languages. *International Journal of Algebra and Computation*, 2026 (online publication). doi:10.1142/S0218196726500463.

[24] O. Rambow and G. Satta. Independent parallelism in finite copying parallel rewriting systems. *Theoretical Computer Science*, 223(1–2):87–120, 1999. doi:10.1016/S0304-3975(97)00190-4.

[25] T. Kuriyama. Distributional learning of context-free languages under fixed finite-semigroup typing. Revised manuscript submitted to *Theoretical Computer Science*, October 2026 (TCS-D-26-00494R1). The cited Proposition 6.4 is in this revision; the earlier, substantively different public preprint is arXiv:1409.6247v4. `https://arxiv.org/abs/1409.6247`.

[26] J.-E. Pin. *Varieties of Formal Languages*. North Oxford Academic, London, and Plenum, New York, 1986.

[27] J. C. Costa. Some pseudovariety joins involving locally trivial semigroups. *Semigroup Forum*, 64(1):12–28, 2002. doi:10.1007/s002330010060.

[28] J. C. Costa. Some pseudovariety joins involving groups and locally trivial semigroups. In *Semigroups, Algorithms, Automata and Languages*, pp. 341–348. World Scientific, 2002.

[29] H. Seki, T. Matsumura, M. Fujii, and T. Kasami. On multiple context-free grammars. *Theoretical Computer Science*, 88(2):191–229, 1991. doi:10.1016/0304-3975(91)90374-B.

[30] L. Kallmeyer. *Parsing Beyond Context-Free Grammars*. Springer, Berlin, Heidelberg, 2010. doi:10.1007/978-3-642-14846-0.

[31] M. Kanazawa. Ogden's lemma, multiple context-free grammars, and the control language hierarchy. *Information and Computation*, 269:104449, 2019. doi:10.1016/j.ic.2019.104449.

[32] C. Gómez-Rodríguez, M. Kuhlmann, G. Satta, and D. Weir. Optimal reduction of rule length in linear context-free rewriting systems. In *Proceedings of Human Language Technologies: The 2009 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 539–547. Association for Computational Linguistics, 2009.

[33] A. Geroldinger and F. Halter-Koch. *Non-Unique Factorizations: Algebraic, Combinatorial and Analytic Theory*, volume 278 of *Pure and Applied Mathematics*. Chapman & Hall/CRC, Boca Raton, 2006.

[34] G. Ramalingam and T. Reps. An incremental algorithm for a generalization of the shortest-path problem. *Journal of Algorithms*, 21(2):267–305, 1996. doi:10.1006/jagm.1996.0046.

[35] M.-J. Nederhof. Weighted deductive parsing and Knuth's algorithm. *Computational Linguistics*, 29(1):135–143, 2003.

[36] J. Almeida, S. Margolis, B. Steinberg, and M. Volkov. Representation theory of finite semigroups, semigroup radicals and formal language theory. *Transactions of the American Mathematical Society*, 361(3):1429–1461, 2009. doi:10.1090/S0002-9947-08-04712-0.

[37] H. Straubing. Finite semigroup varieties of the form $\mathbf{V} * \mathbf{D}$. *Journal of Pure and Applied Algebra*, 36:53–94, 1985. doi:10.1016/0022-4049(85)90062-3.

[38] B. Tilson. Categories as algebra: an essential ingredient in the theory of monoids. *Journal of Pure and Applied Algebra*, 48:83–198, 1987. doi:10.1016/0022-4049(87)90108-3.

[39] C. H. Papadimitriou. On the complexity of integer programming. *Journal of the ACM*, 28(4):765–768, 1981. doi:10.1145/322276.322287.